

## Problem Statement:

The age of abalone is determined by cutting the shell through the cone, staining it, and counting the number of rings through a microscope -- a boring and time-consuming task. Other measurements, which are easier to obtain, are used to predict the age. Further information, such as weather patterns and location (hence food availability) may be required to solve the problem. You have to predict the rings of each abalone which will lead us to the age of that abalone.

## Importing required libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import pickle
import seaborn as sns
from scipy.stats import zscore
import scikitplot as skplt
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import power_transform, StandardScaler
from sklearn.ensemble import AdaBoostClassifier, RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

## Reading data

```
In [2]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\Data Analysis With Python\ML Files\abalone.csv")
df.head()
```

Out[2]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

## Check no of row and column

```
In [7]: print('No of Rows and Columns ----->', df.shape )
```

```
No of Rows and Columns -----> (4177, 9)
```

## Checking for Null values

```
In [9]: print('-----\n')
print(df.isnull().sum())
print('\n-----')
```

```
-----
Sex          0
Length       0
Diameter     0
Height        0
Whole weight  0
Shucked weight 0
Viscera weight 0
Shell weight   0
Rings         0
dtype: int64
-----
```

**There is no null value in our dataset**

## Information about dataset

```
In [11]: print('-----\n')
print(df.info())
print('\n-----')
```

```
-----  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4177 entries, 0 to 4176  
Data columns (total 9 columns):  
 #   Column           Non-Null Count  Dtype     
---  --    
 0   Sex              4177 non-null    object    
 1   Length           4177 non-null    float64   
 2   Diameter         4177 non-null    float64   
 3   Height           4177 non-null    float64   
 4   Whole weight     4177 non-null    float64   
 5   Shucked weight   4177 non-null    float64   
 6   Viscera weight   4177 non-null    float64   
 7   Shell weight     4177 non-null    float64   
 8   Rings            4177 non-null    int64     
dtypes: float64(7), int64(1), object(1)  
memory usage: 293.8+ KB  
None
```

There is sex feature which is object type we have convert into int.

## Checking and converting Sex feature into int

```
In [24]: print('-----')
print('No of Mane, Infant and Female present in our dataset.')
print('-----')
print(df['Sex'].value_counts())
print('-----')
```

No of Mane, Infant and Female present in our dataset.

```
M      1528
I      1342
F      1307
Name: Sex, dtype: int64
```

```
In [25]: df['Sex'] = df['Sex'].replace({'M':0,'F':1,'I':2})  
df.head()
```

Out[25]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	0	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	1	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	0	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	2	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

Replace Male : 0, Female : 1, Infant : 2

## Statistical data

```
In [8]: df.describe()
```

Out[8]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.955470	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594
std	0.827815	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614
min	0.000000	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500
25%	0.000000	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500
50%	1.000000	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000
75%	2.000000	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000
max	2.000000	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000

There is some outliers present in our dataset.

## Relation feature and label

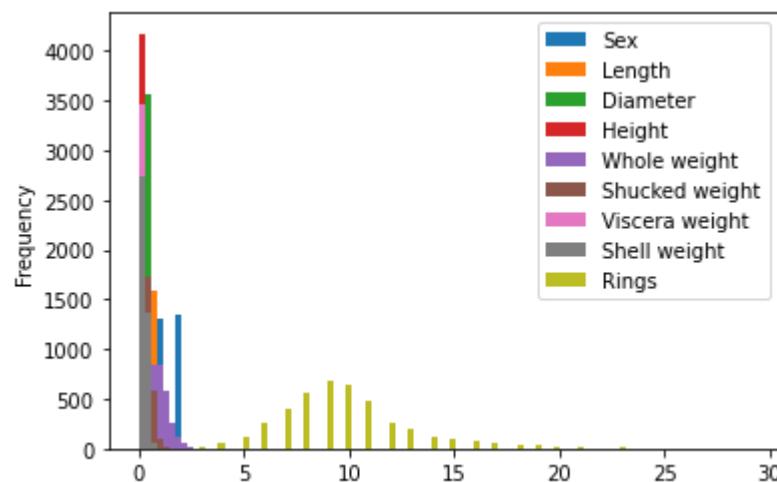
```
In [28]: print('-----')
print(df.groupby('Sex')['Rings'].sum())
print('-----')
```

```
-----  
Sex  
0    16358  
1    14546  
2    10589  
Name: Rings, dtype: int64  
-----
```

Observation : No of rings are vary from sex.

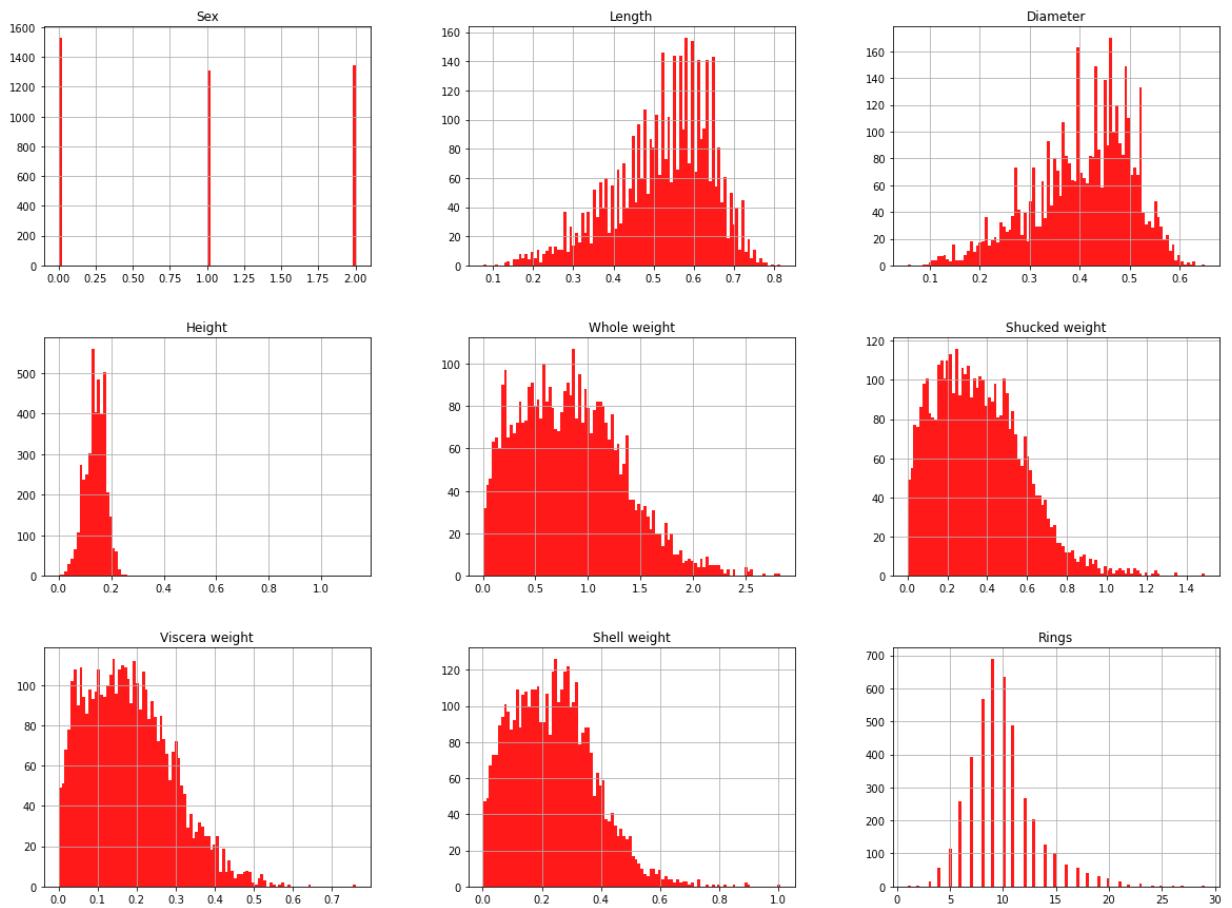
## Understanding our dataset through Visualization

```
In [10]: df.plot(kind = 'hist', bins= 100)
plt.show()
```



Above plot show properties of data distribution scale

```
In [11]: df.hist(color = 'r', alpha = 0.9, figsize = (20,15), bins = 100)  
plt.show()
```



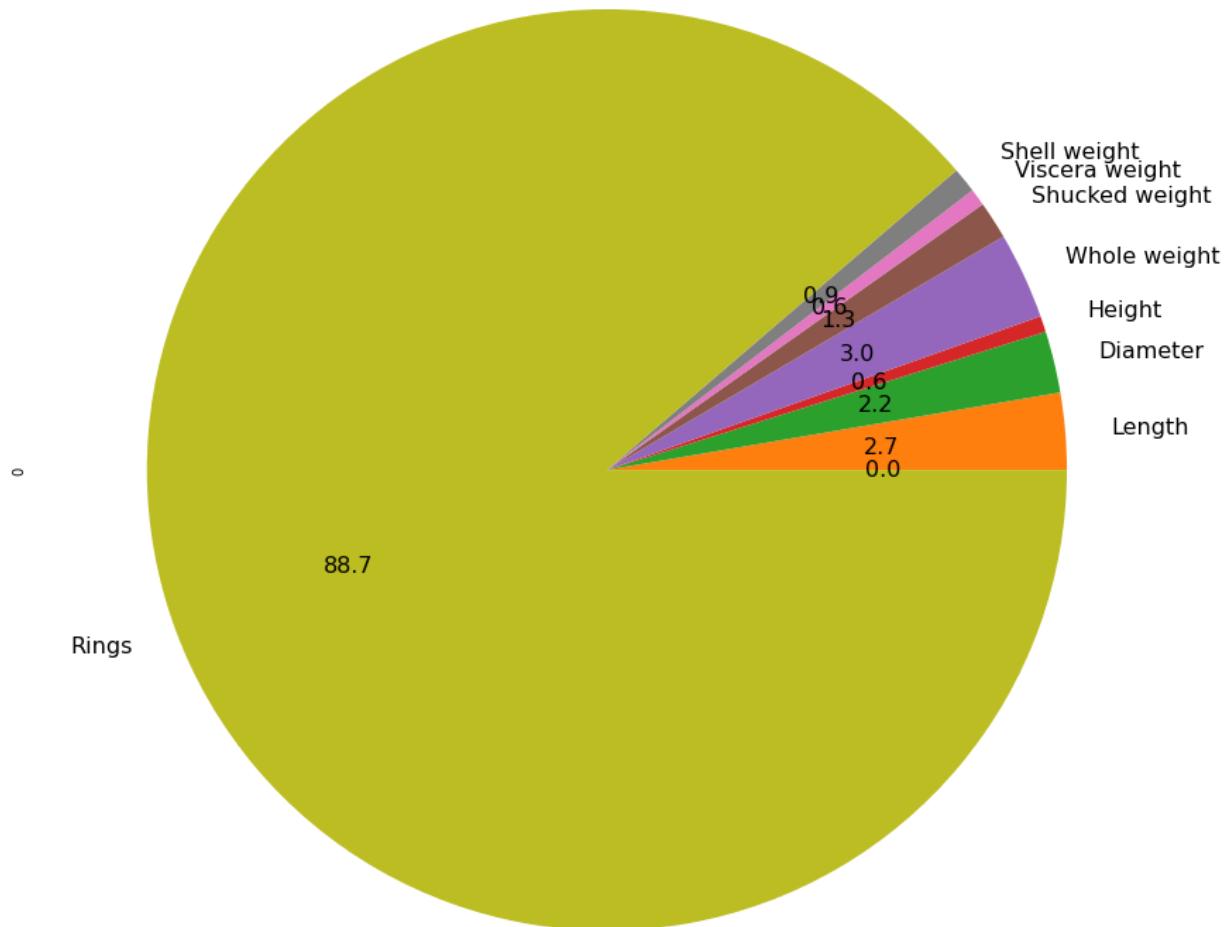
Above plot demonstrate features and how they distributed some are left skewed and some are right skewed

```
In [12]: s = df.iloc[0]
```

```
s
```

```
Out[12]: Sex          0.0000
Length       0.4550
Diameter     0.3650
Height        0.0950
Whole weight  0.5140
Shucked weight 0.2245
Viscera weight 0.1010
Shell weight   0.1500
Rings         15.0000
Name: 0, dtype: float64
```

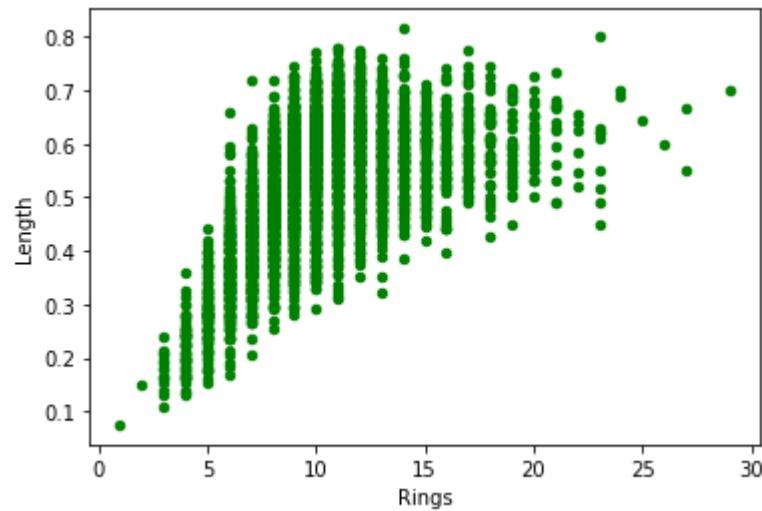
```
In [13]: s.plot.pie(subplots = True, fontsize = 16, autopct = '%.1f', figsize = (20,15))
plt.show()
```



Abalone made of 88.7% of shell.

```
df.plot.scatter(x = 'Rings', y = 'Sex' , color = 'r') plt.show()
```

```
In [14]: df.plot.scatter(x = 'Rings', y = 'Length' , color = 'g')  
plt.show()
```



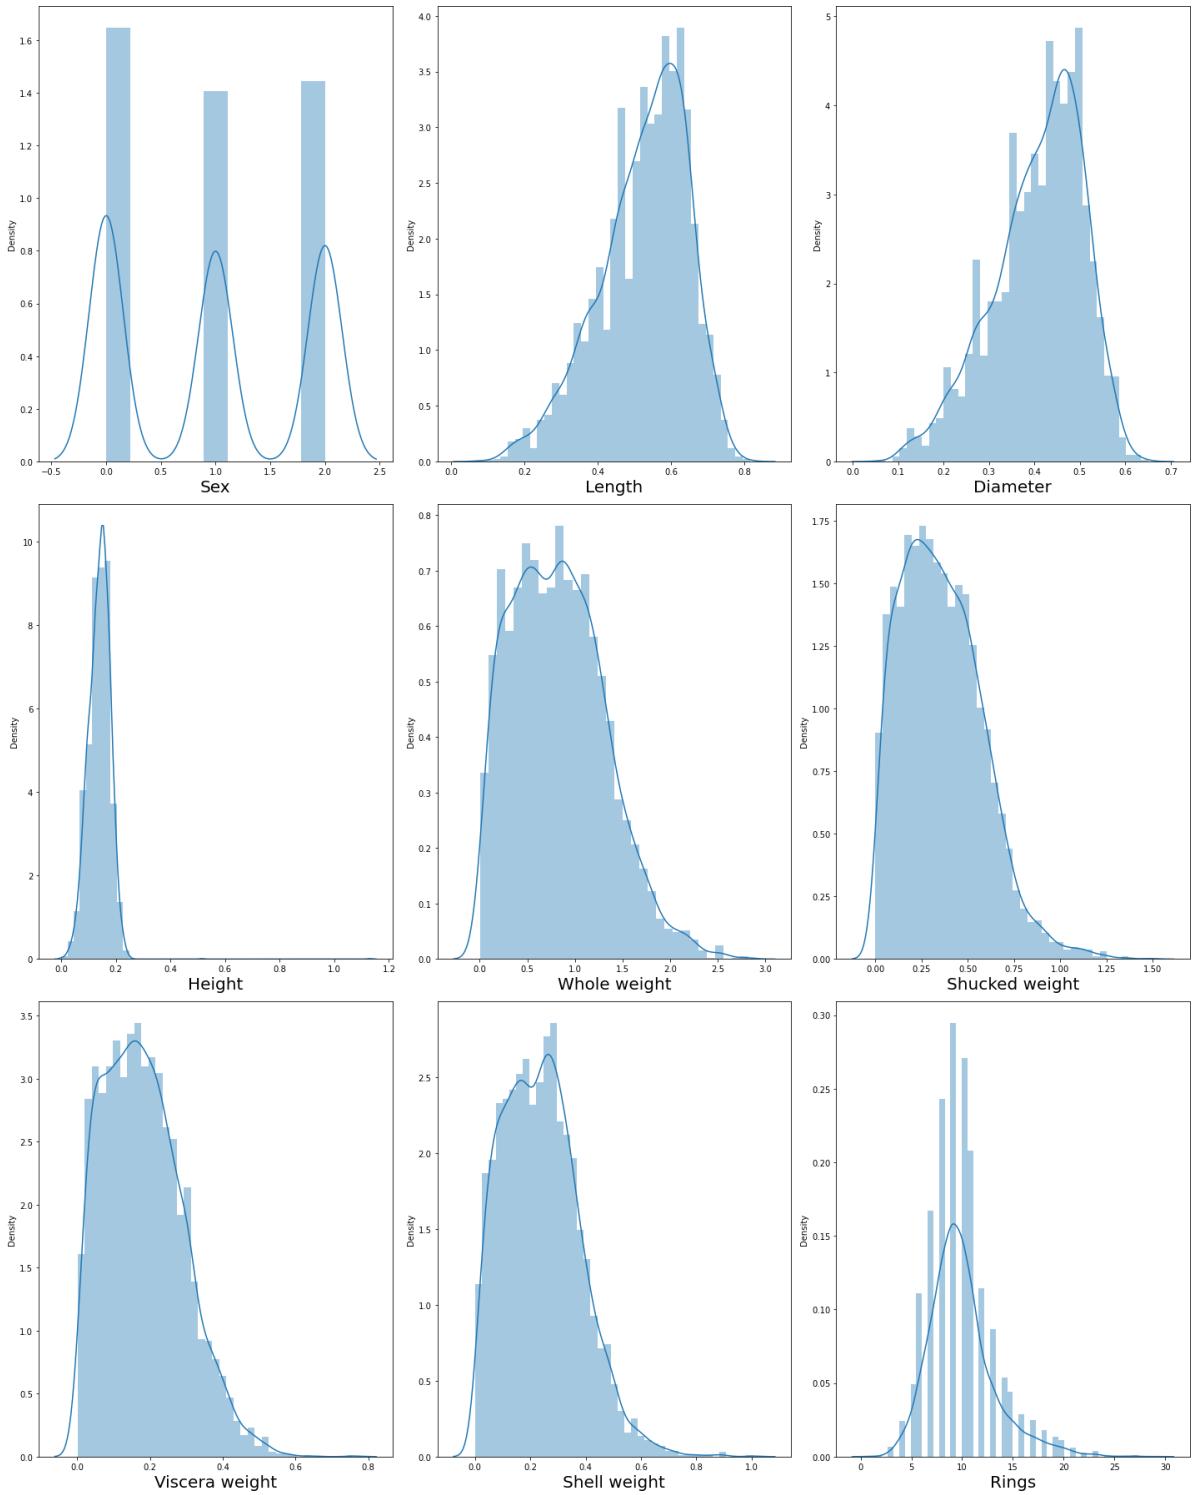
As we see above scatter plot as length increase no of Rings increase

```
In [35]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in df:
    if plotnumber <=9:
        ax = plt.subplot(3,3, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----
Distribution Plot :-
```



**Observations :** Some features is little skewed.

## Checking for Imbalance Label

```
In [37]: print('-----')
print('No of rings present in each class :')
print('-----')
print(df['Rings'].value_counts())
print('-----')
```

No of rings present in each class :

9	689
10	634
8	568
11	487
7	391
12	267
6	259
13	203
14	126
5	115
15	103
16	67
17	58
4	57
18	42
19	32
20	26
3	15
21	14
23	9
22	6
27	2
24	2
26	1
29	1
25	1
1	1
2	1

Name: Rings, dtype: int64

## Corelation of Feature vs Label using Heat map

```
In [39]: print('-----')
```

```
print('Heat Map :-')
```

```
print('-----')
```

```
df_corr = df.corr().abs()
```

```
plt.figure(figsize = (22,16))
```

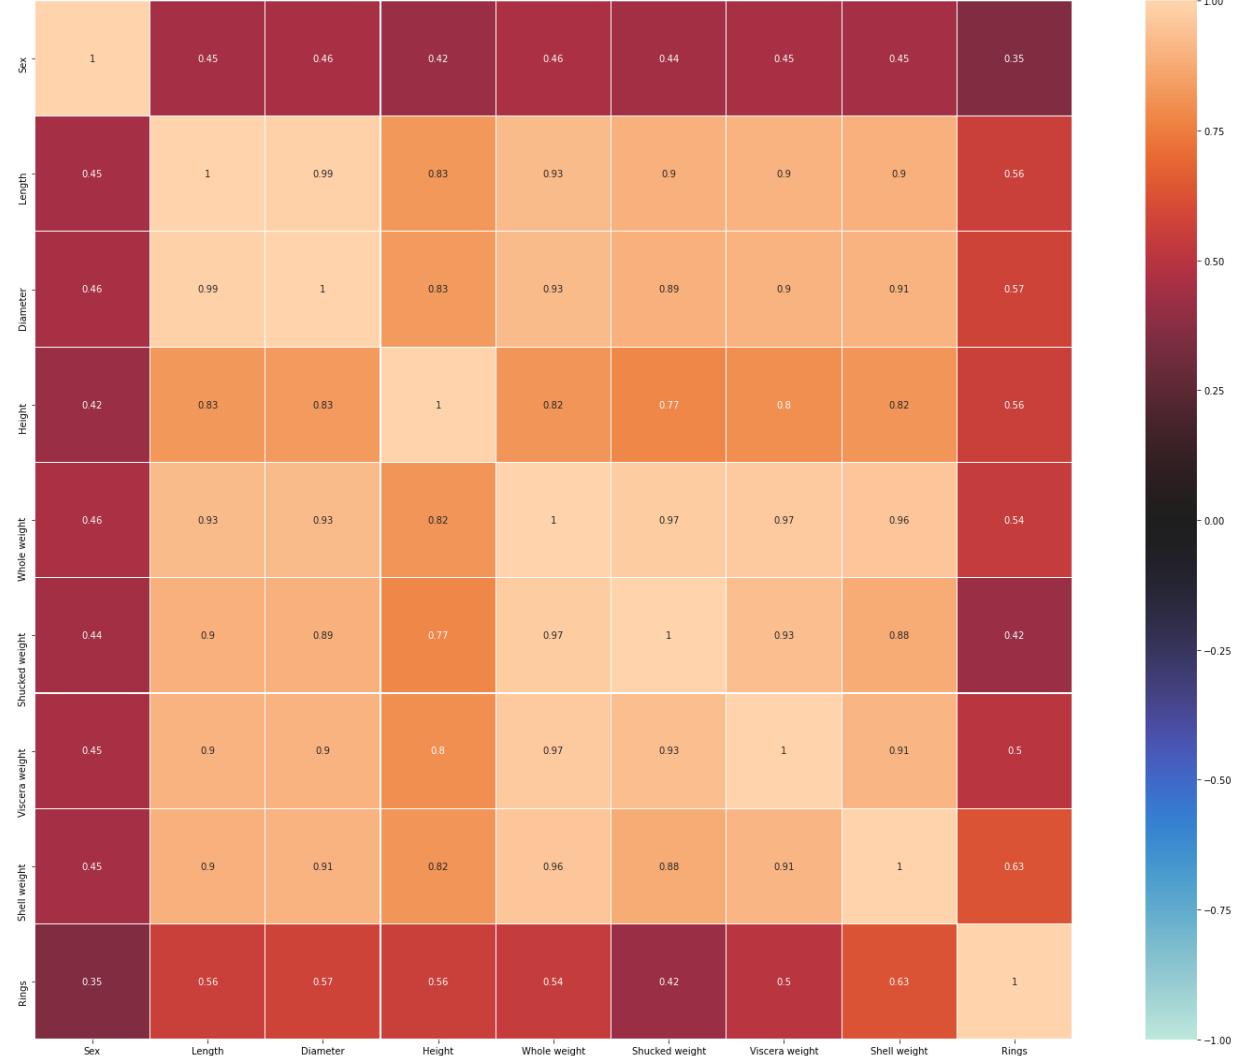
```
sns.heatmap(df_corr, vmin = -1, annot = True, square = True, center = 0, fmt = '.2f')
```

```
plt.tight_layout()
```

-----

Heat Map :-

-----



All features are related

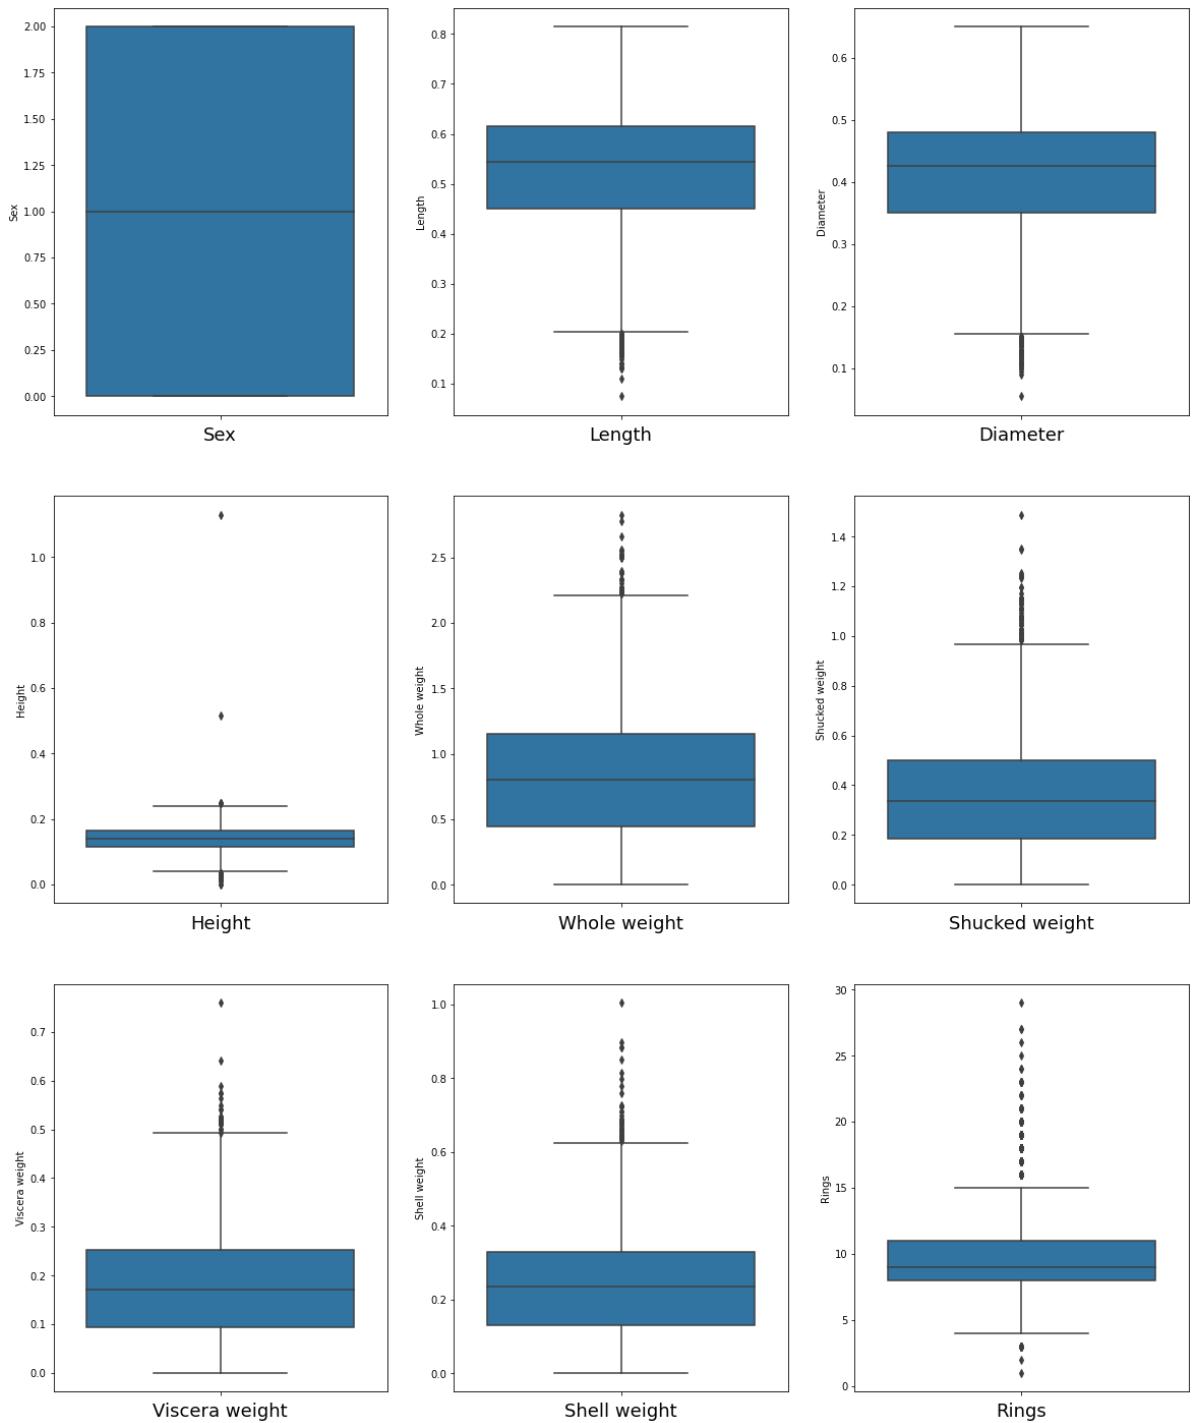
## Checking Outliers

```
In [40]: # Visualize the outliers using boxplot
print('-----')
print('Box Plot :-')
print('-----')

plt.figure(figsize = (20,25))
graph = 1

for column in df:
    if graph <=9:
        ax = plt.subplot(3,3, graph)
        sns.boxplot(y=df[column]) # It is the axis for vertical set as y
        plt.xlabel(column, fontsize = 18)
    graph +=1
plt.show()
```

-----  
Box Plot :-  
-----



There are outliers presents in dataset.

## Removing Outliers using Zscore

```
In [41]: # with std 3 Lets see the stats
```

```
z_score = zscore(df[['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight']])
abs_z_score = np.abs(z_score)

filtering_entry = (abs_z_score < 3).all(axis = 1)

df = df[filtering_entry]

df.describe()
```

```
Out[41]:
```

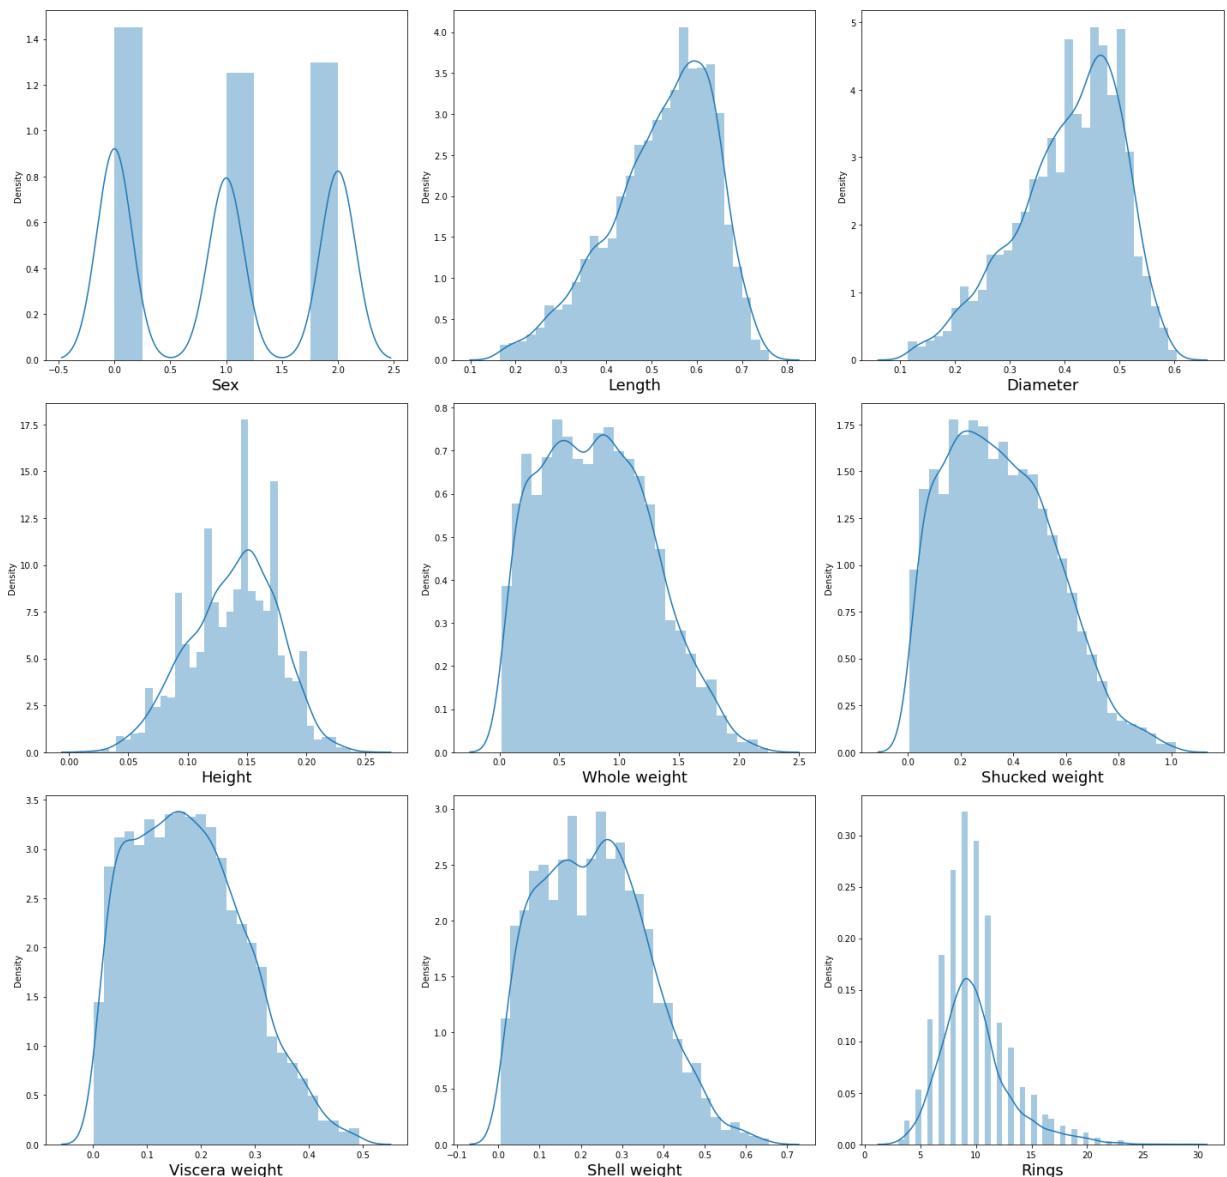
	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
<b>count</b>	4084.000000	4084.000000	4084.000000	4084.000000	4084.000000	4084.000000	4084.000000
<b>mean</b>	0.961557	0.522090	0.406308	0.138530	0.808404	0.350172	0.176579
<b>std</b>	0.828255	0.115709	0.095848	0.037001	0.457321	0.205827	0.103336
<b>min</b>	0.000000	0.165000	0.115000	0.015000	0.014500	0.005500	0.000500
<b>25%</b>	0.000000	0.450000	0.350000	0.115000	0.441000	0.185375	0.092875
<b>50%</b>	1.000000	0.540000	0.420000	0.140000	0.790750	0.332000	0.168500
<b>75%</b>	2.000000	0.610000	0.480000	0.165000	1.134750	0.494000	0.247000
<b>max</b>	2.000000	0.760000	0.605000	0.250000	2.238500	1.017000	0.492500

```
In [42]: # Let's see outliers are removed in columns or not.
```

```
print('-----')
print('Distribution Plot :-')
print('-----')

plt.figure(figsize = (20,25), facecolor = 'white')
plotnumber = 1
for column in df:
    if plotnumber <=12:
        ax = plt.subplot(4,3, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 18)
    plotnumber += 1
plt.tight_layout()
```

-----  
Distribution Plot :-  
-----



**It removes outliers.**

```
In [43]: print('-----')
print('No of Rows and Columns ----->',df.shape)
print('-----')
```

```
-----  
No of Rows and Columns -----> (4084, 9)  
-----
```

**Above we check shape of remaining data after removal of outliers.**

## Checking for Skewness

```
In [46]: print('-----')
print('Skewness in data :-')
print('-----')
print(df.skew())
print('-----')
```

```
-----
Skewness in data :-
-----
Sex           0.071702
Length        -0.633786
Diameter      -0.605450
Height         -0.247192
Whole weight   0.323886
Shucked weight 0.449573
Viscera weight 0.429932
Shell weight   0.358512
Rings          1.136367
dtype: float64
-----
```

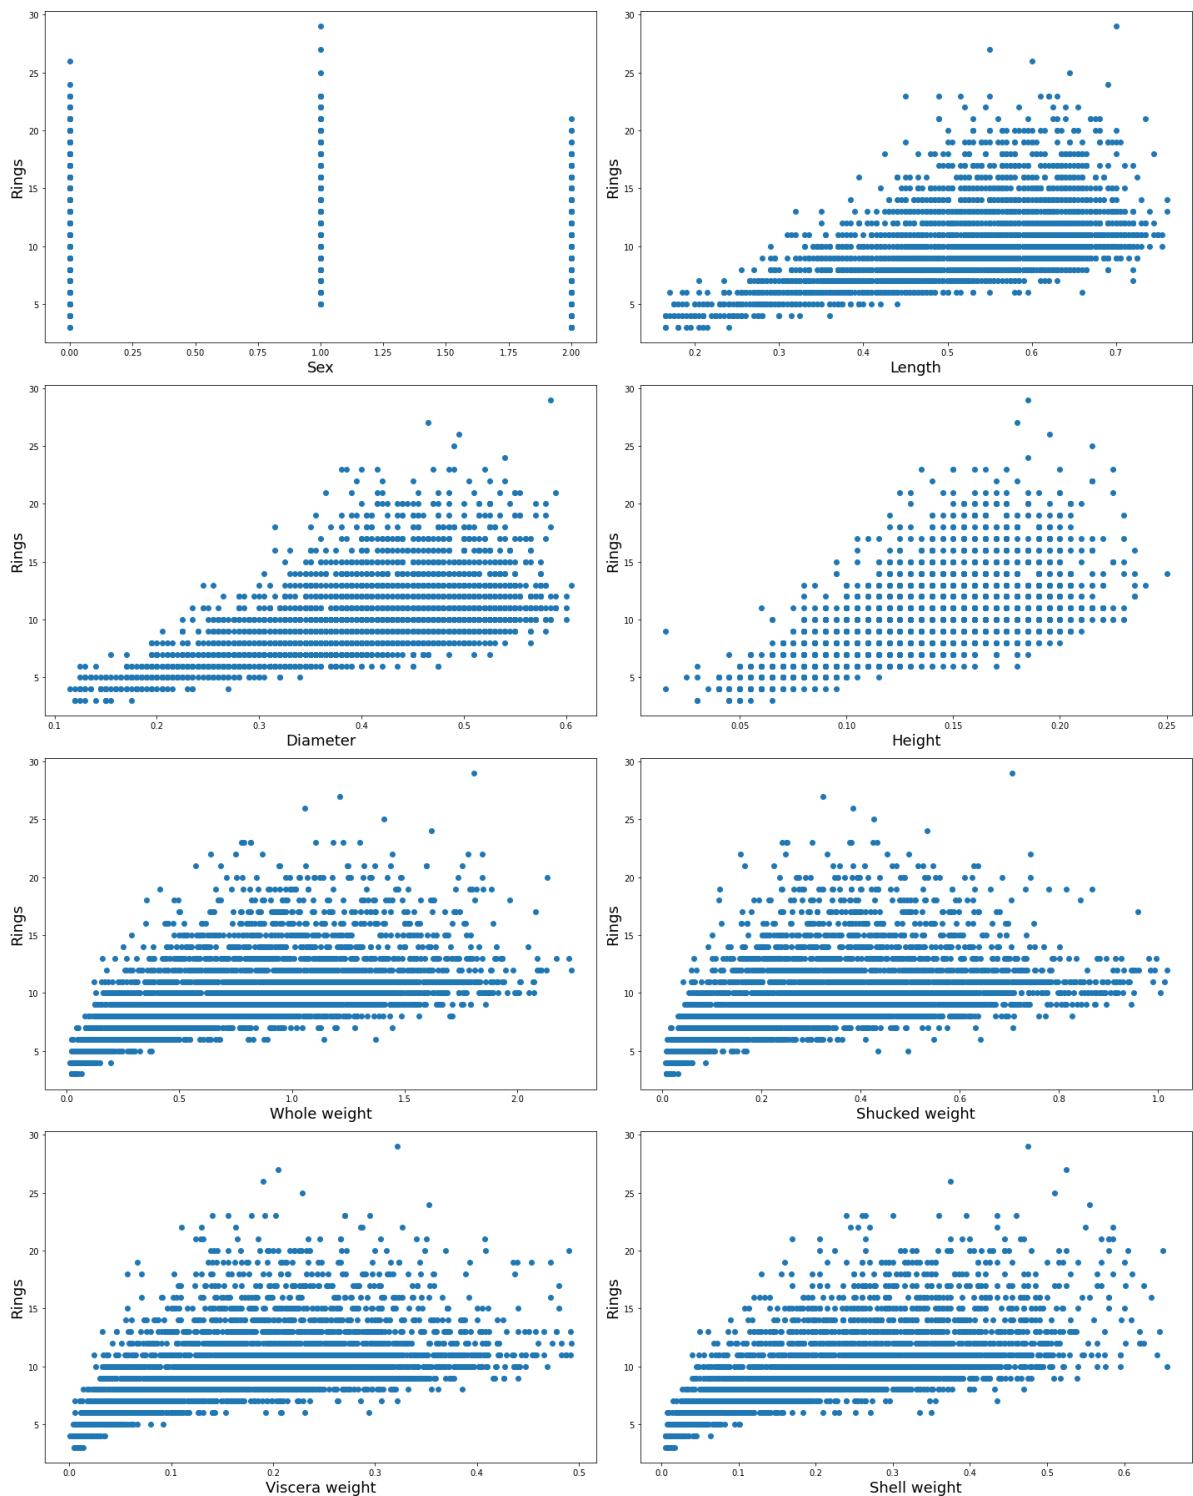
There is some skewness but we kept it because it cause data loss.

## Splitting Dataset into features and labels

```
In [48]: x = df.drop('Rings', axis = 1)
y = df.Rings
```

```
In [50]: # Let's see relation between features and labels.  
print('-----')  
print('Distribution Plot :-')  
print('-----')  
  
plt.figure(figsize = (20,25), facecolor = 'white')  
plotnumber = 1  
for column in x:  
    if plotnumber <=8:  
        ax = plt.subplot(4,2, plotnumber)  
        plt.scatter(x[column],y)  
        plt.xlabel(column, fontsize = 18)  
        plt.ylabel('Rings', fontsize = 18)  
    plotnumber += 1  
plt.tight_layout()
```

-----  
Distribution Plot :-  
-----



Features are related to class

## Handling Class Imbalance

```
In [51]: from imblearn.over_sampling import RandomOverSampler
ros = RandomOverSampler()
x_over, y_over = ros.fit_resample(x, y)
```

```
In [52]: print('-----')
print('Class are balanced :-')
print('-----')
print(y_over.value_counts())
print('-----')
```

```
-----
Class are balanced :-
-----
16    686
17    686
14    686
29    686
13    686
12    686
27    686
11    686
26    686
10    686
25    686
9     686
24    686
8     686
23    686
7     686
22    686
6     686
21    686
5     686
20    686
4     686
19    686
3     686
18    686
15    686
Name: Rings, dtype: int64
-----
```

## Data Scaling

```
In [27]: scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
Out[27]: array([[-1.16108558, -0.5798883 , -0.43102345, ... , -0.61064456,
       -0.73148499, -0.64218694],
      [-1.16108558, -1.48745122, -1.47447322, ... , -1.21802438,
       -1.23959876, -1.26010684],
      [ 0.04641977,  0.06837092,  0.14287393, ... , -0.45515533,
       -0.33951152, -0.17874702],
      ... ,
      [-1.16108558,  0.67341287,  0.7167713 , ... ,  0.85192604,
       1.07352867,  0.57820485],
      [ 0.04641977,  0.88949928,  0.82111628, ... ,  0.87865076,
       0.81705219,  0.48551686],
      [-1.16108558,  1.62419307,  1.55153111, ... ,  2.89272223,
       1.93490248,  2.0225926 ]])
```

Data has been scaled.

**Split data into train and test. Model will be bulit on training data and tested on test data**

```
In [53]: x_train, x_test, y_train, y_test = train_test_split(x_over, y_over, test_size = 0.2)
print('Data has been splited.')
Data has been splited.
```

## Model Building

**AdaBoost model instantiaing, training and evaluating**

```
In [54]: ada = AdaBoostClassifier()
ada.fit(x_train, y_train)
y_pred = ada.predict(x_test)
```

```
In [55]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```

[ 0  0  0  0  0  0  0  0  0  178  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
[ 0  0  0  0  0  0  0  0  0  192  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
[ 0  0  0  0  0  0  0  0  0  167  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 158  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
-----
```

#### Classification Report:

	precision	recall	f1-score	support
3	0.60	1.00	0.75	173
4	0.31	0.46	0.37	171
5	0.00	0.00	0.00	183
6	0.00	0.00	0.00	174
7	0.00	0.00	0.00	188
8	0.00	0.00	0.00	167
9	0.00	0.00	0.00	170
10	0.00	0.00	0.00	173
11	0.04	0.99	0.07	139
12	0.00	0.00	0.00	150
13	0.00	0.00	0.00	177
14	0.00	0.00	0.00	173
15	0.00	0.00	0.00	164
16	0.00	0.00	0.00	171
17	0.00	0.00	0.00	198
18	0.00	0.00	0.00	162
19	0.00	0.00	0.00	158
20	0.00	0.00	0.00	171
21	0.09	0.08	0.09	191
22	0.00	0.00	0.00	167
23	0.00	0.00	0.00	182
24	0.00	0.00	0.00	162
25	0.00	0.00	0.00	178
26	0.00	0.00	0.00	192
27	0.00	0.00	0.00	167
29	0.00	0.00	0.00	158
accuracy			0.09	4459
macro avg	0.04	0.10	0.05	4459
weighted avg	0.04	0.09	0.05	4459

Conclusion : Ada Boost model has 9% score.

Cross Validation score to check if the model is overfitting

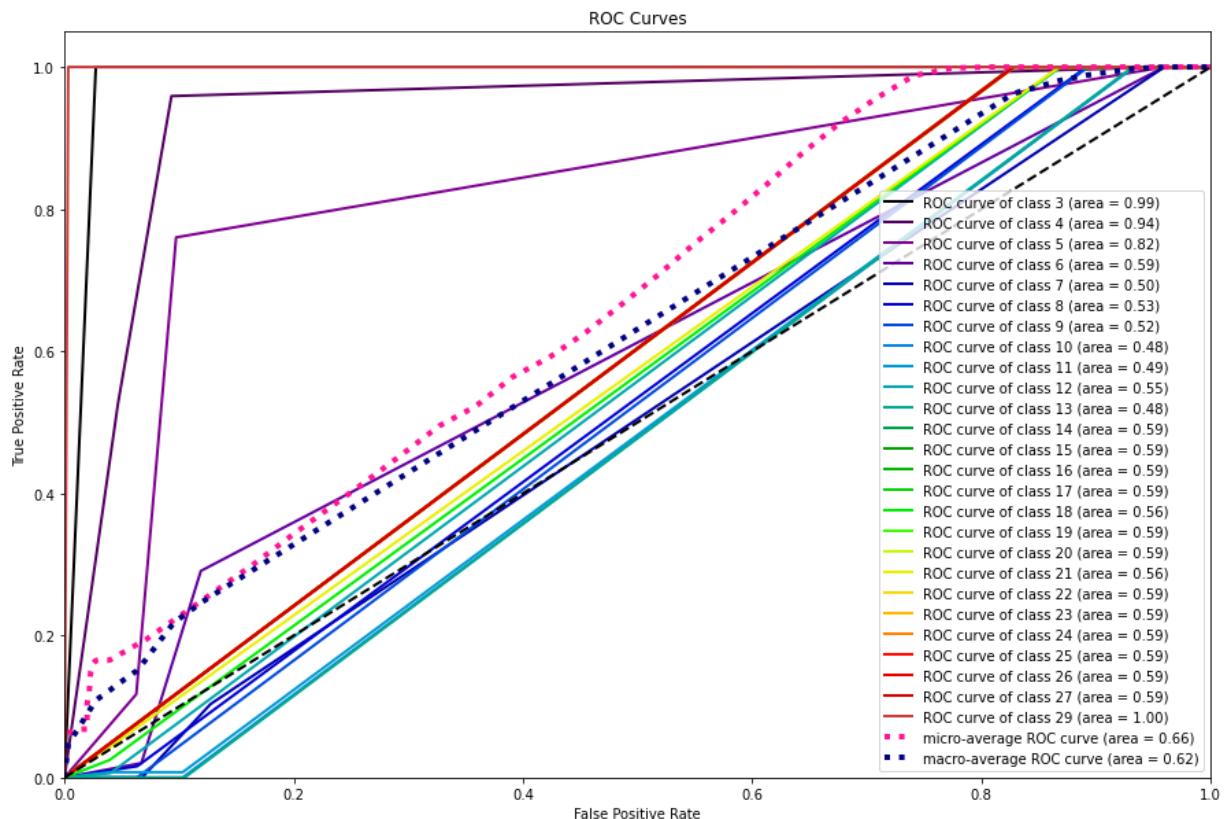
```
In [56]: cv = cross_val_score(ada, x, y, cv = 5)
print('Cross Validation score of Ada Boost model --->', cv.mean())
```

Cross Validation score of Ada Boost model ---> 0.21743886048911606

**Conclusion : Ada Boost model has 21% Cross Validation score**

## ROC, AUC Curve

```
In [51]: prob = ada.predict_proba(x_test) # calculating probability
skplt.metrics.plot_roc(y_pred, prob, figsize = (15,10))
plt.show()
```



## Decision Tree model instantiaing, training and evaluating

```
In [57]: DT = DecisionTreeClassifier()
DT.fit(x_train, y_train)
y_pred = DT.predict(x_test)
```

```
In [58]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 178  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 192  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 167  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 158]]
```

---

#### Classification Report:

	precision	recall	f1-score	support
3	1.00	1.00	1.00	173
4	0.99	1.00	0.99	171
5	0.93	1.00	0.97	183
6	0.88	0.90	0.89	174
7	0.75	0.72	0.74	188
8	0.54	0.47	0.50	167
9	0.34	0.25	0.29	170
10	0.40	0.29	0.34	173
11	0.64	0.60	0.62	139
12	0.80	0.93	0.86	150
13	0.90	0.94	0.92	177
14	0.90	1.00	0.95	173
15	0.91	1.00	0.95	164
16	0.94	1.00	0.97	171
17	0.96	1.00	0.98	198
18	0.98	1.00	0.99	162
19	0.93	1.00	0.97	158
20	0.99	1.00	0.99	171
21	0.99	1.00	0.99	191
22	1.00	1.00	1.00	167
23	1.00	1.00	1.00	182
24	1.00	1.00	1.00	162
25	1.00	1.00	1.00	178
26	1.00	1.00	1.00	192
27	1.00	1.00	1.00	167
29	0.99	1.00	1.00	158
accuracy			0.89	4459
macro avg	0.88	0.89	0.88	4459
weighted avg	0.88	0.89	0.88	4459

---

**Conclusion : Decision Tree model has 89% score**

**Cross Validation score to check if the model is overfitting**

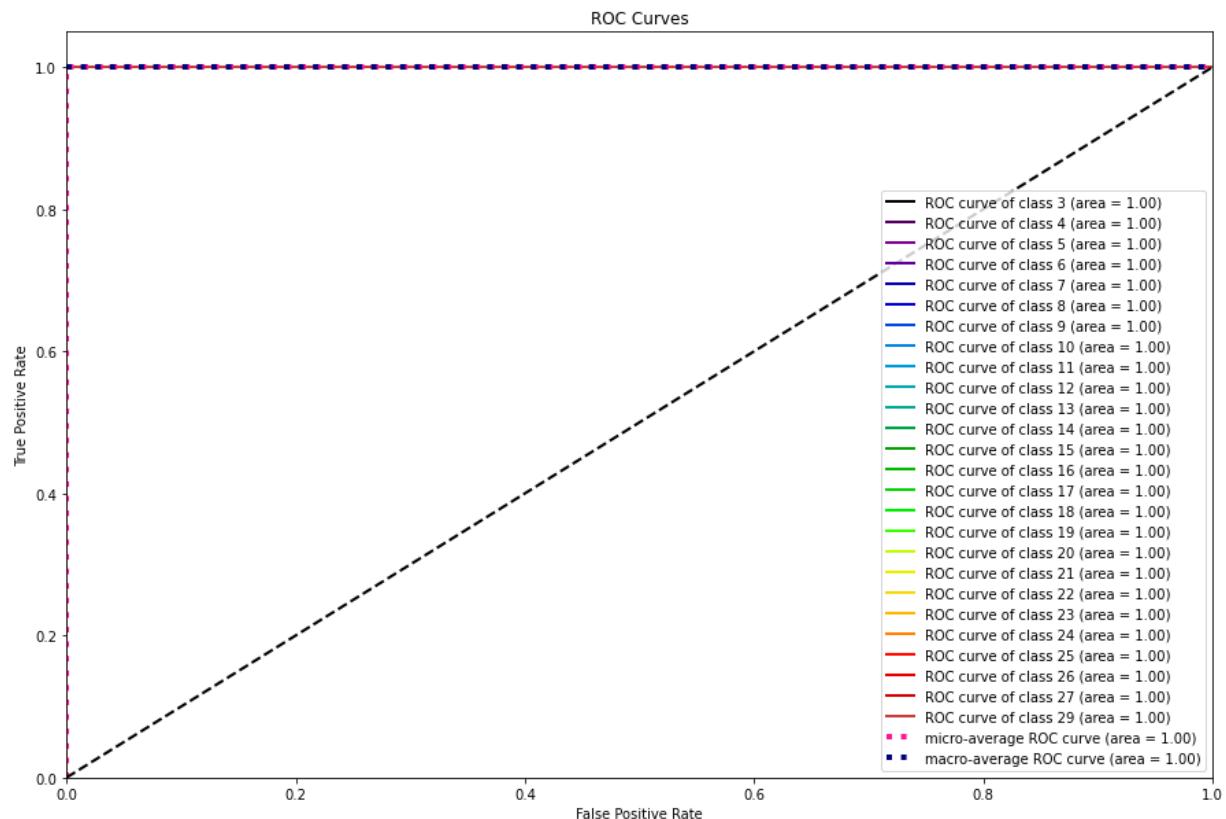
```
In [59]: cv = cross_val_score(DT, x, y, cv = 5)
print('Cross Validation score of Decision Tree model --->', cv.mean())
```

Cross Validation score of Decision Tree model ---> 0.18609361125111

**Conclusion : Decision Tree model has 18% Cross Validation score**

## ROC, AUC Curve

```
In [60]: prob = DT.predict_proba(x_test) # calculating probability
skplt.metrics.plot_roc(y_pred, prob, figsize = (15,10))
plt.show()
```



## Knn model instantiating, training and evaluating

```
In [61]: Knn = KNeighborsClassifier()
Knn.fit(x_train, y_train)
y_pred = Knn.predict(x_test)
```

```
In [62]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

```

[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 178  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 192  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 167  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 158]]
```

---

#### Classification Report:

	precision	recall	f1-score	support
3	0.98	1.00	0.99	173
4	0.90	1.00	0.95	171
5	0.82	0.92	0.86	183
6	0.55	0.62	0.58	174
7	0.46	0.41	0.43	188
8	0.41	0.33	0.36	167
9	0.30	0.23	0.26	170
10	0.26	0.12	0.16	173
11	0.30	0.26	0.28	139
12	0.42	0.43	0.43	150
13	0.67	0.63	0.65	177
14	0.72	0.86	0.79	173
15	0.75	0.94	0.83	164
16	0.90	1.00	0.94	171
17	0.92	0.96	0.94	198
18	0.94	1.00	0.97	162
19	0.93	1.00	0.97	158
20	0.94	1.00	0.97	171
21	0.98	1.00	0.99	191
22	0.98	1.00	0.99	167
23	0.98	1.00	0.99	182
24	1.00	1.00	1.00	162
25	1.00	1.00	1.00	178
26	1.00	1.00	1.00	192
27	1.00	1.00	1.00	167
29	1.00	1.00	1.00	158
accuracy			0.80	4459
macro avg	0.77	0.80	0.78	4459
weighted avg	0.78	0.80	0.79	4459

---

**Conclusion : Knn model has 80% score**

**Cross Validation score to check if the model is overfitting**

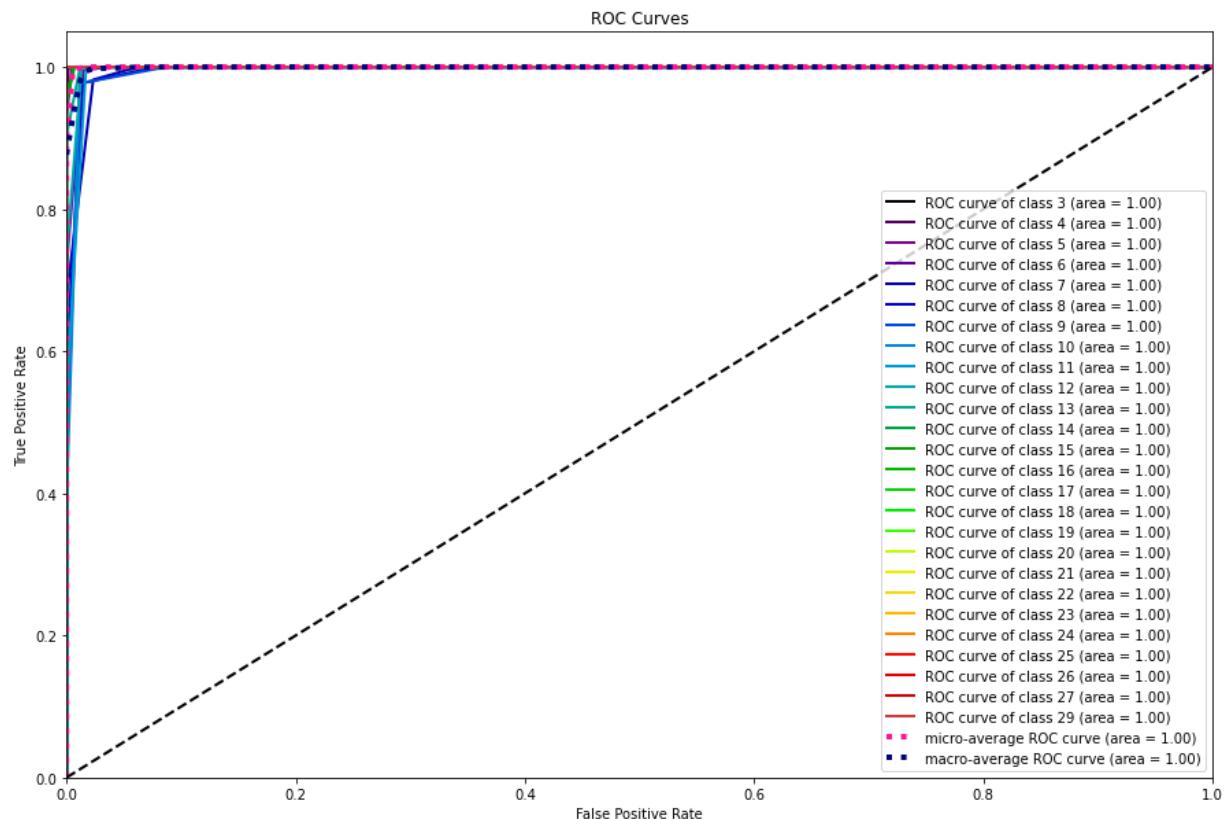
```
In [63]: cv = cross_val_score(Knn, x, y, cv = 5)
print('Cross Validation score of Knn model --->', cv.mean())
```

```
Cross Validation score of Knn model ---> 0.23482492140062877
```

**Conclusion : Knn model has 23% Cross Validation score**

## ROC, AUC Curve

```
In [64]: prob = Knn.predict_proba(x_test) # calculating probability
skplt.metrics.plot_roc(y_pred, prob, figsize = (15,10))
plt.show()
```



## Random Forest model instantiaing, training and evaluating

```
In [65]: Rn = RandomForestClassifier()
Rn.fit(x_train, y_train)
y_pred = Rn.predict(x_test)
```

```
In [66]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

```

[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 178  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 192  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 167  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0 158]]
```

---

#### Classification Report:

	precision	recall	f1-score	support
3	0.99	1.00	0.99	173
4	0.98	1.00	0.99	171
5	0.96	1.00	0.98	183
6	0.87	0.92	0.89	174
7	0.71	0.72	0.72	188
8	0.52	0.53	0.53	167
9	0.35	0.30	0.32	170
10	0.52	0.36	0.42	173
11	0.60	0.61	0.61	139
12	0.86	0.93	0.89	150
13	0.90	0.93	0.92	177
14	0.95	1.00	0.97	173
15	0.95	1.00	0.98	164
16	0.98	1.00	0.99	171
17	0.99	1.00	0.99	198
18	0.99	1.00	1.00	162
19	0.98	1.00	0.99	158
20	0.99	1.00	1.00	171
21	1.00	1.00	1.00	191
22	1.00	1.00	1.00	167
23	1.00	1.00	1.00	182
24	1.00	1.00	1.00	162
25	1.00	1.00	1.00	178
26	1.00	1.00	1.00	192
27	1.00	1.00	1.00	167
29	0.99	1.00	1.00	158
accuracy			0.90	4459
macro avg	0.89	0.90	0.89	4459
weighted avg	0.89	0.90	0.89	4459

---

**Conclusion : Random Forest model has 90% score**

**Cross Validation score to check if the model is overfitting**

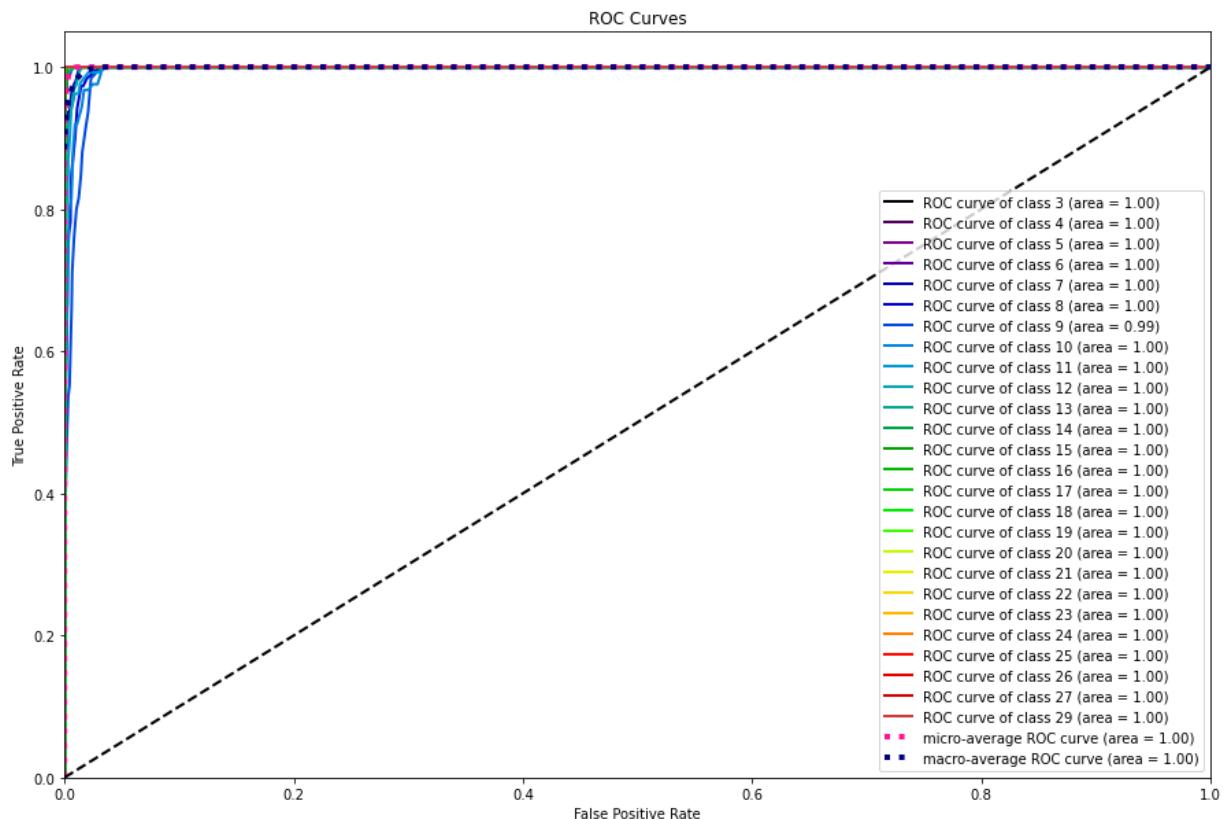
```
In [67]: cv = cross_val_score(Rn, x, y, cv = 5)
print('Cross Validation score of Random Forest model --->', cv.mean())
```

Cross Validation score of Random Forest model ---> 0.253187774497804

**Conclusion : Random Forest model has 25% Cross Validation score**

## ROC, AUC Curve

```
In [54]: prob = Rn.predict_proba(x_test) # calculating probability
skplt.metrics.plot_roc(y_pred,prob, figsize = (15,10))
plt.show()
```



## Ada Boost Hyperparameter Tuning .

```
In [70]: param = {'n_estimators': [50, 100], 'learning_rate' : [0.30,0.40,0.45]}
```

```
In [71]: grid_search = GridSearchCV(estimator = ada, param_grid = param, cv = 5,n_jobs = -1)
```

```
In [72]: grid_search.fit(x_train, y_train)
```

```
Out[72]: GridSearchCV(cv=5, estimator=AdaBoostClassifier(), n_jobs=-1,
param_grid={'learning_rate': [0.3, 0.4, 0.45],
'n_estimators': [50, 100]})
```

```
In [73]: best_parameters = grid_search.best_params_
print(best_parameters)

{'learning_rate': 0.3, 'n_estimators': 50}
```

```
In [74]: hada = AdaBoostClassifier(learning_rate = 0.3, n_estimators= 50)
hada.fit(x_train, y_train)
hada.score(x_test, y_test)
```

Out[74]: 0.12648575913882038

```
In [75]: y_pred = hada.predict(x_test)
```

```
In [76]: print('-----\n')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

Classification Report:				
	precision	recall	f1-score	support
3	0.00	0.00	0.00	173
4	0.00	0.00	0.00	171
5	0.00	0.00	0.00	183
6	0.17	0.80	0.28	174
7	0.00	0.00	0.00	188
8	0.19	0.25	0.22	167
9	0.38	0.11	0.17	170
10	0.33	0.01	0.01	173
11	0.00	0.00	0.00	139
12	0.15	0.04	0.06	150
13	0.05	0.07	0.06	177
14	0.00	0.00	0.00	173
15	0.36	0.02	0.05	164
16	0.08	0.11	0.09	171
17	0.01	0.02	0.01	198
18	0.13	0.22	0.16	162
19	0.15	0.37	0.21	158
20	0.15	0.42	0.22	171
21	0.08	0.20	0.12	191
22	0.00	0.00	0.00	167
23	0.14	0.64	0.23	182
24	0.00	0.00	0.00	162
25	0.00	0.00	0.00	178
26	0.00	0.00	0.00	192
27	0.00	0.00	0.00	167
29	0.00	0.00	0.00	158
accuracy			0.13	4459
macro avg	0.09	0.13	0.07	4459
weighted avg	0.09	0.13	0.07	4459

After Hyperparameter Tuning Ada Boost model accuracy score 13%

## Saving The Model

```
In [77]: # saving the model to the Local file system
filename = 'Abalone Case.pickle'
pickle.dump(hada, open(filename, 'wb'))
```

## Knn Hyperparameter Tuning model

```
In [78]: grid_param = {'leaf_size' : [1,3,5], 'n_neighbors': [5],  
                     'p':[1,2]}
```

```
In [79]: grid_search = GridSearchCV(estimator = Knn, param_grid = grid_param, cv = 5 , n_jobs=-1)
```

```
In [80]: grid_search.fit(x_train, y_train)
```

```
Out[80]: GridSearchCV(cv=5, estimator=KNeighborsClassifier(), n_jobs=-1,  
                      param_grid={'leaf_size': [1, 3, 5], 'n_neighbors': [5],  
                                   'p': [1, 2]})
```

```
In [81]: best_parameters = grid_search.best_params_  
print(best_parameters)
```

```
{'leaf_size': 3, 'n_neighbors': 5, 'p': 1}
```

```
In [82]: hkn = KNeighborsClassifier(leaf_size = 1, n_neighbors = 5, p = 1)  
hkn.fit(x_train, y_train)  
hkn.score(x_test, y_test)
```

```
Out[82]: 0.8004036779546984
```

```
In [83]: y_pred = hkn.predict(x_test)
```

```
In [84]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

	precision	recall	f1-score	support
3	0.98	1.00	0.99	173
4	0.91	1.00	0.96	171
5	0.80	0.91	0.85	183
6	0.54	0.60	0.57	174
7	0.45	0.45	0.45	188
8	0.44	0.35	0.39	167
9	0.27	0.19	0.23	170
10	0.28	0.13	0.18	173
11	0.24	0.20	0.22	139
12	0.42	0.43	0.42	150
13	0.69	0.63	0.66	177
14	0.70	0.87	0.78	173
15	0.79	0.92	0.85	164
16	0.88	1.00	0.93	171
17	0.93	0.96	0.94	198
18	0.93	1.00	0.96	162
19	0.92	1.00	0.96	158
20	0.95	1.00	0.97	171
21	0.98	1.00	0.99	191
22	0.98	1.00	0.99	167
23	0.98	1.00	0.99	182
24	1.00	1.00	1.00	162
25	1.00	1.00	1.00	178
26	1.00	1.00	1.00	192
27	1.00	1.00	1.00	167
29	1.00	1.00	1.00	158
accuracy			0.80	4459
macro avg	0.77	0.79	0.78	4459
weighted avg	0.78	0.80	0.79	4459

After Hyperparameter Tuning model accuracy score increase to 80% .

## Saving The Model

```
In [85]: # saving the model to the Local file system
filename = 'Knn Abalone case.pickle'
pickle.dump(hkn, open(filename, 'wb'))
```

**Final Conclusion : Knn is our best model.**