

Problem Statement:

Sales Channel Prediction Case Study:-

When a company enters a market, the distribution strategy and channel it uses are keys to its success in the market, as well as market know-how and customer knowledge and understanding. Because an effective distribution strategy under efficient supply-chain management opens doors for attaining competitive advantage and strong brand equity in the market, it is a component of the marketing mix that cannot be ignored . The distribution strategy and the channel design have to be right the first time. The case study of Sales channel includes the detailed study of TV, radio and newspaper channel. The predict the total sales generated from all the sales channel.

Importing required libraries

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import pickle
from scipy.stats import zscore
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler
import statsmodels.formula.api as smf
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.linear_model import LinearRegression, LassoCV, Lasso
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

Reading data

```
In [2]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\Data Analysis With Python\Advertising.csv")
df.head()
```

Out[2]:

	Unnamed: 0	TV	radio	newspaper	sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9

Droping Unwanted Column

```
In [3]: df = df.drop(columns = 'Unnamed: 0', axis = 1)
df.head()
```

Out[3]:

	TV	radio	newspaper	sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

Check no of row and column

```
In [4]: print('No of Rows and Columns ----->', df.shape )
```

No of Rows and Columns -----> (200, 4)

Checking for Null values

```
In [5]: print('-----\n')
print(df.isnull().sum())
print('\n-----')
```

```
-----  
TV      0  
radio   0  
newspaper 0  
sales   0  
dtype: int64  
-----
```

There is no null value

Information about dataset

```
In [6]: print('-----\n')
print(df.info())
print('\n-----')
```

```
-----  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 200 entries, 0 to 199  
Data columns (total 4 columns):  
 #   Column     Non-Null Count  Dtype     
---  --          --          --          --  
 0   TV          200 non-null    float64  
 1   radio        200 non-null    float64  
 2   newspaper    200 non-null    float64  
 3   sales        200 non-null    float64  
 dtypes: float64(4)  
 memory usage: 6.4 KB  
None  
-----
```

All features are in float

Statistic of Dataset

```
In [7]: # We use describe command to extracte statistical infomation about dataset.  
df.describe()
```

Out[7]:


	TV	radio	newspaper	sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	14.022500
std	85.854236	14.846809	21.778621	5.217457
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	10.375000
50%	149.750000	22.900000	25.750000	12.900000
75%	218.825000	36.525000	45.100000	17.400000
max	296.400000	49.600000	114.000000	27.000000

Plot graph to visualize how data is distributed in every column

```
In [8]: # Let's see how data is distributed in every columns.
print('\nDistribution Plot :-\n')

plt.figure(figsize = (20,15), facecolor = 'white')
plotnumber = 1
for column in df:
    if plotnumber <=4:
        ax = plt.subplot(2,2, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 18)
    plotnumber += 1
plt.tight_layout()
```

Distribution Plot :-



some column are skewed

Removing outliers and verify it

```
In [9]: # with std 3 Lets see the stats
```

```
z_score = zscore(df[['TV', 'radio', 'newspaper']]) # use only continuous data
abs_z_score = np.abs(z_score)

filtering_entry = (abs_z_score < 3).all(axis = 1)


df = df[filtering_entry]
```

```
In [10]: # Let's see outliers remove or not.
```

```
print('\nDistribution Plot :-\n')

plt.figure(figsize = (20,15), facecolor = 'white')
plotnumber = 1
for column in df:
    if plotnumber <=4:
        ax = plt.subplot(2,2, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 18)
    plotnumber += 1
plt.tight_layout()
```

Distribution Plot :-



```
In [11]: print('Checking No of Rows and Columns After Removing Outliers ----->', df.shape)
```

```
Checking No of Rows and Columns After Removing Outliers -----> (198, 4)
```


Outliers are removed

Visualizing relationship

```
In [12]: TV= df.groupby('sales')['TV'].sum()  
TV
```

```
Out[12]: sales  
1.6      0.7  
3.2      4.1  
4.8      8.6  
5.3     18.5  
5.5      7.3  
...  
24.7    220.3  
25.4    799.8  
25.5    283.6  
26.2    287.6  
27.0    276.9  
Name: TV, Length: 121, dtype: float64
```

```
In [13]: TV.plot.hist(figsize = (20,15), fontsize = 14, bins = 100)
plt.xlabel('TV')
plt.ylabel('Sales')
plt.title('TV vs Sales')
plt.show()
```




Range 10 to 15 highest sales

```
In [14]: radio = df.groupby('radio')[['sales']].sum()
radio
```

```
Out[14]: radio
0.0      8.8
0.3      8.7
0.4      5.3
0.8      9.4
1.3     10.1
...
47.8    16.7
48.9    34.2
49.0    50.1
49.4    38.4
49.6    23.8
Name: sales, Length: 165, dtype: float64
```

```
In [15]: radio.plot.hist(figsize = (20,15), fontsize = 14, bins = 100)
plt.xlabel('Radio')
plt.ylabel('Sales')
plt.title('Radio vs Sales')
plt.show()
```




Range of 5 to 15 highest sales

```
In [16]: newsp= df.groupby('newspaper')['sales'].sum()  
newsp
```

```
Out[16]: newspaper  
0.3      17.4  
0.9      9.3  
1.0      4.8  
1.7     20.7  
1.8     20.7  
...  
75.0      7.2  
75.6     19.2  
79.2     15.9  
84.8     11.9  
89.4      8.7  
Name: sales, Length: 170, dtype: float64
```

```
In [17]: newsp.plot.hist(figsize = (20,15), fontsize = 14, bins = 100)  
plt.xlabel('Newspaper')  
plt.ylabel('Sales')  
plt.title('Newspaper vs Sales')  
plt.show()
```



Range 10 to 12 highest sales


Splitting data into features and label

```
In [18]: x = df.drop('sales', axis = 1)
y = df['sales']
```

```
In [19]: # Let's see how data is related to Label .
print('\nRelationship Plot :-\n')

plt.figure(figsize = (10,5), facecolor = 'white')
plotnumber = 1
for column in x:
    if plotnumber <=4:
        ax = plt.subplot(1,3, plotnumber)
        plt.scatter(x[column], y)
        plt.xlabel(column, fontsize = 18)
    plotnumber += 1
plt.tight_layout()
```

Relationship Plot :-




Newspaper not showing positive trends

Plotting Heatmap

```
In [20]: df_corr = df.corr().abs()
```

```
plt.figure(figsize = (11,8))
sns.heatmap(df_corr, vmin = -1, annot = True, square = True, center = 0, fmt = ".2f")
plt.tight_layout()
```



Observation : TV show maximum relation and newspaper show minimum relation with sales.

Checking for skewness

```
In [21]: x.skew()
```

```
Out[21]: TV      -0.082332
          radio    0.114842
          newspaper 0.650112
          dtype: float64
```

Observation : No skewness is present in dataset

Data Scaling

```
In [22]: scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
Out[22]: array([[ 0.97869734,  0.98952135,  1.93299778],
       [-1.19901165,  1.09070498,  0.75131275],
       [-1.51933199,  1.53591293,  1.93790103],
       [ 0.05645636,  1.22561648,  1.40834924],
       [ 0.40024339, -0.83178391,  1.40344598],
       [-1.61906543,  1.73828018,  2.2173867 ],
       [-1.04647815,  0.6522426 , -0.30779084],
       [-0.31079737, -0.23817331, -0.89127846],
       [-1.62023876, -1.41864895, -1.41102374],
       [ 0.62317696, -1.38492107, -0.42056576],
       [-0.94557138, -1.16906267, -0.27346804],
       [ 0.79800381,  0.05863199, -1.26392602],
       [-1.44189191,  0.80739083,  1.77119028],
       [-0.57714432, -1.04764232, -1.10702179],
       [ 0.67363035,  0.65898817,  0.79544207],
       [ 0.57155024,  1.65733328,  1.13376683],
       [ 1.58061798,  1.1109417 ,  1.27596129],
       [-0.90919801, -0.17746313, -0.56276022],
       [ 0.0071763 ,  0.05188642, -0.52353416],
       [-0.61141718, -0.20001007,  1.15000011]] )
```

OLS Result

```
In [23]: lm = smf.ols(formula = 'sales ~ TV + radio + newspaper', data = df).fit()
lm.summary()
```

Out[23]: OLS Regression Results

Dep. Variable:	sales	R-squared:	0.895			
Model:	OLS	Adj. R-squared:	0.894			
Method:	Least Squares	F-statistic:	553.5			
Date:	Mon, 12 Jul 2021	Prob (F-statistic):	8.35e-95			
Time:	17:13:34	Log-Likelihood:	-383.24			
No. Observations:	198	AIC:	774.5			
Df Residuals:	194	BIC:	787.6			
Df Model:	3					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.9523	0.318	9.280	0.000	2.325	3.580

Split data into train and test. Model will be bulit on training data and tested on test data

```
In [24]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25, random_state = 42)
print('Data has been splited.')
```

Data has been splited.

Model Building

Linear Regression model instantiaing, training and evaluating

```
In [25]: Lr = LinearRegression()
Lr.fit(x_train, y_train)
y_pred = Lr.predict(x_test)
```

```
In [26]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', Lr.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.8532243062942227  
=====  
RMSE of Model -----> 2.1658462808214387  
=====  
MSE of Model -----> 4.690890112148059  
=====  
Score of test data ----> 0.8532243062942227  
=====
```

Here we use Lasso Regularization to avoid overfitting

```
In [27]: # LassoCV will return best alpha after max itratation.  
# Normalize is subtracting the mean and dividing by the L2-norm.  
  
lassocv = LassoCV(alphas = None, max_iter = 1000, normalize = True)  
lassocv.fit(x_train,y_train)
```

```
Out[27]: LassoCV(normalize=True)
```

```
In [28]: alpha = lassocv.alpha_  
alpha
```

```
Out[28]: 0.006607312321213054
```

```
In [29]: lasso_reg = Lasso(alpha)  
lasso_reg.fit(x_train, y_train)
```

```
Out[29]: Lasso(alpha=0.006607312321213054)
```

```
In [30]: print('Lasso Score =====>', lasso_reg.score(x_train, y_train))
```

```
Lasso Score =====> 0.9069044644990196
```

Conclusion : Linear Regression model has 90% score

Knn model instantiaing, training and evaluating

```
In [31]: Knn = KNeighborsRegressor()  
Knn.fit(x_train, y_train)  
y_pred = Knn.predict(x_test)
```

```
In [32]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', Knn.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.9296055918625223  
=====  
RMSE of Model -----> 1.4999253314748706  
=====  
MSE of Model -----> 2.2497760000000007  
=====  
Score of test data ----> 0.9296055918625223  
=====
```

Here we use Lasso Regularization to avoid overfitting

```
In [33]: # LassoCV will return best alpha after max iteration.  
# Normalize is subtracting the mean and dividing by the L2-norm.  
  
lassocv = LassoCV(alphas = None, max_iter = 1000, normalize = True)  
lassocv.fit(x_train,y_train)
```

```
Out[33]: LassoCV(normalize=True)
```

```
In [34]: alpha = lassocv.alpha_  
alpha
```

```
Out[34]: 0.006607312321213054
```

```
In [35]: lasso_reg = Lasso(alpha)  
lasso_reg.fit(x_train, y_train)
```

```
Out[35]: Lasso(alpha=0.006607312321213054)
```

```
In [36]: print('Lasso Score =====>', lasso_reg.score(x_train, y_train))
```

```
Lasso Score =====> 0.9069044644990196
```

Conclusion : Knn model has 90% score

Decision Tree model instantiaing, training and evaluating

```
In [37]: DT = DecisionTreeRegressor()  
DT.fit(x_train, y_train)  
y_pred = DT.predict(x_test)
```

```
In [38]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ---->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ---->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', DT.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.9318451704502787  
=====  
RMSE of Model ----> 1.4758726232300674  
=====  
MSE of Model ----> 2.1782000000000004  
=====  
Score of test data ----> 0.9318451704502787  
=====
```

Here we use Lasso Regularization to avoid overfitting

```
In [39]: # LassoCV will return best alpha after max itratation.  
# Normalize is subtracting the mean and dividing by the L2-norm.  
  
lassocv = LassoCV(alphas = None, max_iter = 1000, normalize = True)  
lassocv.fit(x_train,y_train)
```

```
Out[39]: LassoCV(normalize=True)
```

```
In [40]: alpha = lassocv.alpha_  
alpha
```

```
Out[40]: 0.006607312321213054
```

```
In [41]: lasso_reg = Lasso(alpha)  
lasso_reg.fit(x_train, y_train)
```

```
Out[41]: Lasso(alpha=0.006607312321213054)
```

```
In [42]: print('Lasso Score =====>', lasso_reg.score(x_train, y_train))
```

```
Lasso Score =====> 0.9069044644990196
```

Conclusion : Decision Tree model has 90% score

Random Forest model instantiaing, training and evaluating

```
In [43]: Rn = RandomForestRegressor()  
Rn.fit(x_train, y_train)  
y_pred = Rn.predict(x_test)
```

```
In [44]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ---->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ---->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', Rn.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.9590715210811256  
=====  
RMSE of Model ----> 1.14370326571187  
=====  
MSE of Model ----> 1.3080571599999964  
=====  
Score of test data ----> 0.9590715210811256  
=====
```

Here we use Lasso Regularization to avoid overfitting

```
In [45]: # LassoCV will return best alpha after max itratation.  
# Normalize is subtracting the mean and dividing by the L2-norm.  
  
lassocv = LassoCV(alphas = None, max_iter = 1000, normalize = True)  
lassocv.fit(x_train,y_train)
```

```
Out[45]: LassoCV(normalize=True)
```

```
In [46]: alpha = lassocv.alpha_  
alpha
```

```
Out[46]: 0.006607312321213054
```

```
In [47]: lasso_reg = Lasso(alpha)  
lasso_reg.fit(x_train, y_train)
```

```
Out[47]: Lasso(alpha=0.006607312321213054)
```

```
In [48]: print('Lasso Score =====>', lasso_reg.score(x_train, y_train))
```

```
Lasso Score =====> 0.9069044644990196
```

Conclusion : Random Forest model has 90% score

SVM model instantiaing, training and evaluating

```
In [49]: svr = SVR()
svr.fit(x_train, y_train)
y_pred = svr.predict(x_test)
```

```
In [50]: print('=====')
print('R2 Score ---->', r2_score(y_test, y_pred))
print('=====')
print('RMSE of Model ---->', np.sqrt(mean_squared_error(y_test, y_pred)))
print('=====')
print('MSE of Model ---->', mean_squared_error(y_test, y_pred))
print('=====')
print('Score of test data ---->', svr.score(x_test, y_test))
print('=====')
```

```
=====
R2 Score ----> 0.8431536609343284
=====
RMSE of Model ----> 2.2389157528727637
=====
MSE of Model ----> 5.012743748461815
=====
Score of test data ----> 0.8431536609343284
=====
```

Here we use Lasso Regularization to avoid overfitting

```
In [51]: # LassoCV will return best alpha after max itratation.
# Normalize is subtracting the mean and dividing by the L2-norm.

lassocv = LassoCV(alphas = None, max_iter = 1000, normalize = True)
lassocv.fit(x_train,y_train)
```

```
Out[51]: LassoCV(normalize=True)
```

```
In [52]: alpha = lassocv.alpha_
alpha
```

```
Out[52]: 0.006607312321213054
```

```
In [53]: lasso_reg = Lasso(alpha)
lasso_reg.fit(x_train, y_train)
```

```
Out[53]: Lasso(alpha=0.006607312321213054)
```

```
In [54]: print('Lasso Score =====>', lasso_reg.score(x_train, y_train))
```

```
Lasso Score =====> 0.9069044644990196
```

Conclusion : SVM model has 90% score

Looking R2 score we found Random Forest has best model so we do Hyperparameter Tuning on it.

```
In [55]: param_grid = {'n_estimators': [100, 200, 300, 400, 500],  
                    'max_features': ['auto', 'sqrt'],  
                    'max_depth': [5, 10, 15, 20, 25, 30],  
                    'min_samples_split': [2, 5, 10, 15, 100],  
                    'min_samples_leaf': [1, 2, 5, 10]}
```

```
In [57]: grid_search = GridSearchCV(estimator = Rn, param_grid = param_grid, cv = 5,n_jobs=-1)
```

```
In [58]: grid_search.fit(x_train, y_train)
```

```
Out[58]: GridSearchCV(cv=5, estimator=RandomForestRegressor(), n_jobs=-1,  
                      param_grid={'max_depth': [5, 10, 15, 20, 25, 30],  
                                  'max_features': ['auto', 'sqrt'],  
                                  'min_samples_leaf': [1, 2, 5, 10],  
                                  'min_samples_split': [2, 5, 10, 15, 100],  
                                  'n_estimators': [100, 200, 300, 400, 500]})
```

```
In [59]: best_parameters = grid_search.best_params_  
print(best_parameters)
```

```
{'max_depth': 25, 'max_features': 'auto', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 300}
```

```
In [61]: hRn = RandomForestRegressor(max_depth = 25, max_features = 'auto', min_samples_le  
hRn.fit(x_train, y_train)  
hRn.score(x_test, y_test)
```

```
Out[61]: 0.9591851892273268
```

```
In [62]: y_pred = hRn.predict(x_test)
```

```
In [63]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', hRn.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.9591851892273268  
=====  
RMSE of Model -----> 1.1421139931431339  
=====  
MSE of Model -----> 1.3044243733333547  
=====  
Score of test data ----> 0.9591851892273268  
=====
```

After Hyperparameter Tuning model accuracy score 95%.

Saving The Model

```
In [64]: # saving the model to the Local file system  
filename = 'Advertising Sales Channel Prediction.pickle'  
pickle.dump(hRn, open(filename, 'wb'))
```

Final Conclusion : Random Forest is our best model.

```
In [ ]:
```