

## **Problem Statement:**

**Avocado is a fruit consumed by people heavily in the United States.**

**Content :**

**This data was downloaded from the Hass Avocado Board website in May of 2018 & compiled into a single CSV.**

**The table below represents weekly 2018 retail scan data for National retail volume (units) and price. Retail scan data comes directly from retailers' cash registers based on actual retail sales of Hass avocados.**

**Starting in 2013, the table below reflects an expanded, multi-outlet retail data set. Multi-outlet reporting includes an aggregation of the following channels: grocery, mass, club, drug, dollar and military. The Average Price (of avocados) in the table reflects a per unit (per avocado) cost, even when multiple units (avocados) are sold in bags.**

**The Product Lookup codes (PLU's) in the table are only for Hass avocados. Other varieties of avocados (e.g. greenskins) are not included in this table.**

**Some relevant columns in the dataset:**

**Date - The date of the observation** **AveragePrice - the average price of a single avocado type - conventional or organic** **year - the year** **Region - the city or region of the observation** **Total Volume - Total number of avocados sold** **4046 - Total number of avocados with PLU 4046 sold** **4225 - Total number of avocados with PLU 4225 sold** **4770 - Total number of avocados with PLU 4770 sold**

**Inspiration /Label**

**Your task is to make a mode that can consider the data provided and predict the Average Price.**

## **Importing Required Library**

```
In [45]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pickle
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
import xgboost as xgb
from sklearn.metrics import r2_score, mean_squared_error
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

## Reading Data

```
In [2]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\csv file\avocado.csv")
df.head()
```

Out[2]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags	Large Bags
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62	93.25
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07	97.49
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042.21	103.14
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16	5677.40	133.76
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95	5986.26	197.69

## Check no of row and column

```
In [3]: print('No of Rows and Columns ----->', df.shape)
```

No of Rows and Columns -----> (18249, 14)

## Checking for Null values

```
In [4]: print('-----\n')
print(df.isnull().sum())
print('\n-----')
```

```
-----  
Unnamed: 0      0  
Date            0  
AveragePrice    0  
Total Volume    0  
4046            0  
4225            0  
4770            0  
Total Bags      0  
Small Bags      0  
Large Bags      0  
XLarge Bags     0  
type             0  
year             0  
region           0  
dtype: int64  
-----
```

**There is no null value**

## Information about dataset

```
In [5]: print('-----\n')
print(df.info())
print('\n-----')
```

```
-----  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 18249 entries, 0 to 18248  
Data columns (total 14 columns):  
 #   Column           Non-Null Count  Dtype     
---  --    
 0   Unnamed: 0        18249 non-null   int64    
 1   Date              18249 non-null   object    
 2   AveragePrice      18249 non-null   float64   
 3   Total Volume      18249 non-null   float64   
 4   4046              18249 non-null   float64   
 5   4225              18249 non-null   float64   
 6   4770              18249 non-null   float64   
 7   Total Bags        18249 non-null   float64   
 8   Small Bags        18249 non-null   float64   
 9   Large Bags        18249 non-null   float64   
 10  XLarge Bags       18249 non-null   float64   
 11  type              18249 non-null   object    
 12  year              18249 non-null   int64    
 13  region            18249 non-null   object    
dtypes: float64(9), int64(2), object(3)  
memory usage: 1.9+ MB  
None
```

  
-----

Categorical data present in our data set

## Statistics of Data

```
In [6]: df.describe()
```

Out[6]:

	Unnamed: 0	AveragePrice	Total Volume	4046	4225	4770	To
count	18249.000000	18249.000000	1.824900e+04	1.824900e+04	1.824900e+04	1.824900e+04	1.824
mean	24.232232	1.405978	8.506440e+05	2.930084e+05	2.951546e+05	2.283974e+04	2.396
std	15.481045	0.402677	3.453545e+06	1.264989e+06	1.204120e+06	1.074641e+05	9.862
min	0.000000	0.440000	8.456000e+01	0.000000e+00	0.000000e+00	0.000000e+00	0.000
25%	10.000000	1.100000	1.083858e+04	8.540700e+02	3.008780e+03	0.000000e+00	5.088
50%	24.000000	1.370000	1.073768e+05	8.645300e+03	2.906102e+04	1.849900e+02	3.974
75%	38.000000	1.660000	4.329623e+05	1.110202e+05	1.502069e+05	6.243420e+03	1.107
max	52.000000	3.250000	6.250565e+07	2.274362e+07	2.047057e+07	2.546439e+06	1.937

Outliers are present in data

## Features engineering

### Add Months Column

```
In [7]: df['Months'] = df['Date'].str[5:7]
df['Months'] = df['Months'].astype('int32')
df.head(2)
```

Out[7]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags	Large Bags	XI
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62	93.25	
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07	97.49	

Month column added and convert into int32

```
In [8]: print('-----\n')
print(df.info())
print('\n-----')
```

```
-----  

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        18249 non-null   int64  
 1   Date              18249 non-null   object  
 2   AveragePrice      18249 non-null   float64 
 3   Total Volume      18249 non-null   float64 
 4   4046              18249 non-null   float64 
 5   4225              18249 non-null   float64 
 6   4770              18249 non-null   float64 
 7   Total Bags        18249 non-null   float64 
 8   Small Bags        18249 non-null   float64 
 9   Large Bags        18249 non-null   float64 
 10  XLarge Bags       18249 non-null   float64 
 11  type              18249 non-null   object  
 12  year              18249 non-null   int64  
 13  region             18249 non-null   object  
 14  Months             18249 non-null   int32  
dtypes: float64(9), int32(1), int64(2), object(3)
memory usage: 2.0+ MB
None
```

```
-----
```

**Months column is converted into int32**

## Add Sales column

```
In [9]: df['Sales']= df['Total Volume'] * df['AveragePrice']
df.head(2)
```

Out[9]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags	Large Bags	Xl
0	0	2015-12-27		1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62	93.25
1	1	2015-12-20		1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07	97.49

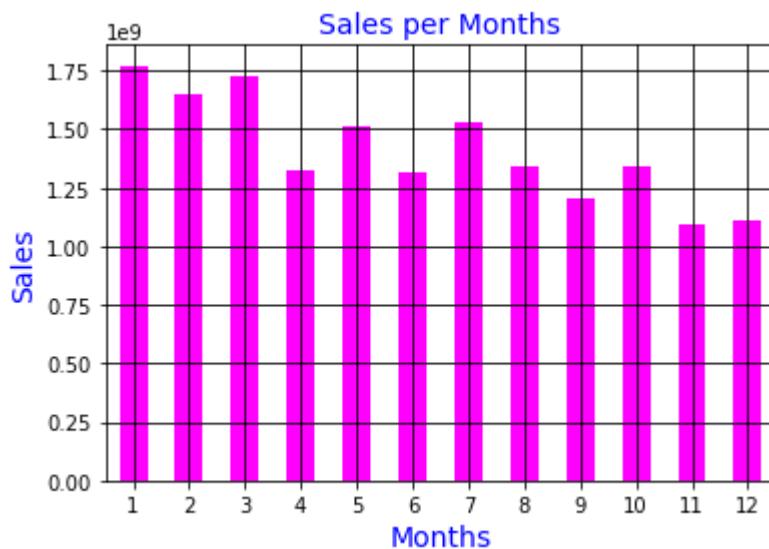
**Sales column added in our dataset**

## To know highest month for sales

```
In [10]: s = df.groupby('Months')['Sales'].sum()  
s
```

```
Out[10]: Months  
1    1.771539e+09  
2    1.648614e+09  
3    1.730009e+09  
4    1.327835e+09  
5    1.515673e+09  
6    1.312736e+09  
7    1.529586e+09  
8    1.343889e+09  
9    1.207088e+09  
10   1.342014e+09  
11   1.097177e+09  
12   1.107970e+09  
Name: Sales, dtype: float64
```

```
In [11]: s.plot.bar(x = 'Months', y = 'Sales', rot = 360,color = 'magenta')  
plt.grid(c = 'black')  
plt.ylabel('Sales',fontsize = 14, color = 'b')  
plt.xlabel('Months',fontsize = 14, color = 'b')  
plt.title('Sales per Months',fontsize = 14, color = 'b')  
plt.show()
```



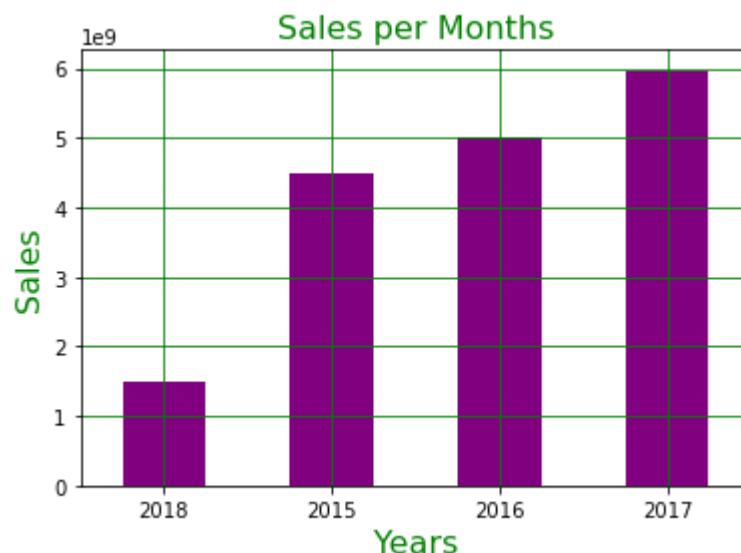
January is best month for sales

## To know highest year of sales

```
In [12]: y = df.groupby('year')['Sales'].sum().sort_values()
y
```

```
Out[12]: year
2018    1.482343e+09
2015    4.488448e+09
2016    4.997590e+09
2017    5.965750e+09
Name: Sales, dtype: float64
```

```
In [13]: y.plot.bar(x = 'year', y = 'Sales', rot = 360,color = 'purple')
plt.grid(c = 'g')
plt.ylabel('Sales',fontsize = 16, color = 'g')
plt.xlabel('Years',fontsize = 16, color = 'g')
plt.title('Sales per Months',fontsize = 16, color = 'g')
plt.show()
```



In 2017 highest sales are done

```
In [14]: df['XLarge Bags'].value_counts()
```

```
Out[14]: 0.00      12048
3.33       29
6.67       16
1.11       15
5.00       12
...
9.95        1
3.47        1
59.32       1
16090.51    1
4920.17     1
Name: XLarge Bags, Length: 5588, dtype: int64
```

Above we check 'XLarge Bags' column we have doubt it is zero value but it is not

## To know which type of avocado highest sales

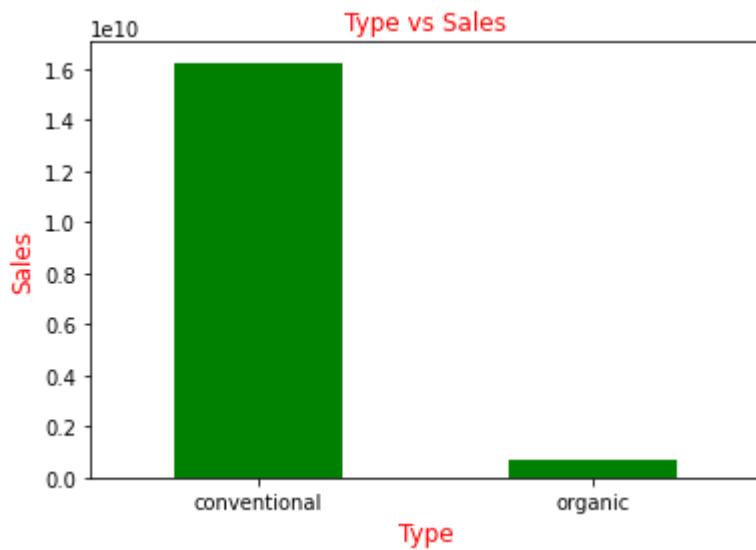
```
In [15]: df['type'].value_counts()
```

```
Out[15]: conventional    9126
          organic        9123
          Name: type, dtype: int64
```

```
In [16]: t = df.groupby('type')['Sales'].sum()
t
```

```
Out[16]: type
          conventional    1.625352e+10
          organic        6.806085e+08
          Name: Sales, dtype: float64
```

```
In [17]: t.plot.bar(x = 'type', y = 'Sales', rot = 360, color = 'g')
plt.ylabel('Sales', fontsize = 12, color = 'r')
plt.xlabel('Type', fontsize = 12, color = 'r')
plt.title('Type vs Sales', fontsize = 12, color = 'r')
plt.show()
```



Conventional type has highest sales

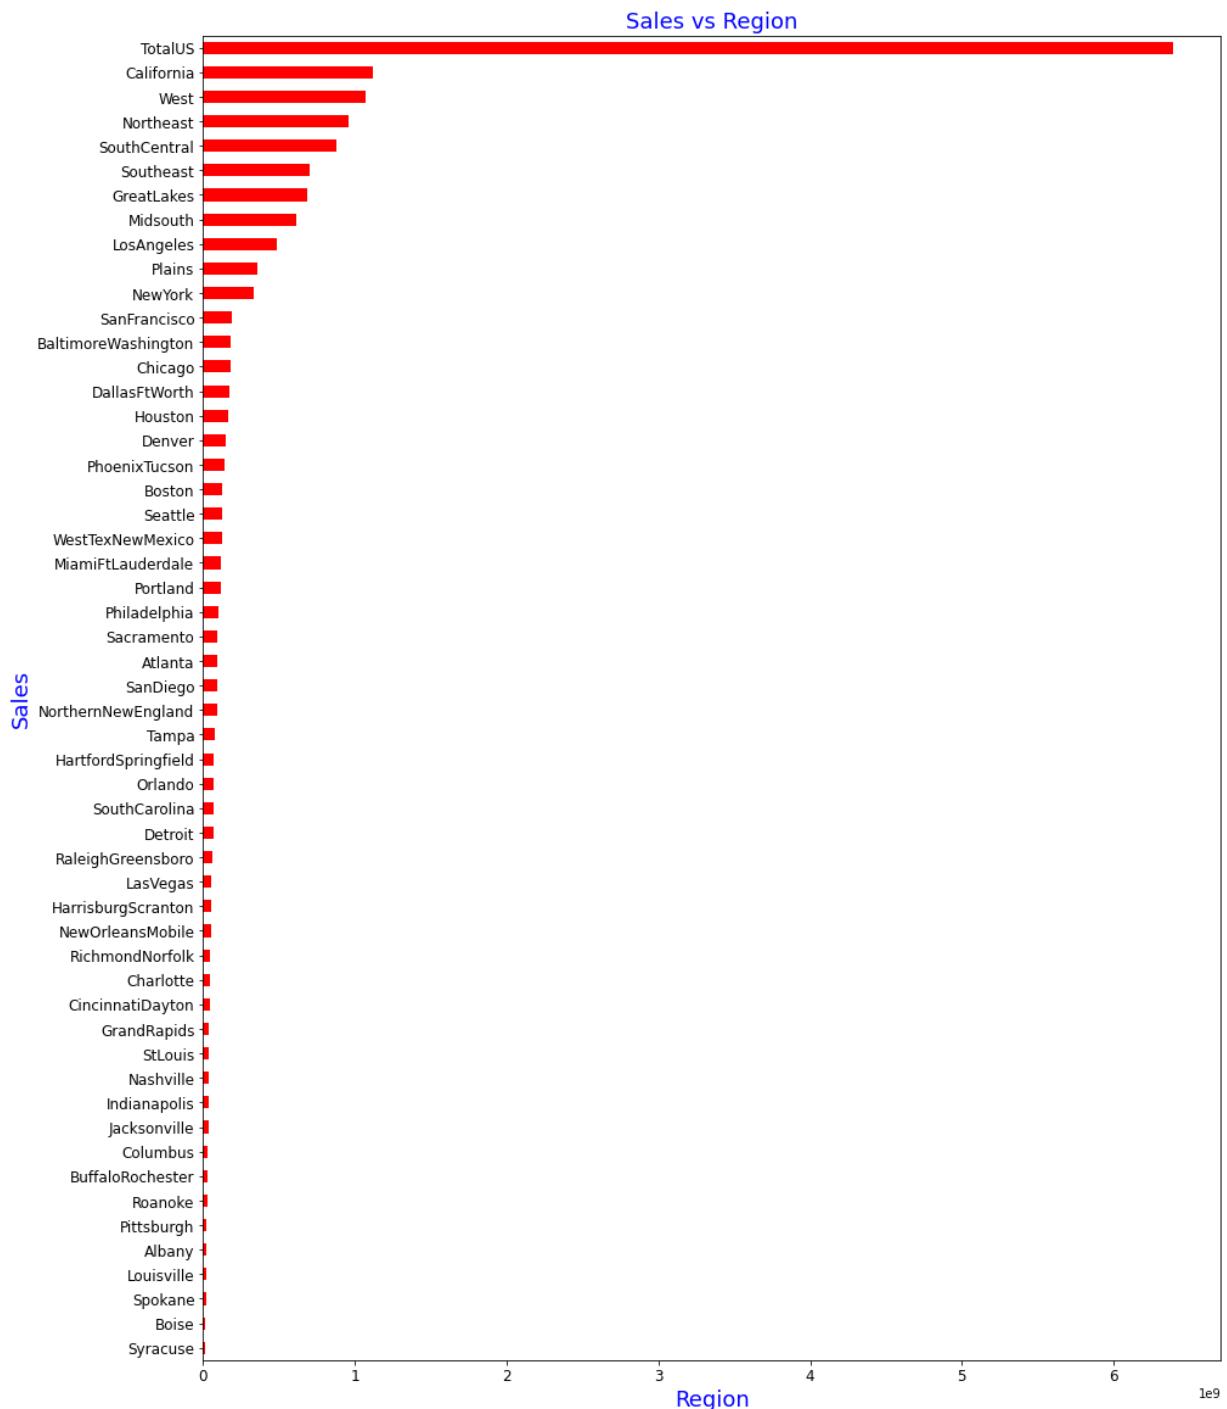
## To know which region has highest sales

```
In [18]: r = df.groupby('region')['Sales'].sum().sort_values()  
r
```

```
Out[18]: region  
Syracuse           1.520519e+07  
Boise              1.534667e+07  
Spokane             1.715649e+07  
Louisville          1.749555e+07  
Albany              2.176672e+07  
Pittsburgh          2.319481e+07  
Roanoke              2.740928e+07  
BuffaloRochester    3.154509e+07  
Columbus             3.156295e+07  
Jacksonville         3.348396e+07  
Indianapolis         3.402474e+07  
Nashville             3.572664e+07  
StLouis              3.743946e+07  
GrandRapids          3.785124e+07  
CincinnatiDayton     4.498958e+07  
Charlotte             4.574304e+07  
RichmondNorfolk       4.734482e+07  
NewOrleansMobile      4.867626e+07  
HarrisburgScranton    5.295472e+07  
LasVegas              5.480905e+07  
RaleighGreensboro     5.983780e+07  
Detroit                6.912624e+07  
SouthCarolina          6.952808e+07  
Orlando                6.961121e+07  
HartfordSpringfield    7.118645e+07  
Tampa                  7.695906e+07  
NorthernNewEngland      8.962525e+07  
SanDiego               9.352710e+07  
Atlanta                 9.379337e+07  
Sacramento              9.480870e+07  
Philadelphia            9.981583e+07  
Portland                 1.128646e+08  
MiamiFtLauderdale        1.185815e+08  
WestTexNewMexico          1.215654e+08  
Seattle                  1.251791e+08  
Boston                   1.265429e+08  
PhoenixTucson             1.384515e+08  
Denver                     1.459828e+08  
Houston                   1.655713e+08  
DallasFtWorth             1.756093e+08  
Chicago                     1.791106e+08  
BaltimoreWashington        1.799084e+08  
SanFrancisco              1.858341e+08  
NewYork                   3.351955e+08  
Plains                     3.600366e+08  
LosAngeles                4.842276e+08  
Midsouth                  6.157238e+08  
GreatLakes                 6.886618e+08  
Southeast                  7.036306e+08  
SouthCentral                8.740593e+08  
Northeast                  9.600079e+08  
West                      1.066834e+09
```

```
California           1.121414e+09
TotalUS            6.387593e+09
Name: Sales, dtype: float64
```

```
In [19]: r.plot.barh(x = 'region', y = 'Sales', rot = 360, figsize = (15,20), color = 'r',
plt.ylabel('Sales', fontsize = 18, color = 'b')
plt.xlabel('Region', fontsize = 18, color = 'b')
plt.title('Sales vs Region', fontsize = 18, color = 'b')
plt.show()
```



US region has highest sales

## Label Encoder

```
In [20]: le = LabelEncoder()
```

```
In [21]: df['type'] = le.fit_transform(df['type'])
```

```
In [22]: df['region'] = le.fit_transform(df['region'])
```

```
In [23]: df.head(2)
```

Out[23]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags	Large Bags	XI
0	0	2015- 12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62	93.25	
1	1	2015- 12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07	97.49	

```
In [24]: print('-----\n')
print(df.info())
print('\n-----')
```

```
-----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        18249 non-null   int64  
 1   Date              18249 non-null   object  
 2   AveragePrice      18249 non-null   float64 
 3   Total Volume     18249 non-null   float64 
 4   4046              18249 non-null   float64 
 5   4225              18249 non-null   float64 
 6   4770              18249 non-null   float64 
 7   Total Bags        18249 non-null   float64 
 8   Small Bags        18249 non-null   float64 
 9   Large Bags        18249 non-null   float64 
 10  XLarge Bags       18249 non-null   float64 
 11  type              18249 non-null   int32  
 12  year              18249 non-null   int64  
 13  region             18249 non-null   int32  
 14  Months             18249 non-null   int32  
 15  Sales              18249 non-null   float64 
dtypes: float64(10), int32(3), int64(2), object(1)
memory usage: 2.0+ MB
None
-----
```

Columns are encoded

## Delete Unwanted Columns

```
In [25]: col = ['Unnamed: 0', 'Date']
```

```
In [26]: df = df.drop(col, axis = 1)
df.head()
```

Out[26]:

	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags	Large Bags	XLarge Bags	type	ye
0	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62	93.25	0.0	0	20
1	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07	97.49	0.0	0	20
2	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042.21	103.14	0.0	0	20
3	1.08	78992.15	1132.00	71976.41	72.58	5811.16	5677.40	133.76	0.0	0	20
4	1.28	51039.60	941.48	43838.39	75.78	6183.95	5986.26	197.69	0.0	0	20

```
In [27]: print('No of Rows and Columns ----->', df.shape )
```

No of Rows and Columns -----> (18249, 14)

Columns has been deleted

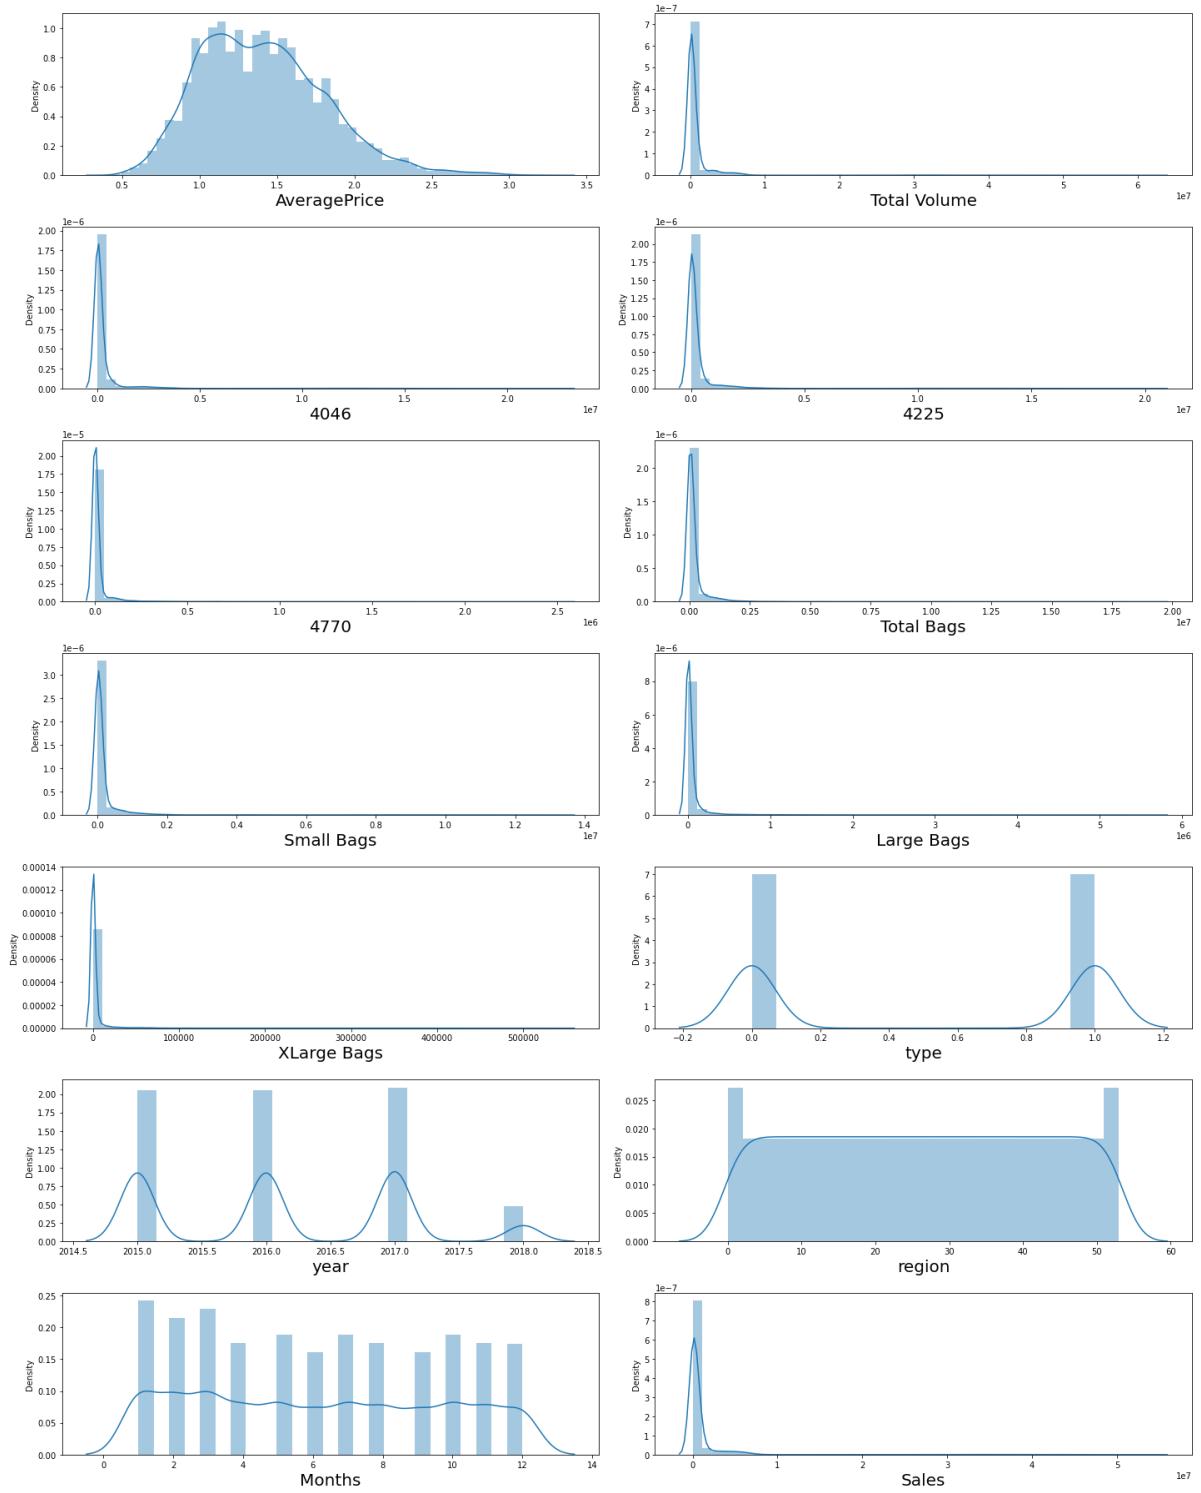
## Data distribution and checking outliers

```
In [28]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in df:
    if plotnumber <=14:
        ax = plt.subplot(7,2, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----
Distribution Plot :-
```



```
In [29]: df.skew()
```

```
Out[29]: AveragePrice      0.580303
Total Volume       9.007687
4046              8.648220
4225              8.942466
4770              10.159396
Total Bags        9.756072
Small Bags        9.540660
Large Bags        9.796455
XLarge Bags       13.139751
type              0.000329
year              0.215339
region            0.000030
Months            0.106617
Sales              8.915413
dtype: float64
```

**Data has outliers and skewed**

**Corelation of Feature vs Label using Heat map**

```
In [30]: print('-----')
```

```
print('Heat Map :-')
```

```
print('-----')
```

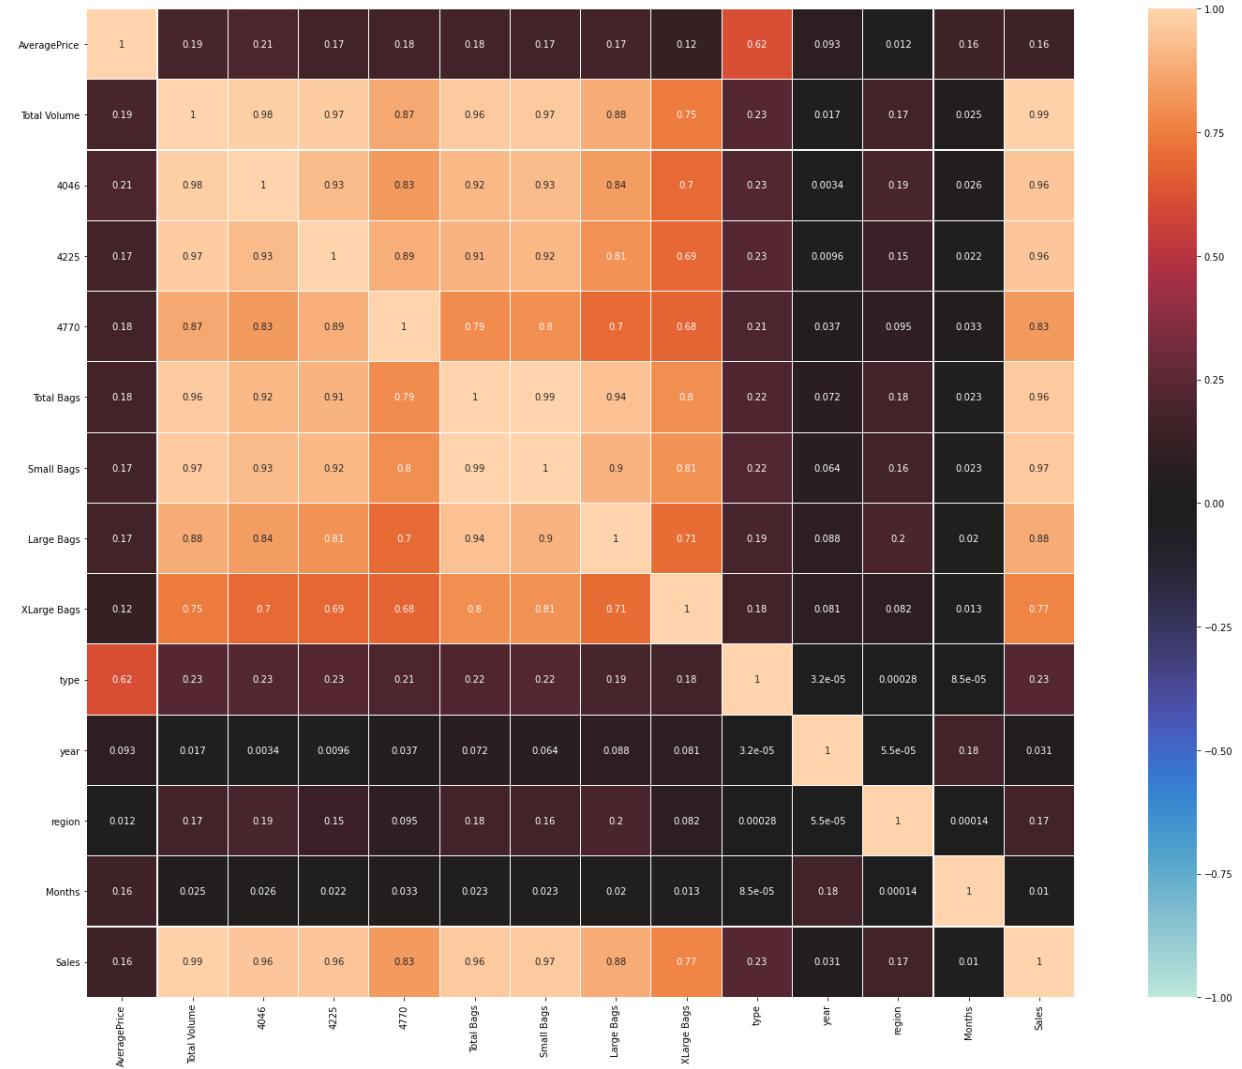
```
df_corr = df.corr().abs()
```

```
plt.figure(figsize = (22,16))
```

```
sns.heatmap(df_corr, vmin = -1, annot = True, square = True, center = 0, fmt = '.2f')
```

```
plt.tight_layout()
```

-----  
Heat Map :-  
-----



Total Volume has highest and months has lowest relation with label

### Checking class imbalance

```
In [31]: df['AveragePrice'].value_counts()
```

```
Out[31]: 1.15      202
1.18      199
1.08      194
1.26      193
1.13      192
...
2.91       1
2.68       1
3.04       1
3.17       1
3.03       1
Name: AveragePrice, Length: 259, dtype: int64
```

Class is not balanced

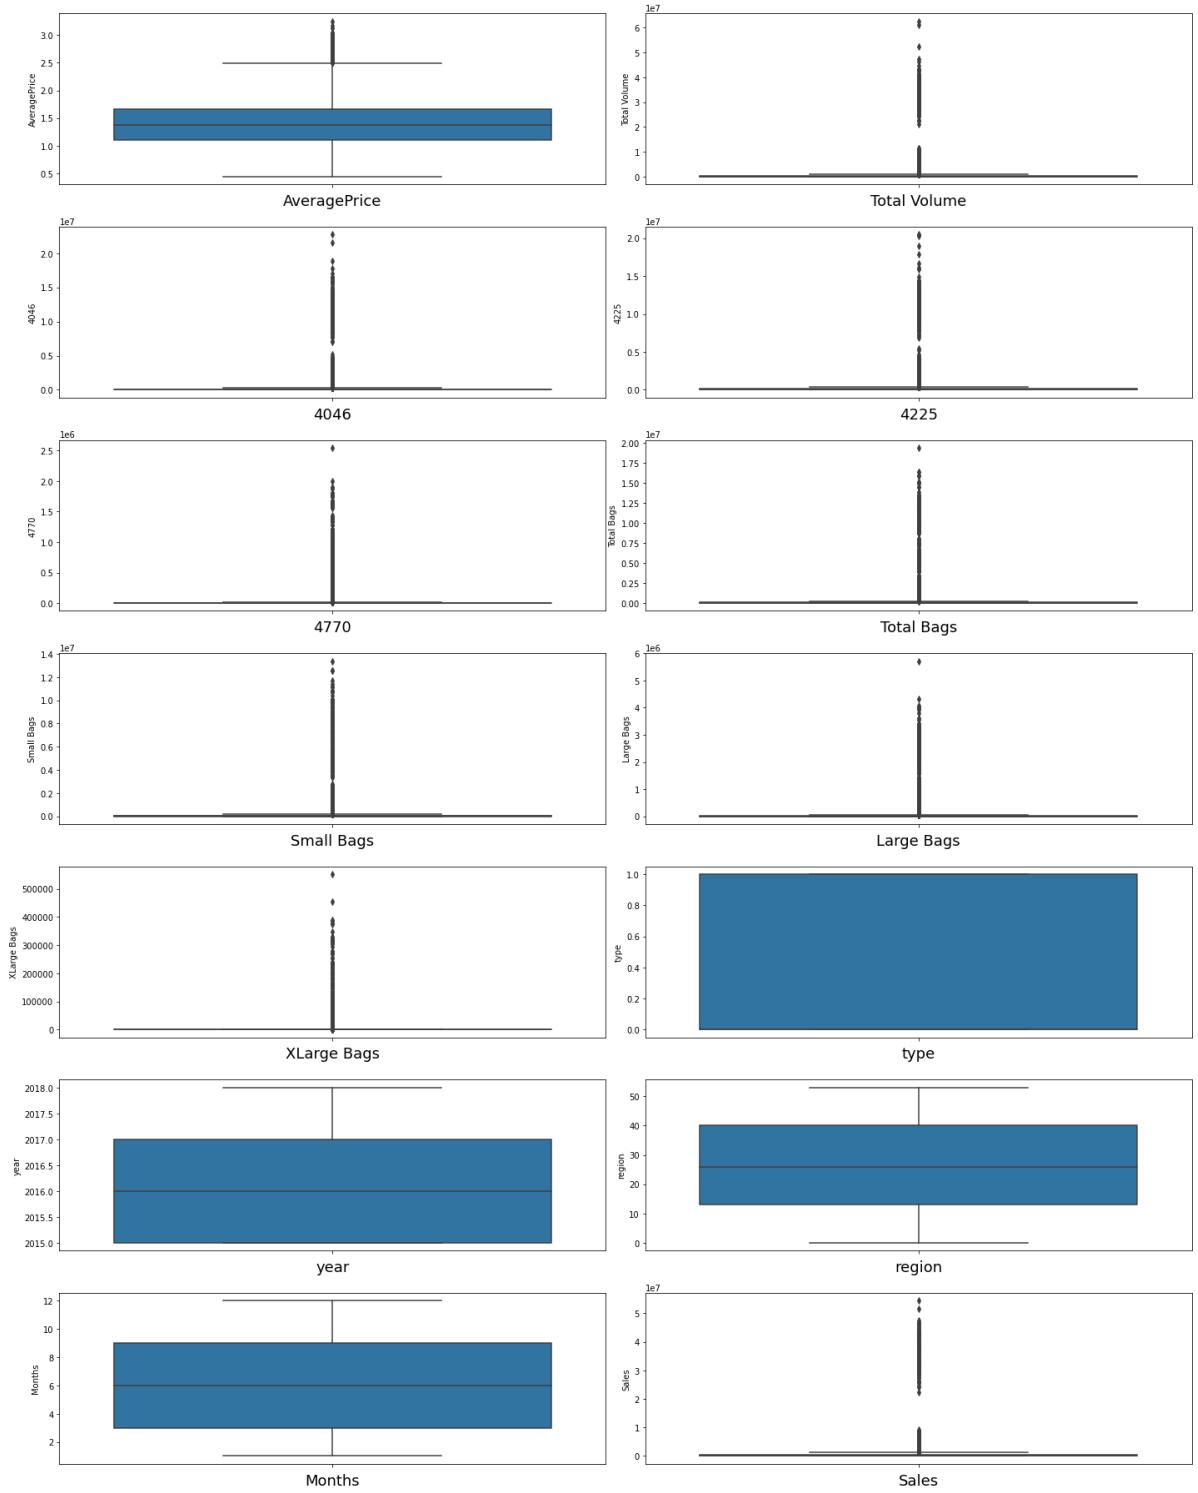
Quntile metthod to removing outliers and skewness.

```
In [32]: # we are removing the top 2% data from the Total Volume column
q = df['Total Volume'].quantile(0.98)
data_cleaned = df[df['Total Volume']<q]
# we are removing the top 2% data from the 4046 column
q = df['4046'].quantile(0.98)
data_cleaned = data_cleaned[data_cleaned['4046']<q]
# we are removing the top 2% data from the 4225 column
q = df['4225'].quantile(0.98)
data_cleaned = data_cleaned[data_cleaned['4225']<q]
# we are removing the top 2% data from the 4770 column
q = df['4770'].quantile(0.98)
data_cleaned = data_cleaned[data_cleaned['4770']<q]
# we are removing the top 2% data from the free Total Bags column
q = df['Total Bags'].quantile(0.98)
data_cleaned = data_cleaned[data_cleaned['Total Bags']<q]
# we are removing the top 2% data from the Small Bags column
q = df['Small Bags'].quantile(0.98)
data_cleaned = data_cleaned[data_cleaned['Small Bags']<q]
# we are removing the top 2% data from the Large Bags column
q = df['Large Bags'].quantile(0.98)
data_cleaned = data_cleaned[data_cleaned['Large Bags']<q]
# we are removing the top 2% data from the XLarge Bags column
q = df['XLarge Bags'].quantile(0.98)
data_cleaned = data_cleaned[data_cleaned['XLarge Bags']<q]
# we are removing the top 2% data from the Sales column
q = df['Sales'].quantile(0.98)
data_cleaned = data_cleaned[data_cleaned['Sales']<q]
```

## Checking Outliers and skewness removed or not

```
In [33]: # Let's see outliers are removed in columns or not.  
print('\nDistribution Plot :-\n')  
  
plt.figure(figsize = (20,25), facecolor = 'white')  
plotnumber = 1  
for column in df:  
    if plotnumber <=14:  
        ax = plt.subplot(7,2, plotnumber)  
        sns.boxplot(y=df[column]) # It is the axis for vertical set as y  
        plt.xlabel(column, fontsize = 18)  
    plotnumber += 1  
plt.tight_layout()
```

Distribution Plot :-



**Outliers are removed**

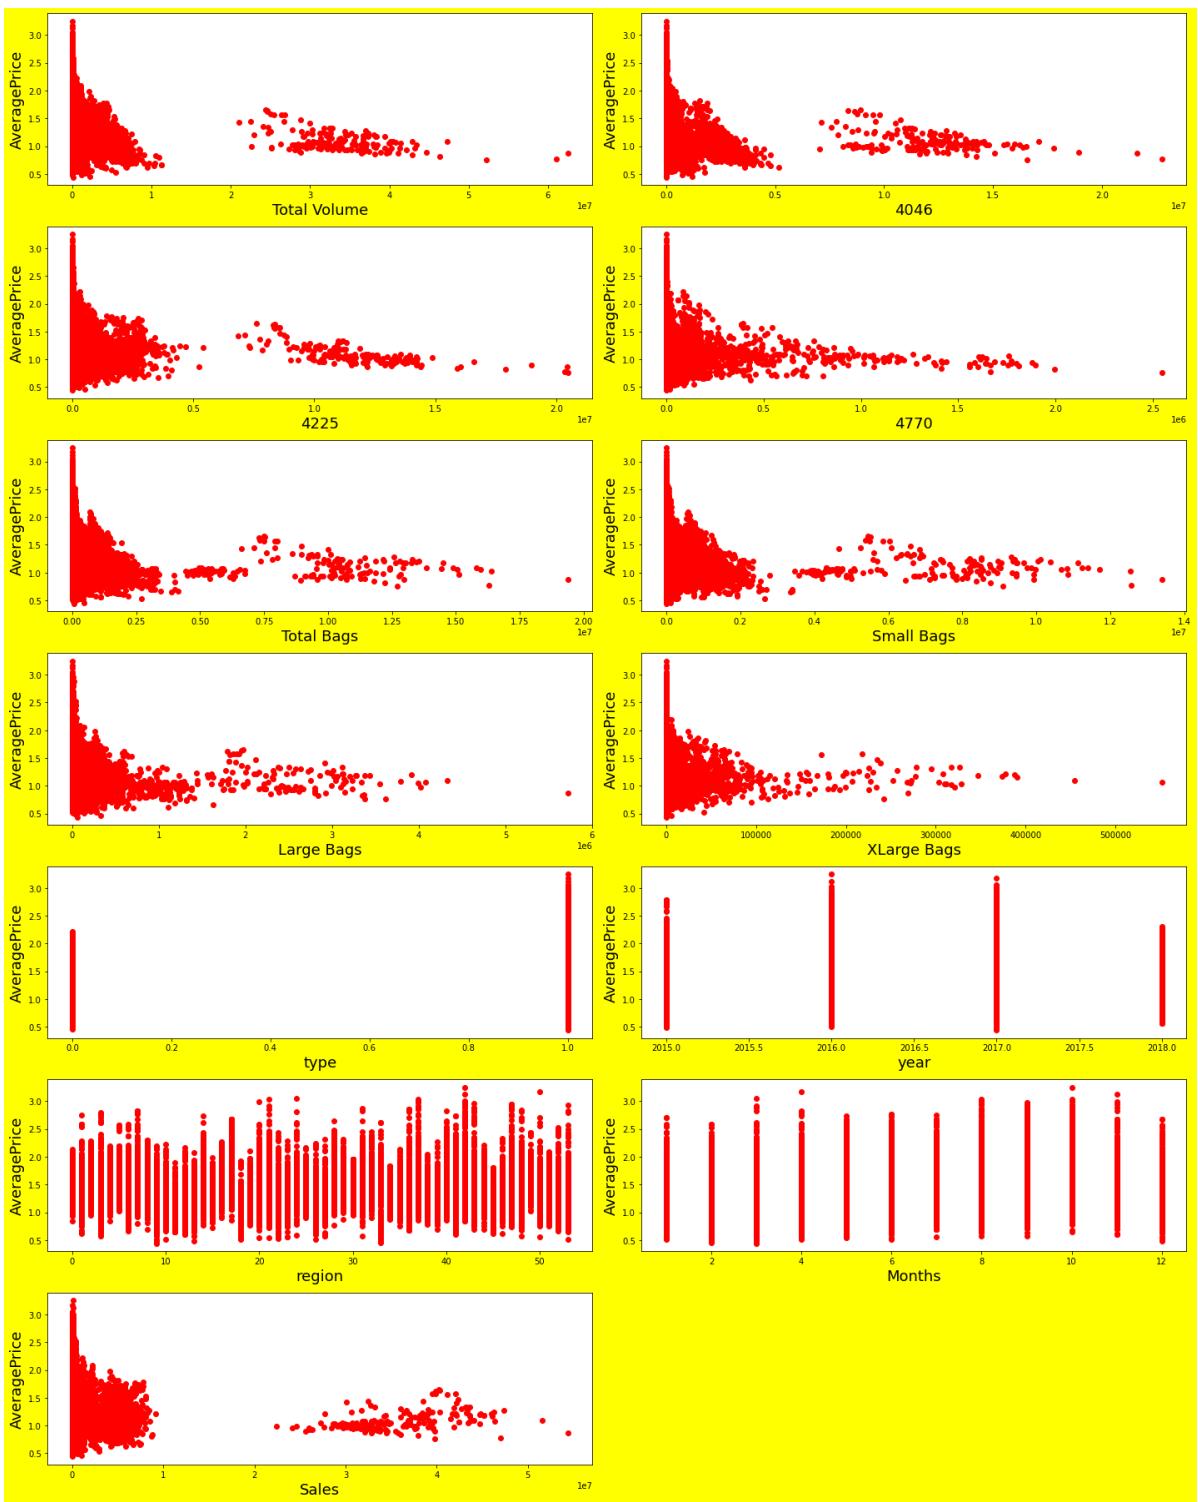
**Splitting Dataset into features and labels**

```
In [34]: x = df.drop('AveragePrice', axis = 1)
y = df. AveragePrice
print('Data has been splited')
```

Data has been splited

```
In [40]: # Let's see relation between features and labels.  
print('-----')  
print('Distribution Plot :-')  
print('-----')  
  
plt.figure(figsize = (20,25), facecolor = 'yellow')  
plotnumber = 1  
for column in x:  
    if plotnumber <=14:  
        ax = plt.subplot(7,2, plotnumber)  
        plt.scatter(x[column],y, c = 'r')  
        plt.xlabel(column, fontsize = 18)  
        plt.ylabel('AveragePrice', fontsize = 18)  
    plotnumber += 1  
plt.tight_layout()
```

```
-----  
Distribution Plot :-  
-----
```



**Features are related to label**

## Data Scaling

```
In [41]: scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
Out[41]: array([[-0.22771641, -0.23081597, -0.1999022 , ..., -1.7002522 ,
       1.64763162, -0.22857453],
      [-0.23042664, -0.23110251, -0.20805446, ..., -1.7002522 ,
       1.64763162, -0.23165401],
      [-0.21208462, -0.23100731, -0.1544779 , ..., -1.7002522 ,
       1.64763162, -0.22192493],
      ...,
      [-0.24233073, -0.2306933 , -0.24309014, ...,  1.70081131,
       -1.46495942, -0.24476871],
      [-0.24162464, -0.2304279 , -0.24265143, ...,  1.70081131,
       -1.46495942, -0.24326781],
      [-0.24125273, -0.22934712, -0.24317042, ...,  1.70081131,
       -1.46495942, -0.24406623]])
```

**Split data into train and test. Model will be bulit on training data and tested on test data**

```
In [42]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25, random_state=42)
print('Data has been splited.')
```

Data has been splited.

## Model Building

**Logistic Regression model instantiaing, training and evaluating**

```
In [46]: Lr = LinearRegression()
Lr.fit(x_train, y_train)
y_pred = Lr.predict(x_test)
```

```
In [47]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', Lr.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.4528681041347923  
=====  
RMSE of Model -----> 0.29939625971037953  
=====  
MSE of Model -----> 0.08963812032856502  
=====  
Score of test data ----> 0.4528681041347923  
=====
```

Conclusion : Linear Regression model has 45% score

## Knn model instantiaing, training and evaluating

```
In [49]: Knn = KNeighborsRegressor()  
Knn.fit(x_train, y_train)  
y_pred = Knn.predict(x_test)
```

```
In [50]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', Knn.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.9642415902557608  
=====  
RMSE of Model -----> 0.07654017521596569  
=====  
MSE of Model -----> 0.0058583984220907295  
=====  
Score of test data ----> 0.9642415902557608  
=====
```

Conclusion : Knn model has 96% score

## Decision Tree model instantiaing, training and evaluating

```
In [52]: DT = DecisionTreeRegressor()  
DT.fit(x_train, y_train)  
y_pred = DT.predict(x_test)
```

```
In [53]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ---->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ---->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', DT.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.8217825605874876  
=====  
RMSE of Model ----> 0.17087379053020385  
=====  
MSE of Model ----> 0.029197852290159986  
=====  
Score of test data ----> 0.8217825605874876  
=====
```

**Conclusion : Decision Tree model has 82% score**

## Random Forest model instantiaing, training and evaluating

```
In [55]: Rn = RandomForestRegressor()  
Rn.fit(x_train, y_train)  
y_pred = Rn.predict(x_test)
```

```
In [56]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', Rn.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.9252366407800727  
=====  
RMSE of Model -----> 0.11067378352550787  
=====  
MSE of Model -----> 0.012248686359850977  
=====  
Score of test data ----> 0.9252366407800727  
=====
```

Conclusion : Random Forest model has 92% score

## XGBoost model instantiaing, training and evaluating

```
In [57]: xgb = xgb.XGBRegressor(eval_metric='mlogloss')  
xgb.fit(x_train, y_train)  
y_pred = xgb.predict(x_test)
```

```
In [58]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', xgb.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.9632619783377186  
=====  
RMSE of Model -----> 0.07758151153188628  
=====  
MSE of Model -----> 0.006018890931572203  
=====  
Score of test data ----> 0.9632619783377186  
=====
```

Conclusion : XGBoost model has 96% score

**Looking R2 score we found KNN has best model so we do Hyperparameter Tuning on it.**

```
In [63]: grid_param = {'leaf_size' : [1,3,5], 'n_neighbors': [3],  
                    'p':[1,2,3,4]}
```

```
In [64]: grid_search = GridSearchCV(estimator = Knn, param_grid = grid_param, cv = 5 , n_
```

```
In [65]: grid_search.fit(x_train, y_train)
```

```
Out[65]: GridSearchCV(cv=5, estimator=KNeighborsRegressor(), n_jobs=-1,  
                      param_grid={'leaf_size': [1, 3, 5], 'n_neighbors': [3],  
                                  'p': [1, 2, 3, 4]})
```

```
In [66]: best_parameters = grid_search.best_params_  
print(best_parameters)
```

```
{'leaf_size': 1, 'n_neighbors': 3, 'p': 3}
```

```
In [68]: hknn = KNeighborsRegressor(leaf_size = 1, n_neighbors = 3, p = 3)  
hknn.fit(x_train, y_train)  
hknn.score(x_test, y_test)
```

```
Out[68]: 0.9663003004702655
```

```
In [69]: print('=====')  
print('R2 Score ----->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ----->', hknn.score(x_test, y_test))  
print('=====')
```

```
=====
```

```
R2 Score -----> 0.9632619783377186
```

```
=====
```

```
RMSE of Model -----> 0.07758151153188628
```

```
=====
```

```
MSE of Model -----> 0.006018890931572203
```

```
=====
```

```
Score of test data -----> 0.9663003004702655
```

```
=====
```

**After Hyperparameter Tuning model accuracy score increase to 96%.**

## Saving The Model

```
In [70]: # saving the model to the Local file system  
filename = 'Avacado Project.pickle'  
pickle.dump(hknn, open(filename, 'wb'))
```

Model has been saved.

**Final Conclusion : Knn is our best model.**

```
In [ ]:
```