

## Problem Statement:

The data scientists at BigMart have collected 2013 sales data for 1559 products across 10 stores in different cities. Also, certain attributes of each product and store have been defined. The aim is to build a predictive model and find out the sales of each product at a particular store. Using this model, BigMart will try to understand the properties of products and stores which play a key role in increasing the sales of their products.

## Importing Required Library

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pickle
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import PowerTransformer
import statsmodels.formula.api as smf
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.svm import SVR
import xgboost as xgb
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

## Reading Data

```
In [2]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\Data Analysis With Python\ML Files\Churn.csv")
df.head()
```

n_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Out
Dairy	249.8092	OUT049	1999	Medium	Tier 1	Su
rt Drinks	48.2692	OUT018	2009	Medium	Tier 3	Su
Meat	141.6180	OUT049	1999	Medium	Tier 1	Su
uits and jetables	182.0950	OUT010	1998	NaN	Tier 3	Su

**Check no of row and column**

```
In [3]: print('No of Rows and Columns ----->', df.shape )
```

No of Rows and Columns ----> (8523, 12)

## Checking for Null values

```
In [4]: print('-----\n')
      print(df.isnull().sum())
      print('\n-----')
```

```
Item_Identifier          0  
Item_Weight              1463  
Item_Fat_Content         0  
Item_Visibility          0  
Item_Type                0  
Item_MRP                 0  
Outlet_Identifier        0  
Outlet_Establishment_Year 0  
Outlet_Size               2410  
Outlet_Location_Type      0  
Outlet_Type                0  
Item_Outlet_Sales          0  
dtype: int64
```

**There is Item weight and Outlet size has Nan value**

## Removing Nan Value

```
In [5]: df['Item_Weight'] = df['Item_Weight'].fillna(df['Item_Weight'].mean())
df['Outlet_Size'] = df['Outlet_Size'].fillna(df['Outlet_Size'].mode()[0])
```

We remove all nan on those features

## Checking for Null values remove or not

```
In [6]: print('-----\n')
print(df.isnull().sum())
print('\n-----')
```

```
-----
Item_Identifier      0
Item_Weight          0
Item_Fat_Content     0
Item_Visibility      0
Item_Type            0
Item_MRP             0
Outlet_Identifier    0
Outlet_Establishment_Year 0
Outlet_Size           0
Outlet_Location_Type 0
Outlet_Type           0
Item_Outlet_Sales    0
dtype: int64
-----
```

```
In [7]: print('No of Rows and Columns Left After Removing NaN ----->', df.shape )
```

No of Rows and Columns Left After Removing NaN -----> (8523, 12)

## Information about dataset

```
In [8]: print('-----\n')
print(df.info())
print('\n-----')
```

```
-----  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 8523 entries, 0 to 8522  
Data columns (total 12 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   Item_Identifier    8523 non-null   object    
 1   Item_Weight        8523 non-null   float64  
 2   Item_Fat_Content   8523 non-null   object    
 3   Item_Visibility    8523 non-null   float64  
 4   Item_Type          8523 non-null   object    
 5   Item_MRP           8523 non-null   float64  
 6   Outlet_Identifier  8523 non-null   object    
 7   Outlet_Establishment_Year 8523 non-null   int64    
 8   Outlet_Size        8523 non-null   object    
 9   Outlet_Location_Type 8523 non-null   object    
 10  Outlet_Type        8523 non-null   object    
 11  Item_Outlet_Sales  8523 non-null   float64  
dtypes: float64(4), int64(1), object(7)  
memory usage: 799.2+ KB  
None  
-----
```

**Some features are in float and some are in object**

## Droping Unwanted Column

```
In [9]: col = ['Item_Identifier', 'Outlet_Identifier']
```

```
In [10]: df = df.drop(columns = col, axis = 1)
df.head()
```

```
Out[10]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Establishment_Year
0	9.30	Low Fat	0.016047	Dairy	249.8092	1999
1	5.92	Regular	0.019278	Soft Drinks	48.2692	2009
2	17.50	Low Fat	0.016760	Meat	141.6180	1999
3	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	1998
4	8.93	Low Fat	0.000000	Household	53.8614	1987

```
In [11]: print('No of Rows and Columns ----->', df.shape )
```

No of Rows and Columns -----> (8523, 10)

## Analysis of data

```
In [12]: df['Item_Fat_Content'].value_counts()
```

```
Out[12]:
```

Low Fat	5089
Regular	2889
LF	316
reg	117
low fat	112

Name: Item\_Fat\_Content, dtype: int64

```
In [13]: df['Item_Type'].value_counts()
```

```
Out[13]:
```

Fruits and Vegetables	1232
Snack Foods	1200
Household	910
Frozen Foods	856
Dairy	682
Canned	649
Baking Goods	648
Health and Hygiene	520
Soft Drinks	445
Meat	425
Breads	251
Hard Drinks	214
Others	169
Starchy Foods	148
Breakfast	110
Seafood	64

Name: Item\_Type, dtype: int64

```
In [14]: df['Outlet_Size'].value_counts()
```

```
Out[14]: Medium      5203  
          Small       2388  
          High        932  
          Name: Outlet_Size, dtype: int64
```

```
In [15]: df['Outlet_Location_Type'].value_counts()
```

```
Out[15]: Tier 3     3350  
          Tier 2     2785  
          Tier 1     2388  
          Name: Outlet_Location_Type, dtype: int64
```

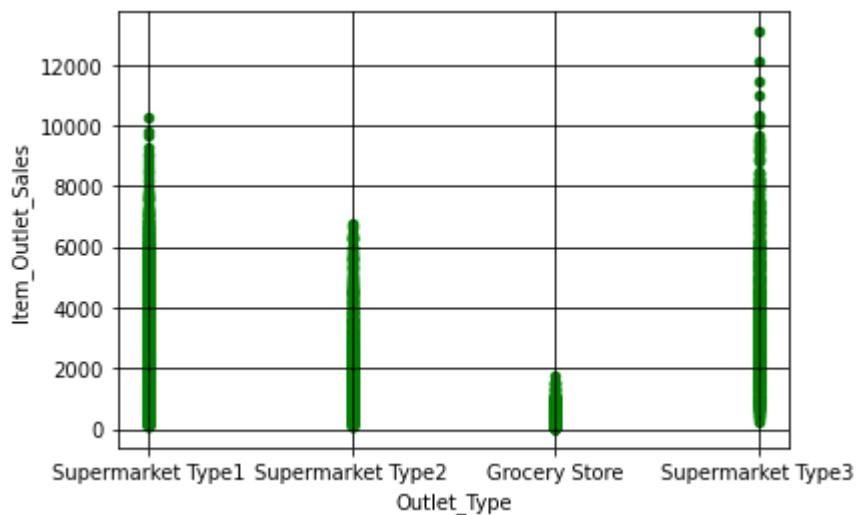
```
In [16]: df['Outlet_Type'].value_counts()
```

```
Out[16]: Supermarket Type1    5577  
          Grocery Store      1083  
          Supermarket Type3   935  
          Supermarket Type2   928  
          Name: Outlet_Type, dtype: int64
```

```
In [17]: df['Outlet_Establishment_Year'].value_counts()
```

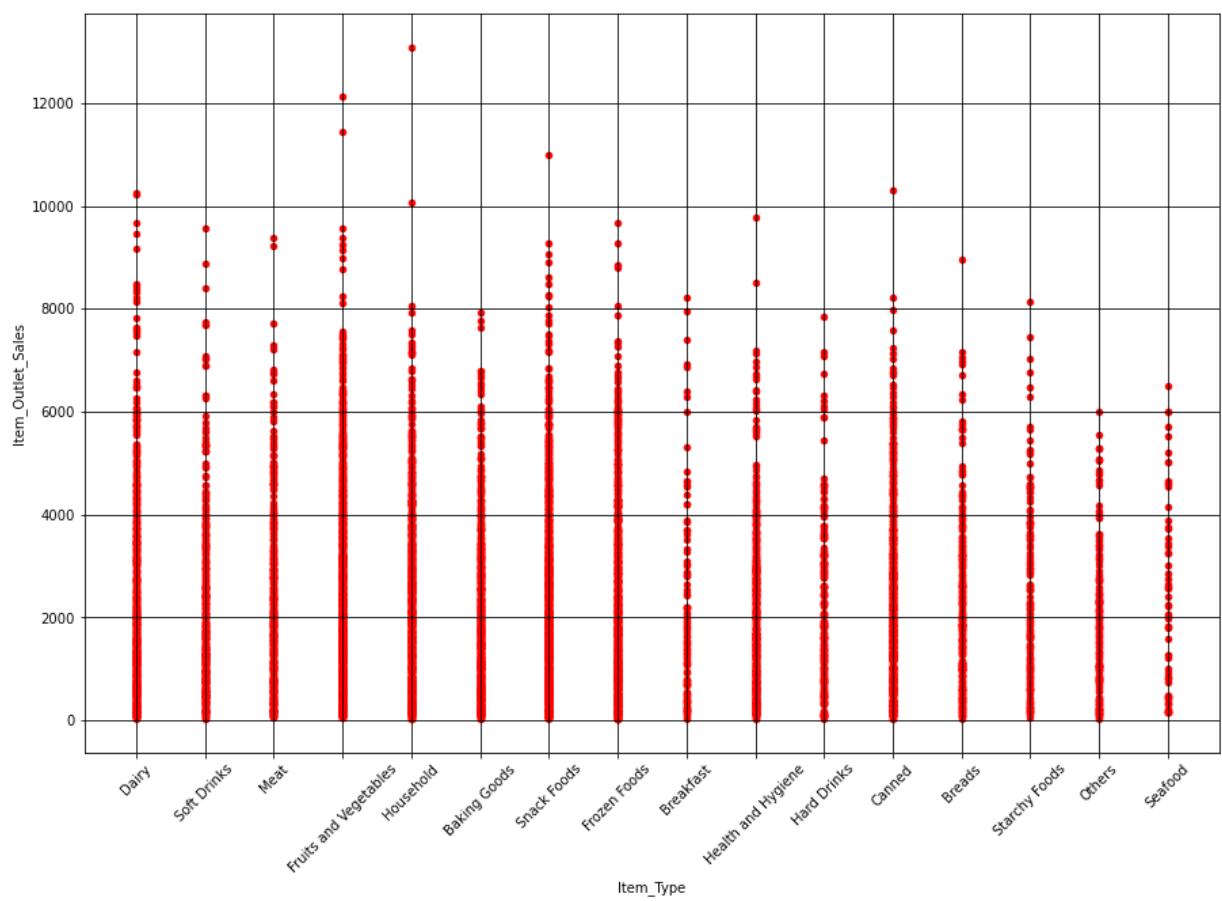
```
Out[17]: 1985      1463  
          1987      932  
          2004      930  
          1997      930  
          1999      930  
          2002      929  
          2009      928  
          2007      926  
          1998      555  
          Name: Outlet_Establishment_Year, dtype: int64
```

```
In [18]: df.plot.scatter('Outlet_Type','Item_Outlet_Sales', c = 'g')
plt.grid(which = 'major', c = 'black')
plt.show()
```



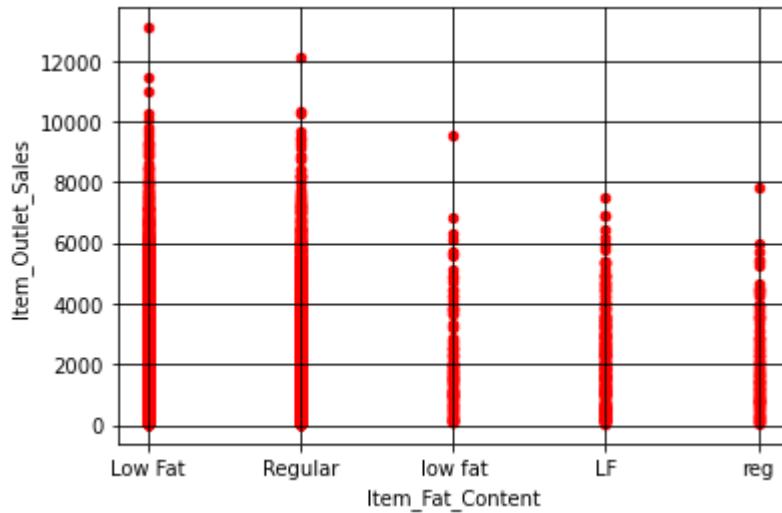
### Supermarket Type3 highest sales

```
In [19]: df.plot.scatter('Item_Type','Item_Outlet_Sales', figsize = (15,10), rot = 45, c =
plt.grid(which = 'major', c = 'black')
plt.show()
```



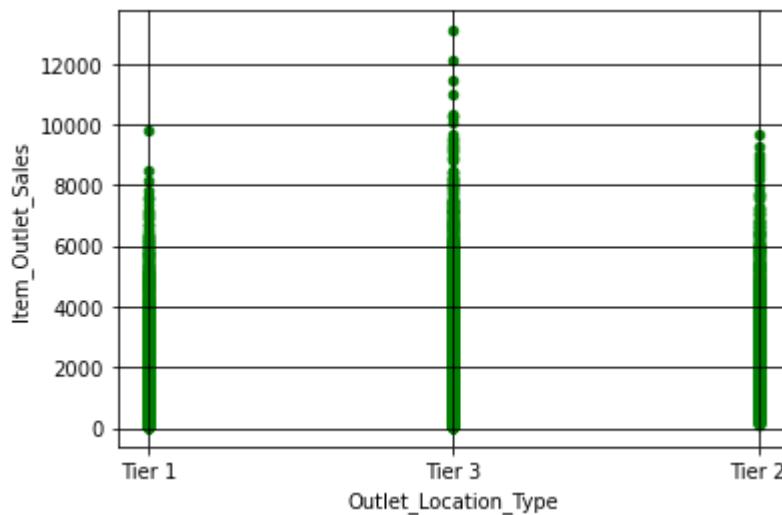
**Household products sales higest and other products least sales**

```
In [20]: df.plot.scatter('Item_Fat_Content','Item_Outlet_Sales', c = 'r')
plt.grid(which = 'major', c = 'black')
plt.show()
```



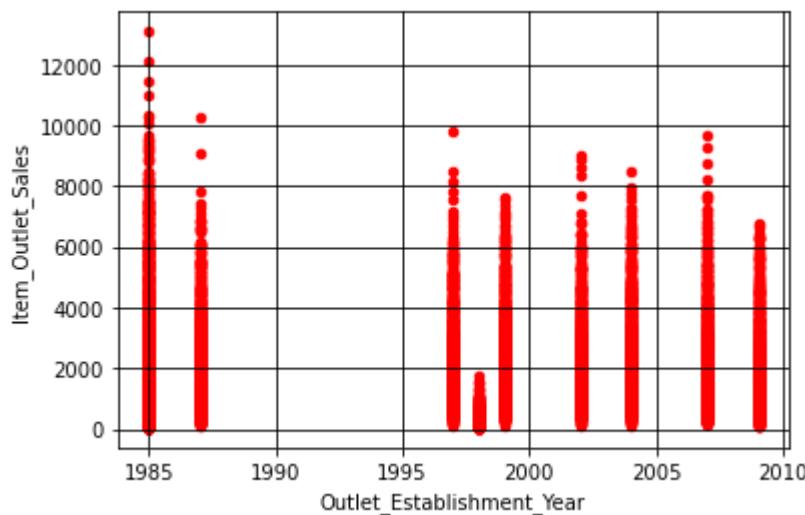
**Low fat content sales higest among all and LF is least**

```
In [21]: df.plot.scatter('Outlet_Location_Type','Item_Outlet_Sales', c = 'g')
plt.grid(which = 'major', c = 'black')
plt.show()
```



**Tier 3 cites sales highest and Tier 2 sales lowest.**

```
In [22]: df.plot.scatter('Outlet_Establishment_Year','Item_Outlet_Sales', c = 'r')
plt.grid(which = 'major', c = 'black')
plt.show()
```



**Older outlets sales higher compare to newly open outlets**

```
In [23]: df.head()
```

```
Out[23]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Establishment_Year
0	9.30	Low Fat	0.016047	Dairy	249.8092	1999
1	5.92	Regular	0.019278	Soft Drinks	48.2692	2009
2	17.50	Low Fat	0.016760	Meat	141.6180	1999
3	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	1998
4	8.93	Low Fat	0.000000	Household	53.8614	1987

```
In [24]: df['No_of_years_outlet_open'] = [2021] - df['Outlet_Establishment_Year']
df.head()
```

```
Out[24]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Establishment_Year
0	9.30	Low Fat	0.016047	Dairy	249.8092	1999
1	5.92	Regular	0.019278	Soft Drinks	48.2692	2009
2	17.50	Low Fat	0.016760	Meat	141.6180	1999
3	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	1998
4	8.93	Low Fat	0.000000	Household	53.8614	1987

Convert Outlet establishment year into No of years outlets open because it easier to encode

```
In [25]: print('-----\n')
print(df.info())
print('\n-----')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8523 entries, 0 to 8522
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Item_Weight      8523 non-null   float64
 1   Item_Fat_Content 8523 non-null   object  
 2   Item_Visibility  8523 non-null   float64
 3   Item_Type        8523 non-null   object  
 4   Item_MRP         8523 non-null   float64
 5   Outlet_Establishment_Year 8523 non-null   int64  
 6   Outlet_Size      8523 non-null   object  
 7   Outlet_Location_Type 8523 non-null   object  
 8   Outlet_Type      8523 non-null   object  
 9   Item_Outlet_Sales 8523 non-null   float64
 10  No_of_years_outlet_open 8523 non-null   int64  
dtypes: float64(4), int64(2), object(5)
memory usage: 732.6+ KB
None
```

```
In [26]: df = df.drop('Outlet_Establishment_Year', axis = 1)
df.head()
```

Out[26]:

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Size	Outlet_Locatio
0	9.30	Low Fat	0.016047	Dairy	249.8092	Medium	
1	5.92	Regular	0.019278	Soft Drinks	48.2692	Medium	
2	17.50	Low Fat	0.016760	Meat	141.6180	Medium	
3	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	Medium	
4	8.93	Low Fat	0.000000	Household	53.8614	High	

Drop Outlet establishment year

## Splitting data into features and label

```
In [27]: x = df.drop('Item_Outlet_Sales', axis = 1)
y = df['Item_Outlet_Sales']
print('Data has been splited')
```

Data has been splited

## Filter Categorical features

```
In [28]: numerics = ['int64', 'float64']
categorical_col = []
features = x.columns.values.tolist()
for col in features:
    if x[col].dtype in numerics:
        continue
    categorical_col.append(col)
```

## Encoding categorical columns using get dummies

```
In [29]: x_dummies = pd.get_dummies(x[categorical_col], drop_first = True)
x_dummies.head()
```

```
Out[29]:
```

	Item_Fat_Content_Low Fat	Item_Fat_Content_Regular	Item_Fat_Content_low fat	Item_Fat_Content_reg	Item_Fat_Content_Very Low Fat
0	1	0	0	0	0
1	0	1	0	0	0
2	1	0	0	0	0
3	0	1	0	0	0
4	1	0	0	0	0

5 rows × 26 columns

```
In [30]: x_dummies.shape
```

```
Out[30]: (8523, 26)
```

```
In [31]: x = x.join(x_dummies)
```

```
In [32]: x.head()
```

```
Out[32]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Size	Outlet_Location
0	9.30	Low Fat	0.016047	Dairy	249.8092	Medium	North
1	5.92	Regular	0.019278	Soft Drinks	48.2692	Medium	North
2	17.50	Low Fat	0.016760	Meat	141.6180	Medium	North
3	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	Medium	South
4	8.93	Low Fat	0.000000	Household	53.8614	High	South

5 rows × 35 columns

```
In [33]: x.drop(columns = categorical_col, axis = 1, inplace = True)
```

In [34]: `x.head()`

Out[34]:

	Item_Weight	Item_Visibility	Item_MRP	No_of_years_outlet_open	Item_Fat_Content_Low Fat	Item_Fat_Content_High Fat
0	9.30	0.016047	249.8092	22	1	0
1	5.92	0.019278	48.2692	12	0	1
2	17.50	0.016760	141.6180	22	1	0
3	19.20	0.000000	182.0950	23	0	1
4	8.93	0.000000	53.8614	34	1	0

5 rows × 30 columns

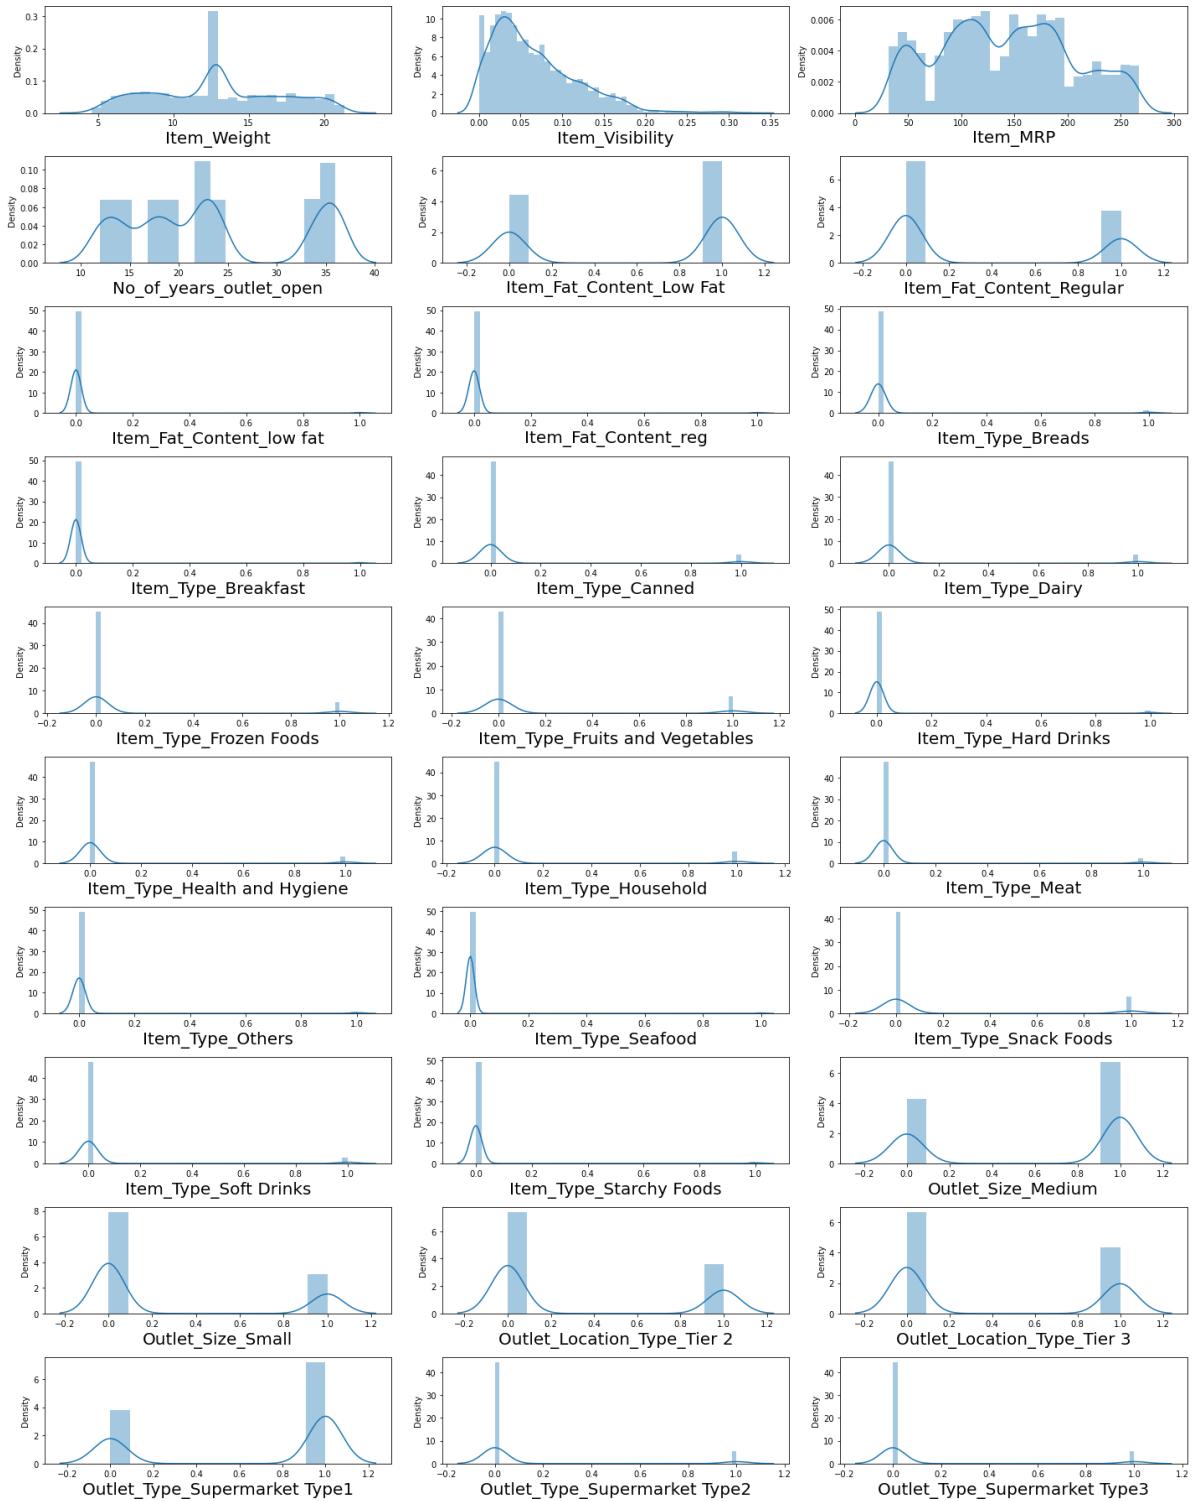
## Checking Outliers

```
In [35]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in x:
    if plotnumber <=30:
        ax = plt.subplot(10,3, plotnumber)
        sns.distplot(x[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----
Distribution Plot :-
```



**Some outliers present in columns**

**Power Transformer to remove outliers**

```
In [36]: scaler = PowerTransformer(method = 'yeo-johnson')
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
Out[36]: array([[-0.82263635, -1.14688621,  1.62367023, ...,  0.72680189,
   -0.34955064, -0.35102831],
 [-1.73380133, -1.03709991, -1.59501643, ..., -1.37589075,
  2.86081583, -0.35102831],
 [ 1.08636176, -1.12237453,  0.08123678, ...,  0.72680189,
  -0.34955064, -0.35102831],
 ...,
 [-0.49655592, -0.54297932, -0.86792986, ...,  0.72680189,
  -0.34955064, -0.35102831],
 [-1.37344321,  1.45998957, -0.54899589, ..., -1.37589075,
  2.86081583, -0.35102831],
 [ 0.49178563, -0.27619496, -1.04722682, ...,  0.72680189,
  -0.34955064, -0.35102831]])
```

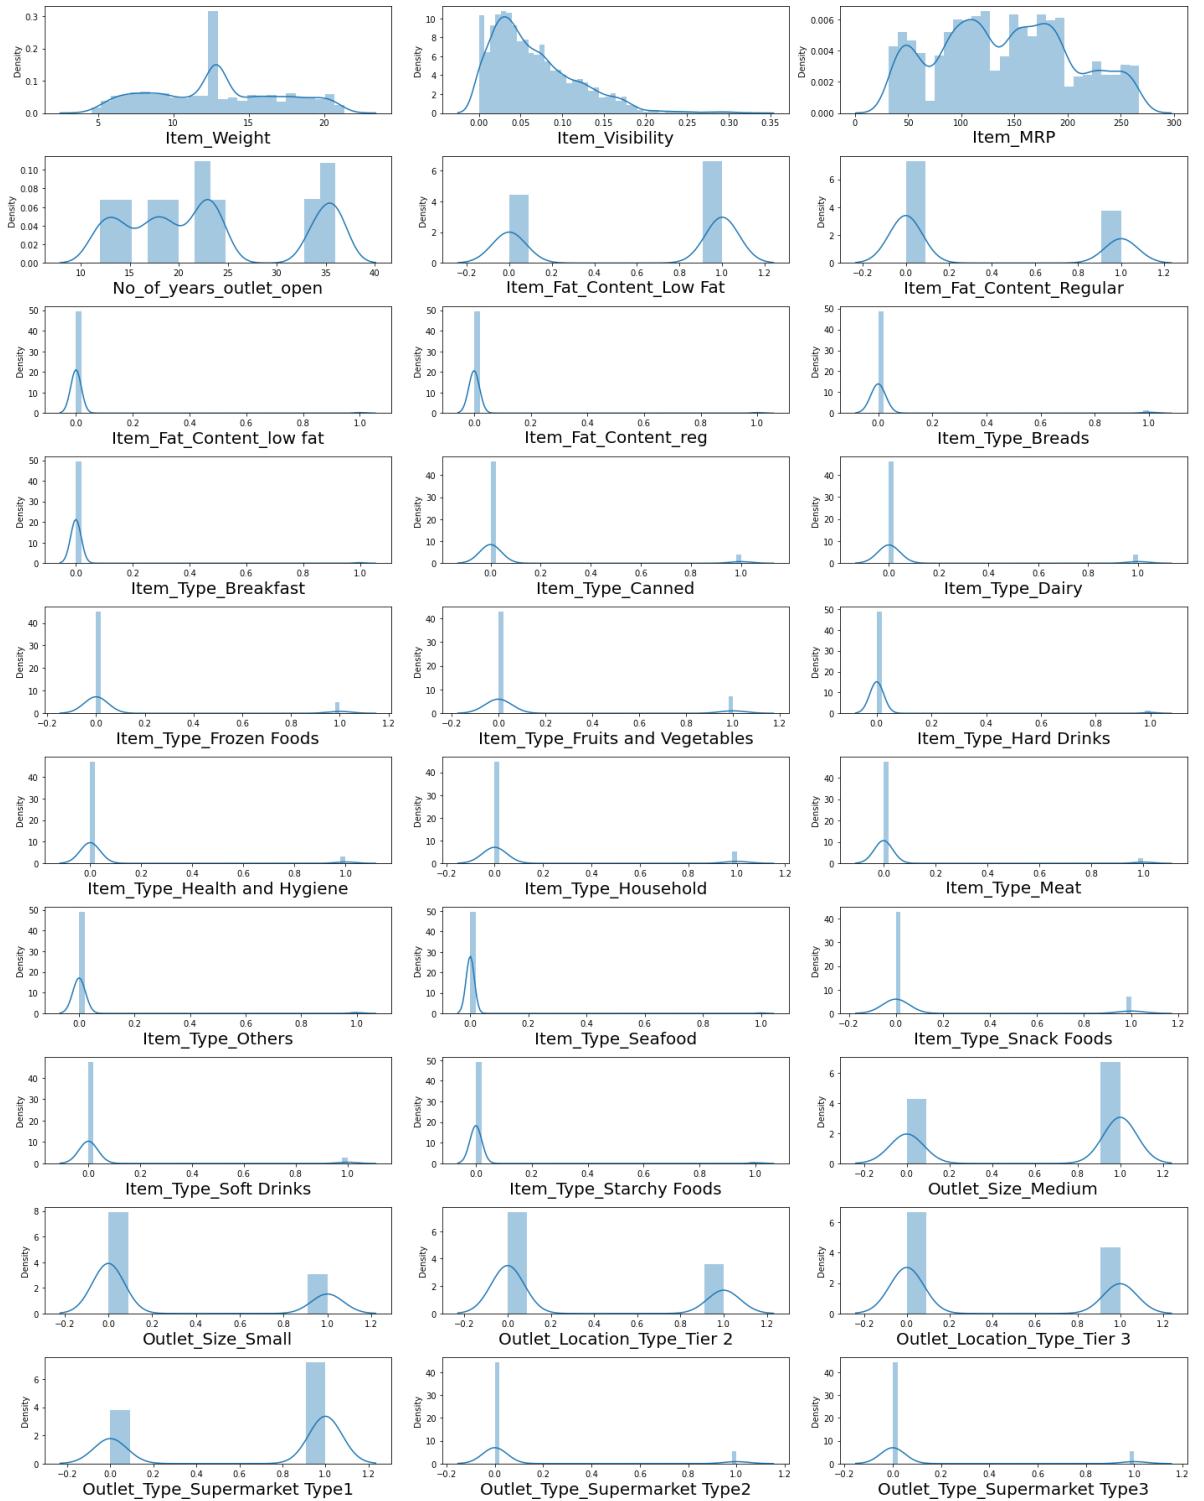
## Checking Outlier remove or not

```
In [37]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in x:
    if plotnumber <=30:
        ax = plt.subplot(10,3, plotnumber)
        sns.distplot(x[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

-----  
Distribution Plot :-  
-----



**Outliers are removed**

**Split data into train and test. Model will be bulit on training data and tested on test data**

```
In [38]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25, random_state = 42)
print('Data has been splitted.')
```

Data has been splitted.

## Model Building

### Xgboost model instantiaing, training and evaluating

```
In [39]: xgb = xgb.XGBRegressor()
xgb.fit(x_train, y_train)
y_pred = xgb.predict(x_test)
```

```
In [40]: print('=====')
print('R2 Score ---->', r2_score(y_test, y_pred))
print('=====')
print('RMSE of Model ---->', np.sqrt(mean_squared_error(y_test, y_pred)))
print('=====')
print('MSE of Model ---->', mean_squared_error(y_test, y_pred))
print('=====')
print('Score of test data ---->', xgb.score(x_test, y_test))
print('=====')
```

```
=====
R2 Score ----> 0.5294371130690214
=====
RMSE of Model ----> 1163.248310421678
=====
MSE of Model ----> 1353146.631698889
=====
Score of test data ----> 0.5294371130690214
=====
```

**Conclusion : XGBoost model has 52% score**

### Knn model instantiaing, training and evaluating

```
In [41]: Knn = KNeighborsRegressor()
Knn.fit(x_train, y_train)
y_pred = Knn.predict(x_test)
```

```
In [42]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', Knn.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.42296758232174914  
=====  
RMSE of Model -----> 1288.1418819087874  
=====  
MSE of Model -----> 1659309.5079275125  
=====  
Score of test data ----> 0.42296758232174914  
=====
```

**Conclusion : Knn model has 42% score**

## Decision Tree model instantiaing, training and evaluating

```
In [43]: DT = DecisionTreeRegressor()  
DT.fit(x_train, y_train)  
y_pred = DT.predict(x_test)
```

```
In [44]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', DT.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.17620759508520245  
=====  
RMSE of Model -----> 1539.1200169339468  
=====  
MSE of Model -----> 2368890.4265267523  
=====  
Score of test data ----> 0.17620759508520245  
=====
```

**Conclusion : Decision Tree model has 19% score**

## Random Forest model instantiaing, training and evaluating

```
In [45]: Rn = RandomForestRegressor()  
Rn.fit(x_train, y_train)  
y_pred = Rn.predict(x_test)
```

```
In [46]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ---->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ---->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', Rn.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.5658710564717431  
=====  
RMSE of Model ----> 1117.3081464916793  
=====  
MSE of Model ----> 1248377.4942166717  
=====  
Score of test data ----> 0.5658710564717431  
=====
```

**Conclusion : Random Forest model has 56% score**

## SVM model instantiaing, training and evaluating

```
In [47]: svr = SVR()  
svr.fit(x_train, y_train)  
y_pred = svr.predict(x_test)
```

```
In [48]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', svr.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.17505733672176238  
=====  
RMSE of Model -----> 1540.1941759675528  
=====  
MSE of Model -----> 2372198.0996843693  
=====  
Score of test data ----> 0.17505733672176238  
=====
```

**Conclusion : SVM model has 17% score**

## Linear Regression model instantiating, training and evaluating

```
In [49]: Lr = LinearRegression()  
Lr.fit(x_train, y_train)  
y_pred = Lr.predict(x_test)
```

```
In [50]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', Lr.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.5761275930961476  
=====  
RMSE of Model -----> 1104.0307445061003  
=====  
MSE of Model -----> 1218883.884814694  
=====  
Score of test data ----> 0.5761275930961476  
=====
```

**Conclusion : Linear Regression model has 57% score**

**Looking R2 score we found Linear Regression has best model so we do Hyperparameter Tuning on it.**

```
In [51]: param_grid = {'fit_intercept' : [True, False], 'normalize' : [True, False], 'copy_X': [True, False]}
```

```
In [52]: grid_search = GridSearchCV(estimator = Lr, param_grid = param_grid, cv = 5, n_jobs=-1)
```

```
In [53]: grid_search.fit(x_train, y_train)
```

```
Out[53]: GridSearchCV(cv=5, estimator=LinearRegression(), n_jobs=-1,
param_grid={'copy_X': [True, False],
'fit_intercept': [True, False],
'normalize': [True, False],
'positive': [True, False]})
```

```
In [54]: best_parameters = grid_search.best_params_
print(best_parameters)
```

```
{'copy_X': True, 'fit_intercept': False, 'normalize': True, 'positive': False}
```

```
In [55]: hlr = LinearRegression(copy_X = True, fit_intercept = False, normalize = True, pco
hlr.fit(x_train, y_train)
hlr.score(x_test, y_test)
```

```
Out[55]: 0.5761141582584479
```

```
In [56]: y_pred = hlr.predict(x_test)
```

```
In [57]: print('=====')
print('R2 Score ----->', r2_score(y_test, y_pred))
print('=====')
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))
print('=====')
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))
print('=====')
print('Score of test data ----->', Lr.score(x_test, y_test))
print('=====')
```

```
=====
R2 Score -----> 0.5761141582584479
=====
RMSE of Model -----> 1104.0482407574457
=====
MSE of Model -----> 1218922.5179196107
=====
Score of test data -----> 0.5761275930961476
=====
```

**After Hyperparameter Tuning model accuracy score 57%.**

## Saving The Model

```
In [58]: # saving the model to the Local file system  
filename = 'Big Data Mart Sales Train.pickle'  
pickle.dump(hlr, open(filename, 'wb'))
```

**Final Conclusion : LinearRegression is our best model.**

```
In [ ]:
```