

## Problem Statement:

This data was extracted from the 1994 Census bureau database by Ronny Kohavi and Barry Becker (Data Mining and Visualization, Silicon Graphics). A set of reasonably clean records was extracted using the following conditions: ((AAGE>16) && (AGI>100) && (AFNLWGT>1) && (HRSWK>0)). The prediction task is to determine whether a person makes over \$50K a year.

Description of fnlwgt (final weight) The weights on the Current Population Survey (CPS) files are controlled to independent estimates of the civilian non-institutional population of the US. These are prepared monthly for us by Population Division here at the Census Bureau. We use 3 sets of controls. These are:

A single cell estimate of the population 16+ for each state.

Controls for Hispanic Origin by age and sex.

Controls by Race, age and sex.

## Importing Required Library

```
In [167]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pickle
import scikitplot as skplt
pd.set_option('display.max_columns', None) # # For display maximum column
from sklearn.preprocessing import StandardScaler, LabelEncoder, PowerTransformer
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, plot_roc_curve
import xgboost as xgb
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

## Reading Data

```
In [2]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\csv file\census_income.csv")
df.head()
```

Out[2]:

|   | Age | Workclass        | Fnlwgt | Education | Education_num | Marital_status | Occupation         | Relationship      | Taxable_inc     |
|---|-----|------------------|--------|-----------|---------------|----------------|--------------------|-------------------|-----------------|
| 0 | 50  | Self-emp-not-inc | 83311  | Bachelors |               | 13             | Married-civ-spouse | Exec-managerial   | Husband \       |
| 1 | 38  | Private          | 215646 | HS-grad   |               | 9              | Divorced           | Handlers-cleaners | Not-in-family \ |
| 2 | 53  | Private          | 234721 | 11th      |               | 7              | Married-civ-spouse | Handlers-cleaners | Husband \ I     |
| 3 | 28  | Private          | 338409 | Bachelors |               | 13             | Married-civ-spouse | Prof-specialty    | Wife \ I        |
| 4 | 37  | Private          | 284582 | Masters   |               | 14             | Married-civ-spouse | Exec-managerial   | Wife \          |

## Check no of row and column

```
In [3]: print('No of Rows and Columns ----->', df.shape )
```

No of Rows and Columns -----> (32560, 15)

## Checking for Null values

```
In [4]: print('=====\\n')
print(df.isnull().sum())
print('\\n=====')
```

```
=====
```

```
Age          0
Workclass    0
Fnlwgt       0
Education    0
Education_num 0
Marital_status 0
Occupation   0
Relationship  0
Race         0
Sex          0
Capital_gain 0
Capital_loss  0
Hours_per_week 0
Native_country 0
Income        0
dtype: int64
```

```
=====
```

**There is no null value**

## Information about dataset

```
In [5]: print('=====\\n')
print(df.info())
print('=====')
```

```
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32560 entries, 0 to 32559
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              32560 non-null   int64  
 1   Workclass        32560 non-null   object  
 2   Fnlwgt           32560 non-null   int64  
 3   Education        32560 non-null   object  
 4   Education_num    32560 non-null   int64  
 5   Marital_status   32560 non-null   object  
 6   Occupation       32560 non-null   object  
 7   Relationship     32560 non-null   object  
 8   Race             32560 non-null   object  
 9   Sex              32560 non-null   object  
 10  Capital_gain    32560 non-null   int64  
 11  Capital_loss    32560 non-null   int64  
 12  Hours_per_week  32560 non-null   int64  
 13  Native_country  32560 non-null   object  
 14  Income           32560 non-null   object  
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
None
=====
```

**Categorical data present in our data set**

## **Analysis of Data Respect To Income**

### **Age column**

```
In [6]: df['Age'].value_counts()
```

```
Out[6]: 36    898
31    888
34    886
23    877
35    876
...
83     6
85     3
88     3
86     1
87     1
Name: Age, Length: 73, dtype: int64
```


```
In [7]: df['Income'].value_counts()
```

```
Out[7]: <=50K    24719
>50K      7841
Name: Income, dtype: int64
```

```
In [8]: b = df.groupby('Age')['Income'].value_counts().sort_values()
b
```

```
Out[8]:   Age  Income
87    <=50K      1
86    <=50K      1
84    >50K       1
19    >50K       2
83    >50K       2
...
20    <=50K    753
27    <=50K    754
24    <=50K    767
25    <=50K    788
23    <=50K    865
Name: Income, Length: 138, dtype: int64
```

```
In [9]: b.plot.barh(figsize = (10,50), color = ['red', 'blue', 'cyan', 'g', 'y'])
plt.xlabel('Frequency', c = 'g', fontsize = 12)
plt.ylabel('Age & Income (<=50K & >50K)', c = 'g', fontsize = 12 )
plt.title('Age vs Income (<=50K & >50K)', c = 'g', fontsize = 12)
plt.show()
```




Above plot shows whoes age is 23 years they have higest income frequency more than 800 people earn (<=50K)

```
In [10]: a = df[ 'Age' ][df[ 'Age' ]<=50].value_counts().sort_values()  
a
```

```
Out[10]: 17    395  
48    543  
18    550  
49    577  
50    602  
47    708  
19    712  
21    720  
44    724  
45    734  
46    737  
20    753  
22    765  
43    770  
42    780  
26    785  
40    794  
24    798  
41    808  
29    813  
39    815  
38    827  
32    828  
27    835  
25    841  
37    858  
30    861  
28    867  
33    875  
35    876  
23    877  
34    886  
31    888  
36    898  
Name: Age, dtype: int64
```

```
In [11]: a.plot.bar(figsize = (10,5), rot = 360, color = ['red','cyan', 'g', 'magenta'])
plt.xlabel('Age', c = 'g', fontsize = 12)
plt.ylabel('Frequency Of Years(<=50)', c = 'g', fontsize = 12 )
plt.title('Age (<=50) Years', c = 'g', fontsize = 12)
plt.show()
```



Abobe plot No of people whoes Age (<=50) years

## Workclass column

```
In [12]: df['Workclass'].value_counts()
```

```
Out[12]: Private           22696  
Self-emp-not-inc      2541  
Local-gov            2093  
?                   1836  
State-gov            1297  
Self-emp-inc          1116  
Federal-gov          960  
Without-pay           14  
Never-worked          7  
Name: Workclass, dtype: int64
```

**Observation : In Workclass column there is '?' present. We solve this problem to replace '?' with nan value**

```
In [13]: # Approach to remove '?' in Workclass column  
df['Workclass'] = df['Workclass'].str.strip().map(lambda x: np.nan if x=="?" else
```

```
In [14]: df.isna().sum()
```

```
Out[14]: Age            0  
Workclass        1836  
Fnlwgt           0  
Education         0  
Education_num     0  
Marital_status    0  
Occupation        0  
Relationship       0  
Race              0  
Sex               0  
Capital_gain      0  
Capital_loss       0  
Hours_per_week    0  
Native_country     0  
Income             0  
dtype: int64
```

**Conclusion : In Workclass column '?' remove and replace with nan value. We fill nan with mode**


```
In [15]: df['Workclass'].fillna(df['Workclass'].mode()[0], inplace = True)
```

```
In [16]: df.isna().sum()
```

```
Out[16]: Age          0  
Workclass      0  
Fnlwgt         0  
Education       0  
Education_num   0  
Marital_status  0  
Occupation      0  
Relationship     0  
Race            0  
Sex              0  
Capital_gain    0  
Capital_loss    0  
Hours_per_week  0  
Native_country  0  
Income           0  
dtype: int64
```

There is no null value in Workclass column we ready to do our analysis

```
In [17]: df['Workclass'].value_counts().plot.bar(figsize = (10,5), rot = 360, color = ['red','cyan','darkgreen','magenta','red','cyan','red','cyan','blue'])  
plt.xlabel('Workclass', c = 'r', fontsize = 12)  
plt.ylabel('Frequency of Jobs', c = 'r', fontsize = 12 )  
plt.title('Workclass vs Frequency of Jobs', c = 'r', fontsize = 12)  
plt.xticks(c = 'b')  
plt.yticks(c = 'b')  
plt.show()
```




Above plot shows Private Job has higest frequency of job working

```
In [18]: w = df.groupby('Workclass')['Income'].value_counts()
w
```

```
Out[18]: Workclass           Income
Federal-gov      <=50K       589
                  >50K        371
Local-gov        <=50K      1476
                  >50K        617
Never-worked     <=50K        7
Private          <=50K    19378
                  >50K      5154
Self-emp-inc     >50K       622
                  <=50K      494
Self-emp-not-inc <=50K      1817
                  >50K       724
State-gov         <=50K      944
                  >50K       353
Without-pay       <=50K       14
Name: Income, dtype: int64
```

```
In [19]: w.plot.barh(figsize = (10,5), color = ['red','cyan', 'g', 'magenta'])
plt.xlabel('Workclass Frequency', c = 'b', fontsize = 12)
plt.ylabel('Income of Jobs', c = 'b', fontsize = 12 )
plt.title('Workclass vs Income Of Jobs', c = 'b', fontsize = 12)
plt.xticks(c = 'g')
plt.yticks(c = 'g')
plt.show()
```




Above plot shows private job which salary <=50K is higest frequency and without-pay has lowest

## Education column

```
In [20]: df['Education'].value_counts()
```

```
Out[20]:   HS-grad      10501
             Some-college    7291
             Bachelors       5354
             Masters         1723
             Assoc-voc        1382
             11th              1175
             Assoc-acdm        1067
             10th              933
             7th-8th           646
             Prof-school        576
             9th                514
             12th              433
             Doctorate          413
             5th-6th            333
             1st-4th            168
             Preschool           51
Name: Education, dtype: int64
```

```
In [21]: df['Education'].value_counts().plot.bar(figsize = (10,5), rot = 45, color = ['red', 'blue', 'green', 'orange', 'purple'])
plt.xlabel('Education', c = 'purple', fontsize = 12)
plt.ylabel('Frequency of Education', c = 'purple', fontsize = 12 )
plt.title('Education vs Frequency of Education', c = 'purple', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```




**Above plot shows HS-grad has more than 10000 counts**

```
In [22]: e = df.groupby('Education')['Income'].value_counts().sort_values()
e
```

```
Out[22]: Education    Income
1st-4th        >50K      6
5th-6th        >50K     16
9th            >50K     27
12th           >50K     33
7th-8th        >50K     40
Preschool      <=50K    51
11th           >50K     60
10th           >50K     62
Doctorate      <=50K   107
Prof-school    <=50K   153
1st-4th        <=50K   162
Assoc-acdm    >50K    265
Doctorate      >50K    306
5th-6th        <=50K   317
Assoc-voc      >50K    361
12th           <=50K   400
Prof-school    >50K    423
9th            <=50K   487
7th-8th        <=50K   606
Masters         <=50K   764
Assoc-acdm    <=50K   802
10th           <=50K   871
Masters         >50K    959
Assoc-voc      <=50K  1021
11th           <=50K  1115
Some-college   >50K  1387
HS-grad         >50K  1675
Bachelors      >50K  2221
                  <=50K  3133
Some-college   <=50K  5904
HS-grad         <=50K  8826
Name: Income, dtype: int64
```


```
In [23]: e.plot.bar(figsize = (15,5),rot= 45,color = ['red','cyan', 'g', 'magenta'])
plt.xlabel('Education with Income', c = 'magenta', fontsize = 12)
plt.ylabel('Frequency', c = 'magenta', fontsize = 12 )
plt.title('Education vs Income', c = 'magenta', fontsize = 12)
plt.xticks(c = 'red')
plt.yticks(c = 'red')
plt.grid()
plt.show()
```



Above plot shows HS-grad who earn (<=50K) has higest frequency

## Marital Status column

```
In [24]: df['Marital_status'].value_counts().plot.bar(figsize = (10,5),rot= 45,color = ['r','g','c','m','b','y','k'],edgecolor = 'black')  
plt.xticks(c = 'g')  
plt.yticks(c = 'g')  
plt.xlabel('Marital Status', c = 'b', fontsize = 12)  
plt.ylabel('Counts', c = 'b', fontsize = 12 )  
plt.title('Marital Status vs Counts', c = 'b', fontsize = 12)  
plt.show()
```




Above plot shows Married-civ-spouse is highest frequency

```
In [25]: m = df.groupby('Marital_status')['Income'].value_counts()
m
```

```
Out[25]: Marital_status           Income
Divorced            <=50K      3980
                  >50K       463
Married-AF-spouse  <=50K       13
                  >50K        10
Married-civ-spouse <=50K     8284
                  >50K    6692
Married-spouse-absent <=50K     384
                  >50K       34
Never-married      <=50K   10191
                  >50K      491
Separated           <=50K      959
                  >50K       66
Widowed             <=50K      908
                  >50K       85
Name: Income, dtype: int64
```

```
In [26]: m.plot.barh(figsize = (10,5),color = ['red','cyan', 'g', 'magenta'])
plt.yticks(c = 'darkgreen', alpha = 1)
plt.xticks(c = '#008000', alpha = 1)
plt.ylabel('Marital Status', c = 'b', fontsize = 12)
plt.xlabel('Counts', c = 'b', fontsize = 12 )
plt.title('Marital Status vs Counts', c = 'b', fontsize = 12)
plt.show()
```



Above plot shows Never married earn (<=50K) is highest counts

## Occupation column

```
In [27]: df['Occupation'].value_counts()
```

```
Out[27]: Prof-specialty      4140  
Craft-repair        4099  
Exec-managerial     4066  
Adm-clerical        3769  
Sales                3650  
Other-service        3295  
Machine-op-inspct    2002  
?                     1843  
Transport-moving     1597  
Handlers-cleaners    1370  
Farming-fishing      994  
Tech-support          928  
Protective-serv       649  
Priv-house-serv       149  
Armed-Forces            9  
Name: Occupation, dtype: int64
```

**Observation : In Occupation column there is '?' present. We solve this problem to replace '?' with nan value**

```
In [28]: # Approach to remove '?' in Occupation column  
df['Occupation'] = df['Occupation'].str.strip().map(lambda x: np.nan if x=="?" else x)
```

```
In [29]: df.isna().sum()
```

```
Out[29]: Age           0  
Workclass      0  
Fnlwgt         0  
Education       0  
Education_num   0  
Marital_status  0  
Occupation     1843  
Relationship    0  
Race           0  
Sex             0  
Capital_gain    0  
Capital_loss    0  
Hours_per_week  0  
Native_country  0  
Income          0  
dtype: int64
```

**Conclusion : In Occupation column '?' remove and replace with nan value. We fill nan with mode**


```
In [30]: df['Occupation'].fillna(df['Occupation'].mode()[0], inplace = True)
```

```
In [31]: df.isna().sum()
```

```
Out[31]: Age          0  
Workclass      0  
Fnlwgt         0  
Education       0  
Education_num   0  
Marital_status  0  
Occupation      0  
Relationship    0  
Race            0  
Sex             0  
Capital_gain    0  
Capital_loss    0  
Hours_per_week  0  
Native_country  0  
Income           0  
dtype: int64
```

There is no null value in Occupation column we ready to do our analysis


```
In [32]: df['Occupation'].value_counts().plot.barh(figsize = (10,5),color = ['red','cyan'],  
plt.ylabel('Occupation', c = 'm', fontsize = 12)  
plt.xlabel('Frequency of Occupation', c = 'm', fontsize = 12 )  
plt.title('Occupation vs Frequency of Occupation', c = 'm', fontsize = 12)  
plt.xticks(c = 'r')  
plt.yticks(c = 'r')  
plt.show()
```



```
In [33]: o = df.groupby('Occupation')['Income'].value_counts().sort_values()  
o
```

```
Out[33]: Occupation      Income  
Armed-Forces        >50K         1  
Priv-house-serv     >50K         1  
Armed-Forces        <=50K        8  
Handlers-cleaners   >50K        86  
Farming-fishing     >50K       115  
Other-service        >50K       137  
Priv-house-serv     <=50K       148  
Protective-serv     >50K       211  
Machine-op-inspct   >50K       250  
Tech-support         >50K       283  
Transport-moving    >50K       320  
Protective-serv     <=50K       438  
Adm-clerical        >50K       507  
Tech-support         <=50K       645  
Farming-fishing     <=50K       879  
Craft-repair         >50K       929  
Sales                >50K       983  
Transport-moving    <=50K      1277  
Handlers-cleaners   <=50K      1284  
Machine-op-inspct   <=50K      1752  
Exec-managerial     >50K      1968  
Prof-specialty      >50K      2050  
Exec-managerial     <=50K      2098  
Sales                <=50K      2667  
Other-service        <=50K      3158  
Craft-repair         <=50K      3170  
Adm-clerical        <=50K      3262  
Prof-specialty      <=50K      3933  
Name: Income, dtype: int64
```

```
In [34]: o.plot.barh(figsize = (12,10),color = ['red','cyan', 'g', 'm'])
plt.ylabel('Occupation Income', c = 'b', fontsize = 12)
plt.xlabel('Frequency of Income', c = 'b', fontsize = 12 )
plt.title('Occupation Income vs Occupation Frequency', c = 'b', fontsize = 12)
plt.xticks(c = 'r')
plt.yticks(c = 'r')
plt.show()
```



**Above plot shows nearly 4000 Prof-specialty people having occupation earn (<=50K) is highest among all**

## Relationship column


```
In [35]: df['Relationship'].value_counts()
```

```
Out[35]: Husband           13193
          Not-in-family    8304
          Own-child         5068
          Unmarried          3446
          Wife                1568
          Other-relative     981
Name: Relationship, dtype: int64
```

```
In [36]: f = df.groupby('Relationship')['Income'].value_counts()
f
```

```
Out[36]: Relationship      Income
          Husband           <=50K    7275
                      >50K    5918
          Not-in-family     <=50K    7448
                      >50K    856
          Other-relative    <=50K    944
                      >50K     37
          Own-child          <=50K    5001
                      >50K     67
          Unmarried          <=50K    3228
                      >50K    218
          Wife                <=50K    823
                      >50K    745
Name: Income, dtype: int64
```

```
In [37]: f=plt.barh(figsize = (12,10),color = ['red','cyan', 'g', 'm'])
plt.ylabel('Relationship Income', c = 'r', fontsize = 12)
plt.xlabel('Frequency of Income', c = 'r', fontsize = 12 )
plt.title('Relationship Income vs Frequency of Income', c = 'r', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```



Above plot shows who has Not in family earn (<=50K) is higest frequency

## Race column


```
In [38]: df['Race'].value_counts()
```

```
Out[38]: White           27815  
Black            3124  
Asian-Pac-Islander  1039  
Amer-Indian-Eskimo  311  
Other             271  
Name: Race, dtype: int64
```

```
In [39]: r = df.groupby('Race')['Income'].value_counts().sort_values()  
r
```

```
Out[39]: Race           Income  
Other          >50K      25  
Amer-Indian-Eskimo >50K     36  
Other          <=50K    246  
Amer-Indian-Eskimo <=50K   275  
Asian-Pac-Islander >50K     276  
Black           >50K    387  
Asian-Pac-Islander <=50K   763  
Black           <=50K   2737  
White           >50K    7117  
                <=50K   20698  
Name: Income, dtype: int64
```

```
In [40]: r.plot.barh(figsize = (10,5),color = ['red','cyan', 'g', 'm'])  
plt.ylabel('Race Income', c = 'r', fontsize = 12)  
plt.xlabel('Frequency of Income', c = 'r', fontsize = 12 )  
plt.title('Race Income vs Frequency of Income', c = 'r', fontsize = 12)  
plt.xticks(c = 'g')  
plt.yticks(c = 'g')  
plt.show()
```



Above plot shows white people earn (<=50K) is highest frequency

## Sex column


```
In [41]: df['Sex'].value_counts()
```

```
Out[41]: Male      21789  
Female     10771  
Name: Sex, dtype: int64
```

```
In [42]: s = df.groupby('Sex')['Income'].value_counts()  
s
```

```
Out[42]: Sex      Income  
Female    <=50K      9592  
          >50K       1179  
Male      <=50K     15127  
          >50K      6662  
Name: Income, dtype: int64
```


```
In [43]: s.plot.bar(figsize = (10,5),rot= 360,color = ['red','cyan', 'g', 'm'])  
plt.xlabel('Sex Income', c = 'g', fontsize = 12)  
plt.ylabel('Frequency of Income', c = 'g', fontsize = 12 )  
plt.title('Sex Income vs Frequency of Income', c = 'g', fontsize = 12)  
plt.xticks(c = 'r')  
plt.yticks(c = 'r')  
plt.show()
```



Above plot shows male earn (<=50K) is highest frequency and female earn (>50K) is lowest frequency

## Hours per week column

```
In [44]: df['Hours_per_week'].value_counts().plot.bar(figsize = (20,10),rot= 360,color = [plt.xlabel('Hours per weeke', c = 'g', fontsize = 16)
plt.ylabel('Counts', c = 'g', fontsize = 16 )
plt.title('Hours per week vs Counts', c = 'g', fontsize = 16)
plt.show()
```





**Above plot shows who work 40 Hours\_per\_week has highest frequency**

```
In [45]: h = df.groupby('Hours_per_week')['Income'].value_counts().sort_values()
h
```

```
Out[45]: Hours_per_week  Income
92          <=50K      1
74          <=50K      1
95          >50K       1
73          >50K       1
17          >50K       1
...
20          <=50K     1146
50          >50K      1276
          <=50K     1543
40          >50K      3247
          <=50K    11969
Name: Income, Length: 173, dtype: int64
```

```
In [46]: h.plot.barh(figsize = (10,30),rot= 360,color = ['red','cyan', 'g', 'm'])
plt.ylabel('Hours_per_week Income', c = 'g', fontsize = 12)
plt.xlabel('Frequency of Income', c = 'g', fontsize = 12 )
plt.title('Hours_per_week Income vs Frequency of Income', c = 'g', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```





Above plot shows who work 40 Hours per week highest frequency

## Native Country column

```
In [47]: df['Native_country'].value_counts()
```

```
Out[47]: United-States      29169
Mexico          643
?
Philippines    198
Germany        137
Canada          121
Puerto-Rico    114
El-Salvador    106
India           100
Cuba            95
England         90
Jamaica         81
South           80
China           75
Italy            73
Dominican-Republic 70
Vietnam          67
Guatemala       64
Japan            62
Poland           60
Columbia         59
Taiwan           51
Haiti            44
Iran             43
Portugal         37
Nicaragua        34
Peru              31
France           29
Greece           29
Ecuador          28
Ireland          24
Hong              20
Trinidad&Tobago 19
Cambodia          19
Thailand          18
Laos              18
Yugoslavia       16
Outlying-US(Guam-USVI-etc) 14
Honduras          13
Hungary           13
Scotland          12
Holand-Netherlands 1
Name: Native_country, dtype: int64
```

**Observation : In Native Country column there is '?' present. We solve this problem to replace '?' with nan value**

```
In [48]: # Approach to remove '?' in Native Country column
```

```
df['Native_country'] = df['Native_country'].str.strip().map(lambda x: np.nan if x == ? else x)
```

```
In [49]: df.isna().sum()
```

```
Out[49]: Age          0  
Workclass      0  
Fnlwgt         0  
Education       0  
Education_num   0  
Marital_status  0  
Occupation      0  
Relationship    0  
Race            0  
Sex             0  
Capital_gain    0  
Capital_loss    0  
Hours_per_week  0  
Native_country  583  
Income           0  
dtype: int64
```

**Conclusion : In Occupation column '?' remove and replace with nan value. We fill nan with mode**


```
In [50]: df['Native_country'].fillna(df['Native_country'].mode()[0], inplace = True)
```

```
In [51]: df.isna().sum()
```

```
Out[51]: Age          0  
Workclass      0  
Fnlwgt         0  
Education       0  
Education_num   0  
Marital_status  0  
Occupation      0  
Relationship    0  
Race            0  
Sex             0  
Capital_gain    0  
Capital_loss    0  
Hours_per_week  0  
Native_country  0  
Income           0  
dtype: int64
```

**There is no null value in Native Country column we ready to do our analysis**

```
In [52]: df['Native_country'].value_counts().plot.barh(figsize = (10,10),rot= 360,color = plt.ylabel('Native Country', c = 'b', fontsize = 16)
plt.xlabel('Counts', c = 'b', fontsize = 16 )
plt.title('Native Country vs Counts', c = 'b', fontsize = 16)
plt.show()
```



Above plot shows US has highest counts

```
In [53]: n = df.groupby('Native_country')['Income'].value_counts().sort_values()  
n
```

```
Out[53]: Native_country      Income  
Holand-Netherlands    <=50K        1  
Honduras              >50K        1  
Laos                  >50K        2  
Peru                  >50K        2  
Dominican-Republic   >50K        2  
                           ...  
Puerto-Rico           <=50K       102  
Philippines           <=50K       137  
Mexico                <=50K       610  
United-States         >50K      7317  
                      <=50K     22435  
Name: Income, Length: 80, dtype: int64
```

```
In [54]: n.plot.barh(figsize = (10,15),color = ['red','cyan', 'g', 'm'])
plt.ylabel('Native Country Income', c = 'r', fontsize = 16)
plt.xlabel('Counts', c = 'r', fontsize = 16 )
plt.title('Native Country Income vs Counts', c = 'r', fontsize = 16)
plt.show()
```



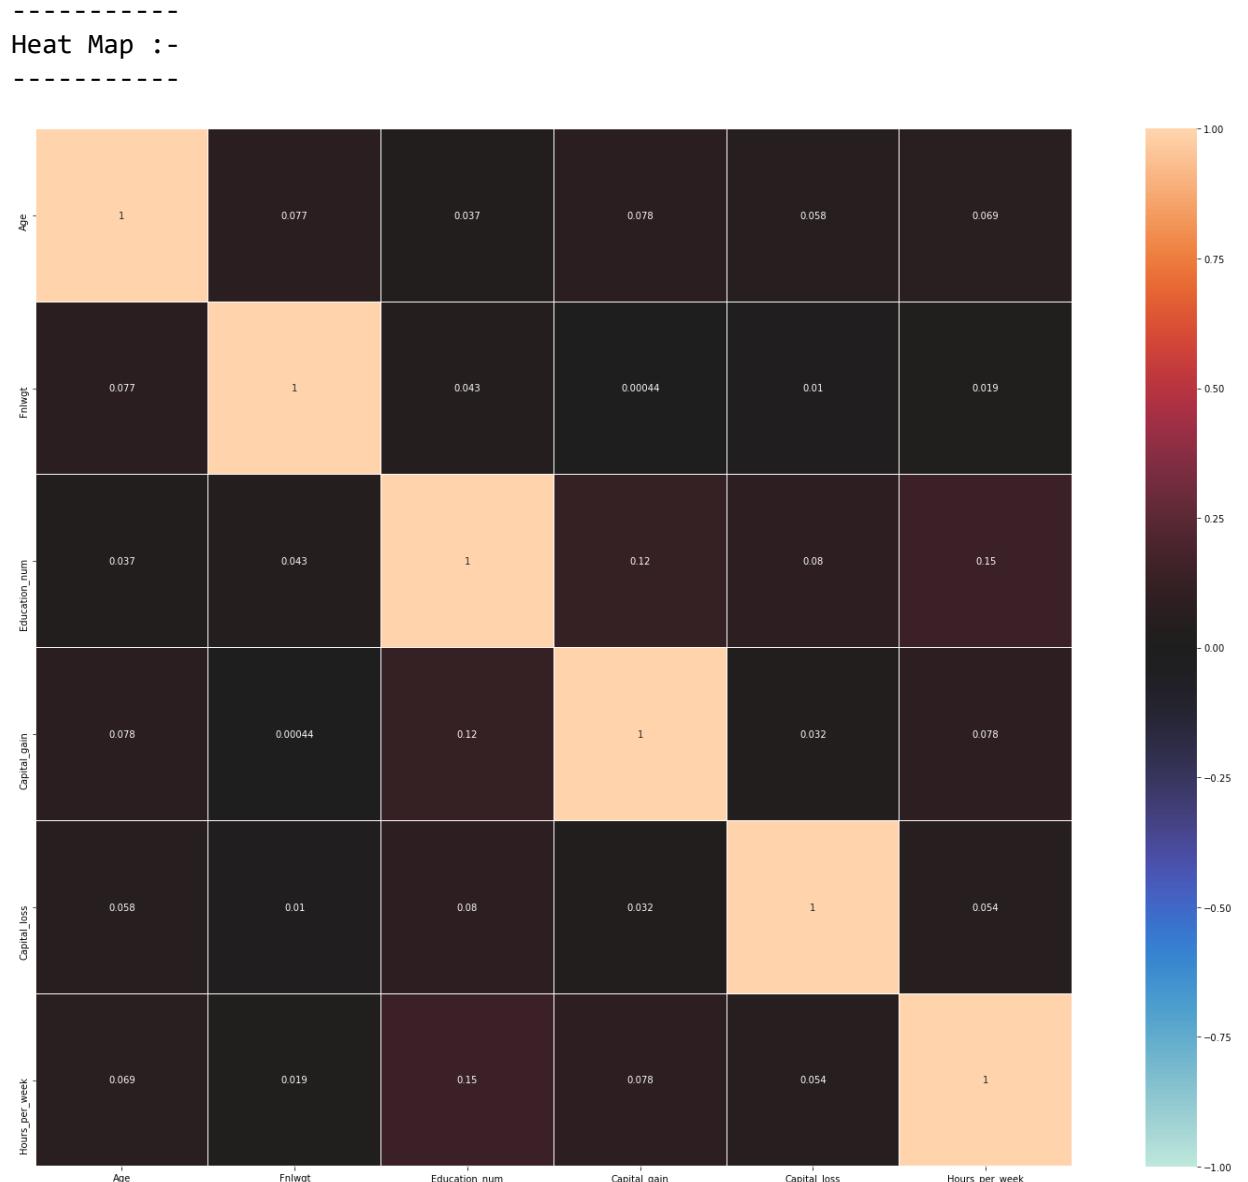


Above plot shows US has highest income

**Corelation of Feature vs Label using Heat map**

```
In [65]: print('-----')
print('Heat Map :-')
print('-----')
df_corr = df.corr().abs()

plt.figure(figsize = (22,16))
sns.heatmap(df_corr, vmin = -1, annot = True, square = True, center = 0, fmt = '.'
plt.tight_layout()
```



**Feature are not related**


## **Spliting Dataset into features and labels**

```
In [104]: x = df.drop('Income', axis = 1)
y = df. Income
print('Data has been splited')
```

Data has been splited

```
In [105]: # Let's see relation between features and labels.  
print('-----')  
print('Scatter Plot :-')  
print('-----')  
  
plt.figure(figsize = (20,25), facecolor = 'white')  
plotnumber = 1  
for column in x:  
    if plotnumber <=14:  
        ax = plt.subplot(5,3, plotnumber)  
        plt.scatter(x[column],y, c = 'g')  
        plt.xlabel(column, fontsize = 18)  
        plt.ylabel('Income', fontsize = 18)  
    plotnumber += 1  
plt.tight_layout()
```

-----  
Scatter Plot :-  
-----



```
In [71]: df['Income'].value_counts()
```

```
Out[71]: <=50K    24719
>50K      7841
Name: Income, dtype: int64
```

**Class are not balanced before balance it we first encode it**

```
In [101]: le = LabelEncoder()
```

```
In [102]: df['Income'] = le.fit_transform(df['Income'])
```

```
In [103]: df['Income'].value_counts()
```

```
Out[103]: 0    24719  
1    7841  
Name: Income, dtype: int64
```

**Class are encoded**

## Filter Categorical features

```
In [108]: numerics = ['int64', 'int32']  
categorical_col = []  
features = x.columns.values.tolist()  
for col in features:  
    if x[col].dtype in numerics:  
        continue  
    categorical_col.append(col)
```

## Encoding categorical columns using get dummies

```
In [109]: x_dummies = pd.get_dummies(x[categorical_col], drop_first = False)  
x_dummies.head()
```

Out[109]:

|   | Workclass_Federal-gov | Workclass_Local-gov | Workclass_Never-worked | Workclass_Private | Workclass_Self-emp-inc | W |
|---|-----------------------|---------------------|------------------------|-------------------|------------------------|---|
| 0 | 0                     | 0                   | 0                      | 0                 | 0                      | 0 |
| 1 | 0                     | 0                   | 0                      | 1                 | 0                      | 0 |
| 2 | 0                     | 0                   | 0                      | 1                 | 0                      | 0 |
| 3 | 0                     | 0                   | 0                      | 1                 | 0                      | 0 |
| 4 | 0                     | 0                   | 0                      | 1                 | 0                      | 0 |


```
In [80]: print('No of Rows and columns of encoded dataset =====>',x_dummies.shape)
```

No of Rows and columns of encoded dataset =====> (32560, 99)

```
In [110]: x = x.join(x_dummies)
x.head()
```

Out[110]:


|   | Age | Workclass        | Fnlwgt | Education | Education_num | Marital_status     | Occupation        | Relationship  | ... |
|---|-----|------------------|--------|-----------|---------------|--------------------|-------------------|---------------|-----|
| 0 | 50  | Self-emp-not-inc | 83311  | Bachelors | 13            | Married-civ-spouse | Exec-managerial   | Husband       | \   |
| 1 | 38  | Private          | 215646 | HS-grad   | 9             | Divorced           | Handlers-cleaners | Not-in-family | \   |
| 2 | 53  | Private          | 234721 | 11th      | 7             | Married-civ-spouse | Handlers-cleaners | Husband       | I   |
| 3 | 28  | Private          | 338409 | Bachelors | 13            | Married-civ-spouse | Prof-specialty    | Wife          | I   |
| 4 | 37  | Private          | 284582 | Masters   | 14            | Married-civ-spouse | Exec-managerial   | Wife          | \   |



```
In [111]: x.drop(columns = categorical_col, axis = 1, inplace = True) # Droping categorical
x.head()
```

Out[111]:

|   | Age | Fnlwgt | Education_num | Capital_gain | Capital_loss | Hours_per_week | Workclass_Federal-gov | ... |
|---|-----|--------|---------------|--------------|--------------|----------------|-----------------------|-----|
| 0 | 50  | 83311  |               | 13           | 0            | 0              | 13                    | 0   |
| 1 | 38  | 215646 |               | 9            | 0            | 0              | 40                    | 0   |
| 2 | 53  | 234721 |               | 7            | 0            | 0              | 40                    | 0   |
| 3 | 28  | 338409 |               | 13           | 0            | 0              | 40                    | 0   |
| 4 | 37  | 284582 |               | 14           | 0            | 0              | 40                    | 0   |



```
In [83]: print('No of Rows and columns of encoded dataset =====>',x.shape)
```

No of Rows and columns of encoded dataset =====> (32560, 105)


## Data distribution and checking outliers and skewness

```
In [85]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in x:
    if plotnumber <=105:
        ax = plt.subplot(15,7, plotnumber)
        sns.distplot(x[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----
Distribution Plot :- 
-----
```



Some outliers present in columns

Power Transformer to remove outliers and skewness

```
In [119]: scaler = PowerTransformer(method = 'yeo-johnson')
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
Out[119]: array([[ 0.89194656, -1.08988899,  1.16479277, ... ,  0.30721362,
       -0.04540906, -0.022173  ],
       [ 0.10271741,  0.39874263, -0.4643296 , ... ,  0.30721362,
       -0.04540906, -0.022173  ],
       [ 1.06438077,  0.56254016, -1.20027437, ... ,  0.30721362,
       -0.04540906, -0.022173  ],
       ... ,
       [ 1.33455599, -0.2197547 , -0.4643296 , ... ,  0.30721362,
       -0.04540906, -0.022173  ],
       [-1.35834147,  0.27166793, -0.4643296 , ... ,  0.30721362,
       -0.04540906, -0.022173  ],
       [ 1.00782289,  0.98216482, -0.4643296 , ... ,  0.30721362,
       -0.04540906, -0.022173 ]])
```

## Checking Outlier remove or not


```
In [89]: # Let's see outliers are removed in columns or not.
```

```
print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in x:
    if plotnumber <=105:
        ax = plt.subplot(15,7, plotnumber)
        sns.distplot(x[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----  
Distribution Plot :-  
-----
```



**Outliers are removed**

## Handling Class Imbalance

```
In [120]: sm = SMOTE()
x_over, y_over = sm.fit_resample(x,y)
```

```
In [121]: print('-----')
print('Class are balanced :-')
print('-----')
print(y_over.value_counts())
print('-----')
```

```
-----
Class are balanced :-  
-----  
0    24719  
1    24719  
Name: Income, dtype: int64  
-----
```

## Data Scaling

```
In [125]: scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
Out[125]: array([[ 0.83709708, -1.0087417 ,  1.13477863, ...,  0.30721362,
       -0.04540906, -0.022173 ],
       [-0.04264043,  0.24504633, -0.42002663, ...,  0.30721362,
       -0.04540906, -0.022173 ],
       [ 1.05703146,  0.42576955, -1.19742926, ...,  0.30721362,
       -0.04540906, -0.022173 ],
       ...,
       [ 1.42358875, -0.3588108 , -0.42002663, ...,  0.30721362,
       -0.04540906, -0.022173 ],
       [-1.21562378,  0.11092744, -0.42002663, ...,  0.30721362,
       -0.04540906, -0.022173 ],
       [ 0.98372 ,  0.92986178, -0.42002663, ...,  0.30721362,
       -0.04540906, -0.022173 ]])
```

Data has been scaled

**Split data into train and test. Model will be bulit on training data and tested on test data**

```
In [126]: x_train, x_test, y_train, y_test = train_test_split(x_over, y_over, test_size = 0.2)
print('Data has been splited.')  
◀ ▶
```

Data has been splited.

## Model Bulding

**Decision Tree model instantiaing, training and evaluating**

```
In [129]: DT = DecisionTreeClassifier()
DT.fit(x_train, y_train)
y_pred = DT.predict(x_test)
```

```
In [130]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

| Classification Report: |           |        |          |         |
|------------------------|-----------|--------|----------|---------|
|                        | precision | recall | f1-score | support |
| 0                      | 0.87      | 0.86   | 0.87     | 6115    |
| 1                      | 0.87      | 0.87   | 0.87     | 6245    |
| accuracy               |           |        | 0.87     | 12360   |
| macro avg              | 0.87      | 0.87   | 0.87     | 12360   |
| weighted avg           | 0.87      | 0.87   | 0.87     | 12360   |

```
-----
```

Conclusion : Decision Tree model has 87% score

## Cross Validation score to check if the model is overfitting


```
In [131]: cv = cross_val_score(DT, x, y, cv = 5)
print('Cross Validation score of Decision Tree model --->', cv.mean())
```

Cross Validation score of Decision Tree model ---> 0.8141584766584767

Conclusion : Decision Tree model has 81% Cross Validation score

## ROC, AUC Curve

```
In [132]: prob = DT.predict_proba(x_test) # calculating probability
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))
plt.show()
```



## XGBoost model instantiating, training and evaluating

```
In [133]: xgb = xgb.XGBClassifier(eval_metric = 'mlogloss')
xgb.fit(x_train, y_train)
y_pred = xgb.predict(x_test)
```

```
In [134]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

| Classification Report: |           |        |          |         |
|------------------------|-----------|--------|----------|---------|
|                        | precision | recall | f1-score | support |
| 0                      | 0.90      | 0.92   | 0.91     | 6115    |
| 1                      | 0.92      | 0.90   | 0.91     | 6245    |
| accuracy               |           |        | 0.91     | 12360   |
| macro avg              | 0.91      | 0.91   | 0.91     | 12360   |
| weighted avg           | 0.91      | 0.91   | 0.91     | 12360   |

```
-----
```

**Conclusion : XGBoost model has 91% score**

## Cross Validation score to check if the model is overfitting


```
In [135]: cv = cross_val_score(xgb, x, y, cv = 5)
print('Cross Validation score of XGBoost model --->', cv.mean())
```

Cross Validation score of XGBoost model ---> 0.8708230958230958

**Conclusion : XGBoost model has 87% Cross Validation score**

## ROC, AUC Curve

```
In [136]: prob = xgb.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



## Knn model instantiaing, training and evaluating

```
In [137]: Knn = KNeighborsClassifier()  
Knn.fit(x_train, y_train)  
y_pred = Knn.predict(x_test)
```

```
In [138]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

| Classification Report: |           |        |          |         |
|------------------------|-----------|--------|----------|---------|
|                        | precision | recall | f1-score | support |
| 0                      | 0.74      | 0.66   | 0.70     | 6115    |
| 1                      | 0.70      | 0.77   | 0.73     | 6245    |
| accuracy               |           |        | 0.72     | 12360   |
| macro avg              | 0.72      | 0.72   | 0.72     | 12360   |
| weighted avg           | 0.72      | 0.72   | 0.72     | 12360   |

```
-----
```

**Conclusion : KNN model has 72% score**

## Cross Validation score to check if the model is overfitting


```
In [139]: cv = cross_val_score(Knn, x, y, cv = 5)
print('Cross Validation score of Knn model --->', cv.mean())
```

Cross Validation score of Knn model ---> 0.7766891891891892

**Conclusion : Knn model has 77% Cross Validation score**

## ROC, AUC Curve

```
In [140]: prob = Knn.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



## Random Forest model instantiaing, training and evaluating

```
In [141]: Rn = RandomForestClassifier()  
Rn.fit(x_train, y_train)  
y_pred = Rn.predict(x_test)
```

```
In [142]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

| Classification Report: |           |        |          |         |
|------------------------|-----------|--------|----------|---------|
|                        | precision | recall | f1-score | support |
| 0                      | 0.89      | 0.92   | 0.90     | 6115    |
| 1                      | 0.92      | 0.89   | 0.90     | 6245    |
| accuracy               |           |        | 0.90     | 12360   |
| macro avg              | 0.90      | 0.90   | 0.90     | 12360   |
| weighted avg           | 0.90      | 0.90   | 0.90     | 12360   |

**Conclusion : Random Forest model has 90% score**

## Cross Validation score to check if the model is overfitting


```
In [143]: cv = cross_val_score(Rn, x, y, cv = 5)
print('Cross Validation score of Random Forest model --->', cv.mean())
```

Cross Validation score of Random Forest model ---> 0.8527334152334152

**Conclusion : Random Forest model has 85% Cross Validation score**

## ROC, AUC Curve

```
In [144]: prob = Rn.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



## Logistic Regression model instantiating, training and evaluating

```
In [146]: Lr = LogisticRegression()  
Lr.fit(x_train, y_train)  
y_pred = Lr.predict(x_test)
```

```
In [147]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

| Classification Report: |           |        |          |         |
|------------------------|-----------|--------|----------|---------|
|                        | precision | recall | f1-score | support |
| 0                      | 0.59      | 0.70   | 0.64     | 6115    |
| 1                      | 0.64      | 0.53   | 0.58     | 6245    |
| accuracy               |           |        | 0.61     | 12360   |
| macro avg              | 0.62      | 0.61   | 0.61     | 12360   |
| weighted avg           | 0.62      | 0.61   | 0.61     | 12360   |

```
-----
```

**Conclusion : Logistic Regression model has 61% score**

## Cross Validation score to check if the model is overfitting


```
In [148]: cv = cross_val_score(Lr, x, y, cv = 5)
print('Cross Validation score of Logistic regression model --->', cv.mean())
```

Cross Validation score of Logistic regression model ---> 0.7976351351351351

**Conclusion : Logistic Regression model has 79% Cross Validation score**

## ROC, AUC Curve

```
In [149]: prob = Lr.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



**Looking CV score we found KNN has best model so we do Hyperparameter Tuning on it**

```
In [150]: param_grid = {'leaf_size': [1, 3, 5], 'n_neighbors': [5], 'p': [1,2]}
```

```
In [151]: grid_search = GridSearchCV(estimator = Knn, param_grid = param_grid, cv = 5,n_jobs
```

```
In [152]: grid_search.fit(x_train, y_train)
```

```
Out[152]: GridSearchCV(cv=5, estimator=KNeighborsClassifier(), n_jobs=-1,  
param_grid={'leaf_size': [1, 3, 5], 'n_neighbors': [5],  
'p': [1, 2]})
```

```
In [153]: best_parameters = grid_search.best_params_  
print(best_parameters)
```

```
{'leaf_size': 1, 'n_neighbors': 5, 'p': 1}
```

```
In [156]: hKnn = KNeighborsClassifier(leaf_size = 1, n_neighbors = 5, p = 1)  
hKnn.fit(x_train, y_train)  
hKnn.score(x_test, y_test)
```

```
Out[156]: 0.7362459546925566
```

```
In [157]: y_pred = hKnn.predict(x_test)
```

```
In [158]: print('-----')  
print('\nClassification Report: ')  
print(classification_report(y_test, y_pred, digits = 2))  
print('-----\n')
```

```
-----  
Classification Report:  
precision    recall   f1-score   support  
  
          0       0.77      0.66      0.71      6115  
          1       0.71      0.81      0.76      6245  
  
accuracy                           0.74      12360  
macro avg       0.74      0.74      0.73      12360  
weighted avg     0.74      0.74      0.73      12360  
  
-----
```

After Hyperparameter Tuning model accuracy score 74%.

## Saving The Model

```
In [166]: # saving the model to the Local file system  
filename = 'Census Income Project.pickle'  
pickle.dump(hKnn, open(filename, 'wb'))
```

**Final Conclusion : KNN is our best model.**

In [ ]: