

## Problem Statement:

Customer churn is when a company's customers stop doing business with that company. Businesses are very keen on measuring churn because keeping an existing customer is far less expensive than acquiring a new customer. New business involves working leads through a sales funnel, using marketing and sales budgets to gain additional customers. Existing customers will often have a higher volume of service consumption and can generate additional customer referrals.

Customer retention can be achieved with good customer service and products. But the most effective way for a company to prevent attrition of customers is to truly know them. The vast volumes of data collected about customers can be used to build churn prediction models. Knowing who is most likely to defect means that a company can prioritise focused marketing efforts on that subset of their customer base.

Preventing customer churn is critically important to the telecommunications sector, as the barriers to entry for switching services are so low.

## Import Required Library

```
In [56]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pickle
import scikitplot as skplt
from imblearn.over_sampling import SMOTE
pd.set_option('display.max_columns', None) # For display maximum columns
from sklearn.preprocessing import OrdinalEncoder, MinMaxScaler
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, plot_roc_curve
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier, BaggingClassifier
from sklearn.linear_model import LogisticRegression
import xgboost as xgb
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

## Rading Data

```
In [2]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\csv file\Telecom_customer_churn.csv")
df.head()
```

Out[2]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	7590-VHVEG	Female	0	Yes	No	1	No	No	No phone service
1	5575-GNVDE	Male	0	No	No	34	Yes	No	No
2	3668-QPYBK	Male	0	No	No	2	Yes	No	No
3	7795-CFOCW	Male	0	No	No	45	No	No	No phone service
4	9237-HQITU	Female	0	No	No	2	Yes	No	No

## Check no of row and column

```
In [3]: print('No of Rows and Columns ----->', df.shape)
```

No of Rows and Columns -----> (7043, 21)

## Checking for Null values

```
In [4]: print('=====\\n')
print(df.isnull().sum())
print('\\n=====')
```

```
=====
```

```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    11
Churn           0
dtype: int64
```

```
=====
```

**There is null value**

## Fill NaN

```
In [5]: df['TotalCharges'].fillna(df['TotalCharges'].mean(), inplace = True)
```

```
In [6]: print('=====\\n')
print(df.isnull().sum())
print('\\n=====')
```

```
=====
```

```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity   0
OnlineBackup     0
DeviceProtection 0
TechSupport      0
StreamingTV     0
StreamingMovies  0
Contract         0
PaperlessBilling 0
PaymentMethod    0
MonthlyCharges   0
TotalCharges     0
Churn            0
dtype: int64
```

```
=====
```

**Nan are removed**

## Information about dataset

```
In [6]: print('=====\\n')
print(df.info())
print('=====')
```

```
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customerID      7043 non-null    object  
 1   gender          7043 non-null    object  
 2   SeniorCitizen   7043 non-null    int64  
 3   Partner         7043 non-null    object  
 4   Dependents     7043 non-null    object  
 5   tenure          7043 non-null    int64  
 6   PhoneService    7043 non-null    object  
 7   MultipleLines   7043 non-null    object  
 8   InternetService 7043 non-null   object  
 9   OnlineSecurity  7043 non-null   object  
 10  OnlineBackup    7043 non-null   object  
 11  DeviceProtection 7043 non-null   object  
 12  TechSupport     7043 non-null   object  
 13  StreamingTV     7043 non-null   object  
 14  StreamingMovies 7043 non-null   object  
 15  Contract        7043 non-null   object  
 16  PaperlessBilling 7043 non-null   object  
 17  PaymentMethod   7043 non-null   object  
 18  MonthlyCharges 7043 non-null   float64 
 19  TotalCharges    7043 non-null   object  
 20  Churn           7043 non-null   object  
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
None
=====
```

Categorical data present in our data set

## Drop unwanted column

```
In [7]: df = df.drop('customerID', axis = 1)
df.head(2)
```

Out[7]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	Female	0	Yes	No	1	No	No phone service	DSL
1	Male	0	No	No	34	Yes	No	DSL

We drop the customerID column

## Analysis of Data with respect to Churn

### Gender column


```
In [8]: df['gender'].value_counts()
```

```
Out[8]: Male      3555  
Female    3488  
Name: gender, dtype: int64
```

```
In [9]: df['Churn'].value_counts()
```

```
Out[9]: No      5174  
Yes     1869  
Name: Churn, dtype: int64
```

```
In [19]: sns.countplot( x="gender", data=df)  
plt.xlabel('Gender', c = 'r', fontsize = 12)  
plt.ylabel('Counts', c = 'r', fontsize = 12)  
plt.title('Gender vs Counts', c = 'r', fontsize = 12)  
plt.show()
```




Male is highest in number

```
In [21]: gn = df.groupby('gender')['Churn'].value_counts()  
gn
```

```
Out[21]: gender Churn  
Female No      2549  
        Yes     939  
Male   No      2625  
        Yes     930  
Name: Churn, dtype: int64
```

```
In [24]: gn.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])  
plt.xlabel('Gender Churn', c = 'purple', fontsize = 12)  
plt.ylabel('Churn Counts', c = 'purple', fontsize = 12 )  
plt.title('Gender Churn vs Churn Counts', c = 'purple', fontsize = 12)  
plt.xticks(c = 'b')  
plt.yticks(c = 'b')  
plt.show()
```




Female has highest no of churn

## Senior Citizen column

```
In [25]: df['SeniorCitizen'].value_counts()
```


```
Out[25]: 0    5901  
1    1142  
Name: SeniorCitizen, dtype: int64
```

```
In [26]: sns.histplot(binwidth=0.5, x="SeniorCitizen", hue="SeniorCitizen", data=df, stat="density")
plt.xlabel('Senior Citizen', c = 'b', fontsize = 12)
plt.ylabel('Counts', c = 'b', fontsize = 12)
plt.title('Senior Citizen vs Counts', c = 'b', fontsize = 12)
plt.show()
```



```
In [27]: sc = df.groupby('SeniorCitizen')['Churn'].value_counts()
sc
```

```
In [28]: sc.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Senior Citizen Churn', c = 'g', fontsize = 12)
plt.ylabel('Churn Counts', c = 'g', fontsize = 12 )
plt.title('Senior Citizen Churn vs Churn Counts', c = 'g', fontsize = 12)
plt.xticks(c = 'r')
plt.yticks(c = 'r')
plt.show()
```




In Senior Citizen highest no of Not Churn ↗

**Partner column**

```
In [29]: df['Partner'].value_counts()
```

```
Out[29]: No      3641  
Yes     3402  
Name: Partner, dtype: int64
```

```
In [31]: sns.countplot( x="Partner", data=df)  
plt.xlabel('Partner', c = 'm', fontsize = 12)  
plt.ylabel('Counts', c = 'm', fontsize = 12)  
plt.title('Partner vs Counts', c = 'm', fontsize = 12)  
plt.show()
```




**Person has No Partner count is high**

```
In [33]: p = df.groupby('Partner')[ 'Churn'].value_counts()  
p
```

```
Out[33]: Partner  Churn  
No        No      2441  
           Yes     1200  
Yes       No      2733  
           Yes     669  
Name: Churn, dtype: int64
```

```
In [35]: p.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Partner Churn', c = 'b', fontsize = 12)
plt.ylabel('Churn Counts', c = 'b', fontsize = 12 )
plt.title('Partner Churn vs Churn Counts', c = 'b', fontsize = 12)
plt.xticks(c = 'k')
plt.yticks(c = 'k')
plt.show()
```




Person which has Partner no of Churn is high ↗

## Dependents column

```
In [36]: df['Dependents'].value_counts()
```


```
Out[36]: No      4933
Yes     2110
Name: Dependents, dtype: int64
```

```
In [37]: sns.histplot(binwidth=0.5, x="Dependents", hue="Dependents", data=df, stat="count")
plt.xlabel('Dependents', c = 'k', fontsize = 12)
plt.ylabel('Counts', c = 'k', fontsize = 12)
plt.title('Dependents vs Counts', c = 'k', fontsize = 12)
plt.show()
```



```
In [38]: d = df.groupby('Dependents')['Churn'].value_counts()  
d
```

```
In [40]: d.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Dependents Churn', c = 'purple', fontsize = 12)
plt.ylabel('Churn Counts', c = 'purple', fontsize = 12 )
plt.title('Dependents Churn vs Churn Counts', c = 'purple', fontsize = 12)
plt.xticks(c = 'k')
plt.yticks(c = 'k')
plt.show()
```



**Dependents person no of Churn is low**

## Tenure column


```
In [42]: df['tenure'].value_counts()
```


```
Out[42]: 1      613
72     362
2      238
3      200
4      176
...
28      57
39      56
44      51
36      50
0       11
Name: tenure, Length: 73, dtype: int64
```

```
In [57]: tn = df.groupby('tenure')['Churn'].value_counts().sort_values()
tn
```

```
Out[57]: tenure  Churn
63      Yes      4
64      Yes      4
62      Yes      5
72      Yes      6
45      Yes      6
...
2       Yes     123
71      No      164
1       No      233
72      No      356
1       Yes     380
Name: Churn, Length: 145, dtype: int64
```

```
In [58]: tn.plot.bahr(figsize = (10,25), color = ['g','r'])
plt.ylabel('Tenure Churn', c = 'purple', fontsize = 12)
plt.xlabel('Churn Counts', c = 'purple', fontsize = 12 )
plt.title('Tenure Churn vs Churn Counts', c = 'purple', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```






Person has year tenure has highest churn

## Phone Service column

```
In [59]: df['PhoneService'].value_counts()
```

```
Out[59]: Yes    6361
No      682
Name: PhoneService, dtype: int64
```

```
In [62]: sns.histplot(binwidth=0.5, x="PhoneService", hue="PhoneService", data=df, stat="c  
plt.xlabel('Phone Service', c = 'r', fontsize = 12)  
plt.ylabel('Counts', c = 'r', fontsize = 12)  
plt.title('Phone Service vs Counts', c = 'r', fontsize = 12)  
plt.show()
```




A person has Phone Service highest counts ↗

```
In [63]: ps = df.groupby('PhoneService')['Churn'].value_counts()  
ps
```

```
Out[63]: PhoneService  Churn  
          No           No      512  
                      Yes     170  
          Yes          No    4662  
                      Yes    1699  
Name: Churn, dtype: int64
```

```
In [64]: ps.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Phone Service Churn', c = 'r', fontsize = 12)
plt.ylabel('Churn Counts', c = 'r', fontsize = 12 )
plt.title('Phone Service Churn vs Churn Counts', c = 'r', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```




Person has Phone Service no of churn is low ↗

## Multiple Lines column

```
In [66]: df['MultipleLines'].value_counts()
```

```
Out[66]: No          3390
Yes         2971
No phone service    682
Name: MultipleLines, dtype: int64
```


```
In [67]: sns.histplot(binwidth=0.5, x="MultipleLines", hue="MultipleLines", data=df, stat="density")
plt.xlabel('MultipleLines', c= 'm', fontsize = 12)
plt.ylabel('Counts', c = 'm', fontsize = 12)
plt.title('MultipleLines vs Counts', c = 'm', fontsize = 12)
plt.show()
```



**Person has no phone service count is highest**

```
In [68]: ml = df.groupby('MultipleLines')['Churn'].value_counts()
ml
```

```
In [70]: ml.plot.bar(figsize = (15,5), rot = 360, color = ['g','r'])
plt.xlabel('Multiple Lines Churn', c = 'purple', fontsize = 12)
plt.ylabel('Churn Counts', c = 'purple', fontsize = 12 )
plt.title('Multiple Lines Churn vs Churn Counts', c = 'purple', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```




A person has Multiple Lines no of churn is low ↴

```
In [71]: df['InternetService'].value_counts()
```

```
Out[71]: Fiber optic    3096
DSL            2421
No             1526
Name: InternetService, dtype: int64
```


```
In [72]: sns.countplot( x="InternetService", data=df)
plt.xlabel('Internet Service', c = 'b', fontsize = 12)
plt.ylabel('Counts', c = 'b', fontsize = 12)
plt.title('Internet Service vs Counts', c = 'b', fontsize = 12)
plt.show()
```



**Fiber optic Internet Service count is highest** 

```
In [73]: ins = df.groupby('InternetService')['Churn'].value_counts()  
ins
```

```
In [74]: ins.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Gender Churn', c = 'r', fontsize = 12)
plt.ylabel('Churn Counts', c = 'r', fontsize = 12 )
plt.title('Gender Churn vs Churn Counts', c = 'r', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```




In fiber optic Internet Service has no of churn is low ⌘

## Online Security column

```
In [75]: df['OnlineSecurity'].value_counts()
```

```
Out[75]: No           3498
Yes          2019
No internet service  1526
Name: OnlineSecurity, dtype: int64
```


```
In [76]: sns.countplot( x="OnlineSecurity", data=df)
plt.xlabel('Online Security', c = 'g', fontsize = 12)
plt.ylabel('Counts', c = 'g', fontsize = 12)
plt.title('Online Security vs Counts', c = 'g', fontsize = 12)
plt.show()
```



**A person has no online security count has highest**

```
In [77]: os = df.groupby('OnlineSecurity')['Churn'].value_counts()
os
```

```
In [80]: os.plot.bar(figsize = (15,5), rot = 360, color = ['g','r'])
plt.xlabel('Online Security Churn', c = 'k', fontsize = 12)
plt.ylabel('Churn Counts', c = 'k', fontsize = 12 )
plt.title('Online Security Churn vs Churn Counts', c = 'k', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```




A person has no online security has highest no of churn

## Online Backup column

```
In [82]: df['OnlineBackup'].value_counts()
```

```
Out[82]: No          3088
Yes         2429
No internet service  1526
Name: OnlineBackup, dtype: int64
```

```
In [83]: sns.histplot(binwidth=0.5, x="OnlineBackup", hue="OnlineBackup", data=df, stat="c  
plt.xlabel('Online Backup', c = 'b', fontsize = 12)  
plt.ylabel('Counts', c = 'b', fontsize = 12)  
plt.title('Online Backup vs Counts', c = 'b', fontsize = 12)  
plt.show()
```




A person has No Online Backup counts is highest

## Device Protection

```
In [84]: df['DeviceProtection'].value_counts()
```

```
Out[84]: No          3095  
Yes         2422  
No internet service  1526  
Name: DeviceProtection, dtype: int64
```

```
In [85]: sns.countplot(x = "DeviceProtection", data=df)
plt.xlabel('Device Protection', c = 'r', fontsize = 12)
plt.ylabel('Counts', c = 'r', fontsize = 12)
plt.title('Device Protection vs Counts', c = 'r', fontsize = 12)
plt.show()
```




**A person has No Device Protection counts is highest**

## Tech Support

```
In [86]: df['TechSupport'].value_counts()
```

```
Out[86]: No          3473
Yes         2044
No internet service  1526
Name: TechSupport, dtype: int64
```

```
In [87]: sns.histplot(binwidth=0.5, x="TechSupport", hue="TechSupport", data=df, stat="counts")
plt.xlabel('Tech Support', c = 'b', fontsize = 12)
plt.ylabel('Counts', c = 'b', fontsize = 12)
plt.title('Tech Support vs Counts', c = 'b', fontsize = 12)
plt.show()
```




A person has No Tech Support counts is highest

## Streaming TV column

```
In [88]: df['StreamingTV'].value_counts()
```

```
Out[88]: No           2810
Yes          2707
No internet service  1526
Name: StreamingTV, dtype: int64
```

```
In [89]: sns.histplot(binwidth=0.5, x="StreamingTV", hue="StreamingTV", data=df, stat="counts")
plt.xlabel('Streaming TV', c = 'g', fontsize = 12)
plt.ylabel('Counts', c = 'g', fontsize = 12)
plt.title('Streaming TV vs Counts', c = 'g', fontsize = 12)
plt.show()
```




A person has Not Streaming TV live counts is highest count

## Streaming Movies column

```
In [91]: df['StreamingMovies'].value_counts()
```

```
Out[91]: No           2785
Yes          2732
No internet service  1526
Name: StreamingMovies, dtype: int64
```

```
In [92]: sns.countplot(x="StreamingMovies", data=df)
plt.xlabel('Streaming Movies', c = 'b', fontsize = 12)
plt.ylabel('Counts', c = 'b', fontsize = 12)
plt.title('Streaming Movies vs Counts', c = 'b', fontsize = 12)
plt.show()
```




A person has Not Streaming Movies is highest count

## Contract columns

```
In [93]: df['Contract'].value_counts()
```

```
Out[93]: Month-to-month    3875
Two year          1695
One year          1473
Name: Contract, dtype: int64
```

```
In [94]: sns.histplot(binwidth=0.5, x="Contract", hue="Contract", data=df, stat="count", r  
plt.xlabel('Contract', c = 'b', fontsize = 12)  
plt.ylabel('Counts', c = 'b', fontsize = 12)  
plt.title('Contract vs Counts', c = 'b', fontsize = 12)  
plt.show()
```




A person has Month to month contract counts is highest

## Paperless Billing column

```
In [95]: df['PaperlessBilling'].value_counts()
```


```
Out[95]: Yes    4171  
No     2872  
Name: PaperlessBilling, dtype: int64
```

```
In [96]: sns.histplot(binwidth=0.5, x="PaperlessBilling", hue="PaperlessBilling", data=df,
plt.xlabel('Paperless Billing', c = 'b', fontsize = 12)
plt.ylabel('Counts', c = 'b', fontsize = 12)
plt.title('Paperless Billing vs Counts', c = 'b', fontsize = 12)
plt.show()
```



```
In [97]: pb = df.groupby('PaperlessBilling')['Churn'].value_counts()  
pb
```

```
In [98]: pb.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Paperless Billing Churn', c = 'purple', fontsize = 12)
plt.ylabel('Churn Counts', c = 'purple', fontsize = 12 )
plt.title('Paperless Billing Churn vs Churn Counts', c = 'purple', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```




A person has Paperless Billing is lowest churn

## Payment Method column

```
In [105]: pm = df['PaymentMethod'].value_counts()
pm
```

```
Out[105]: Electronic check      2365
Mailed check                 1612
Bank transfer (automatic)   1544
Credit card (automatic)     1522
Name: PaymentMethod, dtype: int64
```

```
In [109]: pm.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Payment Method', c = 'purple', fontsize = 12)
plt.ylabel('Churn Counts', c = 'purple', fontsize = 12 )
plt.title('Payment Method vs Churn Counts', c = 'purple', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```




**Electronic check payment method is highest**

```
In [114]: pmg = df.groupby('PaymentMethod')['Churn'].value_counts().sort_values()
pmg
```

```
Out[114]: PaymentMethod      Churn
Credit card (automatic)  Yes    232
Bank transfer (automatic) Yes    258
Mailed check             Yes    308
Electronic check          Yes   1071
Bank transfer (automatic) No     1286
Credit card (automatic)   No     1290
Electronic check          No     1294
Mailed check              No     1304
Name: Churn, dtype: int64
```

```
In [115]: plt.plot.barh(figsize = (10,5), rot = 360, color = ['g','r'])
plt.ylabel('Payment Method', c = 'k', fontsize = 12)
plt.xlabel('Churn Counts', c = 'k', fontsize = 12 )
plt.title('Payment Method vs Churn Counts', c = 'k', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```




A person has No mailed check counts is highest churn

## Monthly Charges column

```
In [119]: df['MonthlyCharges'].value_counts()
```

```
Out[119]: 20.05      61
19.85      45
19.90      44
19.95      44
19.65      43
..
87.65       1
35.30       1
114.85      1
56.50       1
97.25       1
Name: MonthlyCharges, Length: 1585, dtype: int64
```

```
In [133]: sns.distplot(df['MonthlyCharges'])
plt.show()
```




**Above plot normally distributed**

## Total Charges column

```
In [141]: df['TotalCharges'].value_counts()
```

```
Out[141]: 2283.300441    11
20.200000    11
19.750000     9
19.650000     8
19.900000     8
...
1474.350000     1
6668.350000     1
6096.900000     1
140.700000     1
5867.000000     1
Name: TotalCharges, Length: 6531, dtype: int64
```

```
In [142]: sns.distplot(df['TotalCharges'])
plt.show()
```



Above plot is skewed

## Encoding Categorical Column

```
In [8]: oe = OrdinalEncoder()
```

```
In [9]: for i in df.columns:
    if df[i].dtypes == 'object':
        df[i] = oe.fit_transform(df[i].values.reshape(-1,1))
df.head(2)
```

Out[9]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	0.0	0	1.0	0.0	1	0.0	1.0	0.0
1	1.0	0	0.0	0.0	34	1.0	0.0	0.0

```
In [10]: print('=====\\n')
print(df.info())
print('=====')
```

```
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   gender            7043 non-null    float64
 1   SeniorCitizen     7043 non-null    int64  
 2   Partner           7043 non-null    float64
 3   Dependents        7043 non-null    float64
 4   tenure            7043 non-null    int64  
 5   PhoneService      7043 non-null    float64
 6   MultipleLines     7043 non-null    float64
 7   InternetService   7043 non-null    float64
 8   OnlineSecurity    7043 non-null    float64
 9   OnlineBackup       7043 non-null    float64
 10  DeviceProtection  7043 non-null    float64
 11  TechSupport       7043 non-null    float64
 12  StreamingTV       7043 non-null    float64
 13  StreamingMovies   7043 non-null    float64
 14  Contract          7043 non-null    float64
 15  PaperlessBilling  7043 non-null    float64
 16  PaymentMethod     7043 non-null    float64
 17  MonthlyCharges   7043 non-null    float64
 18  TotalCharges      7043 non-null    float64
 19  Churn             7043 non-null    float64
dtypes: float64(18), int64(2)
memory usage: 1.1 MB
None
=====
```


## Data distribution

```
In [149]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in df:
    if plotnumber <=20:
        ax = plt.subplot(5,4, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----
Distribution Plot :-
```



**Some columns are little skewed**

**Check skewness**

```
In [150]: df.skew()
```

```
Out[150]: gender           -0.019031
SeniorCitizen      1.833633
Partner            0.067922
Dependents         0.875199
tenure             0.239540
PhoneService       -2.727153
MultipleLines      0.118719
InternetService   0.205423
OnlineSecurity     0.416985
OnlineBackup        0.182930
DeviceProtection   0.186847
TechSupport         0.402365
StreamingTV        0.028486
StreamingMovies    0.014657
Contract           0.630959
PaperlessBilling   -0.375396
PaymentMethod      -0.170129
MonthlyCharges     -0.220524
TotalCharges       0.962394
Churn              1.063031
dtype: float64
```

**Skewness present in our dataset**

**Corelation of Feature vs Label using Heat map**

```
In [152]: print('-----')
```

```
print('Heat Map :-')
```

```
print('-----')
```


```
df_corr = df.corr().abs()
```

```
plt.figure(figsize = (22,16))
```

```
sns.heatmap(df_corr, vmin = -1, annot = True, square = True, center = 0, fmt = '.2f')
```

```
plt.tight_layout()
```

-----  
Heat Map :-  
-----



**Total charges has highest corelation with label**

## **Spliting Dataset into features and label**

```
In [11]: x = df.drop('Churn', axis = 1)
y = df.Churn
print('Data has been splited')
```


Data has been splited

```
In [156]: # Let's see relation between features and label.
```

```
print('-----')
print('Scatter Plot :-')
print('-----')

plt.figure(figsize = (20,25), facecolor = 'yellow')
plotnumber = 1
for column in x:
    if plotnumber <=20:
        ax = plt.subplot(5,4, plotnumber)
        plt.scatter(x[column],y, c = 'g')
        plt.xlabel(column, fontsize = 18)
        plt.ylabel('Loan_Status', fontsize = 18)
    plotnumber += 1
plt.tight_layout()
```

```
-----
Scatter Plot :-
-----
```



**Positive relation in feature and label**

**Checking for class imbalance**

```
In [12]: df['Churn'].value_counts()
```

```
Out[12]: 0.0    5174  
1.0    1869  
Name: Churn, dtype: int64
```

## Handling Class Imbalance

```
In [13]: sm = SMOTE()  
x_over, y_over = sm.fit_resample(x,y)
```

```
In [14]: print('-----')  
print('Class are balanced :-')  
print('-----')  
print(y_over.value_counts())  
print('-----')
```

```
-----  
Class are balanced :-  
-----  
0.0    5174  
1.0    5174  
Name: Churn, dtype: int64  
-----
```

## Data Scaling

```
In [57]: scaler = MinMaxScaler()  
x_scaled = scaler.fit_transform(x)  
x_scaled
```

```
Out[57]: array([[0.          , 0.          , 1.          , ... , 0.66666667, 0.11542289,  
   0.0012751 ],  
   [1.          , 0.          , 0.          , ... , 1.          , 0.38507463,  
   0.21586661],  
   [1.          , 0.          , 0.          , ... , 1.          , 0.35422886,  
   0.01031041],  
   ... ,  
   [0.          , 0.          , 1.          , ... , 0.66666667, 0.11293532,  
   0.03780868],  
   [1.          , 1.          , 1.          , ... , 1.          , 0.55870647,  
   0.03321025],  
   [1.          , 0.          , 0.          , ... , 0.          , 0.86965174,  
   0.78764136]])
```

**Split data into train and test. Model will be bulit on training data and tested on test data**

```
In [58]: x_train, x_test, y_train, y_test = train_test_split(x_over, y_over, test_size = 0.2)
print('Data has been splited.')
Data has been splited.
```

## Model Bulding

### Decision Tree model instantiaing, training and evaluating

```
In [59]: bag_dt = BaggingClassifier(DecisionTreeClassifier(), n_estimators = 30, max_samples = 100, random_state= 3, oob_score = True)
```

```
In [60]: bag_dt.oob_score
```

```
Out[60]: True
```

```
In [61]: bag_dt.fit(x_train, y_train)
print('Bagging DT score ----->', bag_dt.score(x_test, y_test))
```

```
Bagging DT score -----> 0.8465403942790878
```

```
In [62]: y_pred = bag_dt.predict(x_test)
```

```
In [63]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----  
Classification Report:  
precision    recall    f1-score    support  
  
      0.0        0.83      0.87      0.85      1298  
      1.0        0.86      0.82      0.84      1289  
  
accuracy                           0.85      2587  
macro avg       0.85      0.85      0.85      2587  
weighted avg     0.85      0.85      0.85      2587  
-----
```

Conclusion : Decision Tree model has 85% score

Cross Validation score to check if the model is overfitting


```
In [64]: cv = cross_val_score(bag_dt, x, y, cv = 5)
print('Cross Validation score of Decision Tree model --->', cv.mean())
```

```
Cross Validation score of Decision Tree model ---> 0.78730795131944
```

Conclusion : Decision Tree model has 78% Cross Validation score

## ROC, AUC Curve

```
In [65]: prob = bag_dt.predict_proba(x_test) # calculating probability
skplt.metrics.plot_roc(y_pred, prob, figsize = (10,5))
plt.show()
```



## XGBoost model instantiaing, training and evaluating

```
In [66]: bag_xgb = BaggingClassifier(xgb.XGBClassifier(eval_metric = 'mlogloss'), n_estimators=3, random_state=42, oob_score = True)
```

```
In [67]: bag_xgb.oob_score
```

```
Out[67]: True
```

```
In [68]: bag_xgb.fit(x_train, y_train)
print('Bagging XGBoost score ----->', bag_xgb.score(x_test, y_test))
```

```
Bagging XGBoost score -----> 0.8612292230382682
```

```
In [69]: y_pred = bag_xgb.predict(x_test)
```

```
In [70]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

	precision	recall	f1-score	support
0.0	0.85	0.87	0.86	1298
1.0	0.87	0.85	0.86	1289
accuracy			0.86	2587
macro avg	0.86	0.86	0.86	2587
weighted avg	0.86	0.86	0.86	2587

```
-----
```

**Conclusion : XGBoost model has 86% score**

## Cross Validation score to check if the model is overfitting


```
In [71]: cv = cross_val_score(bag_xgb, x, y, cv = 5)
print('Cross Validation score of XGBoost model --->', cv.mean())
```

Cross Validation score of XGBoost model ---> 0.7968201577521131

**Conclusion : XGBoost model has 79% Cross Validation score**

## ROC, AUC Curve

```
In [72]: prob = bag_xgb.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



## Knn model instantiaing, training and evaluating

```
In [73]: bag_Knn = BaggingClassifier(KNeighborsClassifier(n_neighbors = 5), n_estimators =  
random_state= 3, oob_score = True)
```

```
In [74]: bag_Knn.oob_score
```

```
Out[74]: True
```

```
In [75]: bag_Knn.fit(x_train, y_train)
print('Bagging KNN score ----->', bag_Knn.score(x_test, y_test))
```

Bagging KNN score -----> 0.7529957479706223

```
In [76]: y_pred = bag_dt.predict(x_test)
```

```
In [77]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

	precision	recall	f1-score	support
0.0	0.83	0.87	0.85	1298
1.0	0.86	0.82	0.84	1289
accuracy			0.85	2587
macro avg	0.85	0.85	0.85	2587
weighted avg	0.85	0.85	0.85	2587

Conclusion : KNN model has 85% score

## Cross Validation score to check if the model is overfitting


```
In [78]: cv = cross_val_score(bag_Knn, x, y, cv = 5)
print('Cross Validation score of Knn model --->', cv.mean())
```

Cross Validation score of Knn model ---> 0.7826227700174205

Conclusion : Knn model has 78% Cross Validation score

## ROC, AUC Curve

```
In [79]: prob = bag_Knn.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



## Random Forest model instantiating, training and evaluating

```
In [90]: bag_Rn = BaggingClassifier(RandomForestClassifier(), n_estimators = 10, max_samp[  
random_state= 3, oob_score = True)
```

```
In [91]: bag_Rn.oob_score
```

```
Out[91]: True
```

```
In [92]: bag_Rn.fit(x_train, y_train)  
print('Bagging Random Forest score ----->', bag_Rn.score(x_test, y_test))
```

```
Bagging Random Forest score -----> 0.8554310011596443
```

```
In [93]: y_pred = bag_Rn.predict(x_test)
```

```
In [94]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

Classification Report:				
	precision	recall	f1-score	support
0.0	0.85	0.86	0.86	1298
1.0	0.86	0.85	0.85	1289
accuracy			0.86	2587
macro avg	0.86	0.86	0.86	2587
weighted avg	0.86	0.86	0.86	2587

```
-----
```

**Conclusion : Random Forest model has 86% score**

### Cross Validation score to check if the model is overfitting


```
In [95]: cv = cross_val_score(bag_Rn, x, y, cv = 5)
print('Cross Validation score of Knn model --->', cv.mean())
```

Cross Validation score of Knn model ---> 0.7968203593780244

**Conclusion : Random Forest model has 79% Cross Validation score**

### ROC, AUC Curve

```
In [96]: prob = bag_Rn.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



## Logistic Regression model instantiating, training and evaluating

```
In [116]: bag_Lr = BaggingClassifier(LogisticRegression(), n_estimators = 5, max_samples =  
random_state= 3, oob_score = True)
```

```
In [117]: bag_Lr.oob_score
```

```
Out[117]: True
```

```
In [118]: bag_Lr.fit(x_train, y_train)
print('Bagging Logistic Regression score ----->', bag_Lr.score(x_test, y_test))
```

Bagging Logistic Regression score -----> 0.7634325473521454

```
In [119]: y_pred = bag_Lr.predict(x_test)
```

```
In [120]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----  
Classification Report:  
precision    recall   f1-score   support  
  
      0.0       0.79      0.72      0.75      1298  
      1.0       0.74      0.80      0.77      1289  
  
accuracy                           0.76      2587  
macro avg       0.77      0.76      0.76      2587  
weighted avg     0.77      0.76      0.76      2587  
-----
```

Conclusion : Logistic Regression model has 76% score

## Cross Validation score to check if the model is overfitting


```
In [121]: cv = cross_val_score(bag_Lr, x, y, cv = 5)
print('Cross Validation score of Logistic regression model --->', cv.mean())
```

Cross Validation score of Logistic regression model ---> 0.8025001613007291

Conclusion : Logistic Regression model has 80% Cross Validation score

## ROC, AUC Curve

```
In [122]: prob = bag_Lr.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



## Let's find ROC, AUC score

```
In [128]: # DecisionTreeClassifier  
roc_auc_score(y_test, bag_dt.predict(x_test))
```

Out[128]: 0.8464618838315436

```
In [129]: # XGBoostClassifier  
roc_auc_score(y_test, bag_xgb.predict(x_test))
```

Out[129]: 0.861188544529329

```
In [130]: # KNeighborsClassifier  
roc_auc_score(y_test, bag_Knn.predict(x_test))
```

Out[130]: 0.753136651122871

```
In [126]: # RandomForestClassifier  
roc_auc_score(y_test, bag_Rn.predict(x_test))
```


Out[126]: 0.8553996660135962

```
In [127]: # LogisticRegressionClassifier  
roc_auc_score(y_test, bag_Lr.predict(x_test))
```

Out[127]: 0.7635749216136062

## Let's check ROC, AUC Curve for the fitted model

```
In [131]: dis = plot_roc_curve(bag_dt, x_test, y_test)  
plot_roc_curve(bag_Rn, x_test, y_test, ax = dis.ax_) # ax_ = Axes with confusion  
plot_roc_curve(bag_Knn, x_test, y_test, ax = dis.ax_)  
plot_roc_curve(bag_xgb, x_test, y_test, ax = dis.ax_)  
plot_roc_curve(bag_Lr, x_test, y_test, ax = dis.ax_)  
plt.legend(prop = {'size':11}, loc = 'lower right')  
plt.show()
```



Looking ROC, AUC Curve we found XGBoost has best model so we do Hyperparameter Tuning on it.

```
In [132]: param = {'n_estimators': [50,100,150], 'max_samples': [1.0], 'bootstrap': [True]}
```

```
In [133]: grid_search = GridSearchCV(estimator = bag_xgb, param_grid = param, cv = 5 , n_jc
```

```
In [134]: grid_search.fit(x_train, y_train)
```

```
ty=None),  
                         max_samples=0.5, n_estimators=30,  
                         oob_score=True, random_state=3),  
                         n_jobs=-1,  
                         param_grid={'bootstrap': [True], 'max_samples': [1.0],  
                         'n_estimators': [50, 100, 150]})
```

In [135]: best\_parameters = grid\_search.best\_params\_  
print(best\_parameters)

```
{'bootstrap': True, 'max_samples': 1.0, 'n_estimators': 50}
```

In [137]: hxgb = BaggingClassifier(base\_estimator=xgb.XGBClassifier(eval\_metric = 'mlogloss'))  
hxgb.fit(x\_train, y\_train)  
hxgb.score(x\_test, y\_test)

Out[137]: 0.860456126787785

In [139]: y\_pred = hxgb.predict(x\_test)

In [140]: print('-----')  
print('\nClassification Report: ')  
print(classification\_report(y\_test, y\_pred, digits = 2))  
print('-----\n')

```
-----  
  
Classification Report:  
precision    recall    f1-score    support  
  
      0.0       0.85       0.87       0.86      1298  
      1.0       0.87       0.85       0.86      1289  
  
accuracy                           0.86      2587  
macro avg       0.86       0.86       0.86      2587  
weighted avg     0.86       0.86       0.86      2587  
  
-----
```

After Hyperparameter Tuning model accuracy score 86%.

## Saving The Model

In [148]: # saving the model to the Local file system  
filename = 'Customer Churn Analysis.pickle'  
pickle.dump(hxgb, open(filename, 'wb'))

## Predict Customer Churn

```
In [150]: model = pickle.load(open('Customer Churn Analysis.pickle', 'rb'))
result = model.score(x_test, y_test)
print('Predicted Score ----->', result)
```

Predicted Score -----> 0.860456126787785

```
In [159]: Prediction = pd.DataFrame([model.predict(x_test)[:,], y_test[:,]], index = ['Prediction'])
```

Out[159]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
Predicted	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	1.0	1.0	0.0	1.0
Orginal	0.0	1.0	1.0	0.0	1.0	1.0	1.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	1.0	1.0	1.0	0.0	1.0

## Saving the predicted result in CSV file

```
In [160]: Prediction.to_csv('Customer Churn Analysis Prediction.csv')
```

**Final Conclusion : XGBoost is our best model.**

In [ ]: