

Problem Statement:

Flight ticket prices can be something hard to guess, today we might see a price, check out the price of the same flight tomorrow, it will be a different story. We might have often heard travellers saying that flight ticket prices are so unpredictable. Here you will be provided with prices of flight tickets for various airlines between the months of March and June of 2019 and between various cities. Build a machine learning model to predict the price of the flight ticket.

Importing Required Library

```
In [50]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pickle
import datetime
pd.set_option('Display.max_columns', None)
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor, BaggingRegressor
import xgboost as xgb
from sklearn.metrics import r2_score, mean_squared_error
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

Reading Data

```
In [2]: df = pd.read_excel(r"C:\Users\Kushal Arya\Desktop\csv file\Flight_Ticket_Participant.xlsx")
df.head()
```

Out[2]:

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total
0	IndiGo	24/03/2019	Banglore	New Delhi	BLR → DEL	22:20	01:10 22 Mar	2h 50m	n/a
1	Air India	1/05/2019	Kolkata	Banglore	CCU → IXR → BBI → BLR	05:50	13:15	7h 25m	
2	Jet Airways	9/06/2019	Delhi	Cochin	DEL → LKO → BOM → COK	09:25	04:25 10 Jun	19h	
3	IndiGo	12/05/2019	Kolkata	Banglore	CCU → NAG → BLR	18:05	23:30	5h 25m	
4	IndiGo	01/03/2019	Banglore	New Delhi	BLR → NAG → DEL	16:50	21:35	4h 45m	

Check no of row and column

```
In [3]: print('No of Rows and Columns ----->', df.shape)
```

No of Rows and Columns -----> (10683, 11)

Checking for Null values

```
In [4]: print('=====\\n')
print(df.isnull().sum())
print('\\n=====')
```

```
=====
```

```
Airline      0
Date_of_Journey  0
Source       0
Destination   0
Route        1
Dep_Time     0
Arrival_Time  0
Duration     0
Total_Stops   1
Additional_Info 0
Price        0
dtype: int64
```

```
=====
```

There is null value

Drop NaN

```
In [3]: df = df.dropna()
```

Check NaN remove or not

```
In [4]: print('=====\\n')
print(df.isnull().sum())
print('\\n=====')
```

```
=====
```

```
Airline      0
Date_of_Journey  0
Source       0
Destination   0
Route        0
Dep_Time     0
Arrival_Time  0
Duration     0
Total_Stops   0
Additional_Info 0
Price        0
dtype: int64
```

```
=====
```

Nan are removed

Information about dataset

```
In [7]: print('=====\\n')
print(df.info())
print('=====')
```

```
=====
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10682 entries, 0 to 10682
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Airline          10682 non-null   object  
 1   Date_of_Journey  10682 non-null   object  
 2   Source           10682 non-null   object  
 3   Destination      10682 non-null   object  
 4   Route            10682 non-null   object  
 5   Dep_Time         10682 non-null   object  
 6   Arrival_Time     10682 non-null   object  
 7   Duration         10682 non-null   object  
 8   Total_Stops      10682 non-null   object  
 9   Additional_Info  10682 non-null   object  
 10  Price            10682 non-null   int64  
dtypes: int64(1), object(10)
memory usage: 1001.4+ KB
None
=====
```

Categorical data present in our data set

Checking For Duplicate

```
In [5]: duplicate = df[df.duplicated()]
print('=====')
print("Duplicate Rows :\n\n",duplicate)
print('=====')
```

=====
Duplicate Rows :

	Airline	Date_of_Journey	Source	Destination	\
683	Jet Airways	1/06/2019	Delhi	Cochin	
1061	Air India	21/05/2019	Delhi	Cochin	
1348	Air India	18/05/2019	Delhi	Cochin	
1418	Jet Airways	6/06/2019	Delhi	Cochin	
1674	IndiGo	24/03/2019	Banglore	New Delhi	
...
10594	Jet Airways	27/06/2019	Delhi	Cochin	
10616	Jet Airways	1/06/2019	Delhi	Cochin	
10634	Jet Airways	6/06/2019	Delhi	Cochin	
10672	Jet Airways	27/06/2019	Delhi	Cochin	
10673	Jet Airways	27/05/2019	Delhi	Cochin	

	Route	Dep_Time	Arrival_Time	Duration	Total_Stops	\
683	DEL → NAG → BOM → COK	14:35	04:25 02 Jun	13h 50m	2 stops	
1061	DEL → GOI → BOM → COK	22:00	19:15 22 May	21h 15m	2 stops	
1348	DEL → HYD → BOM → COK	17:15	19:15 19 May	26h	2 stops	
1418	DEL → JAI → BOM → COK	05:30	04:25 07 Jun	22h 55m	2 stops	
1674	BLR → DEL	18:25	21:20	2h 55m	non-stop	
...
10594	DEL → AMD → BOM → COK	23:05	12:35 28 Jun	13h 30m	2 stops	
10616	DEL → JAI → BOM → COK	09:40	12:35 02 Jun	26h 55m	2 stops	
10634	DEL → JAI → BOM → COK	09:40	12:35 07 Jun	26h 55m	2 stops	
10672	DEL → AMD → BOM → COK	23:05	19:00 28 Jun	19h 55m	2 stops	
10673	DEL → AMD → BOM → COK	13:25	04:25 28 May	15h	2 stops	

	Additional_Info	Price
683	No info	13376
1061	No info	10231
1348	No info	12392
1418	In-flight meal not included	10368
1674	No info	7303
...
10594	No info	12819
10616	No info	13014
10634	In-flight meal not included	11733
10672	In-flight meal not included	11150
10673	No info	16704

[220 rows x 11 columns]

=====

Droping Duplicates

```
In [4]: df.drop_duplicates(keep=False,inplace=True)
```

```
In [5]: print('After removing duplicates No of Rows and Columns ---->', df.shape )
```

```
After removing duplicates No of Rows and Columns ----> (10267, 11)
```

Features Engineering

Add Months Column

```
In [6]: df['Date_of_Journey'] = pd.to_datetime(df['Date_of_Journey'])
```

```
In [7]: df['Months'] = df['Date_of_Journey'].dt.month  
df['Months'] = df['Months'].astype('int')  
df.head(2)
```

```
Out[7]:
```

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	22:20	01:10 22 Mar	2h 50m	nc
1	Air India	2019-01-05	Kolkata	Banglore	CCU → IXR → BBI → BLR	05:50	13:15	7h 25m	2

```
◀ ▶
```

```
In [13]: df['Months'].value_counts()
```

```
Out[13]: 6    2403  
3    2129  
5    1984  
9    1349  
1    1043  
12   935  
4    424  
Name: Months, dtype: int64
```

```
In [14]: df['Months'].dtype
```

```
Out[14]: dtype('int32')
```

Add Day Column

```
In [8]: df['Day'] = df['Date_of_Journey'].dt.day  
df['Day'] = df['Day'].astype('int')  
df.head(2)
```

Out[8]:

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	22:20	01:10 22 Mar	2h 50m	nc
1	Air India	2019-01-05	Kolkata	Banglore	CCU → IXR → BBI → BLR	05:50	13:15	7h 25m	2

◀ ▶

```
In [16]: df['Day'].value_counts()
```

```
Out[16]: 6      2064  
5      1355  
3      1328  
21     1063  
27     1057  
24     996  
15     952  
18     799  
4      653  
Name: Day, dtype: int64
```

```
In [17]: df['Day'].dtype
```

```
Out[17]: dtype('int32')
```

Add New Arrival Column

```
In [18]: df['Arrival_Time'].value_counts()
```

```
Out[18]: 19:00          401  
21:00          360  
19:15          333  
16:10          154  
12:35          122  
...  
18:30 22 May      1  
01:10 16 Mar      1  
04:25 02 May      1  
23:00 22 Apr       1  
01:20 28 Jun       1  
Name: Arrival_Time, Length: 1335, dtype: int64
```

```
In [9]: df['Arrival_Time'] = pd.to_datetime(df['Arrival_Time'])
```

```
In [10]: df['NewArrival_Time'] = df['Arrival_Time'].dt.time  
df.head(2)
```

Out[10]:

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_Stops
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	22:20	2021-03-22 01:10:00	2h 50m	nc
1	Air India	2019-01-05	Kolkata	Banglore	CCU → IXR → BBI → BLR	05:50	2021-08-10 13:15:00	7h 25m	2

```
In [21]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 10267 entries, 0 to 10682  
Data columns (total 14 columns):  
 #   Column           Non-Null Count  Dtype     
---  --    
 0   Airline          10267 non-null   object    
 1   Date_of_Journey 10267 non-null   datetime64[ns]   
 2   Source           10267 non-null   object    
 3   Destination      10267 non-null   object    
 4   Route            10267 non-null   object    
 5   Dep_Time         10267 non-null   object    
 6   Arrival_Time     10267 non-null   datetime64[ns]   
 7   Duration         10267 non-null   object    
 8   Total_Stops      10267 non-null   object    
 9   Additional_Info  10267 non-null   object    
 10  Price            10267 non-null   int64     
 11  Months           10267 non-null   int32     
 12  Day              10267 non-null   int32     
 13  NewArrival_Time 10267 non-null   object    
dtypes: datetime64[ns]/int32/int64/object
```

Add New Departure Time column

```
In [11]: df['Dep_Time'] = pd.to_datetime(df['Dep_Time'])
```

```
In [12]: df['NewDep_Time'] = df['Dep_Time'].dt.time  
df.head(2)
```

Out[12]:

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	2021-08- 10 22:20:00	2021-03-22 01:10:00	2h 50m	nc
1	Air India	2019-01-05	Kolkata	Banglore	CCU → IXR → BBI → BLR	2021-08- 10 05:50:00	2021-08-10 13:15:00	7h 25m	2

```
In [24]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 10267 entries, 0 to 10682  
Data columns (total 15 columns):  
 #  Column            Non-Null Count  Dtype     
---  --  
 0   Airline           10267 non-null   object    
 1   Date_of_Journey  10267 non-null   datetime64[ns]  
 2   Source            10267 non-null   object    
 3   Destination       10267 non-null   object    
 4   Route             10267 non-null   object    
 5   Dep_Time          10267 non-null   datetime64[ns]  
 6   Arrival_Time      10267 non-null   datetime64[ns]  
 7   Duration          10267 non-null   object    
 8   Total_Stops       10267 non-null   object    
 9   Additional_Info   10267 non-null   object    
 10  Price             10267 non-null   int64     
 11  Months            10267 non-null   int32     
 12  Day               10267 non-null   int32     
 13  NewArrival_Time   10267 non-null   object    
 14  NewDep_Time       10267 non-null   object    
dtypes: datetime64[ns](3), int32(2), int64(1), object(9)  
memory usage: 1.2+ MB
```

Converting New Arrival Time and New Departure Time column in Minutes

```
In [13]: df['Dep_Hours'] = df['Dep_Time'].apply(lambda x: x.hour * 3600 + x.minute * 60 + df.head(2)
```

Out[13]:

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	2021-08-10 22:20:00	2021-03-22 01:10:00	2h 50m	nc
1	Air India	2019-01-05	Kolkata	Banglore	CCU → IXR → BBI → BLR	2021-08-10 05:50:00	2021-08-10 13:15:00	7h 25m	2

```
In [14]: df['Arrival_Hours'] = df['NewArrival_Time'].apply(lambda x: x.hour * 3600 + x.minute * 60 + df.head(2)
```

Out[14]:

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	2021-08-10 22:20:00	2021-03-22 01:10:00	2h 50m	nc
1	Air India	2019-01-05	Kolkata	Banglore	CCU → IXR → BBI → BLR	2021-08-10 05:50:00	2021-08-10 13:15:00	7h 25m	2

Split Duration column into Hours and Minutes

```
In [15]: new = df['Duration'].str.split(' ', n = 2, expand = True)
```

```
In [16]: df['Hours'] = new[0]
df['Minutes'] = new[1]
df.head(2)
```

Out[16]:

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	2021-08-10 22:20:00	2021-03-22 01:10:00	2h 50m	nc
1	Air India	2019-01-05	Kolkata	Banglore	CCU → IXR → BBI → BLR	2021-08-10 05:50:00	2021-08-10 13:15:00	7h 25m	2

```
In [17]: df['Hours'] = new[0].str.split('h', expand = True)
df['Minutes'] = new[1].str.split('m', expand = True)
df.head(2)
```

Out[17]:

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	2021-08-10 22:20:00	2021-03-22 01:10:00	2h 50m	nc
1	Air India	2019-01-05	Kolkata	Banglore	CCU → IXR → BBI → BLR	2021-08-10 05:50:00	2021-08-10 13:15:00	7h 25m	2

```
In [18]: df['Hours'] = df['Hours'].replace(['5m'], '5')
```

```
In [19]: df['Hours'] = df['Hours'].astype(np.float)
```

```
In [23]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10267 entries, 0 to 10682
Data columns (total 19 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Airline          10267 non-null   object  
 1   Date_of_Journey  10267 non-null   datetime64[ns]
 2   Source           10267 non-null   object  
 3   Destination      10267 non-null   object  
 4   Route            10267 non-null   object  
 5   Dep_Time         10267 non-null   datetime64[ns]
 6   Arrival_Time     10267 non-null   datetime64[ns]
 7   Duration         10267 non-null   object  
 8   Total_Stops      10267 non-null   object  
 9   Additional_Info  10267 non-null   object  
 10  Price            10267 non-null   int64  
 11  Months           10267 non-null   int32  
 12  Day              10267 non-null   int32  
 13  NewArrival_Time 10267 non-null   object  
 14  NewDep_Time     10267 non-null   object  
 15  Dep_Hours        10267 non-null   float64 
 16  Arrival_Hours   10267 non-null   float64 
 17  Hours            10267 non-null   float64 
 18  Minutes          9286 non-null   object  
dtypes: datetime64[ns](3), float64(3), int32(2), int64(1), object(10)
memory usage: 1.5+ MB
```

```
In [20]: df['Minutes'] = df['Minutes'].fillna(df['Minutes'].mode()[0])
```

```
In [21]: df['Minutes'] = df['Minutes'].astype(np.float)
```

In [22]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10267 entries, 0 to 10682
Data columns (total 19 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Airline          10267 non-null   object  
 1   Date_of_Journey  10267 non-null   datetime64[ns]
 2   Source           10267 non-null   object  
 3   Destination      10267 non-null   object  
 4   Route            10267 non-null   object  
 5   Dep_Time         10267 non-null   datetime64[ns]
 6   Arrival_Time     10267 non-null   datetime64[ns]
 7   Duration         10267 non-null   object  
 8   Total_Stops      10267 non-null   object  
 9   Additional_Info  10267 non-null   object  
 10  Price            10267 non-null   int64  
 11  Months           10267 non-null   int32  
 12  Day              10267 non-null   int32  
 13  NewArrival_Time 10267 non-null   object  
 14  NewDep_Time     10267 non-null   object  
 15  Dep_Hours        10267 non-null   float64 
 16  Arrival_Hours   10267 non-null   float64 
 17  Hours            10267 non-null   float64 
 18  Minutes          10267 non-null   float64 
dtypes: datetime64[ns](3), float64(4), int32(2), int64(1), object(9)
memory usage: 1.5+ MB
```

Adding New Duration column

In [23]: df['New_Duration'] = df['Hours']*60 + df['Minutes']
df.head(2)

Out[23]:

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	2021-08-10 22:20:00	2021-03-22 01:10:00	2h 50m	nc
1	Air India	2019-01-05	Kolkata	Banglore	CCU → IXR → BBI → BLR	2021-08-10 05:50:00	2021-08-10 13:15:00	7h 25m	2

```
In [37]: df['New_Duration'].value_counts()
```

```
Out[37]: 170.0      538
90.0       386
165.0      333
155.0      329
175.0      327
...
2135.0      1
2860.0      1
1810.0      1
2850.0      1
815.0       1
Name: New_Duration, Length: 340, dtype: int64
```

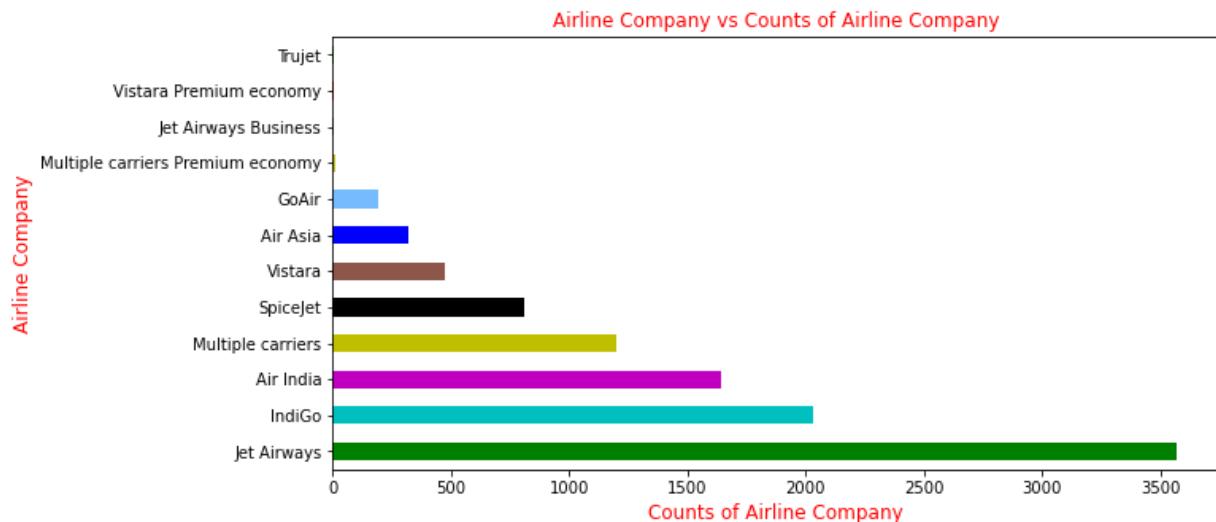
Analysis of Data

Airline column

```
In [38]: ar = df['Airline'].value_counts()
ar
```

```
Out[38]: Jet Airways              3569
IndiGo                     2033
Air India                  1644
Multiple carriers          1196
SpiceJet                   812
Vistara                    477
Air Asia                   319
GoAir                      194
Multiple carriers Premium economy  13
Jet Airways Business        6
Vistara Premium economy    3
Trujet                      1
Name: Airline, dtype: int64
```

```
In [39]: ar.plot.barh(figsize = (10,5), color = ['g','c', 'm', 'y','k','tab:brown','b','xk  
plt.ylabel('Airline Company', c = 'r', fontsize = 12)  
plt.xlabel('Counts of Airline Company', c = 'r', fontsize = 12 )  
plt.title('Airline Company vs Counts of Airline Company', c = 'r', fontsize = 12)  
plt.show()
```

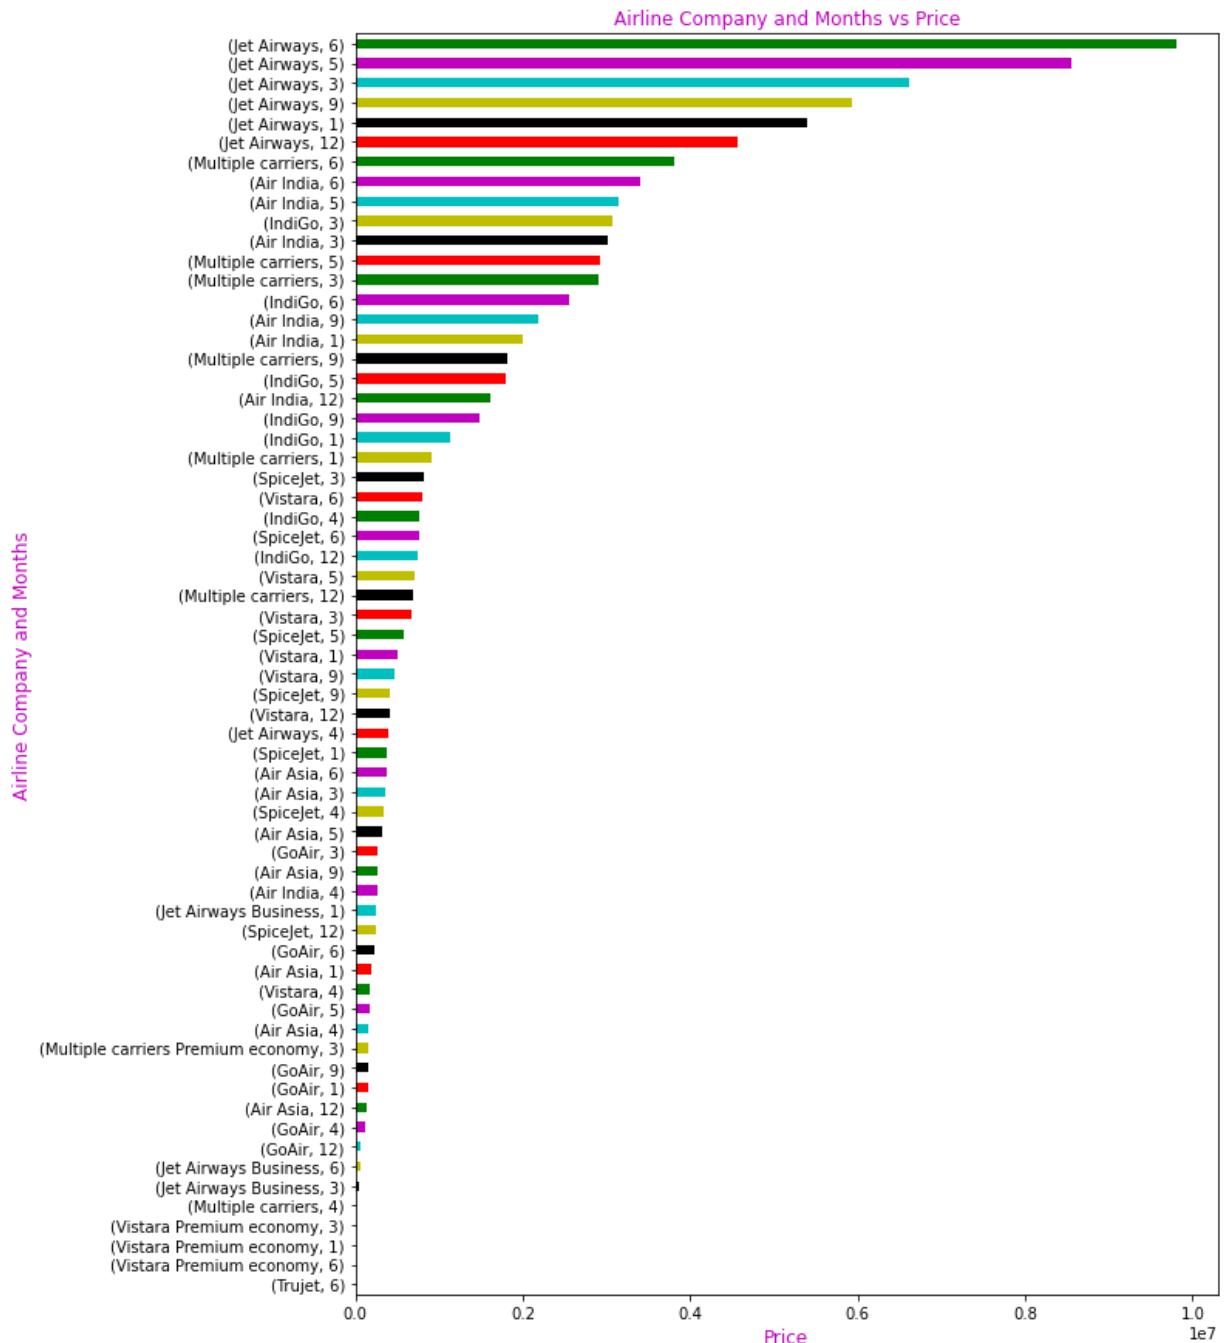


Above plot shows Jet Airways has highest and Trujet has lowest flight counts

```
In [40]: air = df.groupby(['Airline', 'Months'])['Price'].sum().sort_values()  
air
```

```
Out[40]: Airline          Months  
Trujet            6        4140  
Vistara Premium economy    6        5969  
                           1        9125  
                           3       11793  
Multiple carriers      4       12186  
                           ...  
Jet Airways         1       5393640  
                     9       5924281  
                     3       6616650  
                     5       8548865  
                     6       9812957  
Name: Price, Length: 64, dtype: int64
```

```
In [41]: air.plot.barh(figsize = (10, 15),color = ['y','c', 'm', 'g','r','k'])
plt.ylabel('Airline Company and Months', c = 'm', fontsize = 12)
plt.xlabel('Price', c = 'm', fontsize = 12 )
plt.title('Airline Company and Months vs Price', c = 'm', fontsize = 12)
plt.show()
```



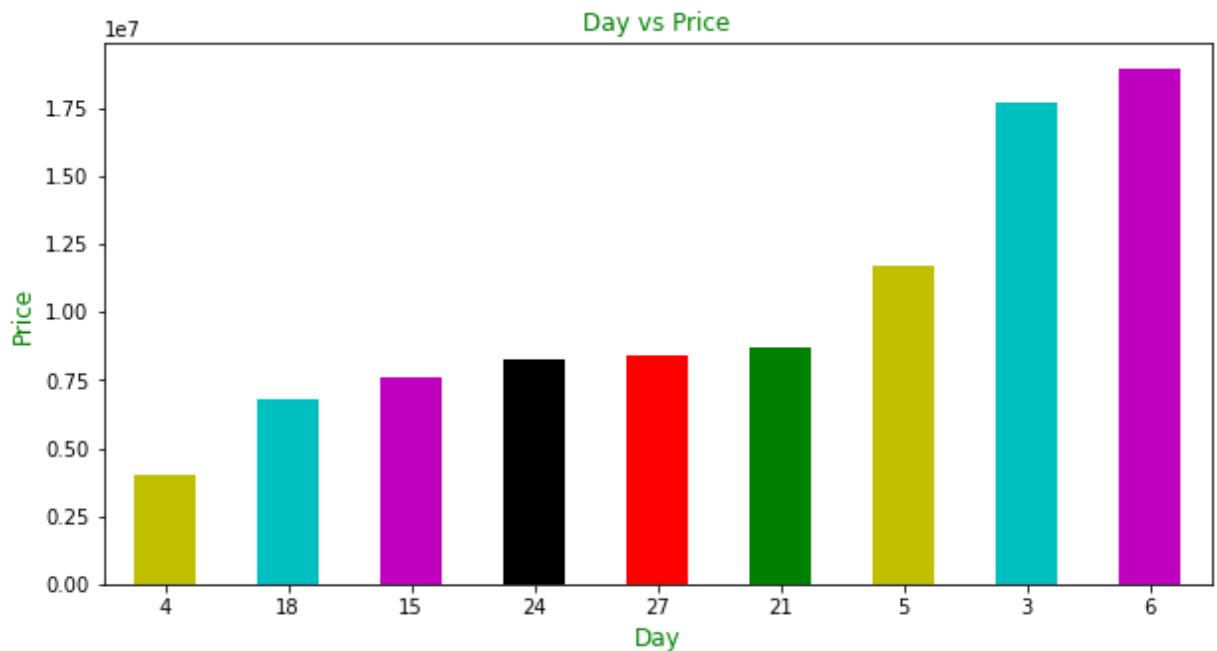
Above plot shows Jet Airways in 6th month has highest Price and Trujet has lowest

Day column

```
In [42]: d = df.groupby(['Day'])['Price'].sum().sort_values()  
d
```

```
Out[42]: Day  
4      3988043  
18     6763741  
15     7618163  
24     8280403  
27     8371368  
21     8727727  
5      11704710  
3      17719024  
6      18950402  
Name: Price, dtype: int64
```

```
In [43]: d.plot.bar(figsize = (10, 5), rot = 360, color = ['y', 'c', 'm', 'k', 'r', 'g'])  
plt.xlabel('Day', c = 'g', fontsize = 12)  
plt.ylabel('Price', c = 'g', fontsize = 12)  
plt.title('Day vs Price', c = 'g', fontsize = 12)  
plt.show()
```

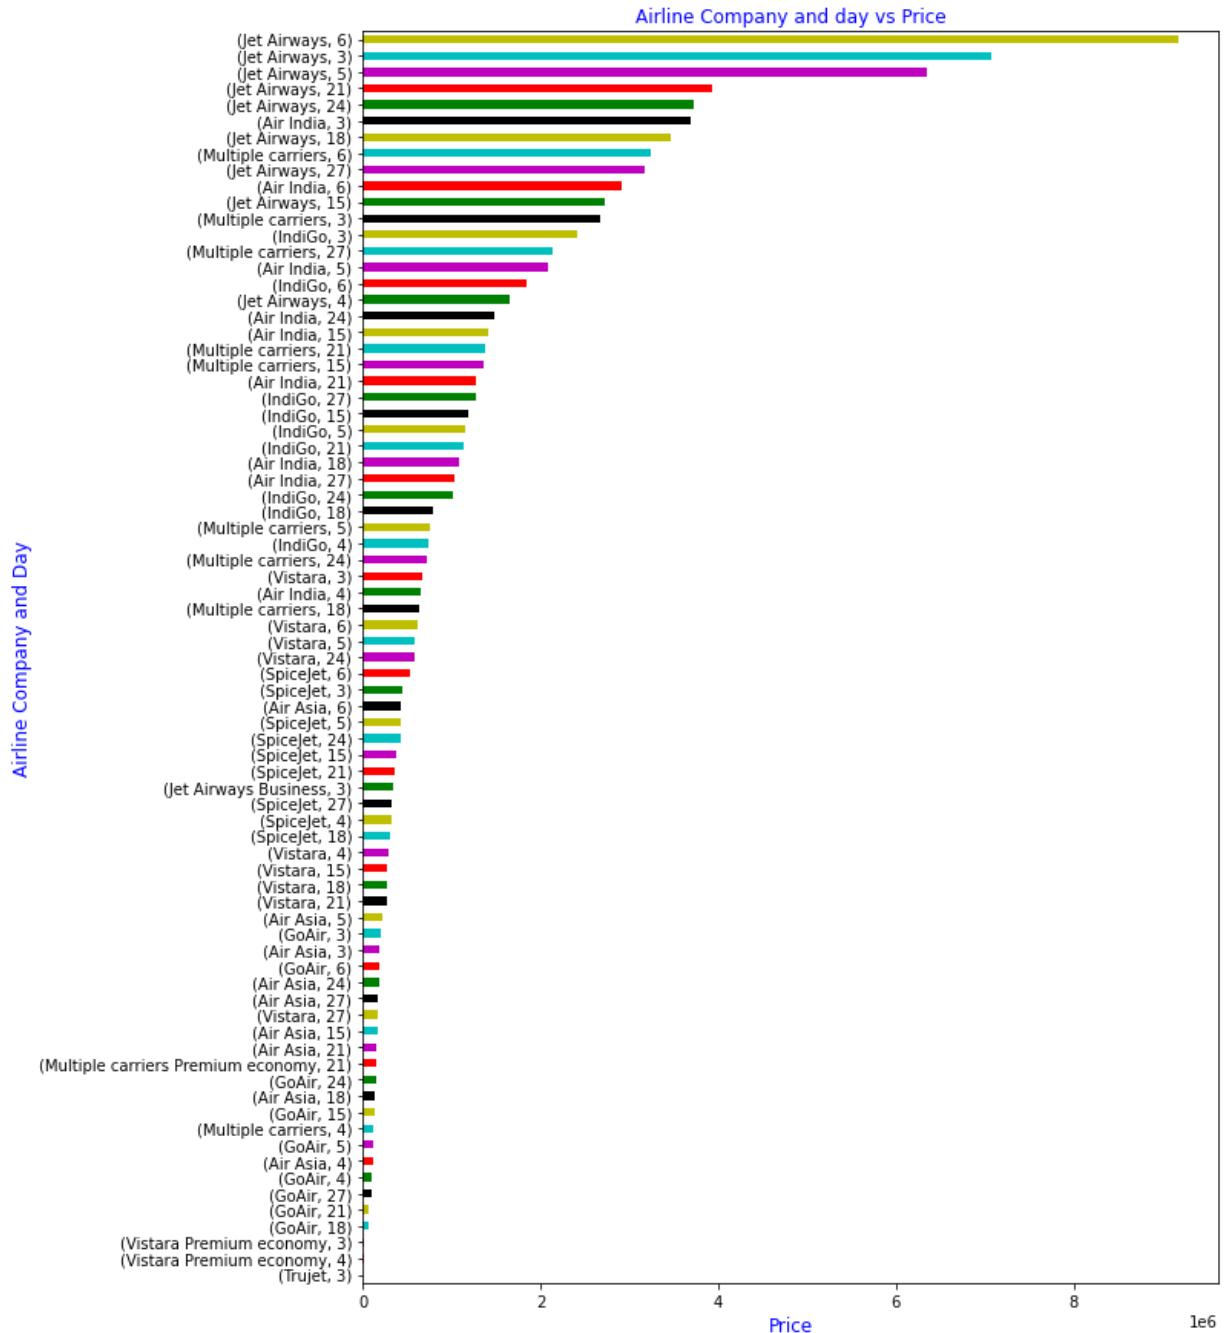


Above plot shows 6th day of every months has highest price and 4th is lowest price

```
In [44]: day = df.groupby(['Airline', 'Day'])['Price'].sum().sort_values()
day
```

```
Out[44]: Airline          Day
Trujet            3      4140
Vistara Premium economy 4      5969
                           3      20918
GoAir             18     63639
                   21     73654
                           ...
Jet Airways       24    3726006
                   21    3925266
                   5     6346158
                   3     7072860
                   6     9171123
Name: Price, Length: 77, dtype: int64
```

```
In [45]: day.plot.barh(figsize = (10, 15),color = ['g','r', 'm', 'c','y','k'])
plt.ylabel('Airline Company and Day', c = 'b', fontsize = 12)
plt.xlabel('Price', c = 'b', fontsize = 12 )
plt.title('Airline Company and day vs Price', c = 'b', fontsize = 12)
plt.show()
```



Above plot shows Jet Airways 6th day of every months has highest price

Route column

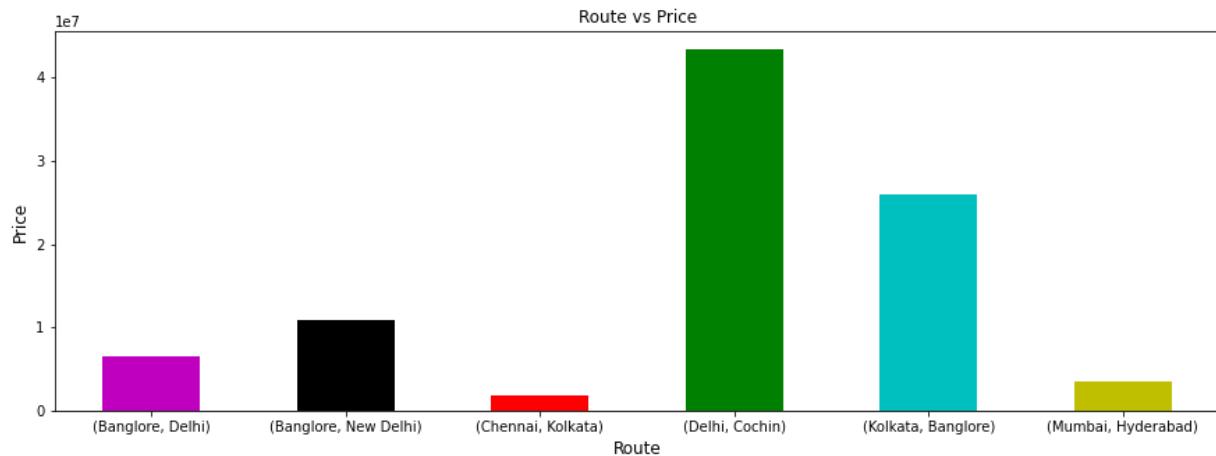
```
In [46]: df['Source'].value_counts()
```

```
Out[46]: Delhi      4178
Kolkata    2850
Banglore   2161
Mumbai     697
Chennai    381
Name: Source, dtype: int64
```

```
In [47]: source = df.groupby(['Source','Destination'])['Price'].sum()
source
```

```
Out[47]: Source    Destination
Banglore   Delhi      6507057
           New Delhi   10842254
Chennai    Kolkata    1824949
Delhi      Cochin    43402838
Kolkata   Banglore   26019866
Mumbai     Hyderabad  3526617
Name: Price, dtype: int64
```

```
In [48]: source.plot.bar(figsize = (15, 5), rot = 360, color = ['m', 'k', 'r', 'g', 'c', 'y'])
plt.ylabel('Price', c = 'k', fontsize = 12)
plt.xlabel('Route', c = 'k', fontsize = 12 )
plt.title('Route vs Price', c = 'k', fontsize = 12)
plt.show()
```



Above plot shows Delhi to Cochin has highest price and Chennai to Kolkata has lowest price

Total Stops column

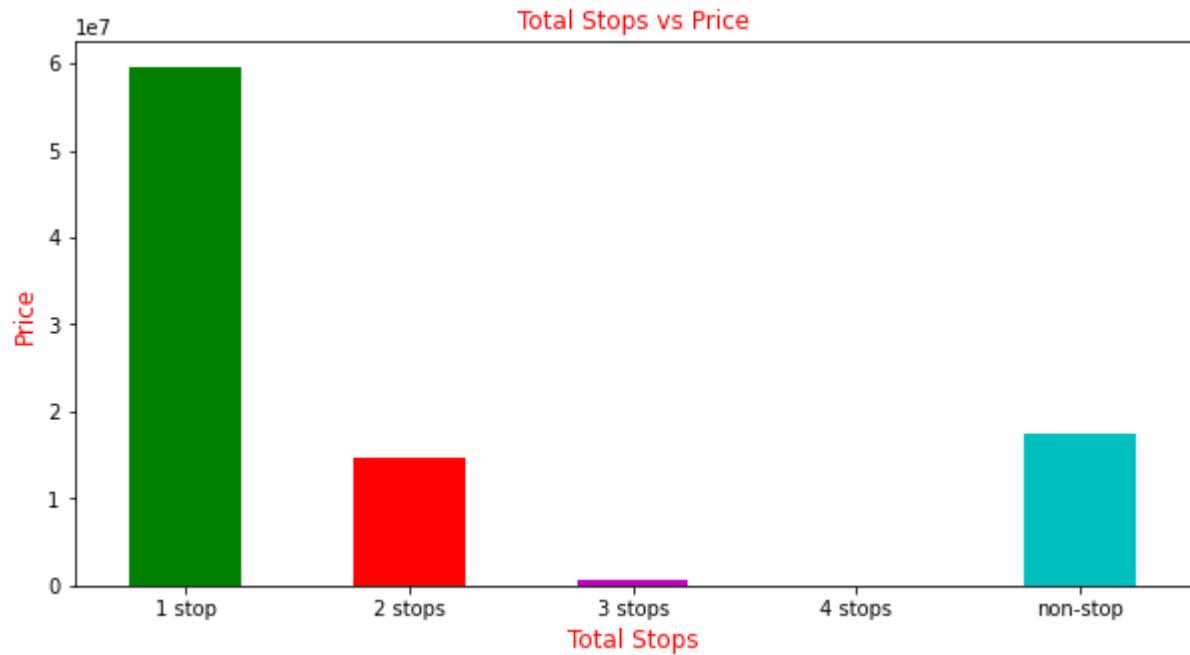
```
In [49]: df['Total_Stops'].value_counts()
```

```
Out[49]: 1 stop      5625
non-stop    3459
2 stops    1141
3 stops     41
4 stops      1
Name: Total_Stops, dtype: int64
```

```
In [50]: tstop = df.groupby('Total_Stops')['Price'].sum()
tstop
```

```
Out[50]: Total_Stops
1 stop      59591945
2 stops    14626877
3 stops     550378
4 stops     17686
non-stop   17336695
Name: Price, dtype: int64
```

```
In [51]: tstop.plot.bar(figsize = (10, 5), rot = 360, color = ['g', 'r', 'm', 'y', 'c'])
plt.xlabel('Total Stops', c = 'r', fontsize = 12)
plt.ylabel('Price', c = 'r', fontsize = 12 )
plt.title('Total Stops vs Price', c = 'r', fontsize = 12)
plt.show()
```



Above plot shows 1 stop flight has highest price and 4 stop has lowest price

New Duration column

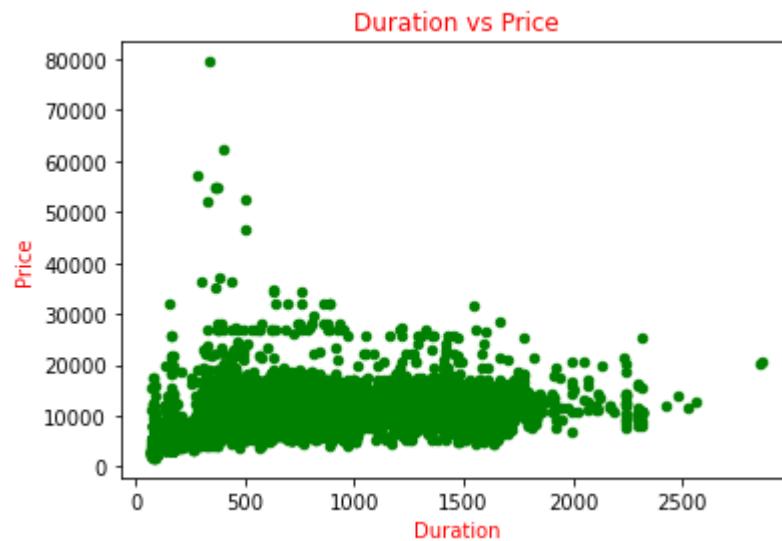
```
In [52]: df['New_Duration'].value_counts()
```

```
Out[52]: 170.0      538
90.0       386
165.0      333
155.0      329
175.0      327
...
2135.0      1
2860.0      1
1810.0      1
2850.0      1
815.0       1
Name: New_Duration, Length: 340, dtype: int64
```

```
In [53]: info = df.groupby('New_Duration')['Price'].sum().sort_values()
info
```

```
Out[53]: New_Duration
250.0      4226
1815.0     7664
1675.0     7932
235.0      8452
820.0      8518
...
155.0      1579634
210.0      1589224
165.0      1843638
175.0      1850293
170.0      2687754
Name: Price, Length: 340, dtype: int64
```

```
In [54]: df.plot.scatter(x = 'New_Duration', y = 'Price', c = 'g')
plt.xlabel('Duration', c = 'r')
plt.ylabel('Price', c = 'r')
plt.title('Duration vs Price', c = 'r')
plt.show()
```



Above plot shows Price increase respect to Duration

Delete Unwanted Columns

```
In [24]: col = ['Date_of_Journey', 'Dep_Time', 'Arrival_Time', 'Duration']
```

```
In [25]: df = df.drop(col, axis = 1)
df.head(2)
```

Out[25]:

	Airline	Source	Destination	Route	Total_Stops	Additional_Info	Price	Months	Day	NewArrival_T
0	IndiGo	Banglore	New Delhi	BLR → DEL	non-stop	No info	3897	3	24	
1	Air India	Kolkata	Banglore	CCU → IXR → BBI → BLR	2 stops	No info	7662	1	5	

```
In [26]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10267 entries, 0 to 10682
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Airline          10267 non-null   object 
 1   Source           10267 non-null   object 
 2   Destination      10267 non-null   object 
 3   Route            10267 non-null   object 
 4   Total_Stops      10267 non-null   object 
 5   Additional_Info  10267 non-null   object 
 6   Price            10267 non-null   int64  
 7   Months           10267 non-null   int32  
 8   Day              10267 non-null   int32  
 9   NewArrival_Time  10267 non-null   object 
 10  NewDep_Time     10267 non-null   object 
 11  Dep_Hours       10267 non-null   float64
 12  Arrival_Hours   10267 non-null   float64
 13  Hours           10267 non-null   float64
 14  Minutes          10267 non-null   float64
 15  New_Duration    10267 non-null   float64
dtypes: float64(5), int32(2), int64(1), object(8)
memory usage: 1.3+ MB
```

Encoding Categorical columns

```
In [27]: le = LabelEncoder()
```

```
In [28]: df['Airline'] = le.fit_transform(df['Airline'])
df['Source'] = le.fit_transform(df['Source'])
df['Destination'] = le.fit_transform(df['Destination'])
df['Route'] = le.fit_transform(df['Route'])
df['Additional_Info'] = le.fit_transform(df['Additional_Info'])
df['NewArrival_Time'] = le.fit_transform(df['NewArrival_Time'])
df['NewDep_Time'] = le.fit_transform(df['NewDep_Time'])
```

```
In [30]: df['Total_Stops'] = le.fit_transform(df['Total_Stops'])
```

```
In [31]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10267 entries, 0 to 10682
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Airline          10267 non-null   int32  
 1   Source           10267 non-null   int32  
 2   Destination      10267 non-null   int32  
 3   Route            10267 non-null   int32  
 4   Total_Stops      10267 non-null   int32  
 5   Additional_Info  10267 non-null   int32  
 6   Price            10267 non-null   int64  
 7   Months           10267 non-null   int32  
 8   Day              10267 non-null   int32  
 9   NewArrival_Time  10267 non-null   int32  
 10  NewDep_Time     10267 non-null   int32  
 11  Dep_Hours       10267 non-null   float64 
 12  Arrival_Hours   10267 non-null   float64 
 13  Hours           10267 non-null   float64 
 14  Minutes          10267 non-null   float64 
 15  New_Duration    10267 non-null   float64 
dtypes: float64(5), int32(10), int64(1)
memory usage: 962.5 KB
```

Columns are encoded

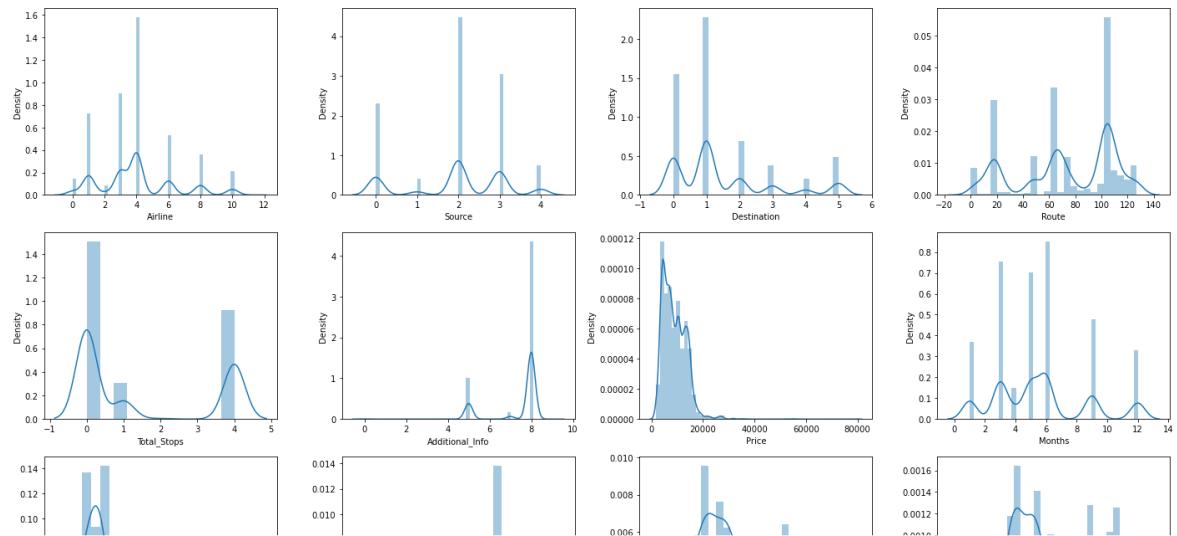
Data distribution

```
In [32]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,15))
plotnumber = 1

for column in df:
    if plotnumber <=16:
        ax = plt.subplot(4,4, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column)
    plotnumber +=1
plt.tight_layout()
```

Distribution Plot :-



Data distribution is fine

Splitting Dataset into features and label

```
In [33]: x = df.drop('Price', axis = 1)
y = df['Price']
print('Data has been splited')
```

Data has been splited

Data Scaling

```
In [34]: scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
Out[34]: array([[-4.18613688e-01, -1.63729319e+00,  2.39487913e+00, ...,
                 -9.47246060e-01,  1.34674268e+00, -9.12290527e-01],
                 [-1.26043266e+00,  8.74799840e-01, -9.70420339e-01, ...,
                  -3.42577058e-01, -4.22910364e-01, -3.55942283e-01],
                 [ 2.29579537e-03,  3.74354957e-02, -2.97360444e-01, ...,
                  1.10862855e+00, -6.89797554e-02,  1.11079400e+00],
                 ...,
                 [ 2.29579537e-03, -1.63729319e+00,  3.75699450e-01, ...,
                  -8.26312260e-01, -6.89797554e-02, -8.31367146e-01],
                 [ 2.52775270e+00, -1.63729319e+00,  2.39487913e+00, ...,
                  -9.47246060e-01,  6.38881463e-01, -9.32521372e-01],
                 [-1.26043266e+00,  3.74354957e-02, -2.97360444e-01, ...,
                  -2.21643257e-01, -7.76840974e-01, -2.44672634e-01]])
```

Data has been scaled

Split data into train and test. Model will be bulit on training data and tested on test data

```
In [35]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25, random_state=3)
print('Data has been splited.')
```

Data has been splited.

Model Bulding

Decision Tree model instantiaing, training and evaluating

```
In [69]: bag_dt = BaggingRegressor(DecisionTreeRegressor(), n_estimators = 50, max_samples = 20, max_features = 3, random_state=3, oob_score = True)
```

```
In [70]: bag_dt.oob_score
```

```
Out[70]: True
```

```
In [71]: bag_dt.fit(x_train, y_train)
print('Bagging DT score ----->', bag_dt.score(x_test, y_test))
```

```
Bagging DT score -----> 0.8569681026273173
```

```
In [72]: y_pred = bag_dt.predict(x_test)
```

```
In [73]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', bag_dt.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.8569681026273173  
=====  
RMSE of Model -----> 1774.210733289188  
=====  
MSE of Model -----> 3147823.726118558  
=====  
Score of test data ----> 0.8569681026273173  
=====
```

Conclusion : Decision Tree model has 85% score

XGBoost model instantiaing, training and evaluating

```
In [74]: bag_xgb = BaggingRegressor(xgb.XGBRegressor(eval_metric = 'mlogloss'), n_estimators=100,  
                                 random_state= 3, oob_score = True)
```

```
In [75]: bag_xgb.oob_score
```

```
Out[75]: True
```

```
In [76]: bag_xgb.fit(x_train, y_train)  
print('Bagging XGBoost score ----->', bag_xgb.score(x_test, y_test))
```

```
Bagging XGBoost score -----> 0.8778938096639761
```

```
In [77]: y_pred = bag_xgb.predict(x_test)
```

```
In [78]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', bag_xgb.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.8778938096639761  
=====  
RMSE of Model -----> 1639.2968077693845  
=====  
MSE of Model -----> 2687294.0239628945  
=====  
Score of test data ----> 0.8778938096639761  
=====
```

Conclusion : XGBoost model has 87% score

Knn model instantiaing, training and evaluating

```
In [79]: bag_Knn = BaggingRegressor(KNeighborsRegressor(n_neighbors = 5), n_estimators = 3,  
                                 random_state= 3, oob_score = True)
```

```
In [80]: bag_Knn.oob_score
```

```
Out[80]: True
```

```
In [81]: bag_Knn.fit(x_train, y_train)  
print('Bagging KNN score ----->', bag_Knn.score(x_test, y_test))
```

```
Bagging KNN score -----> 0.5529532461895169
```

```
In [82]: y_pred = bag_dt.predict(x_test)
```

```
In [83]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', bag_Knn.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.8569681026273173  
=====  
RMSE of Model -----> 1774.210733289188  
=====  
MSE of Model -----> 3147823.726118558  
=====  
Score of test data ----> 0.5529532461895169  
=====
```

Conclusion : KNN model has 85% score

Random Forest model instantiaing, training and evaluating

```
In [84]: bag_Rn = BaggingRegressor(RandomForestRegressor(), n_estimators = 50, max_samples  
random_state= 3, oob_score = True)
```

```
In [85]: bag_Rn.oob_score
```

```
Out[85]: True
```

```
In [86]: bag_Rn.fit(x_train, y_train)  
print('Bagging Random Forest score ----->', bag_Rn.score(x_test, y_test))
```

```
Bagging Random Forest score -----> 0.8491420044591111
```

```
In [87]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', bag_Rn.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.8569681026273173  
=====  
RMSE of Model -----> 1774.210733289188  
=====  
MSE of Model -----> 3147823.726118558  
=====  
Score of test data ----> 0.8491420044591111  
=====
```

Conclusion : Random Forest model has 85% score

Looking RMSE score we found XGBoost has best model so we do Hyperparameter Tuning on it.

```
In [91]: param = {'n_estimators': [50,100], 'max_samples': [1.0], 'bootstrap': [True]}
```

```
In [92]: grid_search = GridSearchCV(estimator = bag_xgb, param_grid = param, cv = 5 , n_jc
```

```
In [93]: grid_search.fit(x_train, y_train)
```

```
=None),  
                         max_samples=0.5, n_estimators=30,  
                         oob_score=True, random_state=3),  
                         n_jobs=-1,  
                         param_grid={'bootstrap': [True], 'max_samples': [1.0],  
                         'n_estimators': [50, 100]})
```

In [94]: best_parameters = grid_search.best_params_
print(best_parameters)

```
{'bootstrap': True, 'max_samples': 1.0, 'n_estimators': 100}
```

In [95]: hxgb = BaggingRegressor(base_estimator=xgb.XGBRegressor(), max_samples = 1.0, bootstrap=True)
hxgb.fit(x_train, y_train)
hxgb.score(x_test, y_test)

Out[95]: 0.8805171134539985

In [96]: y_pred = hxgb.predict(x_test)

In [97]: print('=====')
print('R2 Score ----->', r2_score(y_test, y_pred))
print('=====')
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))
print('=====')
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))
print('=====')
print('Score of test data ----->', hxgb.score(x_test, y_test))
print('=====')

```
=====  
R2 Score -----> 0.8805171134539985  
=====  
RMSE of Model -----> 1621.592045682996  
=====  
MSE of Model -----> 2629560.762622364  
=====  
Score of test data -----> 0.8805171134539985  
=====
```

Saving The Model

In [98]: # saving the model to the Local file system
filename = 'Flight Price Prediction Train.pickle'
pickle.dump(hxgb, open(filename, 'wb'))

Final Conclusion : XGBoost is our best model.

In []:

