

## Problem Statement:

The Global Power Plant Database is a comprehensive, open source database of power plants around the world. It centralizes power plant data to make it easier to navigate, compare and draw insights for one's own analysis. The database covers approximately 35,000 power plants from 167 countries and includes thermal plants (e.g. coal, gas, oil, nuclear, biomass, waste, geothermal) and renewables (e.g. hydro, wind, solar). Each power plant is geolocated and entries contain information on plant capacity, generation, ownership, and fuel type. It will be continuously updated as data becomes available.

## Importing Required Library

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pickle
from scipy.stats import zscore
import scikitplot as skplt
pd.set_option('display.max_columns', None) # # For display maximum column
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score, mean_squared_error
import xgboost as xgb
%matplotlib inline

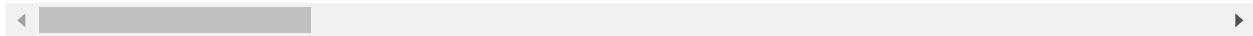
import warnings
warnings.filterwarnings('ignore')
```

## Reading Data

```
In [2]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\csv file\database_IND.csv")
df.head()
```

Out[2]:

	country	country_long		name	gppd_idnr	capacity_mw	latitude	longitude	primary_fuel
0	IND	India		ACME Solar Tower	WRI1020239	2.5	28.1839	73.2407	Solar
1	IND	India		ADITYA CEMENT WORKS	WRI1019881	98.0	24.7663	74.6090	Coal
2	IND	India	AES Saurashtra Windfarms		WRI1026669	39.2	21.9038	69.3732	Wind
3	IND	India	AGARTALA GT		IND0000001	135.0	23.8712	91.3602	Gas
4	IND	India	AKALTARA TPP		IND0000002	1800.0	21.9603	82.4091	Coal



## Check no of row and column

```
In [3]: print('No of Rows and Columns ----->', df.shape )
```

No of Rows and Columns -----> (908, 25)

## Checking for Null values

```
In [4]: print('-----\n')
print(df.isnull().sum())
print('\n-----')
```

```
-----  
country 0  
country_long 0  
name 0  
gppd_idnr 0  
capacity_mw 0  
latitude 46  
longitude 46  
primary_fuel 0  
other_fuel1 709  
other_fuel2 907  
other_fuel3 908  
commissioning_year 380  
owner 566  
source 0  
url 0  
geolocation_source 19  
wepp_id 908  
year_of_capacity_data 388  
generation_gwh_2013 524  
generation_gwh_2014 507  
generation_gwh_2015 483  
generation_gwh_2016 471  
generation_gwh_2017 465  
generation_data_source 458  
estimated_generation_gwh 908  
dtype: int64
```

-----

There is null value

Information about dataset

```
In [5]: print('-----\n')
print(df.info())
print('-----')
```

```
-----  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 908 entries, 0 to 907  
Data columns (total 25 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   country          908 non-null    object    
 1   country_long     908 non-null    object    
 2   name              908 non-null    object    
 3   gppd_idnr        908 non-null    object    
 4   capacity_mw      908 non-null    float64  
 5   latitude          862 non-null    float64  
 6   longitude         862 non-null    float64  
 7   primary_fuel      908 non-null    object    
 8   other_fuel1       199 non-null    object    
 9   other_fuel2       1 non-null     object    
 10  other_fuel3      0 non-null     float64  
 11  commissioning_year 528 non-null    float64  
 12  owner              342 non-null    object    
 13  source             908 non-null    object    
 14  url                908 non-null    object    
 15  geolocation_source 889 non-null    object    
 16  wepp_id            0 non-null     float64  
 17  year_of_capacity_data 520 non-null    float64  
 18  generation_gwh_2013 384 non-null    float64  
 19  generation_gwh_2014 401 non-null    float64  
 20  generation_gwh_2015 425 non-null    float64  
 21  generation_gwh_2016 437 non-null    float64  
 22  generation_gwh_2017 443 non-null    float64  
 23  generation_data_source 450 non-null    object    
 24  estimated_generation_gwh 0 non-null     float64  
dtypes: float64(13), object(12)  
memory usage: 177.5+ KB  
None  
-----
```

Categorical data present in our data set

## Statistics of Data

In [6]: df.describe()

Out[6]:

	capacity_mw	latitude	longitude	other_fuel3	commissioning_year	wepp_id	year_of_c
count	908.000000	862.000000	862.000000	0.0	528.000000	0.0	
mean	321.046378	21.196189	77.447848	NaN	1996.876894	NaN	
std	580.221767	6.248627	4.907260	NaN	17.047817	NaN	
min	0.000000	8.168900	68.644700	NaN	1927.000000	NaN	
25%	16.837500	16.771575	74.258975	NaN	1988.000000	NaN	
50%	60.000000	21.778300	76.719250	NaN	2000.000000	NaN	
75%	388.125000	25.516375	79.441475	NaN	2011.250000	NaN	
max	4760.000000	34.649000	95.408000	NaN	2018.000000	NaN	

Outliers are present in our data set

## Analysis of Null value

```
In [7]: print('-----\n')
print(df.isnull().sum())
print('\n-----')
```

```
-----  
country 0  
country_long 0  
name 0  
gppd_idnr 0  
capacity_mw 0  
latitude 46  
longitude 46  
primary_fuel 0  
other_fuel1 709  
other_fuel2 907  
other_fuel3 908  
commissioning_year 380  
owner 566  
source 0  
url 0  
geolocation_source 19  
wepp_id 908  
year_of_capacity_data 388  
generation_gwh_2013 524  
generation_gwh_2014 507  
generation_gwh_2015 483  
generation_gwh_2016 471  
generation_gwh_2017 465  
generation_data_source 458  
estimated_generation_gwh 908  
dtype: int64
```

```
In [8]: df['other_fuel1'].value_counts()
```

```
Out[8]: Oil      196
Gas      2
Cogeneration      1
Name: other_fuel1, dtype: int64
```

```
In [9]: df['other_fuel2'].value_counts()
```

```
Out[9]: Oil      1
Name: other_fuel2, dtype: int64
```

```
In [10]: df['other_fuel3'].value_counts()
```

```
Out[10]: Series([], Name: other_fuel3, dtype: int64)
```

Approach : We drop above other\_fuel2 and other\_fuel3 column because maximum null value and keep other\_fuel1 and fill null value with mode

```
In [11]: df['latitude'].value_counts()
```

```
Out[11]: 24.1917      3  
19.0004      3  
16.5697      2  
23.4639      2  
13.2450      2  
..  
20.9099      1  
17.2387      1  
23.5594      1  
27.3426      1  
16.5973      1  
Name: latitude, Length: 837, dtype: int64
```

```
In [12]: df['longitude'].value_counts()
```

```
Out[12]: 71.6917      4  
71.6918      3  
75.8988      3  
72.8983      3  
81.2875      3  
..  
80.1264      1  
76.1137      1  
74.6447      1  
86.0970      1  
79.5748      1  
Name: longitude, Length: 828, dtype: int64
```

**Approach : We fill above longitude and latitude columns with mean**

```
In [13]: df['commissioning_year'].value_counts()
```

```
Out[13]: 2013.0      28  
2015.0      26  
2012.0      23  
2016.0      21  
2014.0      17  
..  
1958.0      1  
1949.0      1  
1954.0      1  
1956.0      1  
1927.0      1  
Name: commissioning_year, Length: 73, dtype: int64
```

**Approach: We fill commissioning\_year column with mode**

```
In [14]: df['owner'].value_counts()
```

```
Out[14]: Acc Acc ltd          4  
Sterling Agro Industries ltd. 4  
Jk Cement ltd                4  
Tata Power Solar Systems Limited (TPREL) 3  
Ujaas Energy Limited         3  
..  
Gm Energy ltd                1  
Dcm & chem                  1  
REI Agro Limited              1  
National And paper            1  
Spr Pvt ltd                  1  
Name: owner, Length: 280, dtype: int64
```

**Approach : We fill owner column with mode**

```
In [15]: df['geolocation_source'].value_counts()
```

```
Out[15]: WRI                   766  
Industry About                119  
National Renewable Energy Laboratory 4  
Name: geolocation_source, dtype: int64
```

**Approach : We fill geolocation\_source column with mode**

```
In [16]: df['wepp_id'].value_counts()
```

```
Out[16]: Series([], Name: wepp_id, dtype: int64)
```

**Approach : We drop wepp\_id column because maximum null value**

```
In [17]: df['year_of_capacity_data'].value_counts()
```

```
Out[17]: 2018.0      520  
Name: year_of_capacity_data, dtype: int64
```

**Approach : We fill year\_of\_capacity\_data column with bfill beacuse only one data in it**

```
In [18]: df['generation_gwh_2013'].value_counts()
```

```
Out[18]: 0.00000    21
14881.88000     1
42.49645       1
2036.00000     1
97.73885       1
..
7229.33000     1
657.21740       1
507.89775       1
8556.42400     1
8211.00000     1
Name: generation_gwh_2013, Length: 364, dtype: int64
```

```
In [19]: df['generation_gwh_2014'].value_counts()
```

```
Out[19]: 0.00000    28
6803.31250     1
4735.13000     1
145.81400      1
2022.57000     1
..
6224.00000     1
268.48085      1
1255.73200     1
164.32425      1
1153.65300     1
Name: generation_gwh_2014, Length: 374, dtype: int64
```

```
In [20]: df['generation_gwh_2015'].value_counts()
```

```
Out[20]: 0.00000    28
5837.76600     1
1297.97750     1
8076.81050     1
1.09395        1
..
2636.86400     1
665.19730      1
1516.36010     1
741.86205      1
7130.50700     1
Name: generation_gwh_2015, Length: 398, dtype: int64
```

```
In [21]: df['generation_gwh_2016'].value_counts()
```

```
Out[21]: 0.00000    31
8470.57000     2
1511.00000     2
7.31325        1
94.85500        1
...
433.84800      1
283.74811      1
259.94375      1
403.96000      1
307.87290      1
Name: generation_gwh_2016, Length: 405, dtype: int64
```

```
In [22]: df['generation_gwh_2017'].value_counts()
```

```
Out[22]: 0.00000    33
170.08530      2
344.35955      1
2265.47000      1
59.43135      1
...
214.48220      1
272.73945      1
2887.00000      1
12.73600      1
158.73235      1
Name: generation_gwh_2017, Length: 410, dtype: int64
```

**Approach : We fill above generation\_gwh columns with mean**

```
In [23]: df['generation_data_source'].value_counts()
```

```
Out[23]: Central Electricity Authority    450
Name: generation_data_source, dtype: int64
```

**Approach : We fill above generation\_data\_source columns with bfill because only one data in it**

```
In [24]: df['estimated_generation_gwh'].value_counts()
```

```
Out[24]: Series([], Name: estimated_generation_gwh, dtype: int64)
```

**Approach : We drop estimated\_generation\_gwh columns because no data in it**

**Drop Unwanted column**

```
In [7]: col = ['estimated_generation_gwh', 'wepp_id', 'other_fuel2', 'other_fuel3', 'url']
```

```
In [8]: df = df.drop(col, axis = 1)
df.head(2)
```

Out[8]:

		country_long	name	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commission
0		India	ACME Solar Tower	2.5	28.1839	73.2407	Solar		NaN
1		India	ADITYA CEMENT WORKS	98.0	24.7663	74.6090	Coal		NaN

```
In [9]: print('After droping no of Rows and Columns ---->', df.shape )
```

After droping no of Rows and Columns ----> (908, 18)

## Fill NaN

```
In [10]: df = df.apply(lambda x:x.fillna(x.mean()))if x.dtype == 'float' else x.fillna(x.v
```

```
In [11]: print('-----\n')
print(df.isnull().sum())
print('\n-----')
```

```
-----
country_long          0
name                  0
capacity_mw           0
latitude              0
longitude             0
primary_fuel          0
other_fuel1           0
commissioning_year    0
owner                 0
source                0
geolocation_source    0
year_of_capacity_data 0
generation_gwh_2013    0
generation_gwh_2014    0
generation_gwh_2015    0
generation_gwh_2016    0
generation_gwh_2017    0
generation_data_source 0
dtype: int64
```

There is no null value left

## Analysis of data respect to Capacity MW

```
In [12]: df.head(2)
```

Out[12]:

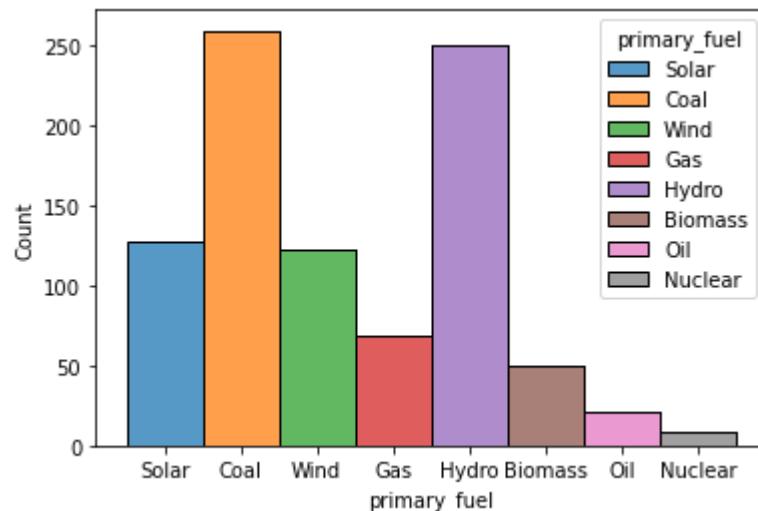
		country_long	name	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commission
0	India	ACME Solar Tower		2.5	28.1839	73.2407	Solar	Oil	201
1	India	ADITYA CEMENT WORKS		98.0	24.7663	74.6090	Coal	Oil	199

## Fule Type column

```
In [31]: df['primary_fuel'].value_counts()
```

```
Out[31]: Coal      259  
Hydro     250  
Solar     127  
Wind      123  
Gas       69  
Biomass    50  
Oil        21  
Nuclear     9  
Name: primary_fuel, dtype: int64
```

```
In [32]: sns.histplot(binwidth=0.5, x="primary_fuel", hue="primary_fuel", data=df, stat="count")  
plt.show()
```

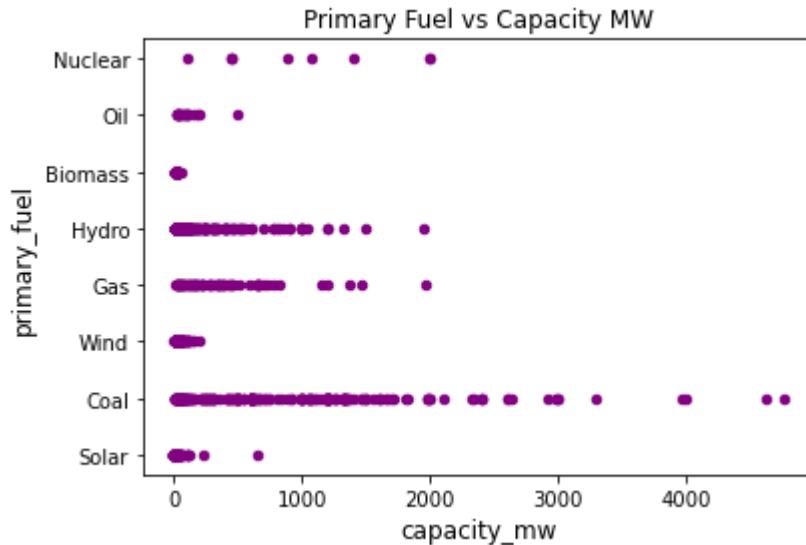


Cole type fuel used highest and Nuclear type fuel used least

```
In [13]: df['capacity_mw'].value_counts()
```

```
Out[13]: 5.0      39  
10.0     22  
600.0    21  
15.0     20  
1200.0   19  
..  
26.4     1  
68.8     1  
91.8     1  
1.8      1  
816.4    1  
Name: capacity_mw, Length: 365, dtype: int64
```

```
In [19]: df.plot.scatter(x = 'capacity_mw', y = 'primary_fuel', c = 'purple')
plt.xlabel('capacity_mw', fontsize = 12, c = 'black')
plt.ylabel('primary_fuel', fontsize = 12, c = 'black')
plt.title('Primary Fuel vs Capacity MW', fontsize = 12, c = 'black')
plt.show()
```



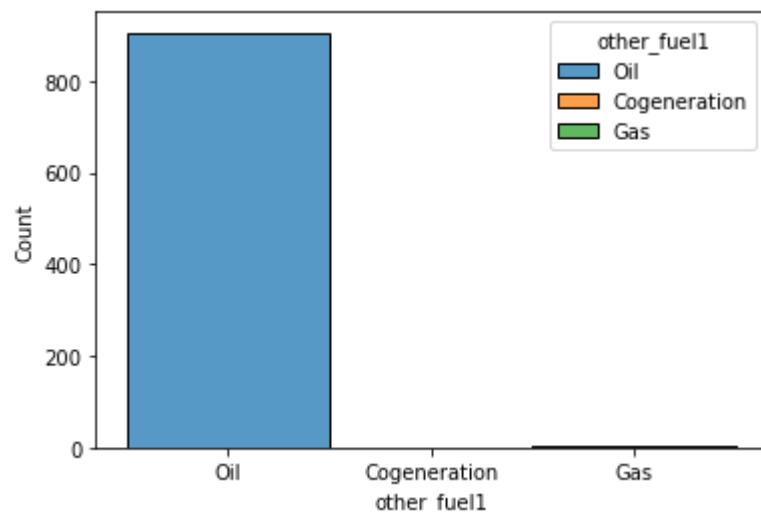
Above plot shows Coal fuel give higest energy and Biomass give lowest energy

## Other Fuel Type column

```
In [33]: df['other_fuel1'].value_counts()
```

```
Out[33]: Oil           905
          Gas            2
          Cogeneration    1
          Name: other_fuel1, dtype: int64
```

```
In [34]: sns.histplot(binwidth=0.5, x="other_fuel1", hue="other_fuel1", data=df, stat="count")
plt.show()
```



Oil fuel used most of all fuel

## Commission Year column

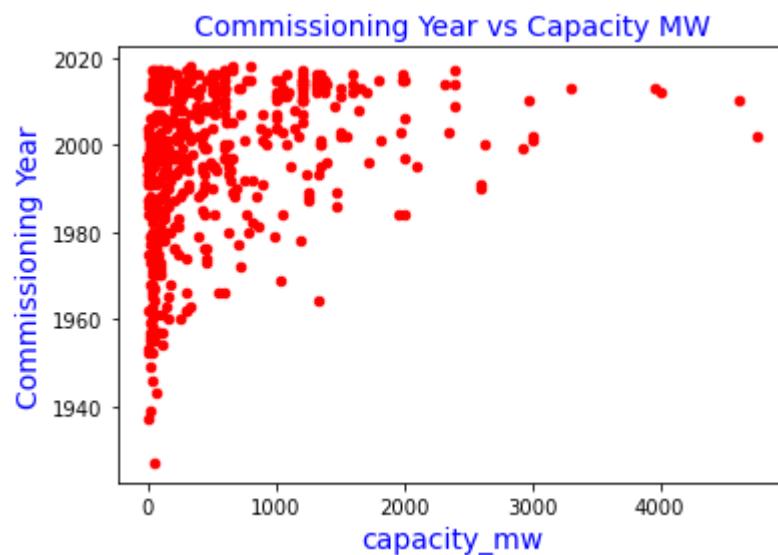
```
In [35]: df['commissioning_year'].value_counts()
```

```
Out[35]: 1996.876894      380
2013.000000       28
2015.000000       26
2012.000000       23
2016.000000       21
...
1949.000000        1
1958.000000        1
1954.000000        1
1956.000000        1
1927.000000        1
Name: commissioning_year, Length: 74, dtype: int64
```

```
In [23]: df.groupby('commissioning_year')['capacity_mw'].sum().sort_values()
```

```
Out[23]: commissioning_year
1953.0      4.00
1937.0      5.00
1949.0     9.30
1956.0    10.00
1959.0    15.00
...
2010.0  16198.00
2012.0  16801.13
2014.0  17468.00
2013.0  21953.48
2015.0  23185.50
Name: capacity_mw, Length: 74, dtype: float64
```

```
In [20]: df.plot.scatter(x = 'capacity_mw', y= 'commissioning_year', c = 'r')
plt.ylabel('Commissioning Year', fontsize = 14, color = 'b')
plt.xlabel('capacity_mw', fontsize = 14, color = 'b')
plt.title('Commissioning Year vs Capacity MW', fontsize = 14, color = 'b')
plt.show()
```



Above plot shows early years production of energy less than 1000 MW but Now a days it produce more than 4000 MW

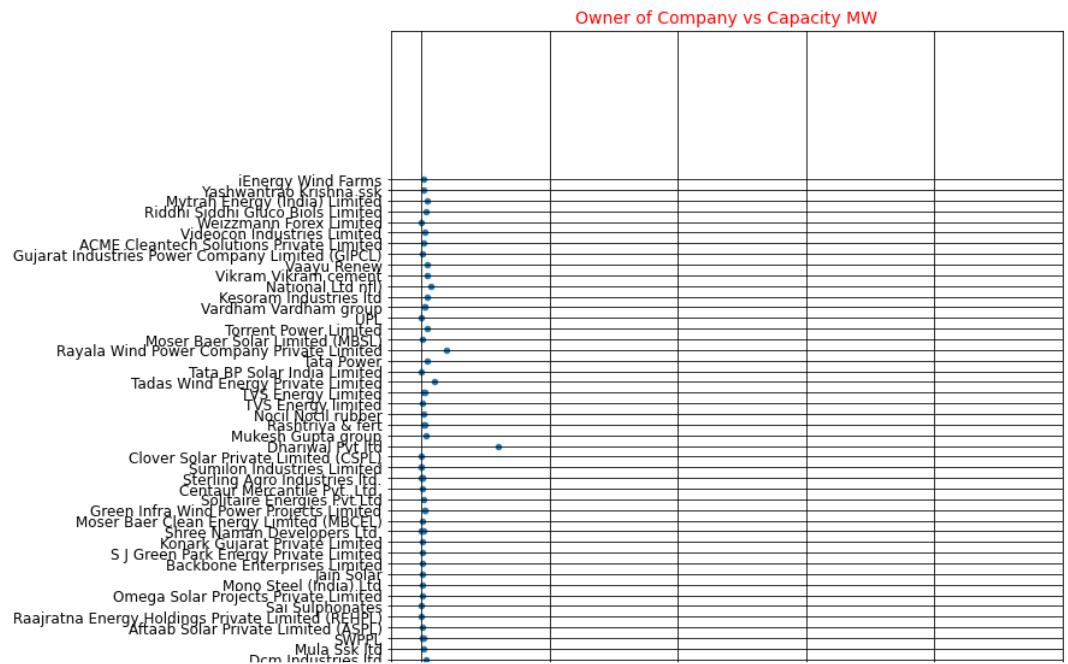
```
In [38]: df['owner'].value_counts()
```

```
Out[38]: Acc Acc ltd                               570
Sterling Agro Industries ltd.                      4
Jk Cement ltd                                     4
Tata Power Solar Systems Limited (TPREL)          3
Ujaas Energy Limited                            3
...
Gm Energy ltd                                    1
Dcm & chem                                      1
REI Agro Limited                                1
National And paper                            1
Spr Pvt ltd                                     1
Name: owner, Length: 280, dtype: int64
```

```
In [39]: o = df.groupby('owner')['primary_fuel'].sum()
o
```

```
Out[39]: owner
ACME Cleantech Solutions Private Limited        Solar
ACME Solar Energy                               Solar
AES                                         Wind
AEW Infratech Private Limited                  Solar
Abellon CleanEnergy Limited                    Solar
...
West Coast Paper Mills Ltd.                     Gas
Yashwantrao Krishna ssk                         Biomass
Ym Ssk ltd                                     Biomass
Zamil New Delhi Infrastructure Private Limited Solar
iEnergy Wind Farms                            Wind
Name: primary_fuel, Length: 280, dtype: object
```

```
In [25]: df.plot.scatter(y = 'owner', x = 'capacity_mw', figsize = (10,50), rot = 360, font
plt.grid(c = 'black')
plt.ylabel('Owner', fontsize = 14, color = 'r')
plt.xlabel('capacity_mw', fontsize = 14, color = 'r')
plt.title('Owner of Company vs Capacity MW', fontsize = 14, color = 'r')
plt.show()
```



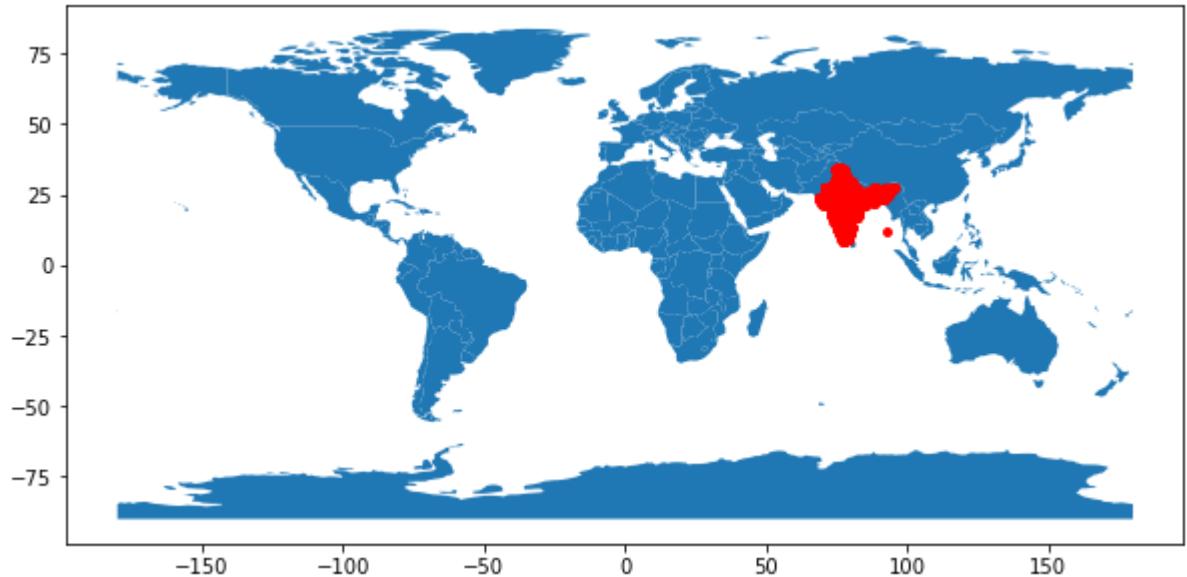
Above plot show JK Cement company consumption energy is more than 4000 MW and other company consume energy less than 1000 MW

## Co-ordinates

```
In [26]: import geopandas as gpd
from shapely.geometry import Point, Polygon
import descartes
from geopandas import GeoDataFrame
from pyproj import CRS
```

```
In [27]: geometry = [Point(xy) for xy in zip(df['longitude'], df['latitude'])]
gdf = GeoDataFrame(df, geometry=geometry)

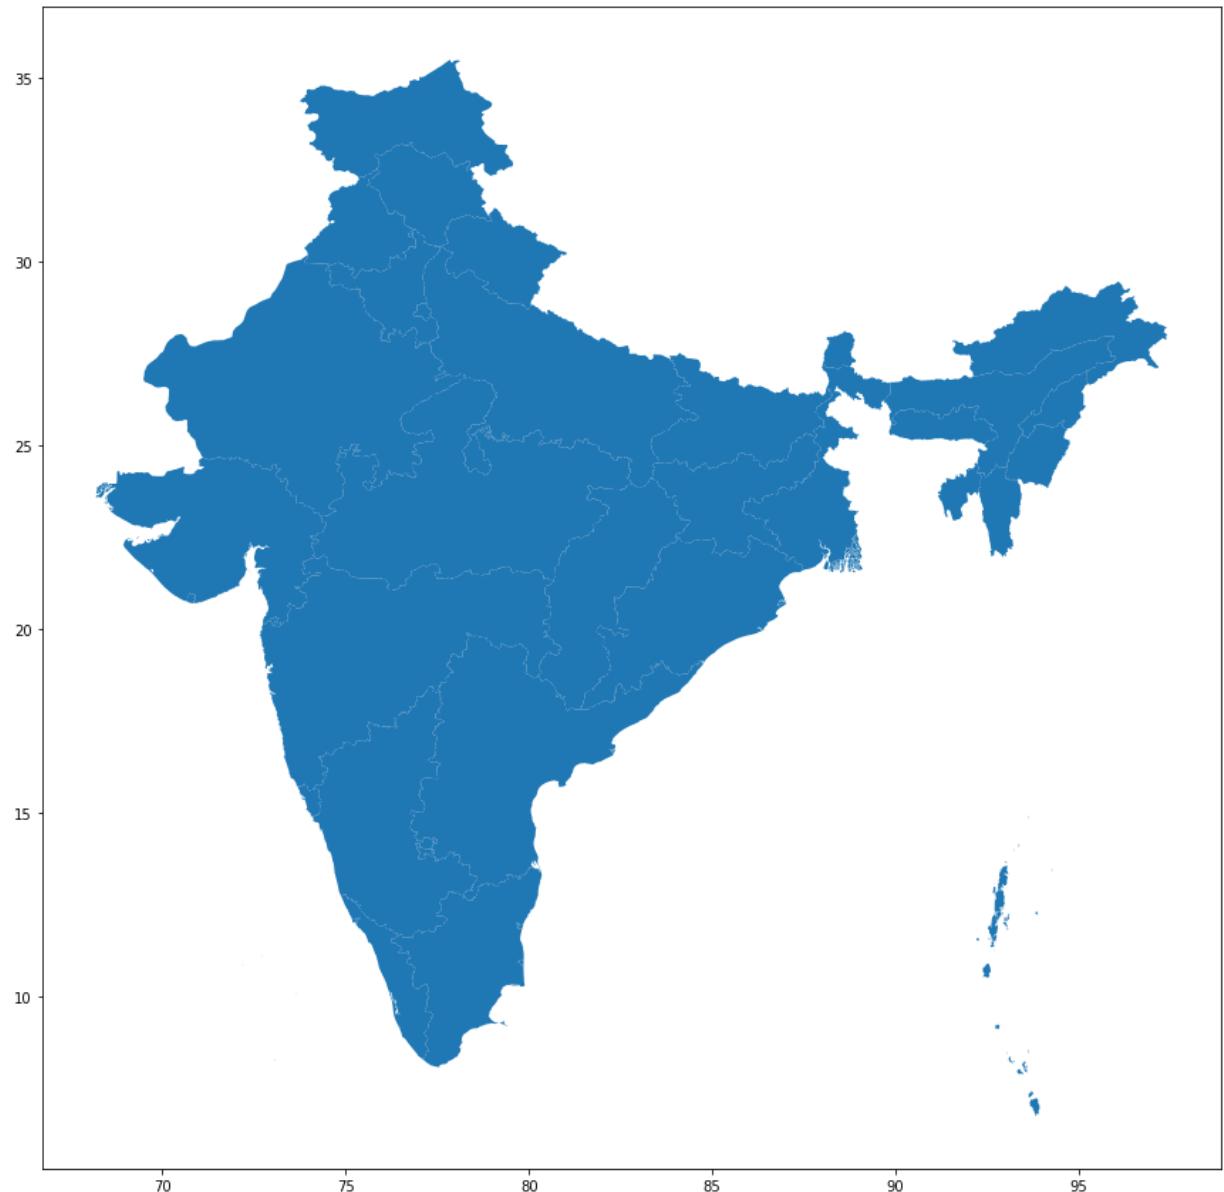
#this is a simple map that goes with geopandas
world = gpd.read_file(gpd.datasets.get_path('naturalearth_lowres'))
gdf.plot(ax=world.plot(figsize=(10, 6)), marker='o', color='red', markersize=15);
```



Above plot shows all coordinate belong to indian

```
In [28]: indmap = gpd.read_file(r"C:\Users\Kushal Arya\Desktop\csv file\map.shx")
fig,ax = plt.subplots(figsize = (15,15))
indmap.plot(ax = ax)
```

Out[28]: <AxesSubplot:>



```
In [29]:
```

```
crs=CRS('EPSG:4326').to_proj4()
```

```
In [30]:
```

```
geometry = [Point(xy) for xy in zip(df['longitude'], df['latitude'])]
geometry[:3]
```

```
Out[30]:
```

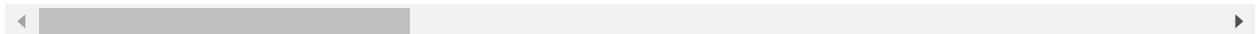
```
[<shapely.geometry.point.Point at 0x200d8d86fa0>,
 <shapely.geometry.point.Point at 0x200d94afb50>,
 <shapely.geometry.point.Point at 0x200d94af370>]
```

```
In [31]:
```

```
geo_df = gpd.GeoDataFrame(df, crs = crs, geometry = geometry)
geo_df.head()
```

```
Out[31]:
```

	country_long	name	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commissioning
0	India	ACME Solar Tower	2.5	28.1839	73.2407	Solar	Oil	2010
1	India	ADITYA CEMENT WORKS	98.0	24.7663	74.6090	Coal	Oil	1998
2	India	AES Saurashtra Windfarms	39.2	21.9038	69.3732	Wind	Oil	1998
3	India	AGARTALA GT	135.0	23.8712	91.3602	Gas	Oil	2000
4	India	AKALTARA TPP	1800.0	21.9603	82.4091	Coal	Oil	2000



In [32]: `try:`

```
fig, ax = plt.subplots(figsize = (15,15))
indmap.plot(ax = ax, alpha = 0.4, color = 'grey')
geo_df[geo_df['capacity_mw']].plot(ax = ax, markersize = 20, color = 'b', mar
geo_df[geo_df['capacity_mw']].plot(ax = ax, markersize = 20, color = '^', mar
plt.legend(prop = {'size': 15})
except:
    pass
```



```
In [34]: import plotly  
import plotly.graph_objects as go  
import chart_studio.plotly as py  
from plotly.offline import iplot
```

```
In [35]: data = dict(type = 'choropleth',
                  locations = df['geometry'],
                  locationmode = 'country names',
                  z = df['capacity_mw'],
                  text = df['geometry'],
                  colorscale = 'YlGnBu',
                  autocolorscale = False,
                  )
layout = dict(geo = {'scope':'asia'})

diagram = go.Figure(data = [data], layout = layout)
iplot(diagram)
```



```
In [52]: from geopy.geocoders import Nominatim
```

```
In [53]: # initialize Nominatim API
geolocator = Nominatim(user_agent="geoapiExercises")
```

```
In [54]: # Latitude & Longitude input
Latitude = "28.1839"
Longitude = "73.2407"

location = geolocator.reverse(Latitude+","+Longitude)

address = location.raw['address']

# traverse the data
city = address.get('city', '')
state = address.get('state', '')
country = address.get('country', '')
code = address.get('country_code')
zipcode = address.get('postcode')
print('City : ', city)
print('State : ', state)
print('Country : ', country)
print('Zip Code : ', zipcode)
```

City :  
 State : Rajasthan  
 Country : India  
 Zip Code : None

## Drop Columns

```
In [36]: col = ['geometry', 'generation_data_source', 'source', 'owner', 'name', 'country_
```

```
In [37]: df = df.drop(col, axis = 1)
df.head(2)
```

Out[37]:

	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commissioning_year	year_of_capacity
0	2.5	28.1839	73.2407	Solar	Oil	2011.000000	
1	98.0	24.7663	74.6090	Coal	Oil	1996.876894	

```
In [38]: print('After dropping no of Rows and Columns ---->', df.shape )
```

After dropping no of Rows and Columns ----> (908, 12)

## Encoding columns

```
In [39]: le = LabelEncoder()
```

```
In [40]: df['primary_fuel'] = le.fit_transform(df['primary_fuel'])
```

```
In [41]: df['other_fuel1'] = le.fit_transform(df['other_fuel1'])
```

```
In [42]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 908 entries, 0 to 907
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   capacity_mw      908 non-null    float64
 1   latitude          908 non-null    float64
 2   longitude         908 non-null    float64
 3   primary_fuel      908 non-null    int32  
 4   other_fuel1       908 non-null    int32  
 5   commissioning_year 908 non-null    float64
 6   year_of_capacity_data 908 non-null    float64
 7   generation_gwh_2013 908 non-null    float64
 8   generation_gwh_2014 908 non-null    float64
 9   generation_gwh_2015 908 non-null    float64
 10  generation_gwh_2016 908 non-null    float64
 11  generation_gwh_2017 908 non-null    float64
dtypes: float64(10), int32(2)
memory usage: 78.2 KB
```

```
In [62]: df.head(2)
```

```
Out[62]:
```

	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commissioning_year	year_of_capacity
0	2.5	28.1839	73.2407		6	2	2011.000000
1	98.0	24.7663	74.6090		1	2	1996.876894

All columns are encoded

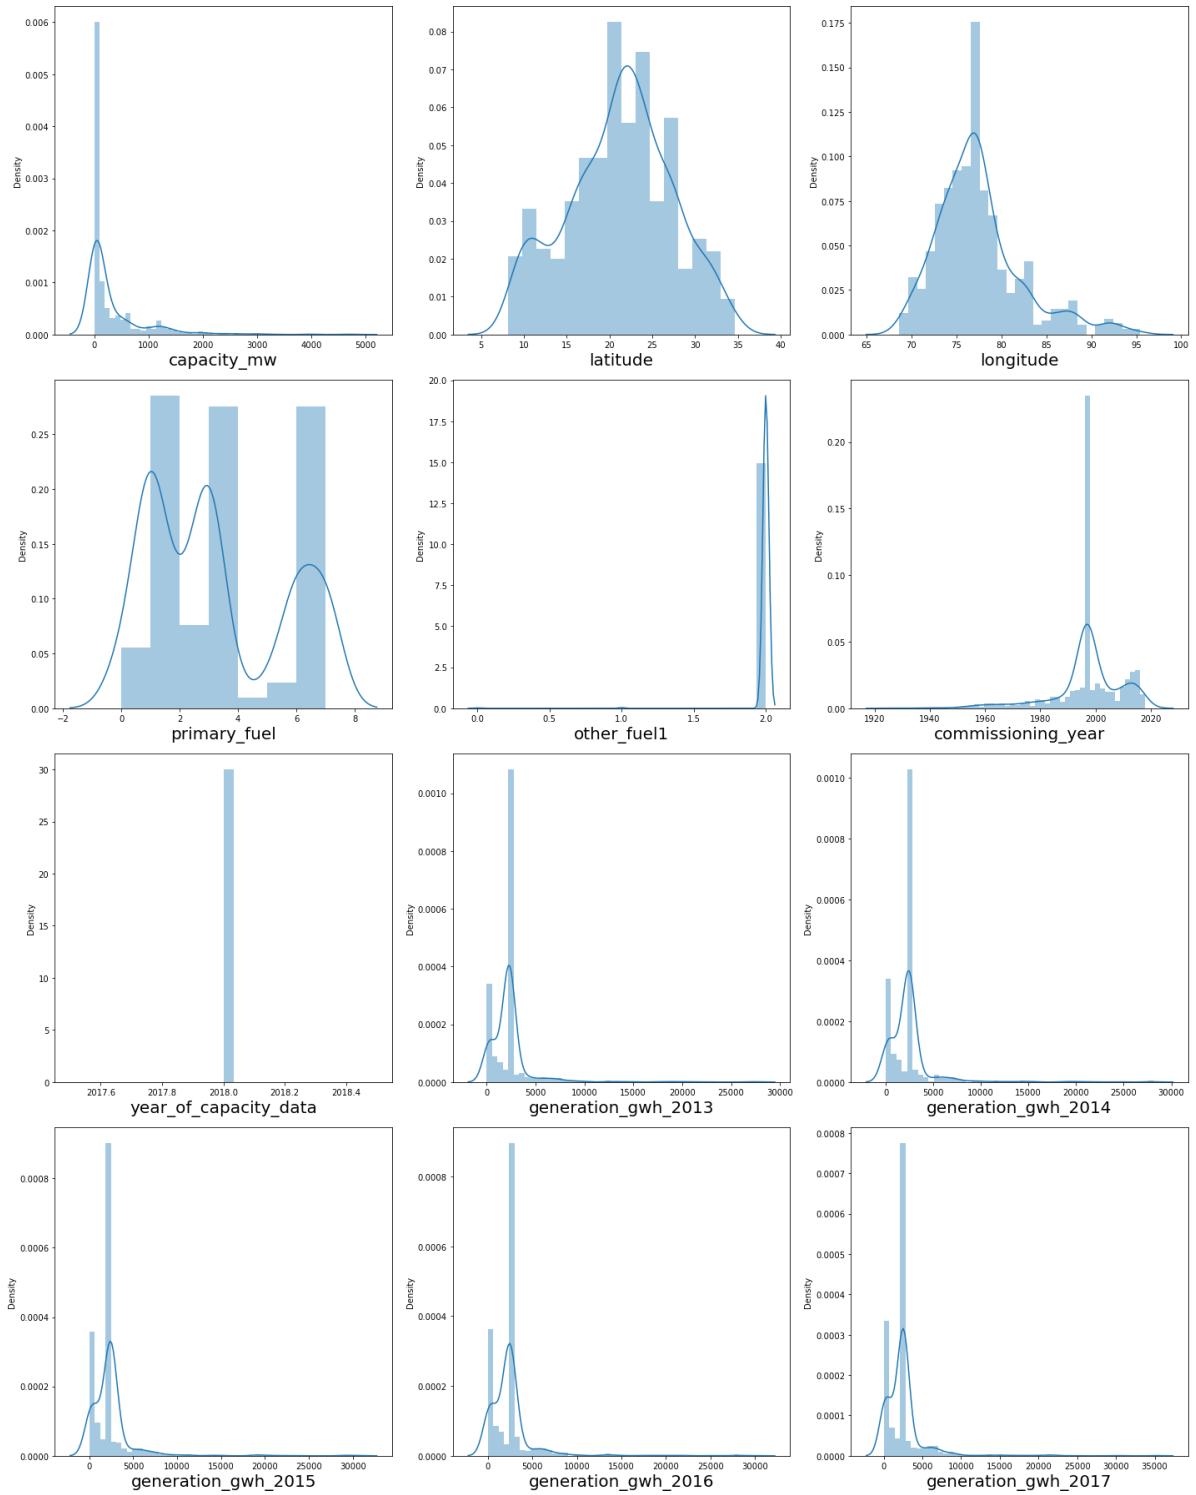
## Data distribution and checking outliers and skewness

```
In [43]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in df:
    if plotnumber <=12:
        ax = plt.subplot(4,3, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----
Distribution Plot :-
```



```
In [44]: df.skew()
```

```
Out[44]: capacity_mw            3.193257
latitude                  -0.147391
longitude                 1.129836
primary_fuel              0.471141
other_fuel1              -20.464435
commissioning_year        -1.383330
year_of_capacity_data    0.000000
generation_gwh_2013       5.241491
generation_gwh_2014       5.041961
generation_gwh_2015       5.367370
generation_gwh_2016       5.071758
generation_gwh_2017       5.111938
dtype: float64
```

Data has outliers and skewness

Corelation of Feature vs Label using Heat map

```
In [65]: print('-----')
```

```
print('Heat Map :-')
```

```
print('-----')
```

```
df_corr = df.corr().abs()
```

```
plt.figure(figsize = (22,16))
```

```
sns.heatmap(df_corr, vmin = -1, annot = True, square = True, center = 0, fmt = '.2f')
```

```
plt.tight_layout()
```

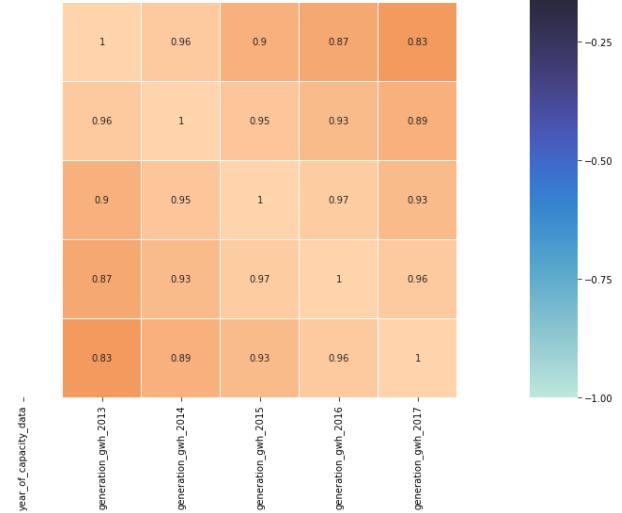
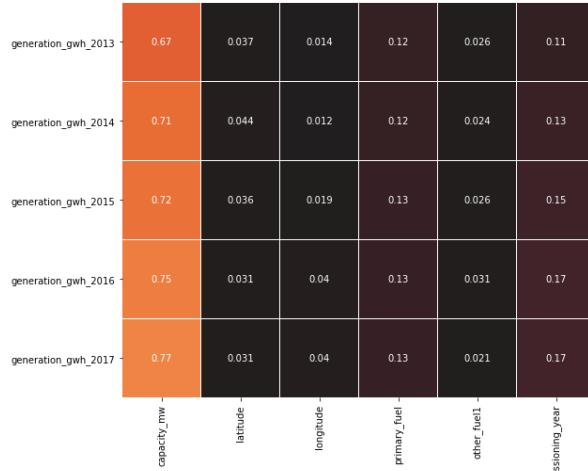
```
-----
```

```
Heat Map :-
```

```
-----
```



```
year_of_capacity_data ~
```



**generation\_gwh\_2013,generation\_gwh\_2014 higest relation**

## Removing Outliers using Zscore

In [45]: *# with std 3 Lets see the stats*

```
z_score = zscore(df[['longitude', 'generation_gwh_2013', 'generation_gwh_2014', 'abs_z_score = np.abs(z_score)

filtering_entry = (abs_z_score < 3).all(axis = 1)

df = df[filtering_entry]

df.describe()
```

Out[45]:

	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commissioning_year	year_c
<b>count</b>	871.000000	871.000000	871.000000	871.000000	871.000000	871.000000	871.000000
<b>mean</b>	272.583767	21.089931	77.036419	3.257176	1.995408	1996.855732	
<b>std</b>	439.479660	6.145269	4.197722	2.303519	0.082918	13.017081	
<b>min</b>	0.000000	8.168900	68.644700	0.000000	0.000000	1927.000000	
<b>25%</b>	16.500000	16.899050	74.329400	1.000000	2.000000	1996.876894	
<b>50%</b>	50.700000	21.196189	76.760600	3.000000	2.000000	1996.876894	
<b>75%</b>	330.000000	25.114850	78.922500	6.000000	2.000000	2003.000000	
<b>max</b>	2400.000000	34.649000	91.565000	7.000000	2.000000	2018.000000	

**Checking Outliers and skewness removed or not**

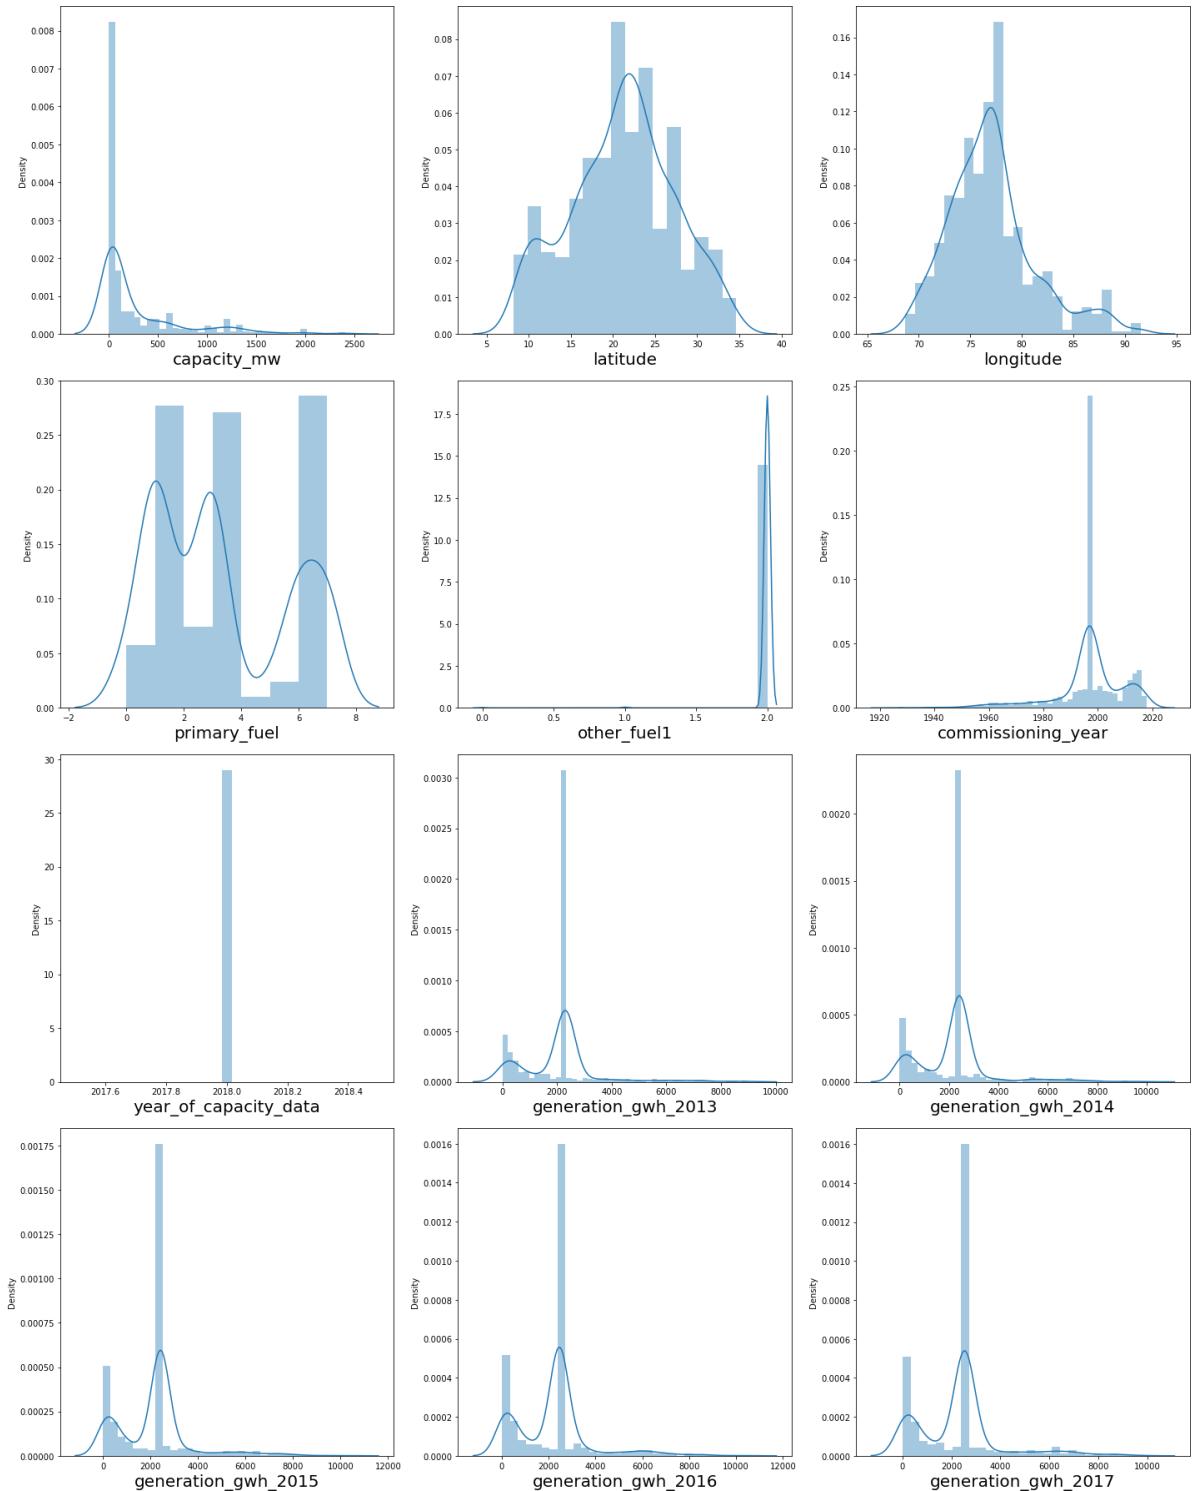
```
In [46]: # Let's see outliers are removed in columns or not.
```

```
print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in df:
    if plotnumber <=12:
        ax = plt.subplot(4,3, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----  
Distribution Plot :-  
-----
```



**Outliers are removed**

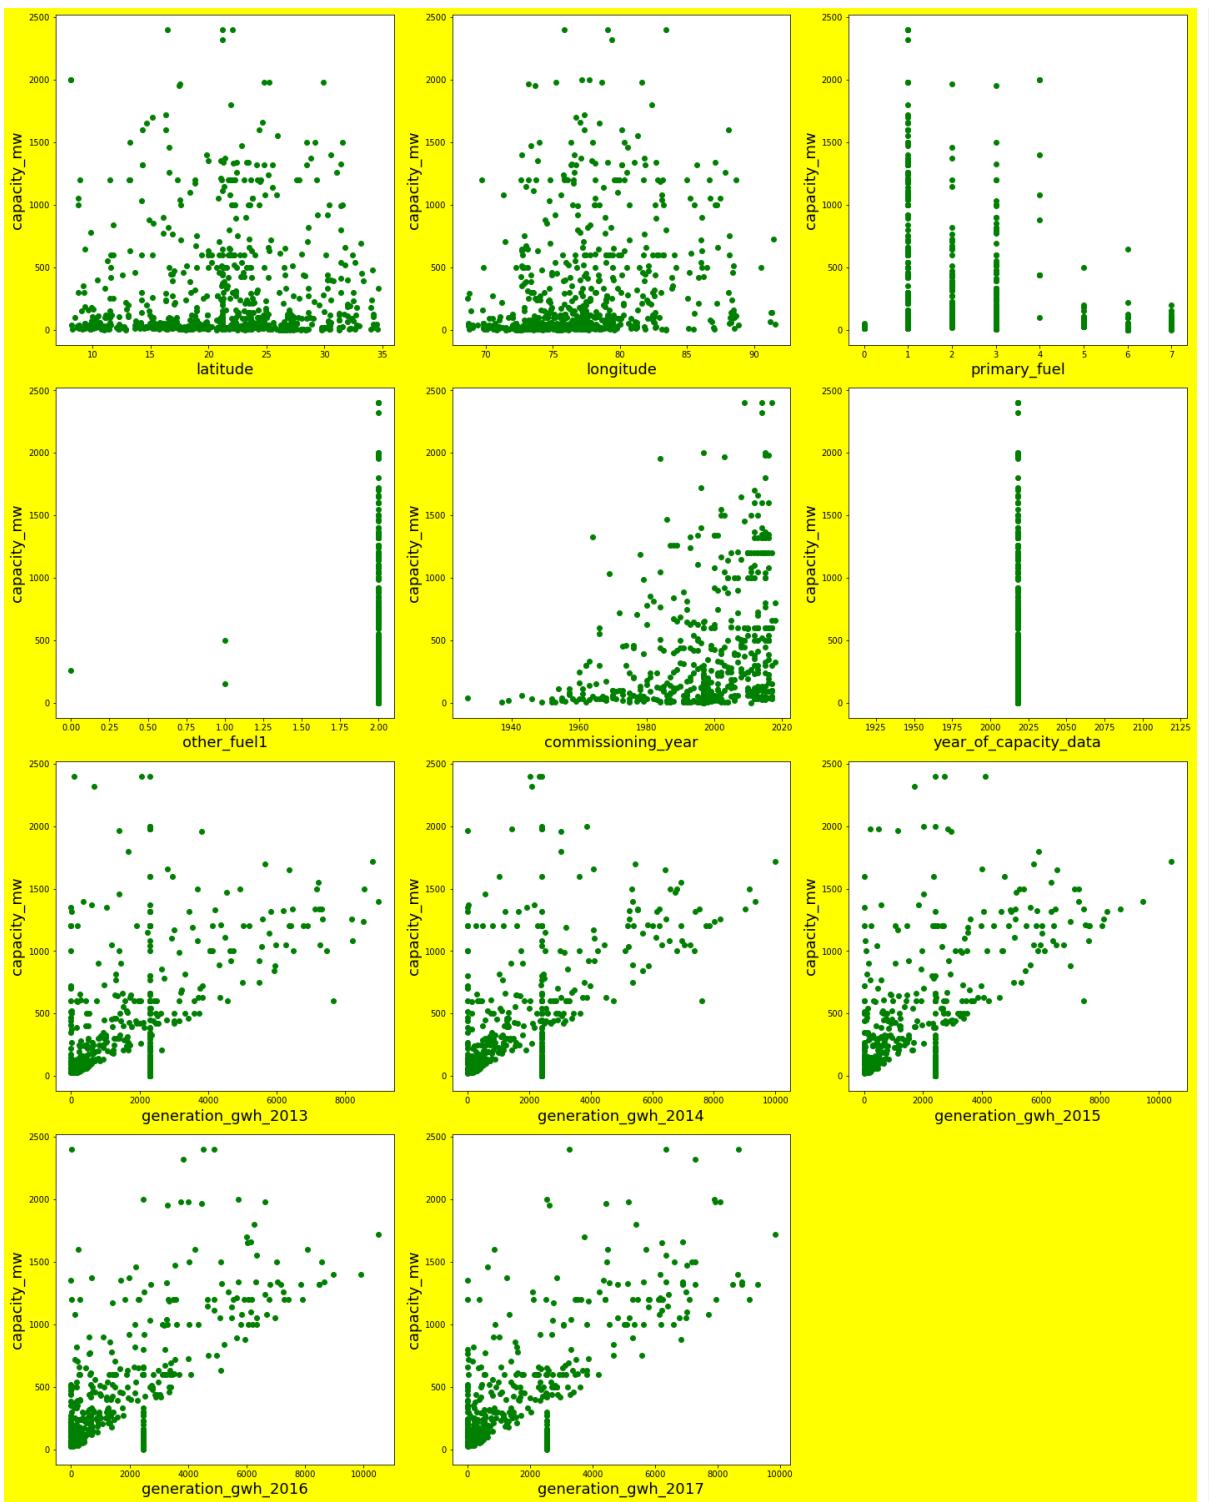
**Splitting Dataset into features and labels**

```
In [47]: x = df.drop('capacity_mw', axis = 1)
y = df['capacity_mw']
print('Data has been splited')
```

Data has been splited

```
In [49]: # Let's see relation between features and labels.  
print('-----')  
print('Scatter Plot :-')  
print('-----')  
  
plt.figure(figsize = (20,25), facecolor = 'yellow')  
plotnumber = 1  
for column in x:  
    if plotnumber <=12:  
        ax = plt.subplot(4,3, plotnumber)  
        plt.scatter(x[column],y, c = 'g')  
        plt.xlabel(column, fontsize = 18)  
        plt.ylabel('capacity_mw', fontsize = 18)  
    plotnumber += 1  
plt.tight_layout()
```

-----  
Scatter Plot :-  
-----



**Features are related to label**

## Data Scaling

```
In [52]: scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
Out[52]: array([[ 1.15504214, -0.90475257,  1.1913948 , ...,  0.20019784,
   0.1919957 ,  0.1785893 ],
   [ 0.5985875 , -0.57860276, -0.98044461, ...,  0.20019784,
   0.1919957 ,  0.1785893 ],
   [ 0.13251441, -1.82661503,  1.62576269, ...,  0.20019784,
   0.1919957 ,  0.1785893 ],
   ...,
   [-0.9466585 , -0.34689171,  1.62576269, ...,  0.20019784,
   0.1919957 ,  0.1785893 ],
   [ 0.53080541, -0.78390338, -0.98044461, ...,  0.20019784,
   0.1919957 ,  0.1785893 ],
   [-1.81634681,  0.10496987,  1.62576269, ...,  0.20019784,
   0.1919957 ,  0.1785893 ]])
```

**Data has been scaled**

**Split data into train and test. Model will be bulit on training data and tested on test data**

```
In [53]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25, random_state = 42)
print('Data has been splited.')
```

Data has been splited.

## Model Bulding

**Decision Tree model instantiaing, training and evaluating**

```
In [54]: DT = DecisionTreeRegressor()
DT.fit(x_train, y_train)
y_pred = DT.predict(x_test)
```

```
In [57]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', DT.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.6111258730830611  
=====  
RMSE of Model -----> 274.22435732040987  
=====  
MSE of Model -----> 75198.99814779183  
=====  
Score of test data ----> 0.6111258730830611  
=====
```

Conclusion : Decision Tree model has 61% score

## XGBoost model instantiaing, training and evaluating

```
In [58]: xgb = xgb.XGBRegressor(eval_metric = 'mlogloss')  
xgb.fit(x_train, y_train)  
y_pred = xgb.predict(x_test)
```

```
In [59]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', xgb.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.752973794775826  
=====  
RMSE of Model -----> 218.56117471317214  
=====  
MSE of Model -----> 47768.98709200176  
=====  
Score of test data ----> 0.752973794775826  
=====
```

Conclusion : XGBoost model has 75% score

## Knn model instantiaing, training and evaluating

```
In [60]: Knn = KNeighborsRegressor()  
Knn.fit(x_train, y_train)  
y_pred = Knn.predict(x_test)
```

```
In [61]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ---->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ---->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', Knn.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.7713684825660243  
=====  
RMSE of Model ----> 210.2662417010793  
=====  
MSE of Model ----> 44211.8923990967  
=====  
Score of test data ----> 0.7713684825660243  
=====
```

Conclusion : KNN model has 77% score

## Random Forest model instantiaing, training and evaluating

```
In [62]: Rn = RandomForestRegressor()  
Rn.fit(x_train, y_train)  
y_pred = Rn.predict(x_test)
```

```
In [63]: print('=====')  
print('R2 Score ---->', r2_score(y_test, y_pred))  
print('=====')  
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))  
print('=====')  
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))  
print('=====')  
print('Score of test data ---->', Rn.score(x_test, y_test))  
print('=====')
```

```
=====  
R2 Score ----> 0.8146813211243855  
=====  
RMSE of Model -----> 189.30457093961715  
=====  
MSE of Model -----> 35836.22057863254  
=====  
Score of test data ----> 0.8146813211243855  
=====
```

Conclusion : Random Forest model has 81% score

Looking RSME score we found Random Forest has best model so we do Hyperparameter Tuning on it

```
In [64]: param_grid = {'n_estimators': [100, 200, 300, 400, 500],  
                    'max_features': ['auto', 'sqrt'],  
                    'max_depth': [5, 10, 15, 20, 25, 30],  
                    'min_samples_split': [2, 5, 10, 15, 100],  
                    'min_samples_leaf': [1, 2, 5, 10]}
```

```
In [65]: grid_search = GridSearchCV(estimator = Rn, param_grid = param_grid, cv = 5,n_jobs=-1)
```

```
In [66]: grid_search.fit(x_train, y_train)
```

```
Out[66]: GridSearchCV(cv=5, estimator=RandomForestRegressor(), n_jobs=-1,  
param_grid={'max_depth': [5, 10, 15, 20, 25, 30],  
           'max_features': ['auto', 'sqrt'],  
           'min_samples_leaf': [1, 2, 5, 10],  
           'min_samples_split': [2, 5, 10, 15, 100],  
           'n_estimators': [100, 200, 300, 400, 500]})
```

```
In [67]: best_parameters = grid_search.best_params_  
print(best_parameters)
```

```
{'max_depth': 30, 'max_features': 'auto', 'min_samples_leaf': 5, 'min_samples_split': 5, 'n_estimators': 200}
```

```
In [68]: hRn = RandomForestRegressor(max_depth = 30, max_features = 'auto', min_samples_leaf=1, n_estimators=100)
hRn.fit(x_train, y_train)
hRn.score(x_test, y_test)
```

```
Out[68]: 0.8034611402413743
```

```
In [69]: y_pred = hRn.predict(x_test)
```

```
In [70]: print('=====')
print('R2 Score ---->', r2_score(y_test, y_pred))
print('=====')
print('RMSE of Model ----->', np.sqrt(mean_squared_error(y_test, y_pred)))
print('=====')
print('MSE of Model ----->', mean_squared_error(y_test, y_pred))
print('=====')
print('Score of test data ---->', hRn.score(x_test, y_test))
print('=====')
```

```
=====
R2 Score ----> 0.8034611402413743
=====
RMSE of Model -----> 194.95111295013075
=====
MSE of Model -----> 38005.93644049463
=====
Score of test data ----> 0.8034611402413743
=====
```

After Hyperparameter Tuning model accuracy score 80%.

## Saving The Model

```
In [71]: # saving the model to the Local file system
filename = 'Global Power Plant Project (Capacity MW).pickle'
pickle.dump(hRn, open(filename, 'wb'))
```

Final Conclusion : Random Forest is our best model.

```
In [ ]:
```