

Problem Statement:

The Global Power Plant Database is a comprehensive, open source database of power plants around the world. It centralizes power plant data to make it easier to navigate, compare and draw insights for one's own analysis. The database covers approximately 35,000 power plants from 167 countries and includes thermal plants (e.g. coal, gas, oil, nuclear, biomass, waste, geothermal) and renewables (e.g. hydro, wind, solar). Each power plant is geolocated and entries contain information on plant capacity, generation, ownership, and fuel type. It will be continuously updated as data becomes available.

Importing Required Library

```
In [125]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pickle
from scipy.stats import zscore
import scikitplot as skplt
pd.set_option('display.max_columns', None) # # For display maximum column
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, plot_roc_curve
import xgboost as xgb
%matplotlib inline

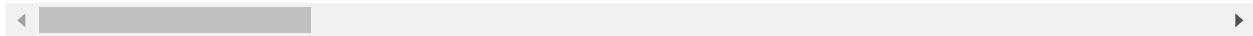
import warnings
warnings.filterwarnings('ignore')
```

Reading Data

```
In [2]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\csv file\database_IND.csv")
df.head()
```

Out[2]:

	country	country_long		name	gppd_idnr	capacity_mw	latitude	longitude	primary_fuel
0	IND	India		ACME Solar Tower	WRI1020239	2.5	28.1839	73.2407	Solar
1	IND	India		ADITYA CEMENT WORKS	WRI1019881	98.0	24.7663	74.6090	Coal
2	IND	India	AES Saurashtra Windfarms		WRI1026669	39.2	21.9038	69.3732	Wind
3	IND	India	AGARTALA GT		IND0000001	135.0	23.8712	91.3602	Gas
4	IND	India	AKALTARA TPP		IND0000002	1800.0	21.9603	82.4091	Coal



Check no of row and column

```
In [3]: print('No of Rows and Columns ----->', df.shape )
```

No of Rows and Columns -----> (908, 25)

Checking for Null values

```
In [4]: print('-----\n')
print(df.isnull().sum())
print('\n-----')
```

```
-----  
country 0  
country_long 0  
name 0  
gppd_idnr 0  
capacity_mw 0  
latitude 46  
longitude 46  
primary_fuel 0  
other_fuel1 709  
other_fuel2 907  
other_fuel3 908  
commissioning_year 380  
owner 566  
source 0  
url 0  
geolocation_source 19  
wepp_id 908  
year_of_capacity_data 388  
generation_gwh_2013 524  
generation_gwh_2014 507  
generation_gwh_2015 483  
generation_gwh_2016 471  
generation_gwh_2017 465  
generation_data_source 458  
estimated_generation_gwh 908  
dtype: int64
```

There is null value

Information about dataset

```
In [5]: print('-----\n')
print(df.info())
print('-----')
```

```
-----  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 908 entries, 0 to 907  
Data columns (total 25 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   country          908 non-null    object    
 1   country_long     908 non-null    object    
 2   name              908 non-null    object    
 3   gppd_idnr        908 non-null    object    
 4   capacity_mw      908 non-null    float64  
 5   latitude          862 non-null    float64  
 6   longitude         862 non-null    float64  
 7   primary_fuel      908 non-null    object    
 8   other_fuel1       199 non-null    object    
 9   other_fuel2       1 non-null     object    
 10  other_fuel3      0 non-null     float64  
 11  commissioning_year 528 non-null    float64  
 12  owner              342 non-null    object    
 13  source             908 non-null    object    
 14  url                908 non-null    object    
 15  geolocation_source 889 non-null    object    
 16  wepp_id            0 non-null     float64  
 17  year_of_capacity_data 520 non-null    float64  
 18  generation_gwh_2013 384 non-null    float64  
 19  generation_gwh_2014 401 non-null    float64  
 20  generation_gwh_2015 425 non-null    float64  
 21  generation_gwh_2016 437 non-null    float64  
 22  generation_gwh_2017 443 non-null    float64  
 23  generation_data_source 450 non-null    object    
 24  estimated_generation_gwh 0 non-null     float64  
dtypes: float64(13), object(12)  
memory usage: 177.5+ KB  
None  
-----
```

Categorical data present in our data set

Statistics of Data

In [6]: df.describe()

Out[6]:

	capacity_mw	latitude	longitude	other_fuel3	commissioning_year	wepp_id	year_of_c
count	908.000000	862.000000	862.000000	0.0	528.000000	0.0	
mean	321.046378	21.196189	77.447848	NaN	1996.876894	NaN	
std	580.221767	6.248627	4.907260	NaN	17.047817	NaN	
min	0.000000	8.168900	68.644700	NaN	1927.000000	NaN	
25%	16.837500	16.771575	74.258975	NaN	1988.000000	NaN	
50%	60.000000	21.778300	76.719250	NaN	2000.000000	NaN	
75%	388.125000	25.516375	79.441475	NaN	2011.250000	NaN	
max	4760.000000	34.649000	95.408000	NaN	2018.000000	NaN	

Outliers are present in our data set

Analysis of Null value

```
In [7]: print('-----\n')
print(df.isnull().sum())
print('\n-----')
```

```
-----  
country 0  
country_long 0  
name 0  
gppd_idnr 0  
capacity_mw 0  
latitude 46  
longitude 46  
primary_fuel 0  
other_fuel1 709  
other_fuel2 907  
other_fuel3 908  
commissioning_year 380  
owner 566  
source 0  
url 0  
geolocation_source 19  
wepp_id 908  
year_of_capacity_data 388  
generation_gwh_2013 524  
generation_gwh_2014 507  
generation_gwh_2015 483  
generation_gwh_2016 471  
generation_gwh_2017 465  
generation_data_source 458  
estimated_generation_gwh 908  
dtype: int64
```

```
In [8]: df['other_fuel1'].value_counts()
```

```
Out[8]: Oil      196
Gas      2
Cogeneration      1
Name: other_fuel1, dtype: int64
```

```
In [9]: df['other_fuel2'].value_counts()
```

```
Out[9]: Oil      1
Name: other_fuel2, dtype: int64
```

```
In [10]: df['other_fuel3'].value_counts()
```

```
Out[10]: Series([], Name: other_fuel3, dtype: int64)
```

Approach : We drop above other_fuel2 and other_fuel3 column because maximum null value and keep other_fuel1 and fill null value with mode

```
In [11]: df['latitude'].value_counts()
```

```
Out[11]: 24.1917      3  
19.0004      3  
16.5697      2  
23.4639      2  
13.2450      2  
..  
20.9099      1  
17.2387      1  
23.5594      1  
27.3426      1  
16.5973      1  
Name: latitude, Length: 837, dtype: int64
```

```
In [12]: df['longitude'].value_counts()
```

```
Out[12]: 71.6917      4  
71.6918      3  
75.8988      3  
72.8983      3  
81.2875      3  
..  
80.1264      1  
76.1137      1  
74.6447      1  
86.0970      1  
79.5748      1  
Name: longitude, Length: 828, dtype: int64
```

Approach : We fill above longitude and latitude columns with mean

```
In [13]: df['commissioning_year'].value_counts()
```

```
Out[13]: 2013.0      28  
2015.0      26  
2012.0      23  
2016.0      21  
2014.0      17  
..  
1958.0      1  
1949.0      1  
1954.0      1  
1956.0      1  
1927.0      1  
Name: commissioning_year, Length: 73, dtype: int64
```

Approach: We fill commissioning_year column with mode

```
In [14]: df['owner'].value_counts()
```

```
Out[14]: Acc Acc ltd          4
Sterling Agro Industries ltd. 4
Jk Cement ltd                 4
Tata Power Solar Systems Limited (TPREL) 3
Ujaas Energy Limited          3
..
Gm Energy ltd                 1
Dcm & chem                     1
REI Agro Limited               1
National And paper             1
Spr Pvt ltd                   1
Name: owner, Length: 280, dtype: int64
```

Approach : We fill owner column with mode

```
In [15]: df['geolocation_source'].value_counts()
```

```
Out[15]: WRI                  766
Industry About                119
National Renewable Energy Laboratory 4
Name: geolocation_source, dtype: int64
```

Approach : We fill geolocation_source column with mode

```
In [16]: df['wepp_id'].value_counts()
```

```
Out[16]: Series([], Name: wepp_id, dtype: int64)
```

Approach : We drop wepp_id column because maximum null value

```
In [17]: df['year_of_capacity_data'].value_counts()
```

```
Out[17]: 2018.0      520
Name: year_of_capacity_data, dtype: int64
```

Approach : We fill year_of_capacity_data column with bfill beacuse only one data in it

```
In [18]: df['generation_gwh_2013'].value_counts()
```

```
Out[18]: 0.00000    21
14881.88000     1
42.49645       1
2036.00000     1
97.73885       1
..
7229.33000     1
657.21740       1
507.89775       1
8556.42400     1
8211.00000     1
Name: generation_gwh_2013, Length: 364, dtype: int64
```

```
In [19]: df['generation_gwh_2014'].value_counts()
```

```
Out[19]: 0.00000    28
6803.31250     1
4735.13000     1
145.81400      1
2022.57000     1
..
6224.00000     1
268.48085      1
1255.73200     1
164.32425      1
1153.65300     1
Name: generation_gwh_2014, Length: 374, dtype: int64
```

```
In [20]: df['generation_gwh_2015'].value_counts()
```

```
Out[20]: 0.00000    28
5837.76600     1
1297.97750     1
8076.81050     1
1.09395        1
..
2636.86400     1
665.19730      1
1516.36010     1
741.86205      1
7130.50700     1
Name: generation_gwh_2015, Length: 398, dtype: int64
```

```
In [21]: df['generation_gwh_2016'].value_counts()
```

```
Out[21]: 0.00000    31
8470.57000     2
1511.00000     2
7.31325        1
94.85500        1
...
433.84800      1
283.74811      1
259.94375      1
403.96000      1
307.87290      1
Name: generation_gwh_2016, Length: 405, dtype: int64
```

```
In [22]: df['generation_gwh_2017'].value_counts()
```

```
Out[22]: 0.00000    33
170.08530      2
344.35955      1
2265.47000      1
59.43135      1
...
214.48220      1
272.73945      1
2887.00000      1
12.73600      1
158.73235      1
Name: generation_gwh_2017, Length: 410, dtype: int64
```

Approach : We fill above generation_gwh columns with mean

```
In [23]: df['generation_data_source'].value_counts()
```

```
Out[23]: Central Electricity Authority    450
Name: generation_data_source, dtype: int64
```

Approach : We fill above generation_data_source columns with bfill because only one data in it

```
In [24]: df['estimated_generation_gwh'].value_counts()
```

```
Out[24]: Series([], Name: estimated_generation_gwh, dtype: int64)
```

Approach : We drop estimated_generation_gwh columns because no data in it

Drop Unwanted column

```
In [25]: col = ['estimated_generation_gwh', 'wepp_id', 'other_fuel2', 'other_fuel3', 'url']
```

```
In [26]: df = df.drop(col, axis = 1)
df.head(2)
```

Out[26]:

	country_long	name	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commission
0	India	ACME Solar Tower	2.5	28.1839	73.2407	Solar		NaN
1	India	ADITYA CEMENT WORKS	98.0	24.7663	74.6090	Coal		NaN

```
In [27]: print('After droping no of Rows and Columns ---->', df.shape )
```

After droping no of Rows and Columns ----> (908, 18)

Fill NaN

```
In [28]: df = df.apply(lambda x:x.fillna(x.mean()))if x.dtype == 'float' else x.fillna(x.v
```

```
In [29]: print('-----\n')
print(df.isnull().sum())
print('\n-----')
```

```
-----  
country_long          0  
name                  0  
capacity_mw           0  
latitude              0  
longitude             0  
primary_fuel          0  
other_fuel1           0  
commissioning_year    0  
owner                 0  
source                0  
geolocation_source    0  
year_of_capacity_data 0  
generation_gwh_2013   0  
generation_gwh_2014   0  
generation_gwh_2015   0  
generation_gwh_2016   0  
generation_gwh_2017   0  
generation_data_source 0  
dtype: int64  
-----
```

There is no null value left

Analysis of data respect to fuel type

```
In [30]: df.head(2)
```

Out[30]:


		country_long	name	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commission
0	India	ACME Solar Tower		2.5	28.1839	73.2407	Solar	Oil	201
1	India	ADITYA CEMENT WORKS		98.0	24.7663	74.6090	Coal	Oil	199

Fule Type column

```
In [31]: df['primary_fuel'].value_counts()
```

```
Out[31]: Coal      259  
Hydro     250  
Solar     127  
Wind      123  
Gas       69  
Biomass    50  
Oil        21  
Nuclear     9  
Name: primary_fuel, dtype: int64
```

```
In [32]: sns.histplot(binwidth=0.5, x="primary_fuel", hue="primary_fuel", data=df, stat="count")  
plt.show()
```




Cole type fuel used highest and Nuclear type fuel used least

Other Fuel Type column

```
In [33]: df['other_fuel1'].value_counts()
```

```
Out[33]: Oil          905  
Gas           2  
Cogeneration   1  
Name: other_fuel1, dtype: int64
```

```
In [34]: sns.histplot(binwidth=0.5, x="other_fuel1", hue="other_fuel1", data=df, stat="count")
plt.show()
```



Oil fuel used most of all fuel

Commission Year column


```
In [35]: df['commissioning_year'].value_counts()
```

```
Out[35]: 1996.876894    380
2013.000000     28
2015.000000     26
2012.000000     23
2016.000000     21
...
1949.000000      1
1958.000000      1
1954.000000      1
1956.000000      1
1927.000000      1
Name: commissioning_year, Length: 74, dtype: int64
```

```
In [36]: y = df.groupby('commissioning_year')['primary_fuel'].sum()  
y
```

```
Out[36]: commissioning_year  
1927.0                               Hydro  
1937.0                               Hydro  
1939.0                               Hydro  
1943.0                               Hydro  
1946.0                               Hydro  
...  
2014.0   CoalCoalCoalCoalGasSolarSolarCoalCoalSolar...  
2015.0   CoalCoalSolarHydroCoalCoalGasCoalHydroHydr...  
2016.0   CoalCoalCoalGasCoalCoalHydroGasGasCoalCoal...  
2017.0   CoalCoalHydroHydroCoalCoalHydroHydroHydroCoalC...  
2018.0                               HydroCoalCoal  
Name: primary_fuel, Length: 74, dtype: object
```

```
In [37]: df.plot.scatter(x = 'primary_fuel', y= 'commissioning_year', c = 'r')  
plt.ylabel('Commissioning Year', fontsize = 14, color = 'b')  
plt.xlabel('Primary Fuel', fontsize = 14, color = 'b')  
plt.title('Commissioning Year vs Primary Fuel', fontsize = 14, color = 'b')  
plt.show()
```



Above plot shows early years 1927 Hydro and Coal Fuel used but Now a days we used Solar and Gas Fuel most


```
In [38]: df['owner'].value_counts()
```

```
Out[38]: Acc Acc ltd                570
Sterling Agro Industries ltd.        4
Jk Cement ltd                      4
Tata Power Solar Systems Limited (TPREL) 3
Ujaas Energy Limited                3
...
Gm Energy ltd                      1
Dcm & chem                         1
REI Agro Limited                   1
National And paper                 1
Spr Pvt ltd                        1
Name: owner, Length: 280, dtype: int64
```

```
In [39]: o = df.groupby('owner')['primary_fuel'].sum()
o
```

```
Out[39]: owner
ACME Cleantech Solutions Private Limited      Solar
ACME Solar Energy                            Solar
AES                                         Wind
AEW Infratech Private Limited                Solar
Abellon CleanEnergy Limited                  Solar
...
West Coast Paper Mills Ltd.                  Gas
Yashwantrao Krishna ssk                     Biomass
Ym Ssk ltd                                  Biomass
Zamil New Delhi Infrastructure Private Limited Solar
iEnergy Wind Farms                          Wind
Name: primary_fuel, Length: 280, dtype: object
```

```
In [40]: df.plot.scatter(y = 'owner', x = 'primary_fuel', figsize = (10,50), rot = 360, for  
plt.grid(c = 'black')  
plt.show()
```




Above plot show most of the company used solar energy

Geolocation Source column

```
In [41]: g = df['geolocation_source'].value_counts()
```

```
In [42]: sns.histplot(binwidth=0.5, x="geolocation_source", data=df, stat="count", multip]plt.show()
```




Above plot shows WRI highest among all

Co-ordinates

```
In [43]: import geopandas as gpd
from shapely.geometry import Point, Polygon
import descartes
from geopandas import GeoDataFrame
from pyproj import CRS
```

```
In [44]: geometry = [Point(xy) for xy in zip(df['longitude'], df['latitude'])]
gdf = GeoDataFrame(df, geometry=geometry)


#this is a simple map that goes with geopandas
world = gpd.read_file(gpd.datasets.get_path('naturalearth_lowres'))
gdf.plot(ax=world.plot(figsize=(10, 6)), marker='o', color='red', markersize=15);
```



Above plot shows all coordinate belong to indian

```
In [45]: indmap = gpd.read_file(r"C:\Users\Kushal Arya\Desktop\csv file\map.shx")
fig,ax = plt.subplots(figsize = (15,15))
indmap.plot(ax = ax)
```

Out[45]: <AxesSubplot:>



```
In [46]:
```

```
crs=CRS('EPSG:4326').to_proj4()
```

```
In [47]:
```

```
geometry = [Point(xy) for xy in zip(df['longitude'], df['latitude'])]
geometry[:3]
```

```
Out[47]:
```


```
[<shapely.geometry.point.Point at 0x1733acc1130>,
 <shapely.geometry.point.Point at 0x1733acc1250>,
 <shapely.geometry.point.Point at 0x1733acc1340>]
```

```
In [48]:
```

```
geo_df = gpd.GeoDataFrame(df, crs = crs, geometry = geometry)
geo_df.head()
```


```
Out[48]:
```

	country_long	name	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commissioning
0	India	ACME Solar Tower	2.5	28.1839	73.2407	Solar	Oil	2010
1	India	ADITYA CEMENT WORKS	98.0	24.7663	74.6090	Coal	Oil	1998
2	India	AES Saurashtra Windfarms	39.2	21.9038	69.3732	Wind	Oil	1998
3	India	AGARTALA GT	135.0	23.8712	91.3602	Gas	Oil	2010
4	India	AKALTARA TPP	1800.0	21.9603	82.4091	Coal	Oil	2010



In [49]: **try:**

```
fig, ax = plt.subplots(figsize = (15,15))
indmap.plot(ax = ax, alpha = 0.4, color = 'grey')
geo_df[geo_df['primary_fuel']].plot(ax = ax, markersize = 20, color = 'b', ma
geo_df[geo_df['primary_fuel']].plot(ax = ax, markersize = 20, color = '^', ma
plt.legend(prop = {'size': 15})
except:
    pass
```




In [50]:

```
import plotly
import plotly.graph_objects as go
import chart_studio.plotly as py
from plotly.offline import iplot
```

```
In [51]: data = dict(type = 'choropleth',
                  locations = df['geometry'],
                  locationmode = 'country names',
                  z = df['primary_fuel'],
                  text = df['geometry'],
                  colorscale = 'YlGnBu',
                  autocolorscale = False,
                  )
layout = dict(geo = {'scope':'asia'})

diagram = go.Figure(data = [data], layout = layout)
iplot(diagram)
```



```
In [52]: from geopy.geocoders import Nominatim
```

```
In [53]: # initialize Nominatim API
geolocator = Nominatim(user_agent="geoapiExercises")
```

```
In [54]: # Latitude & Longitude input
Latitude = "28.1839"
Longitude = "73.2407"

location = geolocator.reverse(Latitude+","+Longitude)

address = location.raw['address']

# traverse the data
city = address.get('city', '')
state = address.get('state', '')
country = address.get('country', '')
code = address.get('country_code')
zipcode = address.get('postcode')
print('City : ', city)
print('State : ', state)
print('Country : ', country)
print('Zip Code : ', zipcode)
```

City :
 State : Rajasthan
 Country : India
 Zip Code : None

Drop Columns

```
In [55]: col = ['geometry', 'generation_data_source', 'source', 'owner', 'name', 'country_
```

```
In [56]: df = df.drop(col, axis = 1)
df.head(2)
```

Out[56]:

	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commissioning_year	year_of_capacity
0	2.5	28.1839	73.2407	Solar	Oil	2011.000000	
1	98.0	24.7663	74.6090	Coal	Oil	1996.876894	

```
In [57]: print('After dropping no of Rows and Columns ---->', df.shape )
```

After dropping no of Rows and Columns ----> (908, 12)

Encoding columns

```
In [58]: le = LabelEncoder()
```

```
In [59]: df['primary_fuel'] = le.fit_transform(df['primary_fuel'])
```

```
In [60]: df['other_fuel1'] = le.fit_transform(df['other_fuel1'])
```

```
In [61]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 908 entries, 0 to 907
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   capacity_mw      908 non-null    float64
 1   latitude          908 non-null    float64
 2   longitude         908 non-null    float64
 3   primary_fuel      908 non-null    int32  
 4   other_fuel1       908 non-null    int32  
 5   commissioning_year 908 non-null    float64
 6   year_of_capacity_data 908 non-null    float64
 7   generation_gwh_2013 908 non-null    float64
 8   generation_gwh_2014 908 non-null    float64
 9   generation_gwh_2015 908 non-null    float64
 10  generation_gwh_2016 908 non-null    float64
 11  generation_gwh_2017 908 non-null    float64
dtypes: float64(10), int32(2)
memory usage: 78.2 KB
```

```
In [62]: df.head(2)
```

Out[62]:

	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commissioning_year	year_of_capacity_data
0	2.5	28.1839	73.2407		6	2	2011.000000
1	98.0	24.7663	74.6090		1	2	1996.876894

All columns are encoded


Data distribution and checking outliers and skewness

```
In [63]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in df:
    if plotnumber <=12:
        ax = plt.subplot(4,3, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----
Distribution Plot :-
```



```
In [64]: df.skew()
```

```
Out[64]: capacity_mw            3.193257
latitude                  -0.147391
longitude                 1.129836
primary_fuel              0.471141
other_fuel1              -20.464435
commissioning_year        -1.383330
year_of_capacity_data    0.000000
generation_gwh_2013       5.241491
generation_gwh_2014       5.041961
generation_gwh_2015       5.367370
generation_gwh_2016       5.071758
generation_gwh_2017       5.111938
dtype: float64
```

Data has outliers and skewness

Corelation of Feature vs Label using Heat map

```
In [65]: print('-----')
```

```
print('Heat Map :-')
```

```
print('-----')
```

```
df_corr = df.corr().abs()
```

```
plt.figure(figsize = (22,16))
```


```
sns.heatmap(df_corr, vmin = -1, annot = True, square = True, center = 0, fmt = '.2f')
```

```
plt.tight_layout()
```


```
-----
```

```
Heat Map :-
```

```
-----
```



```
year_of_capacity_data ~
```



generation_gwh_2013,generation_gwh_2014 higest relation

Removing Outliers using Zscore

In [66]: # with std 3 Lets see the stats

```
z_score = zscore(df[['capacity_mw', 'longitude', 'generation_gwh_2013', 'generation_gwh_2014']])
abs_z_score = np.abs(z_score)

filtering_entry = (abs_z_score < 3).all(axis = 1)

df = df[filtering_entry]

df.describe()
```

Out[66]:

	capacity_mw	latitude	longitude	primary_fuel	other_fuel1	commissioning_year	year_constructed
count	867.000000	867.000000	867.000000	867.000000	867.000000	867.000000	867.000000
mean	262.860970	21.093784	77.025306	3.267589	1.995386	1996.778942	1996.778942
std	416.438061	6.157370	4.200192	2.303707	0.083109	12.996286	12.996286
min	0.000000	8.168900	68.644700	0.000000	0.000000	1927.000000	1927.000000
25%	16.500000	16.899050	74.318150	1.000000	2.000000	1996.876894	1996.876894
50%	50.400000	21.196189	76.737000	3.000000	2.000000	1996.876894	1996.876894
75%	317.500000	25.134100	78.906250	6.000000	2.000000	2002.000000	2002.000000
max	2000.000000	34.649000	91.565000	7.000000	2.000000	2018.000000	2018.000000

Checking Outliers and skewness removed or not


```
In [67]: # Let's see outliers are removed in columns or not.
```

```
print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in df:
    if plotnumber <=12:
        ax = plt.subplot(4,3, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----  
Distribution Plot :-  
-----
```



Outliers are removed


Splitting Dataset into features and labels

```
In [68]: x = df.drop('primary_fuel', axis = 1)
y = df. primary_fuel
print('Data has been splited')
```

Data has been splited

```
In [69]: # Let's see relation between features and labels.  
print('-----')  
print('Scatter Plot :-')  
print('-----')  
  
plt.figure(figsize = (20,25), facecolor = 'yellow')  
plotnumber = 1  
for column in x:  
    if plotnumber <=12:  
        ax = plt.subplot(4,3, plotnumber)  
        plt.scatter(x[column],y, c = 'g')  
        plt.xlabel(column, fontsize = 18)  
        plt.ylabel('Primary Fuel', fontsize = 18)  
    plotnumber += 1  
plt.tight_layout()
```

Scatter Plot :-



Features are related to label

Checking for class imbalance

```
In [70]: df['primary_fuel'].value_counts()
```

```
Out[70]: 1    237  
3    236  
6    126  
7    123  
2     65  
0     50  
5     21  
4      9  
Name: primary_fuel, dtype: int64
```

Class are not balanced

Handling Class Imbalance

```
In [71]: from imblearn.over_sampling import RandomOverSampler  
ros = RandomOverSampler()  
x_over, y_over = ros.fit_resample(x, y)
```

```
In [72]: print('-----')  
print('Class are balanced :-')  
print('-----')  
print(y_over.value_counts())  
print('-----')
```

```
-----  
Class are balanced :-  
-----  
0    237  
1    237  
2    237  
3    237  
4    237  
5    237  
6    237  
7    237  
Name: primary_fuel, dtype: int64  
-----
```

Data Scaling

```
In [73]: scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
Out[73]: array([[-0.62557022,  1.15214913, -0.90157571, ...,  0.20188562,
       0.19589199,  0.19330316],
      [-0.39611203,  0.5967866 , -0.57561684, ...,  0.20188562,
       0.19589199,  0.19330316],
      [-0.537391 ,  0.13162823, -1.82289847, ...,  0.20188562,
       0.19589199,  0.19330316],
      ...,
      [-0.57030804, -0.94542669, -0.34404144, ...,  0.20188562,
       0.19589199,  0.19330316],
      [-0.4393607 ,  0.52913754, -0.78079727, ...,  0.20188562,
       0.19589199,  0.19330316],
      [-0.59193237, -1.81340815,  0.1075556 , ...,  0.20188562,
       0.19589199,  0.19330316]])
```

Data has been scaled

Split data into train and test. Model will be bulit on training data and tested on test data

```
In [74]: x_train, x_test, y_train, y_test = train_test_split(x_over, y_over, test_size = 0.2)
print('Data has been splited.')
Data has been splited.
```

Model Bulding

Decision Tree model instantiaing, training and evaluating

```
In [76]: DT = DecisionTreeClassifier()
DT.fit(x_train, y_train)
y_pred = DT.predict(x_test)
```

```
In [77]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

	precision	recall	f1-score	support
0	0.95	1.00	0.98	62
1	0.82	0.66	0.73	61
2	0.90	1.00	0.95	53
3	0.85	0.70	0.76	56
4	0.94	1.00	0.97	67
5	0.93	1.00	0.96	52
6	0.94	0.99	0.96	68
7	0.89	0.93	0.91	55
accuracy			0.91	474
macro avg	0.90	0.91	0.90	474
weighted avg	0.90	0.91	0.90	474

```
-----
```

Conclusion : Decision Tree model has 91% score

Cross Validation score to check if the model is overfitting


```
In [81]: cv = cross_val_score(DT, x, y, cv = 5)
print('Cross Validation score of Decision Tree model --->', cv.mean())
```

Cross Validation score of Decision Tree model ---> 0.7047172945319249

Conclusion : Decision Tree model has 70% Cross Validation score

ROC, AUC Curve

```
In [85]: prob = DT.predict_proba(x_test) # calculating probability
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))
plt.show()
```



XGBoost model instantiaing, training and evaluating

```
In [96]: xgb = xgb.XGBClassifier(eval_metric = 'mlogloss')
xgb.fit(x_train, y_train)
y_pred = xgb.predict(x_test)
```

```
In [97]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

	precision	recall	f1-score	support
0	0.93	1.00	0.96	62
1	0.82	0.75	0.79	61
2	0.86	0.96	0.91	53
3	0.92	0.79	0.85	56
4	0.99	1.00	0.99	67
5	0.98	1.00	0.99	52
6	0.97	0.96	0.96	68
7	0.93	0.95	0.94	55
accuracy			0.93	474
macro avg	0.92	0.93	0.92	474
weighted avg	0.93	0.93	0.92	474

```
-----
```

Conclusion : XGBoost model has 93% score

Cross Validation score to check if the model is overfitting


```
In [98]: cv = cross_val_score(xgb, x, y, cv = 5)
print('Cross Validation score of Ada Boost model --->', cv.mean())
```

Cross Validation score of Ada Boost model ---> 0.778566208225367

Conclusion : XGBoost model has 77% Cross Validation score

ROC, AUC Curve

```
In [100]: prob = xgb.predict_proba(x_test) # calculating probability
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))
plt.show()
```



Knn model instantiaing, training and evaluating

```
In [101]: Knn = KNeighborsClassifier()
Knn.fit(x_train, y_train)
y_pred = Knn.predict(x_test)
```

```
In [102]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

	precision	recall	f1-score	support
0	0.75	0.97	0.85	62
1	0.67	0.51	0.58	61
2	0.69	0.81	0.75	53
3	0.72	0.46	0.57	56
4	0.91	1.00	0.95	67
5	0.85	1.00	0.92	52
6	0.91	0.63	0.75	68
7	0.54	0.67	0.60	55
accuracy			0.76	474
macro avg	0.76	0.76	0.74	474
weighted avg	0.76	0.76	0.75	474

```
-----
```

Conclusion : KNN model has 76% score

Cross Validation score to check if the model is overfitting


```
In [103]: cv = cross_val_score(Knn, x, y, cv = 5)
print('Cross Validation score of Ada Boost model --->', cv.mean())
```

Cross Validation score of Ada Boost model ---> 0.604391734768454

Conclusion : Knn model has 60% Cross Validation score

ROC, AUC Curve

```
In [104]: prob = Knn.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



Random Forest model instantiating, training and evaluating

```
In [105]: Rn = RandomForestClassifier()  
Rn.fit(x_train, y_train)  
y_pred = Rn.predict(x_test)
```

```
In [106]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	62
1	0.82	0.80	0.81	61
2	0.89	0.96	0.93	53
3	0.89	0.75	0.82	56
4	1.00	1.00	1.00	67
5	0.96	1.00	0.98	52
6	0.97	0.97	0.97	68
7	0.91	0.91	0.91	55
accuracy			0.93	474
macro avg	0.92	0.92	0.92	474
weighted avg	0.93	0.93	0.92	474

```
-----
```

Conclusion : Random Forest model has 93% score

Cross Validation score to check if the model is overfitting


```
In [107]: cv = cross_val_score(Rn, x, y, cv = 5)
print('Cross Validation score of Ada Boost model --->', cv.mean())
```

Cross Validation score of Ada Boost model ---> 0.7854893362567272

Conclusion : Random Forest model has 78% Cross Validation score

ROC, AUC Curve

```
In [109]: prob = Rn.predict_proba(x_test) # calculating probability
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))
plt.show()
```



Looking CV score we found Random Forest has best model so we do Hyperparameter Tuning on it

```
In [115]: param_grid = {'n_estimators': [100,200,300,400,500],
                  'max_features': ['auto', 'sqrt'],
                  'max_depth': [5, 10, 15, 20, 25, 30],
                  'min_samples_split': [2, 5, 10, 15, 100],
                  'min_samples_leaf': [1, 2, 5, 10]}
```

```
In [116]: grid_search = GridSearchCV(estimator = Rn, param_grid = param_grid, cv = 5,n_jobs
```

```
In [117]: grid_search.fit(x_train, y_train)
```

```
Out[117]: GridSearchCV(cv=5, estimator=RandomForestClassifier(), n_jobs=-1,
                       param_grid={'max_depth': [5, 10, 15, 20, 25, 30],
                                   'max_features': ['auto', 'sqrt'],
                                   'min_samples_leaf': [1, 2, 5, 10],
                                   'min_samples_split': [2, 5, 10, 15, 100],
                                   'n_estimators': [100, 200, 300, 400, 500]})
```

```
In [118]: best_parameters = grid_search.best_params_
print(best_parameters)
```

```
{'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 400}
```

```
In [120]: hRn = RandomForestClassifier(max_depth = 20, max_features = 'auto', min_samples_
```

hRn.fit(x_train, y_train)
hRn.score(x_test, y_test)

```
Out[120]: 0.919831223628692
```

```
In [122]: y_pred = hRn.predict(x_test)
```

```
In [123]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

Classification Report:				
	precision	recall	f1-score	support
0	0.93	1.00	0.96	62
1	0.80	0.77	0.78	61
2	0.88	0.96	0.92	53
3	0.91	0.73	0.81	56
4	1.00	1.00	1.00	67
5	0.96	1.00	0.98	52
6	0.97	0.97	0.97	68
7	0.89	0.91	0.90	55
accuracy			0.92	474
macro avg	0.92	0.92	0.92	474
weighted avg	0.92	0.92	0.92	474

```
-----
```

After Hyperparameter Tuning model accuracy score 92%.

Saving The Model

```
In [124]: # saving the model to the Local file system  
filename = 'Global Power Plant Project (Fuel Type).pickle'  
pickle.dump(hRn, open(filename, 'wb'))
```

Final Conclusion : Random Forest is our best model.

```
In [ ]:
```