

Problem Statement:

Business case: Insurance fraud is a huge problem in the industry. It's difficult to identify fraud claims. Machine Learning is in a unique position to help the Auto Insurance industry with this problem. In this project, you are provided a dataset which has the details of the insurance policy along with the customer details. It also has the details of the accident on the basis of which the claims have been made.

Importing Required Library

```
In [204]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.stats import zscore
import scikitplot as skplt
from imblearn.over_sampling import SMOTE
import pickle
pd.set_option('display.max_columns',None) # For display maximum columns
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.preprocessing import MinMaxScaler, OrdinalEncoder
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, plot_roc_curve
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier, BaggingClassifier
from sklearn.linear_model import LogisticRegression
import xgboost as xgb
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

Reading Data

```
In [2]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\csv file\Automobile_insurance_fraud.csv")
df.head()
```

Out[2]:

	months_as_customer	age	policy_number	policy_bind_date	policy_state	policy_csl	policy_deductible
0	328	48	521585	17-10-2014	OH	250/500	
1	228	42	342868	27-06-2006	IN	250/500	
2	134	29	687698	06-09-2000	OH	100/300	
3	256	41	227811	25-05-1990	IL	250/500	
4	228	44	367455	06-06-2014	IL	500/1000	

Check no of row and column

```
In [3]: print('No of Rows and Columns ----->', df.shape )
```

No of Rows and Columns -----> (1000, 40)

Checking for Null values

```
In [4]: print('=====\\n')
print(df.isnull().sum())
print('\\n=====')
```

```
=====
months_as_customer          0
age                          0
policy_number                0
policy_bind_date              0
policy_state                  0
policy_csl                     0
policy_deductable              0
policy_annual_premium            0
umbrella_limit                  0
insured_zip                      0
insured_sex                      0
insured_education_level            0
insured_occupation                  0
insured_hobbies                  0
insured_relationship                  0
capital-gains                      0
capital-loss                      0
incident_date                      0
incident_type                      0
collision_type                   178
incident_severity                  0
authorities_contacted                  0
incident_state                      0
incident_city                      0
incident_location                  0
incident_hour_of_the_day            0
number_of_vehicles_involved            0
property_damage                   360
bodily_injuries                      0
witnesses                        0
police_report_available            343
total_claim_amount                  0
injury_claim                      0
property_claim                      0
vehicle_claim                      0
auto_make                          0
auto_model                          0
auto_year                          0
fraud_reported                      0
_c39                           1000
dtype: int64
=====
```

There is null value

Information about dataset

```
In [5]: print('=====\\n')
print(df.info())
print('=====')
```

```
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 40 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   months_as_customer    1000 non-null   int64  
 1   age                  1000 non-null   int64  
 2   policy_number        1000 non-null   int64  
 3   policy_bind_date     1000 non-null   object  
 4   policy_state         1000 non-null   object  
 5   policy_csl           1000 non-null   object  
 6   policy_deductable   1000 non-null   int64  
 7   policy_annual_premium 1000 non-null   float64 
 8   umbrella_limit       1000 non-null   int64  
 9   insured_zip          1000 non-null   int64  
 10  insured_sex          1000 non-null   object  
 11  insured_education_level 1000 non-null   object  
 12  insured_occupation   1000 non-null   object  
 13  insured_hobbies      1000 non-null   object  
 14  insured_relationship 1000 non-null   object  
 15  capital-gains       1000 non-null   int64  
 16  capital-loss         1000 non-null   int64  
 17  incident_date        1000 non-null   object  
 18  incident_type        1000 non-null   object  
 19  collision_type       822 non-null    object  
 20  incident_severity    1000 non-null   object  
 21  authorities_contacted 1000 non-null   object  
 22  incident_state        1000 non-null   object  
 23  incident_city         1000 non-null   object  
 24  incident_location     1000 non-null   object  
 25  incident_hour_of_the_day 1000 non-null   int64  
 26  number_of_vehicles_involved 1000 non-null   int64  
 27  property_damage       640 non-null    object  
 28  bodily_injuries       1000 non-null   int64  
 29  witnesses             1000 non-null   int64  
 30  police_report_available 657 non-null    object  
 31  total_claim_amount    1000 non-null   int64  
 32  injury_claim          1000 non-null   int64  
 33  property_claim         1000 non-null   int64  
 34  vehicle_claim          1000 non-null   int64  
 35  auto_make              1000 non-null   object  
 36  auto_model              1000 non-null   object  
 37  auto_year               1000 non-null   int64  
 38  fraud_reported         1000 non-null   object  
 39  _c39                   0 non-null    float64 
dtypes: float64(2), int64(17), object(21)
memory usage: 312.6+ KB
None
=====
```

Categorical data present in our data set

Drop unwanted column

```
In [6]: df = df.drop('_c39', axis = 1)  
df.head(2)
```

Out[6]:

	months_as_customer	age	policy_number	policy_bind_date	policy_state	policy_csl	policy_ded
0	328	48	521585	17-10-2014	OH	250/500	
1	228	42	342868	27-06-2006	IN	250/500	

We drop the '_c39' column

Fill NaN

```
In [7]: df = df.apply(lambda x:x.fillna(x.mean()))if x.dtype == 'int64' else x.fillna(x.v
```

```
In [8]: print('=====\\n')
print(df.isnull().sum())
print('\\n=====')
```

```
=====
```

```
months_as_customer          0
age                          0
policy_number                0
policy_bind_date              0
policy_state                  0
policy_csl                     0
policy_deductable              0
policy_annual_premium           0
umbrella_limit                  0
insured_zip                      0
insured_sex                      0
insured_education_level           0
insured_occupation                 0
insured_hobbies                   0
insured_relationship                 0
capital-gains                      0
capital-loss                      0
incident_date                      0
incident_type                      0
collision_type                      0
incident_severity                   0
authorities_contacted                 0
incident_state                      0
incident_city                      0
incident_location                   0
incident_hour_of_the_day             0
number_of_vehicles_involved           0
property_damage                      0
bodily_injuries                      0
witnesses                         0
police_report_available                 0
total_claim_amount                   0
injury_claim                         0
property_claim                         0
vehicle_claim                         0
auto_make                           0
auto_model                            0
auto_year                            0
fraud_reported                        0
dtype: int64
```

```
=====
```

There is no null value left

Statistics of Data

```
In [9]: df.describe()
```

Out[9]:

	months_as_customer	age	policy_number	policy_deductable	policy_annual_premium
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	203.954000	38.948000	546238.648000	1136.000000	1256.406150
std	115.113174	9.140287	257063.005276	611.864673	244.167398
min	0.000000	19.000000	100804.000000	500.000000	433.330000
25%	115.750000	32.000000	335980.250000	500.000000	1089.607500
50%	199.500000	38.000000	533135.000000	1000.000000	1257.200000
75%	276.250000	44.000000	759099.750000	2000.000000	1415.695000
max	479.000000	64.000000	999435.000000	2000.000000	2047.590000

Features Engineering

Incident Date column

```
In [10]: df.head(2)
```

Out[10]:

	months_as_customer	age	policy_number	policy_bind_date	policy_state	policy_csl	policy_ded
0	328	48	521585	17-10-2014	OH	250/500	
1	228	42	342868	27-06-2006	IN	250/500	

```
In [11]: df['incident_date'].value_counts()
```

```
Out[11]: 02-02-2015    28  
17-02-2015    26  
07-01-2015    25  
10-01-2015    24  
04-02-2015    24  
24-01-2015    24  
19-01-2015    23  
08-01-2015    22  
30-01-2015    21  
13-01-2015    21  
22-02-2015    20  
31-01-2015    20  
12-02-2015    20  
06-02-2015    20  
01-01-2015    19  
14-01-2015    19  
21-01-2015    19  
12-01-2015    19  
23-02-2015    19  
21-02-2015    19  
01-02-2015    18  
03-01-2015    18  
14-02-2015    18  
20-01-2015    18  
25-02-2015    18  
18-01-2015    18  
28-02-2015    18  
08-02-2015    17  
24-02-2015    17  
26-02-2015    17  
06-01-2015    17  
09-01-2015    17  
16-02-2015    16  
15-02-2015    16  
16-01-2015    16  
13-02-2015    16  
05-02-2015    16  
28-01-2015    15  
15-01-2015    15  
17-01-2015    15  
18-02-2015    15  
20-02-2015    14  
22-01-2015    14  
27-02-2015    14  
09-02-2015    13  
03-02-2015    13  
27-01-2015    13  
23-01-2015    13  
01-03-2015    12  
04-01-2015    12  
26-01-2015    11  
29-01-2015    11  
02-01-2015    11  
11-02-2015    10
```

```
07-02-2015    10
10-02-2015    10
25-01-2015    10
19-02-2015    10
11-01-2015     9
05-01-2015     7
Name: incident_date, dtype: int64
```

```
In [12]: df['IncidentMonth&Year'] = df['incident_date'].str[3:]
df.head(2)
```

Out[12]:

	months_as_customer	age	policy_number	policy_bind_date	policy_state	policy_csl	policy_ded
0	328	48	521585	17-10-2014	OH	250/500	
1	228	42	342868	27-06-2006	IN	250/500	

```
In [13]: df['IncidentMonth&Year'].value_counts()
```

```
Out[13]: 01-2015      516
02-2015      472
03-2015       12
Name: IncidentMonth&Year, dtype: int64
```

Policy Bind Date column

```
In [14]: df['policy_bind_date'].value_counts()
```

```
Out[14]: 05-08-1992    3
28-04-1992    3
01-01-2006    3
22-08-1991    2
07-07-1996    2
..
11-06-2008    1
11-12-1994    1
19-06-2008    1
16-03-1998    1
19-04-2002    1
Name: policy_bind_date, Length: 951, dtype: int64
```

```
In [15]: df['PolicyBindMonth&year'] = df['policy_bind_date'].str[3:]
df.head(2)
```

Out[15]:

	months_as_customer	age	policy_number	policy_bind_date	policy_state	policy_csl	policy_ded
0	328	48	521585	17-10-2014	OH	250/500	
1	228	42	342868	27-06-2006	IN	250/500	

```
In [16]: df['PolicyBindMonth&year'].value_counts()
```

```
Out[16]: 11-1991    9
07-1996    8
03-2007    8
08-1994    8
12-1995    8
..
02-2012    1
07-1998    1
04-1997    1
01-2000    1
09-1998    1
Name: PolicyBindMonth&year, Length: 286, dtype: int64
```

```
In [17]: df['Incident_Pincode'] = df['incident_location'].str[0:4]
df['Incident_Pincode'] = df['Incident_Pincode'].astype('int64')
df.head(2)
```

Out[17]:

	months_as_customer	age	policy_number	policy_bind_date	policy_state	policy_csl	policy_ded
0	328	48	521585	17-10-2014	OH	250/500	
1	228	42	342868	27-06-2006	IN	250/500	

To Get Details With Pincode 

```
In [18]: import pgeocode
```

```
nomi = pgeocode.Nominatim('in')
nomi.query_postal_code('801503')
```

```
Out[18]: postal_code          801503
         country_code           IN
         place_name      Jamsaut, Sherpur, Ganghara, Sadikpur, Dalip Ch...
         state_name                Bihar
         state_code                  34
         county_name               Patna
         county_code                 230.0
         community_name            Danapur
         community_code              NaN
         latitude                      25.61
         longitude                     84.978713
         accuracy                         4
Name: 0, dtype: object
```


Analysis of Data

Months as customer column

```
In [19]: df['months_as_customer'].value_counts()
```

```
Out[19]: 194     8
254     7
210     7
101     7
140     7
..
312     1
62      1
309     1
308     1
0       1
Name: months_as_customer, Length: 391, dtype: int64
```

```
In [20]: sns.distplot(df['months_as_customer'], kde = True, hist = True)  
plt.show()
```



Months as customer has skewed


```
In [21]: df['fraud_reported'].value_counts()
```

```
Out[21]: N    753  
Y    247  
Name: fraud_reported, dtype: int64
```

```
In [22]: m = df.groupby('months_as_customer')['fraud_reported'].value_counts()  
m
```

```
Out[22]: months_as_customer  fraud_reported  
0                  N          1  
1                  N          3  
2                  N          2  
3                  N          1  
                   Y          1  
..  
475                 N          2  
476                 N          1  
478                 N          1  
                   Y          1  
479                 N          2  
Name: fraud_reported, Length: 532, dtype: int64
```

```
In [23]: sns.catplot(y = 'months_as_customer', x = 'fraud_reported', data = df)  
plt.show()
```




Above plot shows who is less period of time customer is changes of fraud is high

Age column

```
In [24]: df['age'].value_counts()
```

```
Out[24]: 43    49  
39    48  
41    45  
34    44  
38    42  
30    42  
31    42  
37    41  
33    39  
40    38  
32    38  
29    35  
46    33  
35    32  
44    32  
36    32  
42    32  
28    30  
45    26  
26    26  
48    25  
47    24  
27    24  
57    16  
55    14  
25    14  
49    14  
50    13  
53    13  
24    10  
54    10  
61    10  
51     9  
60     9  
56     8  
58     8  
23     7  
21     6  
59     5  
52     4  
62     4  
63     2  
64     2  
20     1  
22     1  
19     1  
Name: age, dtype: int64
```

```
In [25]: sns.histplot( x="age", data=df)
plt.xlabel('Age', c = 'r', fontsize = 12)
plt.ylabel('Counts', c = 'r', fontsize = 12)
plt.title('Age vs Counts', c = 'r', fontsize = 12)
plt.show()
```



Age 39 to 43 is highest counts


```
In [26]: a = df.groupby('age')['fraud_reported'].value_counts().sort_values()
a
```

```
Out[26]: age  fraud_reported
19    Y          1
20    N          1
63    N          1
22    Y          1
62    Y          1
...
34    N         31
40    N         34
30    N         36
39    N         38
43    N         38
Name: fraud_reported, Length: 86, dtype: int64
```

```
In [27]: sns.catplot(y = 'age', x = 'fraud_reported', data = df)
plt.ylabel('Age', c = 'b', fontsize = 12)
plt.xlabel('Fraud Reported', c = 'b', fontsize = 12 )
plt.title('Age vs Fraud Reported', c = 'b', fontsize = 12)
plt.show()
```



```
In [28]: a.plot.barh(figsize = (10,15), rot = 360, color = ['g','r','y','k','c','b'])
plt.ylabel('Age', c = 'g', fontsize = 12)
plt.xlabel('Fraud Reported', c = 'g', fontsize = 12 )
plt.title('Age vs Fraud Reported', c = 'g', fontsize = 12)
plt.show()
```



Age group 31 to 41 highest Fraud Reported


Policy State

In [29]: `df['policy_state'].value_counts()`

Out[29]:

OH	352
IL	338
IN	310
Name:	policy_state, dtype: int64

```
In [30]: sns.countplot( x="policy_state", data=df)
plt.xlabel('Policy State', c = 'brown', fontsize = 12)
plt.ylabel('Counts', c = 'brown', fontsize = 12)
plt.title('Policy State vs Counts', c = 'brown', fontsize = 12)
plt.show()
```




OH highest counts

Policy Deductable

```
In [31]: df['policy_deductable'].value_counts()
```

```
Out[31]: 1000    351
500     342
2000    307
Name: policy_deductable, dtype: int64
```

```
In [32]: sns.countplot( x="policy_deductable", data=df)
plt.xlabel('Policy Deductable', c = 'g', fontsize = 12)
plt.ylabel('Counts', c = 'g', fontsize = 12)
plt.title('Policy Deductable vs Counts', c = 'g', fontsize = 12)
plt.show()
```




1000 policy deductible highest counts

Insured Sex Columns

```
In [33]: df['insured_sex'].value_counts()
```

```
Out[33]: FEMALE    537
MALE      463
Name: insured_sex, dtype: int64
```

```
In [34]: sns.histplot(binwidth=0.5, x="insured_sex", hue="insured_sex", data=df, stat="count")
plt.xlabel('Insured Sex', c = 'b', fontsize = 12)
plt.ylabel('Counts', c = 'b', fontsize = 12)
plt.title('Insured Sex vs Counts', c = 'b', fontsize = 12)
plt.show()
```




Female is highest in counts

```
In [35]: s = df.groupby('insured_sex')['fraud_reported'].value_counts()
s
```

```
Out[35]: insured_sex  fraud_reported
FEMALE      N          411
              Y          126
MALE        N          342
              Y          121
Name: fraud_reported, dtype: int64
```

```
In [36]: s.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Insured Sex', c = 'g', fontsize = 12)
plt.ylabel('Fraud Reported Counts', c = 'g', fontsize = 12 )
plt.title('Insured Sex vs Fraud Reported Counts', c = 'g', fontsize = 12)
plt.xticks(c = 'r')
plt.yticks(c = 'r')
plt.show()
```




In female highest no of fraud reported

Insured Education Level counts

```
In [37]: df['insured_education_level'].value_counts()
```

```
Out[37]: JD      161  
High School  160  
Associate    145  
MD          144  
Masters     143  
PhD         125  
College     122  
Name: insured_education_level, dtype: int64
```

```
In [38]: sns.countplot( x="insured_education_level", data=df)  
plt.xlabel('Education', c = 'b', fontsize = 12)  
plt.ylabel('Counts', c = 'b', fontsize = 12)  
plt.title('Education vs Counts', c = 'b', fontsize = 12)  
plt.show()
```



In education JD is highest counts


```
In [39]: e = df.groupby('insured_education_level')['fraud_reported'].value_counts()  
e
```

```
Out[39]: insured_education_level  fraud_reported
```

Associate	N	111
	Y	34
College	N	90
	Y	32
High School	N	124
	Y	36
JD	N	119
	Y	42
MD	N	106
	Y	38
Masters	N	111
	Y	32
PhD	N	92
	Y	33

```
Name: fraud_reported, dtype: int64
```

```
In [40]: e.plot.barh(figsize = (10,15), color = ['g','r'])
plt.ylabel('Education', c = 'r', fontsize = 12)
plt.xlabel('Fraud Reported Counts', c = 'r', fontsize = 12 )
plt.title('Education vs Fraud Reported Counts', c = 'r', fontsize = 12)
plt.xticks(c = 'm')
plt.yticks(c = 'm')
plt.show()
```




JD has highest fraud reported

Insured Occupation Columns

```
In [41]: o = df['insured_occupation'].value_counts()  
o
```

```
Out[41]: machine-op-inspct    93  
prof-specialty      85  
tech-support        78  
sales                76  
exec-managerial     76  
craft-repair         74  
transport-moving     72  
other-service        71  
priv-house-serv      71  
armed-forces         69  
adm-clerical         65  
protective-serv       63  
handlers-cleaners     54  
farming-fishing       53  
Name: insured_occupation, dtype: int64
```

```
In [42]: o.plot.barh( figsize = (15,10), color = ['g','r','c','k','y', 'brown', 'm'])  
plt.xlabel('Counts', c = 'g', fontsize = 12)  
plt.ylabel('Insured Occupation', c = 'g', fontsize = 12 )  
plt.title('Insured Occupation vs Counts', c = 'g', fontsize = 12)  
plt.xticks(c = 'r')  
plt.yticks(c = 'r')  
plt.show()
```




Machine-op-inspct highest counts

```
In [43]: op = df.groupby('insured_occupation')['fraud_reported'].value_counts()
```

```
Out[43]: insured_occupation  fraud_reported
adm-clerical            N          54
                           Y          11
armed-forces             N          52
                           Y          17
craft-repair              N          52
                           Y          22
exec-managerial           N          48
                           Y          28
farming-fishing            N          37
                           Y          16
handlers-cleaners          N          43
                           Y          11
machine-op-inspct           N          71
                           Y          22
other-service              N          59
                           Y          12
priv-house-serv             N          59
                           Y          12
prof-specialty              N          67
                           Y          18
protective-serv              N          49
                           Y          14
sales                      N          55
                           Y          21
tech-support                 N          56
                           Y          22
transport-moving              N          51
                           Y          21
Name: fraud_reported, dtype: int64
```

```
In [44]: op.plot.barh(figsize = (15,10),color = ['g','r','c','k','y', 'brown', 'm'])
plt.ylabel('Insured Occupation', c = 'b', fontsize = 12)
plt.xlabel('Fraud Reported Counts', c = 'b', fontsize = 12 )
plt.title('Insured Occupation vs Fraud Reported Counts', c = 'b', fontsize = 12)
plt.show()
```




Exec-managerial occupation has highest fraud reported counts

Insured Hobbies Columns

```
In [45]: h = df['insured_hobbies'].value_counts()  
h
```

```
Out[45]: reading      64  
exercise      57  
paintball      57  
bungie-jumping 56  
movies         55  
camping        55  
golf           55  
kayaking       54  
yachting       53  
hiking          52  
video-games    50  
skydiving      49  
base-jumping    49  
board-games    48  
polo            47  
chess           46  
dancing         43  
sleeping        41  
cross-fit       35  
basketball      34  
Name: insured_hobbies, dtype: int64
```

```
In [46]: h.plot.bar(figsize = (10,5), rot = 45, color = [ 'g', 'r', 'c', 'k', 'y', 'brown', 'm' ]  
plt.xlabel('Hobbies', c = 'g', fontsize = 12)  
plt.ylabel('Counts', c = 'g', fontsize = 12 )  
plt.title('Hobbies vs Fraud Reported Counts', c = 'g', fontsize = 12)  
plt.xticks(c = 'r')  
plt.yticks(c = 'r')  
plt.show()
```




Reading is highest hobbies

Insured Relationship Column

```
In [47]: r = df['insured_relationship'].value_counts()  
r
```

```
Out[47]: own-child      183  
other-relative    177  
not-in-family     174  
husband           170  
wife              155  
unmarried         141  
Name: insured_relationship, dtype: int64
```

```
In [48]: r.plot.bar(figsize = (12,5), rot = 360, color = ['g','r','c','k','y', 'brown', 'm'])  
plt.xlabel('Relationship', c = 'r', fontsize = 12)  
plt.ylabel('Counts', c = 'r', fontsize = 12)  
plt.title('Relationship vs Counts', c = 'r', fontsize = 12)  
plt.show()
```




Own child has highest counts

```
In [49]: re = df.groupby('insured_relationship')['fraud_reported'].value_counts()
```

```
Out[49]: insured_relationship  fraud_reported
husband                  N          135
                           Y           35
not-in-family            N          129
                           Y           45
other-relative           N          125
                           Y           52
own-child                N          144
                           Y           39
unmarried                N          107
                           Y           34
wife                     N          113
                           Y           42
Name: fraud_reported, dtype: int64
```

```
In [50]: re.plot.bar(figsize = (18,8), rot = 360, color = ['g','r','c','k','y', 'brown',
plt.xlabel('Relationship Fraud ', c = 'r', fontsize = 12)
plt.ylabel('Counts', c = 'r', fontsize = 12)
plt.title('Relationship Fraud vs Counts', c = 'r', fontsize = 12)
plt.show()
```




Other relative highest fraud counts

Capital Gains Column

```
In [51]: df['capital-gains'].value_counts()
```

```
Out[51]: 0      508  
46300     5  
68500     4  
51500     4  
45500     3  
...  
54700     1  
40100     1  
33200     1  
37300     1  
72700     1  
Name: capital-gains, Length: 338, dtype: int64
```

```
In [52]: sns.distplot(df['capital-gains'], kde = True, hist = True)  
plt.show()
```




Capital Gains has skewed

Capital Loss Column

```
In [53]: df['capital-loss'].value_counts()
```

```
Out[53]: 0      475  
-53700    5  
-50300    5  
-31700    5  
-51000    4  
...  
-43300    1  
-66100    1  
-55900    1  
-66500    1  
-48000    1  
Name: capital-loss, Length: 354, dtype: int64
```

```
In [54]: sns.distplot(df['capital-loss'], kde = True, hist = True)  
plt.show()
```




Capital loss has skewed

Incident Type Column

```
In [55]: i = df['incident_type'].value_counts()  
i
```

```
Out[55]: Multi-vehicle Collision      419  
Single Vehicle Collision            403  
Vehicle Theft                      94  
Parked Car                         84  
Name: incident_type, dtype: int64
```

```
In [56]: i.plot.pie(figsize = (8,8), autopct = '%.1f')
plt.title('Incident Type', c = 'r', fontsize = 18)
plt.show()
```




Multi Vehicle Collision Highest Counts

```
In [57]: ig = df.groupby('incident_type')['fraud_reported'].value_counts()
ig
```

```
Out[57]: incident_type      fraud_reported
          Multi-vehicle Collision    N        305
                           Y        114
          Parked Car                 N        76
                           Y         8
          Single Vehicle Collision   N        286
                           Y        117
          Vehicle Theft              N        86
                           Y         8
Name: fraud_reported, dtype: int64
```

```
In [58]: ig.plot.bar(figsize = (20,8), rot = 360, color = ['g','r'])
plt.xlabel('Incident Type Fraud', c = 'g', fontsize = 18)
plt.ylabel('Counts', c = 'g', fontsize = 18 )
plt.title('Incident Type Fraud vs Fraud Reported Counts', c = 'g', fontsize = 18)
plt.xticks(c = 'r', fontsize = 10)
plt.yticks(c = 'r', fontsize = 12)
plt.show()
```




Single Vehicle Collision Highest Fraud Counts

Collision Type Counts

```
In [59]: df['collision_type'].value_counts()
```

```
Out[59]: Rear Collision      470
Side Collision       276
Front Collision      254
Name: collision_type, dtype: int64
```

```
In [60]: sns.histplot(binwidth=0.5, x="collision_type", hue="collision_type", data=df, stat="count")
plt.xlabel('Collision Type', c = 'b', fontsize = 12)
plt.ylabel('Counts', c = 'b', fontsize = 12)
plt.title('Collision Type vs Counts', c = 'b', fontsize = 12)
plt.show()
```




Rear Collision Has Highest Counts

```
In [61]: ct = df.groupby('collision_type')['fraud_reported'].value_counts()
ct
```

```
Out[61]: collision_type  fraud_reported
Front Collision    N          184
                  Y           70
Rear Collision     N          363
                  Y           107
Side Collision      N          206
                  Y            70
Name: fraud_reported, dtype: int64
```

```
In [62]: ct.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Collision Type Fraud', c = 'k', fontsize = 12)
plt.ylabel('Fraud Reported Counts', c = 'k', fontsize = 12 )
plt.title('Collision Type Fraud vs Fraud Reported Counts', c = 'k', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```




Rear Collision Has Highest Fraud Reported Counts

Authorities Contacted Column

```
In [63]: df['authorities_contacted'].value_counts()
```

```
Out[63]: Police      292  
Fire        223  
Other       198  
Ambulance   196  
None        91  
Name: authorities_contacted, dtype: int64
```

```
In [64]: sns.countplot(x="authorities_contacted", data=df)  
plt.xlabel('Authorities Contacted', c = 'r', fontsize = 12)  
plt.ylabel('Counts', c = 'r', fontsize = 12)  
plt.title('Authorities Contacted vs Counts', c = 'r', fontsize = 12)  
plt.show()
```




Police Has Highest Counts

```
In [65]: au = df.groupby('authorities_contacted')['fraud_reported'].value_counts()  
au
```

```
Out[65]: authorities_contacted  fraud_reported  
Ambulance           N          139  
                      Y          57  
Fire                N          163  
                      Y          60  
None               N          85  
                      Y          6  
Other               N          135  
                      Y          63  
Police              N         231  
                      Y          61  
Name: fraud_reported, dtype: int64
```

```
In [66]: au.plot.bar(figsize = (15,5), rot = 360, color = ['g','r'])
plt.xlabel('Authorities Contacted Fraud', c = 'm', fontsize = 12)
plt.ylabel('Fraud Reported Counts', c = 'm', fontsize = 12 )
plt.title('Authorities Contacted vs Fraud Reported Counts', c = 'm', fontsize = 12)
plt.show()
```




Other has highest no of fraud reported

Incident State Column

```
In [67]: ist =df['incident_state'].value_counts()
ist
```

```
Out[67]: NY    262
          SC    248
          WV    217
          VA    110
          NC    110
          PA     30
          OH     23
Name: incident_state, dtype: int64
```

```
In [68]: ist.plot.pie(figsize = (8,8), autopct = '%.1f')
plt.title('Incident State', c ='b', fontsize = 18)
plt.show()
```




NY state has highest counts

```
In [69]: sf = df.groupby('incident_state')['fraud_reported'].value_counts()  
sf
```

```
Out[69]: incident_state  fraud_reported  
NC           N            76  
              Y            34  
NY           N           204  
              Y            58  
OH           N            13  
              Y            10  
PA           N            22  
              Y             8  
SC           N           175  
              Y            73  
VA           N            85  
              Y            25  
WV           N           178  
              Y            39  
Name: fraud_reported, dtype: int64
```

```
In [70]: sf.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])  
plt.xlabel('Incident State Fraud Reported', c = 'g', fontsize = 12)  
plt.ylabel('Fraud Reported Counts', c = 'g', fontsize = 12 )  
plt.title('Incident State Fraud Reported vs Fraud Reported Counts', c = 'g', fonti  
plt.xticks(c = 'b')  
plt.yticks(c = 'b')  
plt.show()
```




SC state highest fraud reported counts

Incident City Column

```
In [71]: ic = df['incident_city'].value_counts()
ic
```

```
Out[71]: Springfield    157
Arlington      152
Columbus       149
Northbend      145
Hillsdale      141
Riverwood      134
Northbrook     122
Name: incident_city, dtype: int64
```

```
In [72]: ic.plot.bar(figsize = (10,5), rot = 360, color = ['g','r','c','m','y'])
plt.xlabel('Incident City', c = 'g', fontsize = 12)
plt.ylabel('Counts', c = 'g', fontsize = 12 )
plt.title('Incident City vs Counts', c = 'g', fontsize = 12)
plt.show()
```




Springfield City has highest counts

```
In [73]: c = df.groupby('incident_city')['fraud_reported'].value_counts()
c
```

```
Out[73]: incident_city  fraud_reported
Arlington      N          108
                  Y           44
Columbus       N          110
                  Y           39
Hillsdale      N          106
                  Y           35
Northbend      N          111
                  Y           34
Northbrook     N           95
                  Y           27
Riverwood      N          104
                  Y           30
Springfield    N          119
                  Y           38
Name: fraud_reported, dtype: int64
```

```
In [74]: c.plot.bar(figsize = (20,5), rot = 360, color = ['g','r'])
plt.xlabel('Incident City Fraud Counts', c = 'r', fontsize = 18)
plt.ylabel('Fraud Reported Counts', c = 'r', fontsize = 18 )
plt.title('Incident City Fraud Counts vs Fraud Reported Counts', c = 'r', fontsize = 18)
plt.xticks(c = 'b', fontsize = 10)
plt.yticks(c = 'b', fontsize = 10)
plt.show()
```



Arlington city highest fraud reported counts

Incident Hour Of The Day Column

```
In [75]: df['incident_hour_of_the_day'].value_counts()
```

```
Out[75]: 17    54
3     53
0     52
23    51
16    49
4     46
10    46
13    46
6     44
14    43
9     43
21    42
18    41
7     40
19    40
12    40
15    39
22    38
8     36
20    34
5     33
2     31
11    30
1     29
Name: incident_hour_of_the_day, dtype: int64
```


```
In [76]: hr = df.groupby('incident_hour_of_the_day')['fraud_reported'].value_counts().sort_index()
```

```
Out[76]: incident_hour_of_the_day  fraud_reported
```

9	Y	4
22	Y	5
8	Y	6
12	Y	8
7	Y	8
13	Y	8
17	Y	8
4	Y	8
1	Y	8
19	Y	9
20	Y	9
5	Y	9
2	Y	10
21	Y	10
0	Y	12
3	Y	12
11	Y	12
15	Y	13
18	Y	14
23	Y	14
6	Y	14
10	Y	15
14	Y	15
16	Y	16
11	N	18
1	N	21
2	N	21
5	N	24
20	N	25
15	N	26
18	N	27
14	N	28
8	N	30
6	N	30
19	N	31
10	N	31
7	N	32
21	N	32
12	N	32
16	N	33
22	N	33
23	N	37
4	N	38
13	N	38
9	N	39
0	N	40
3	N	41
17	N	46

```
Name: fraud_reported, dtype: int64
```

```
In [77]: hr.plot.bahr(figsize = (10,15), rot = 360, color = ['g','r','m','k','y','c','b'])
plt.ylabel('Incident Hour Of The Day Fraud ', c = 'b', fontsize = 18)
plt.xlabel('Fraud Reported Counts', c = 'b', fontsize = 18 )
plt.title('Incident Hour Of The Day Fraud vs Fraud Reported Counts', c = 'b', fontweight = 'bold')
plt.show()
```




In a day highest 16 fraud case reported

Number Of Vehicles Involved Column

```
In [78]: df['number_of_vehicles_involved'].value_counts()
```

```
Out[78]: 1    581
3    358
4    31
2    30
Name: number_of_vehicles_involved, dtype: int64
```

```
In [79]: sns.countplot( x="number_of_vehicles_involved", data=df)
plt.xlabel('Number of vehicles involved', c = 'r', fontsize = 12)
plt.ylabel('Counts', c = 'r', fontsize = 12)
plt.title('Number of vehicles involved vs Counts', c = 'r', fontsize = 12)
plt.show()
```




One vehicles involved fraud case highest

```
In [80]: v = df.groupby('number_of_vehicles_involved')['fraud_reported'].value_counts()  
v
```

```
Out[80]: number_of_vehicles_involved  fraud_reported  
1                           N      448  
                           Y      133  
2                           N       21  
                           Y        9  
3                           N      264  
                           Y      94  
4                           N      20  
                           Y      11  
Name: fraud_reported, dtype: int64
```

```
In [81]: v.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])  
plt.xlabel('Number of vehicles involved fraud case', c = 'g', fontsize = 12)  
plt.ylabel('Fraud Reported Counts', c = 'g', fontsize = 12 )  
plt.title('Number of vehicles involved fraud case vs Fraud Reported Counts', c =  
plt.xticks(c = 'r')  
plt.yticks(c = 'r')  
plt.show()
```




One vehicles involved fraud case highest

Property Damage

```
In [82]: df['property_damage'].value_counts()
```

```
Out[82]: NO      698  
YES     302  
Name: property_damage, dtype: int64
```

```
In [83]: sns.histplot(binwidth=0.5, x="property_damage", hue="property_damage", data=df, s  
plt.xlabel('Property Damage', c = 'b', fontsize = 12)  
plt.ylabel('Counts', c = 'b', fontsize = 12)  
plt.title('Property Damage vs Counts', c = 'b', fontsize = 12)  
plt.show()
```



Property Damage cases are low

Bodily Injuries Column

```
In [84]: df['bodily_injuries'].value_counts()
```

```
Out[84]: 0    340  
2    332  
1    328  
Name: bodily_injuries, dtype: int64
```

```
In [85]: sns.countplot( x="bodily_injuries", data=df)
plt.xlabel('Bodily Injuries', c = 'b', fontsize = 12)
plt.ylabel('Counts', c = 'b', fontsize = 12)
plt.title('Bodily Injuries vs Counts', c = 'b', fontsize = 12)
plt.show()
```




Zero has highest counts

Witnesses Column

```
In [86]: df['witnesses'].value_counts()
```

```
Out[86]: 1    258
2    250
0    249
3    243
Name: witnesses, dtype: int64
```

```
In [87]: sns.countplot( x="witnesses", data=df)
plt.xlabel('Witnesses', c = 'k', fontsize = 12)
plt.ylabel('Counts', c = 'k', fontsize = 12)
plt.title('Witnesses vs Counts', c = 'k', fontsize = 12)
plt.show()
```



One witnesses counts is highest

Total Claim Amount Column

```
In [88]: print('Maximum Claim Amount----->',df['total_claim_amount'].max())
```

Maximum Claim Amount-----> 114920

```
In [89]: print('Minimum Claim Amount----->',df['total_claim_amount'].min())
```

Minimum Claim Amount-----> 100

Injury Claim Column

```
In [90]: print('Maximum Injury Claim Amount----->',df['injury_claim'].max())
```

Maximum Injury Claim Amount-----> 21450

```
In [91]: print('Minimum Injury Claim Amount----->',df['injury_claim'].min())
```

Minimum Injury Claim Amount-----> 0

Property Claim Column

```
In [92]: print('Maximum Property Claim ----->',df['property_claim'].max())
```

```
Maximum Property Claim -----> 23670
```

```
In [93]: print('Minimum Property Claim ----->',df['property_claim'].min())
```

```
Minimum Property Claim -----> 0
```

Vehicle Claim Column

```
In [94]: print('Maximum Vehicle Claim ----->',df['vehicle_claim'].max())
```

```
Maximum Vehicle Claim -----> 79560
```

```
In [95]: print('Minimum Vehicle Claim ----->',df['vehicle_claim'].min())
```


```
Minimum Vehicle Claim -----> 70
```

Auto Make Column

```
In [96]: au = df['auto_make'].value_counts()  
au
```

```
Out[96]: Dodge      80  
          Saab       80  
          Subaru     80  
          Nissan     78  
          Chevrolet  76  
          Ford       72  
          BMW        72  
          Toyota     70  
          Audi       69  
          Accura     68  
          Volkswagen 68  
          Jeep        67  
          Mercedes   65  
          Honda      55  
          Name: auto_make, dtype: int64
```

```
In [97]: au.plot.bar(figsize = (15,5), rot = 360, color = ['g','r','b','c','y','m'])
plt.xlabel('Vehicles Company', c = 'g', fontsize = 12)
plt.ylabel('Counts', c = 'g', fontsize = 12 )
plt.title('Vehicles Company vs Counts', c = 'g', fontsize = 12)
plt.xticks(c = 'r')
plt.yticks(c = 'r')
plt.show()
```




Dodge is highest counts and Honda is lowest

```
In [98]: am = df.groupby('auto_make')['fraud_reported'].value_counts().sort_values()
```

```
Out[98]: auto_make    fraud_reported
Jeep          Y            11
Accura        Y            13
Toyota        Y            13
Honda         Y            14
Nissan        Y            14
Saab          Y            18
Suburu        Y            19
Volkswagen   Y            19
Dodge          Y            20
BMW           Y            20
Audi           Y            21
Chevrolet    Y            21
Mercedes      Y            22
Ford           Y            22
Honda         N            41
Mercedes      N            43
Audi           N            48
Volkswagen   N            49
Ford           N            50
BMW           N            52
Accura         N            55
Chevrolet    N            55
Jeep           N            56
Toyota         N            57
Dodge          N            60
Suburu        N            61
Saab           N            62
Nissan         N            64
Name: fraud_reported, dtype: int64
```

```
In [99]: am.plot.barh(figsize = (10,15), rot = 360, color = ['g','r'])
plt.ylabel('Vehicle Company Fraud', c = 'g', fontsize = 12)
plt.xlabel('Fraud Reported Counts', c = 'g', fontsize = 12 )
plt.title('Vehhicle Company Fraud vs Fraud Reported Counts', c = 'g', fontsize =
plt.show()
```




Mercedes and Ford has highest fraud reported

Auto Year Column

```
In [100]: y = df['auto_year'].value_counts()  
y
```

```
Out[100]: 1995      56  
1999      55  
2005      54  
2006      53  
2011      53  
2007      52  
2003      51  
2010      50  
2009      50  
2013      49  
2002      49  
2015      47  
1997      46  
2012      46  
2008      45  
2014      44  
2001      42  
2000      42  
1998      40  
2004      39  
1996      37  
Name: auto_year, dtype: int64
```

```
In [101]: y.plot.bar(figsize = (15,5), rot = 360, color = ['g','r','k','c','b','y','m'])
plt.xlabel('Auto Year', c = 'r', fontsize = 12)
plt.ylabel('Counts', c = 'r', fontsize = 12 )
plt.title('Auto Year vs Counts', c = 'r', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```



1995 highest sell of auto


```
In [102]: ay = df.groupby('auto_year')['fraud_reported'].value_counts().sort_values()
```

```
Out[102]: auto_year  fraud_reported
```

2010	Y	7
1998	Y	7
2012	Y	9
2003	Y	9
2001	Y	9
2002	Y	10
2008	Y	10
1999	Y	10
2015	Y	11
2000	Y	11
2009	Y	11
2014	Y	12
1997	Y	12
2005	Y	12
1995	Y	13
1996	Y	14
2006	Y	14
2013	Y	15
2004	Y	16
2011	Y	17
2007	Y	18
2004	N	23
1996	N	23
2000	N	31
2014	N	32
2001	N	33
1998	N	33
2007	N	34
1997	N	34
2013	N	34
2008	N	35
2011	N	36
2015	N	36
2012	N	37
2006	N	39
2002	N	39
2009	N	39
2005	N	42
2003	N	42
2010	N	43
1995	N	43
1999	N	45

```
Name: fraud_reported, dtype: int64
```

```
In [103]: ay.plot.barh(figsize = (10,15), rot = 360, color = ['g','r'])
plt.ylabel('Auto Year Fraud', c = 'g', fontsize = 12)
plt.xlabel('Fraud Reported Counts', c = 'g', fontsize = 12 )
plt.title('Auto Year Fraud vs Fraud Reported Counts', c = 'g', fontsize = 12)
plt.xticks(c = 'r')
plt.yticks(c = 'r')
plt.show()
```




Year 2007 highest fraud reported counts

Incident Month & Year Column

```
In [104]: df['IncidentMonth&Year'].value_counts()
```

```
Out[104]: 01-2015    516
           02-2015    472
           03-2015     12
Name: IncidentMonth&Year, dtype: int64
```

```
In [105]: sns.countplot( x="IncidentMonth&Year", data=df)
plt.xlabel('Incident Month & Year', c = 'r', fontsize = 12)
plt.ylabel('Counts', c = 'r', fontsize = 12)
plt.title('Incident Month & Year vs Counts', c = 'r', fontsize = 12)
plt.show()
```




In 2015 first month highest incident counts

```
In [106]: imy = df.groupby('IncidentMonth&Year')['fraud_reported'].value_counts()  
imy
```

```
Out[106]: IncidentMonth&Year    fraud_reported  
01-2015          N            381  
                  Y            135  
02-2015          N            364  
                  Y            108  
03-2015          N             8  
                  Y             4  
Name: fraud_reported, dtype: int64
```

```
In [107]: imy.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])  
plt.xlabel('Incident Month & Year Fraud', c = 'g', fontsize = 12)  
plt.ylabel('Fraud Reported Counts', c = 'g', fontsize = 12 )  
plt.title('Incident Month & Year Counts vs Fraud Reported Counts', c = 'g', fontst  
plt.xticks(c = 'r')  
plt.yticks(c = 'r')  
plt.show()
```



In 2015 first month highest fraud reported counts


Drop Columns

```
In [108]: col = ['policy_bind_date', 'incident_date', 'incident_location']
```

```
In [109]: df = df.drop(col, axis = 1)
df.head(2)
```

Out[109]:

	months_as_customer	age	policy_number	policy_state	policy_csl	policy_deductable	policy_an
0	328	48	521585	OH	250/500	1000	
1	228	42	342868	IN	250/500	2000	



```
In [110]: print('No of Rows and Columns after removing columns ----->', df.shape )
```

No of Rows and Columns after removing columns -----> (1000, 39)

Encoding Categorical Column

```
In [111]: oe = OrdinalEncoder()
```

```
In [112]: for i in df.columns:
    if df[i].dtypes == 'object':
        df[i] = oe.fit_transform(df[i].values.reshape(-1,1))
df.head(2)
```

Out[112]:

	months_as_customer	age	policy_number	policy_state	policy_csl	policy_deductable	policy_an
0	328	48	521585	2.0	1.0	1000	
1	228	42	342868	1.0	1.0	2000	



```
In [113]: print('=====\\n')
print(df.info())
print('=====')
```

```
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 39 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   months_as_customer    1000 non-null   int64  
 1   age                  1000 non-null   int64  
 2   policy_number        1000 non-null   int64  
 3   policy_state         1000 non-null   float64 
 4   policy_csl           1000 non-null   float64 
 5   policy_deductable   1000 non-null   int64  
 6   policy_annual_premium 1000 non-null   float64 
 7   umbrella_limit       1000 non-null   int64  
 8   insured_zip          1000 non-null   int64  
 9   insured_sex          1000 non-null   float64 
 10  insured_education_level 1000 non-null   float64 
 11  insured_occupation   1000 non-null   float64 
 12  insured_hobbies      1000 non-null   float64 
 13  insured_relationship 1000 non-null   float64 
 14  capital-gains       1000 non-null   int64  
 15  capital-loss         1000 non-null   int64  
 16  incident_type        1000 non-null   float64 
 17  collision_type       1000 non-null   float64 
 18  incident_severity    1000 non-null   float64 
 19  authorities_contacted 1000 non-null   float64 
 20  incident_state        1000 non-null   float64 
 21  incident_city         1000 non-null   float64 
 22  incident_hour_of_the_day 1000 non-null   int64  
 23  number_of_vehicles_involved 1000 non-null   int64  
 24  property_damage       1000 non-null   float64 
 25  bodily_injuries       1000 non-null   int64  
 26  witnesses             1000 non-null   int64  
 27  police_report_available 1000 non-null   float64 
 28  total_claim_amount    1000 non-null   int64  
 29  injury_claim          1000 non-null   int64  
 30  property_claim        1000 non-null   int64  
 31  vehicle_claim         1000 non-null   int64  
 32  auto_make              1000 non-null   float64 
 33  auto_model             1000 non-null   float64 
 34  auto_year              1000 non-null   int64  
 35  fraud_reported        1000 non-null   float64 
 36  IncidentMonth&Year    1000 non-null   float64 
 37  PolicyBindMonth&year   1000 non-null   float64 
 38  Incident_Pincode      1000 non-null   int64  
dtypes: float64(21), int64(18)
memory usage: 304.8 KB
None
=====
```

All categorical columns are encoded


Data distribution

```
In [114]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in df:
    if plotnumber <=40:
        ax = plt.subplot(8,5, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----
Distribution Plot :-
```



Some columns are little skewed

Check skewness

```
In [115]: print('=====')  
print(df.skew())  
print('=====')
```

```
=====  
months_as_customer      0.362177  
age                     0.478988  
policy_number           0.038991  
policy_state            -0.026177  
policy_csl              0.088928  
policy_deductable       0.477887  
policy_annual_premium   0.004402  
umbrella_limit          1.806712  
insured_zip             0.816554  
insured_sex              0.148630  
insured_education_level -0.000148  
insured_occupation      -0.058881  
insured_hobbies          -0.061563  
insured_relationship     0.077488  
capital-gains           0.478850  
capital-loss             -0.391472  
incident_type            0.101507  
collision_type           -0.033682  
incident_severity        0.279016  
authorities_contacted   -0.121744  
incident_state           -0.148865  
incident_city             0.049531  
incident_hour_of_the_day -0.035584  
number_of_vehicles_involved 0.502664  
property_damage           0.863806  
bodily_injuries          0.014777  
witnesses                0.019636  
police_report_available  0.802728  
total_claim_amount        -0.594582  
injury_claim              0.264811  
property_claim            0.378169  
vehicle_claim             -0.621098  
auto_make                 -0.018797  
auto_model                -0.080773  
auto_year                  -0.048289  
fraud_reported            1.175051  
IncidentMonth&Year       0.267378  
PolicyBindMonth&year      -0.001113  
Incident_Pincode          0.002897  
dtype: float64  
=====
```


Skewness present in our dataset

Corelation of Feature vs Label using Heat map

```
In [116]: print('-----')
print('Heat Map :-')
print('-----')
df_corr = df.corr().abs()

plt.figure(figsize = (22,16))
sns.heatmap(df_corr, vmin = -1, annot = True, square = True, center = 0, fmt = '.'
plt.tight_layout()
```

Heat Map :-




Vehicle Claim highest corelation

Checking Outliers

```
In [117]: print('=====')  
print('Box Plot :-')  
print('=====')  
  
plt.figure(figsize = (20,25), facecolor = 'white')  
plotnumber = 1  
for column in df:  
    if plotnumber <=40:  
        ax = plt.subplot(8,5, plotnumber)  
        sns.boxplot(y=df[column]) # It is the axis for vertical set as y  
        plt.xlabel(column, fontsize = 18)  
    plotnumber += 1  
plt.tight_layout()
```

```
=====  
Box Plot :-  
=====
```



There are outliers presents in dataset

Removing Outliers

```
In [118]: # with std 3 Lets see the stats
```

```
z_score = zscore(df[['age', 'umbrella_limit', 'total_claim_amount', 'property_cla  
abs_z_score = np.abs(z_score)  
  
filtering_entry = (abs_z_score < 3).all(axis = 1)  
  
df = df[filtering_entry]  
df.describe()
```


```
Out[118]:
```

	months_as_customer	age	policy_number	policy_state	policy_csl	policy_deductable
count	981.000000	981.000000	981.000000	981.000000	981.000000	981.000000
mean	204.594292	38.972477	545930.544343	1.020387	0.945973	1138.124363
std	115.362802	9.179406	257515.314276	0.830289	0.804412	611.567191
min	0.000000	19.000000	100804.000000	0.000000	0.000000	500.000000
25%	116.000000	32.000000	335780.000000	0.000000	0.000000	500.000000
50%	200.000000	38.000000	533940.000000	1.000000	1.000000	1000.000000
75%	278.000000	44.000000	760179.000000	2.000000	2.000000	2000.000000
max	479.000000	64.000000	999435.000000	2.000000	2.000000	2000.000000

```
In [119]: # Let's see outliers are removed in columns or not.
```

```
print('=====')  
print('Box Plot :-')  
print('=====')  
  
plt.figure(figsize = (20,25), facecolor = 'white')  
plotnumber = 1  
for column in df:  
    if plotnumber <=40:  
        ax = plt.subplot(8,5, plotnumber)  
        sns.boxplot(y=df[column]) # It is the axis for vertical set as y  
        plt.xlabel(column, fontsize = 18)  
    plotnumber += 1  
plt.tight_layout()
```

```
=====  
Box Plot :-  
=====
```



In [120]: `df.shape # Here we check shape of remaining data after removal of outliers.`

Out[120]: (981, 39)

Outliers are removed


Splitting Dataset into features and label

```
In [121]: x = df.drop('fraud_reported', axis = 1)
y = df. fraud_reported
print('Data has been splited')
```

Data has been splited

```
In [122]: # Let's see relation between features and label.  
print('-----')  
print('Scatter Plot :-')  
print('-----')  
  
plt.figure(figsize = (20,25), facecolor = 'white')  
plotnumber = 1  
for column in x:  
    if plotnumber <=40:  
        ax = plt.subplot(8,5, plotnumber)  
        plt.scatter(x[column],y, c = 'r')  
        plt.xlabel(column, fontsize = 18)  
    plotnumber += 1  
plt.tight_layout()
```

Scatter Plot :-



Positive relation in feature and label

Checking for class imbalance

```
In [123]: df['fraud_reported'].value_counts()
```

```
Out[123]: 0.0    741  
1.0    240  
Name: fraud_reported, dtype: int64
```

Class are not balance

Handling Class Imbalance

```
In [124]: sm = SMOTE()  
x_over, y_over = sm.fit_resample(x,y)
```

```
In [125]: print('-----')  
print('Class are balanced :-')  
print('-----')  
print(y_over.value_counts())  
print('-----')
```

```
-----  
Class are balanced :-  
-----  
0.0    741  
1.0    741  
Name: fraud_reported, dtype: int64  
-----
```

Data Scaling

```
In [170]: scaler = MinMaxScaler()  
x_scaled = scaler.fit_transform(x)  
x_scaled
```

```
Out[170]: array([[0.68475992, 0.64444444, 0.4682467 , ..., 0.         , 0.82807018,  
                  0.99409537],  
                 [0.47599165, 0.51111111, 0.26936974, ..., 0.         , 0.47368421,  
                  0.62344029],  
                 [0.27974948, 0.22222222, 0.65309788, ..., 0.5       , 0.69122807,  
                  0.68059269],  
                 ...,  
                 [0.27139875, 0.33333333, 0.90995303, ..., 0.         , 0.13333333,  
                  0.73718806],  
                 [0.95615866, 0.95555556, 0.48199539, ..., 0.5       , 0.90175439,  
                  0.56996435],  
                 [0.9519833 , 0.91111111, 0.50663287, ..., 0.5       , 0.85263158,  
                  0.04500891]])
```

Data has been scaled

Split data into train and test. Model will be bulit on training data and tested on test data

```
In [171]: x_train, x_test, y_train, y_test = train_test_split(x_over, y_over, test_size = 0.2)
print('Data has been splited.')
```

Data has been splited.

Model Bulding

Decision Tree model instantiaing, training and evaluating

```
In [172]: bag_dt = BaggingClassifier(DecisionTreeClassifier(), n_estimators = 30, max_samples = 5,
                                 random_state= 3, oob_score = True)
```

```
In [173]: bag_dt.oob_score
```

Out[173]: True

```
In [174]: bag_dt.fit(x_train, y_train)
print('Bagging DT score ----->', bag_dt.score(x_test, y_test))
```

Bagging DT score -----> 0.8598382749326146

```
In [175]: y_pred = bag_dt.predict(x_test)
```

```
In [176]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

	precision	recall	f1-score	support
0.0	0.85	0.88	0.87	191
1.0	0.87	0.83	0.85	180
accuracy			0.86	371
macro avg	0.86	0.86	0.86	371
weighted avg	0.86	0.86	0.86	371

Conclusion : Decision Tree model has 86% score

Cross Validation score to check if the model is overfitting


```
In [177]: cv = cross_val_score(bag_dt, x, y, cv = 5)
print('Cross Validation score of Decision Tree model --->', cv.mean())
```

Cross Validation score of Decision Tree model ---> 0.8318398425359993

Conclusion : Decision Tree model has 83% Cross Validation score

ROC, AUC Curve

```
In [178]: prob = bag_dt.predict_proba(x_test) # calculating probability
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))
plt.show()
```



XGBoost model instantiaing, training and evaluating

```
In [179]: bag_xgb = BaggingClassifier(xgb.XGBClassifier(eval_metric = 'mlogloss'), n_estimators=10, random_state= 3, oob_score = True)
```

```
In [180]: bag_xgb.oob_score
```

Out[180]: True

```
In [181]: bag_xgb.fit(x_train, y_train)
print('Bagging XGBoost score ----->', bag_xgb.score(x_test, y_test))
```

Bagging XGBoost score -----> 0.8814016172506739

```
In [182]: y_pred = bag_xgb.predict(x_test)
```

```
In [183]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

	precision	recall	f1-score	support
0.0	0.90	0.86	0.88	191
1.0	0.86	0.90	0.88	180
accuracy			0.88	371
macro avg	0.88	0.88	0.88	371
weighted avg	0.88	0.88	0.88	371

Conclusion : XGBoost model has 88% score

Cross Validation score to check if the model is overfitting


```
In [184]: cv = cross_val_score(bag_xgb, x, y, cv = 5)
print('Cross Validation score of XGBoost model --->', cv.mean())
```

Cross Validation score of XGBoost model ---> 0.8256863151351912

Conclusion : XGBoost model has 82% Cross Validation score

ROC, AUC Curve

```
In [141]: prob = bag_xgb.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



Knn model instantiaing, training and evaluating

```
In [185]: bag_Knn = BaggingClassifier(KNeighborsClassifier(n_neighbors = 5), n_estimators =  
random_state= 3, oob_score = True)
```

```
In [186]: bag_Knn.oob_score
```

```
Out[186]: True
```

```
In [187]: bag_Knn.fit(x_train, y_train)  
print('Bagging KNN score ----->', bag_Knn.score(x_test, y_test))
```

```
Bagging KNN score -----> 0.6630727762803235
```

```
In [188]: y_pred = bag_dt.predict(x_test)
```

```
In [189]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

	precision	recall	f1-score	support
0.0	0.85	0.88	0.87	191
1.0	0.87	0.83	0.85	180
accuracy			0.86	371
macro avg	0.86	0.86	0.86	371
weighted avg	0.86	0.86	0.86	371

Conclusion : KNN model has 86% score

Cross Validation score to check if the model is overfitting


```
In [190]: cv = cross_val_score(bag_Knn, x, y, cv = 5)
print('Cross Validation score of Knn model --->', cv.mean())
```

Cross Validation score of Knn model ---> 0.7451362270796643

Conclusion : Knn model has 74% Cross Validation score

ROC, AUC Curve

```
In [148]: prob = bag_Knn.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



Random Forest model instantiaing, training and evaluating

```
In [191]: bag_Rn = BaggingClassifier(RandomForestClassifier(), n_estimators = 30, max_samples = 200,  
random_state= 3, oob_score = True)
```

```
In [192]: bag_Rn.oob_score
```

```
Out[192]: True
```

```
In [193]: bag_Rn.fit(x_train, y_train)
print('Bagging Random Forest score ----->', bag_Rn.score(x_test, y_test))
```

Bagging Random Forest score -----> 0.8706199460916442

```
In [194]: y_pred = bag_Rn.predict(x_test)
```

```
In [195]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

	precision	recall	f1-score	support
0.0	0.87	0.87	0.87	191
1.0	0.87	0.87	0.87	180
accuracy			0.87	371
macro avg	0.87	0.87	0.87	371
weighted avg	0.87	0.87	0.87	371

Conclusion : Random Forest model has 87% score

Cross Validation score to check if the model is overfitting


```
In [196]: cv = cross_val_score(bag_Rn, x, y, cv = 5)
print('Cross Validation score of Rn model --->', cv.mean())
```

Cross Validation score of Rn model ---> 0.7573604060913706

Conclusion : Random Forest model has 75% Cross Validation score

ROC, AUC Curve

```
In [155]: prob = bag_Rn.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



Logistic Regression model instantiating, training and evaluating

```
In [197]: bag_Lr = BaggingClassifier(LogisticRegression(), n_estimators = 30, max_samples =  
random_state= 3, oob_score = True)
```

```
In [198]: bag_Lr.oob_score
```

```
Out[198]: True
```

```
In [199]: bag_Lr.fit(x_train, y_train)  
print('Bagging Logistic Regression score ----->', bag_Lr.score(x_test, y_test))
```

```
Bagging Logistic Regression score -----> 0.5498652291105122
```

```
In [200]: y_pred = bag_Lr.predict(x_test)
```

```
In [201]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

	precision	recall	f1-score	support
0.0	0.57	0.52	0.54	191
1.0	0.53	0.58	0.55	180
accuracy			0.55	371
macro avg	0.55	0.55	0.55	371
weighted avg	0.55	0.55	0.55	371

Conclusion : Logistic Regression model has 55% score

Cross Validation score to check if the model is overfitting


```
In [202]: cv = cross_val_score(bag_Lr, x, y, cv = 5)
print('Cross Validation score of Logistic regression model --->', cv.mean())
```

Cross Validation score of Logistic regression model ---> 0.7553506681860561

Conclusion : Logistic Regression model has 75% Cross Validation score

ROC, AUC Curve

```
In [203]: prob = bag_lr.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred, prob, figsize = (10,5))  
plt.show()
```



Let's find ROC, AUC score

```
In [205]: # DecisionTreeClassifier  
roc_auc_score(y_test, bag_dt.predict(x_test))
```

Out[205]: 0.8590750436300175

```
In [206]: # XGBoostClassifier  
roc_auc_score(y_test, bag_xgb.predict(x_test))
```

Out[206]: 0.881937172774869

```
In [207]: # KNeighborsClassifier  
roc_auc_score(y_test, bag_Knn.predict(x_test))
```

```
Out[207]: 0.6660558464223385
```

```
In [208]: # RandomForestClassifier  
roc_auc_score(y_test, bag_Rn.predict(x_test))
```


```
Out[208]: 0.8705061082024432
```

```
In [209]: # LogisticRegressionClassifier  
roc_auc_score(y_test, bag_Lr.predict(x_test))
```

```
Out[209]: 0.5506689936009308
```

Let's check ROC, AUC Curve for the fitted model

```
In [210]: dis = plot_roc_curve(bag_dt, x_test, y_test)  
plot_roc_curve(bag_Rn, x_test, y_test, ax = dis.ax_) # ax_ = Axes with confusion  
plot_roc_curve(bag_Knn, x_test, y_test, ax = dis.ax_)  
plot_roc_curve(bag_xgb, x_test, y_test, ax = dis.ax_)  
plot_roc_curve(bag_Lr, x_test, y_test, ax = dis.ax_)  
plt.legend(prop = {'size':11}, loc = 'lower right')  
plt.show()
```



Looking ROC, AUC Curve we found Random Forest has best model so we do Hyperparameter Tuning on it.

```
In [213]: param = {'n_estimators': [50,100], 'max_samples': [1.0], 'bootstrap': [True]}
```

```
In [216]: grid_search = GridSearchCV(estimator = bag_Rn, param_grid = param, cv = 5 , n_jobs
```

```
In [217]: grid_search.fit(x_train, y_train)
```

```
Out[217]: GridSearchCV(cv=5,
                        estimator=BaggingClassifier(base_estimator=RandomForestClassifier(),
                        (),
                        max_samples=0.5, n_estimators=30,
                        oob_score=True, random_state=3),
                        n_jobs=-1,
                        param_grid={'bootstrap': [True], 'max_samples': [1.0],
                        'n_estimators': [50, 100]})
```

```
In [218]: best_parameters = grid_search.best_params_
print(best_parameters)
```

```
{'bootstrap': True, 'max_samples': 1.0, 'n_estimators': 50}
```

```
In [219]: hRn = BaggingClassifier(base_estimator=RandomForestClassifier(),max_samples = 1.0)
hRn.fit(x_train, y_train)
hRn.score(x_test, y_test)
```

```
Out[219]: 0.8679245283018868
```

```
In [220]: y_pred = hRn.predict(x_test)
```

```
In [221]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
Classification Report:
      precision    recall  f1-score   support

       0.0       0.88      0.86      0.87      191
       1.0       0.86      0.87      0.87      180

  accuracy                           0.87      371
   macro avg       0.87      0.87      0.87      371
weighted avg       0.87      0.87      0.87      371
-----
```

After Hyperparameter Tuning model accuracy score 87%.

Saving The Model

```
In [222]: # saving the model to the Local file system  
filename = 'Insurance Claims Fraud Detection.pickle'  
pickle.dump(hRn, open(filename, 'wb'))
```

Predict Insurance Claims Fraud Detection

```
In [223]: model = pickle.load(open('Insurance Claims Fraud Detection.pickle', 'rb'))  
result = model.score(x_test, y_test)  
print('Predicted Score ----->', result)
```

Predicted Score -----> 0.8679245283018868

```
In [224]: Prediction = pd.DataFrame([model.predict(x_test)[:,], y_test[:,]], index = ['Prediction'])
```

Out[224]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
Predicted	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	1.0	1.0	0.0	1.0	1.0	0.0	1.0	0.0	0.0
Orginal	0.0	0.0	1.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	1.0	0.0	0.0	1.0	1.0	0.0	1.0	0.0	0.0

Saving the predicted result in CSV file

```
In [225]: Prediction.to_csv('Insurance Claims Fraud Detection.csv')
```

Final Conclusion : Random Forest is our best model.

In []: