

Problem Statement:

This dataset includes details of applicants who have applied for loan. The dataset includes details like credit history, loan amount, their income, dependents etc. You have to build a model that can predict whether the loan of the applicant will be approved or not on the basis of the details provided in the dataset

Importing Required Library

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pickle
import scikitplot as skplt
from imblearn.over_sampling import SMOTE
pd.set_option('display.max_columns', None) # # For display maximum column
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, plot_roc_curve
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier, BaggingClassifier
from sklearn.linear_model import LogisticRegression
import xgboost as xgb
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

Reading Data

```
In [2]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\csv file\loan_prediction.csv")
df.head()
```

Out[2]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	

Check no of row and column

```
In [3]: print('No of Rows and Columns ----->', df.shape )
```

```
No of Rows and Columns -----> (614, 13)
```

Checking for Null values

```
In [4]: print('=====\\n')
print(df.isnull().sum())
print('=====')
```

```
=====
```

```
Loan_ID          0
Gender           13
Married          3
Dependents       15
Education         0
Self_Employed    32
ApplicantIncome   0
CoapplicantIncome 0
LoanAmount        22
Loan_Amount_Term 14
Credit_History    50
Property_Area     0
Loan_Status        0
dtype: int64
```

```
=====
```

There is null value

Information about dataset

```
In [5]: print('=====\\n')
print(df.info())
print('=====')
```

```
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Loan_ID          614 non-null    object  
 1   Gender           601 non-null    object  
 2   Married          611 non-null    object  
 3   Dependents       599 non-null    object  
 4   Education         614 non-null    object  
 5   Self_Employed    582 non-null    object  
 6   ApplicantIncome  614 non-null    int64  
 7   CoapplicantIncome 614 non-null    float64 
 8   LoanAmount        592 non-null    float64 
 9   Loan_Amount_Term  600 non-null    float64 
 10  Credit_History   564 non-null    float64 
 11  Property_Area    614 non-null    object  
 12  Loan_Status       614 non-null    object  
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
None
=====
```

Categorical data present in our data set

Statistics of Data

```
In [6]: df.describe()
```

Out[6]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	614.000000	614.000000	592.000000	600.000000	564.000000
mean	5403.459283	1621.245798	146.412162	342.000000	0.842199
std	6109.041673	2926.248369	85.587325	65.12041	0.364878
min	150.000000	0.000000	9.000000	12.000000	0.000000
25%	2877.500000	0.000000	100.000000	360.000000	1.000000
50%	3812.500000	1188.500000	128.000000	360.000000	1.000000
75%	5795.000000	2297.250000	168.000000	360.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.000000	1.000000

Outliers are present in our data set

Drop Unwanted Columns

```
In [7]: df = df.drop('Loan_ID', axis = 1)
df.head(2)
```

Out[7]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	Male	No	0	Graduate	No	5849	0.0
1	Male	Yes	1	Graduate	No	4583	1508.0

Analysis of Null value

```
In [8]: print('=====\\n')
print(df.isnull().sum())
print('\\n=====')
```

```
=====
Gender           13
Married          3
Dependents       15
Education         0
Self_Employed    32
ApplicantIncome   0
CoapplicantIncome 0
LoanAmount        22
Loan_Amount_Term 14
Credit_History    50
Property_Area     0
Loan_Status        0
dtype: int64
=====
```

Gender column

```
In [9]: df['Gender'].value_counts()
```

```
Out[9]: Male      489
Female    112
Name: Gender, dtype: int64
```

```
In [10]: df['Gender'].dtype
```

```
Out[10]: dtype('O')
```

Approach : Gender column is categorical value so, we fill null value with mode

Married column

```
In [11]: df['Married'].value_counts()
```

```
Out[11]: Yes    398  
No     213  
Name: Married, dtype: int64
```

```
In [12]: df['Married'].dtype
```

```
Out[12]: dtype('O')
```

Approach : Married column is categorical value so, we fill null value with mode

Dependents columns

```
In [13]: df['Dependents'].value_counts()
```

```
Out[13]: 0    345  
1    102  
2    101  
3+   51  
Name: Dependents, dtype: int64
```

```
In [14]: df['Dependents'].dtype
```

```
Out[14]: dtype('O')
```

```
In [15]: df['Dependents'] = df['Dependents'].replace(to_replace= '3+', value = 3)
```

```
In [16]: df['Dependents'].value_counts()
```

```
Out[16]: 0    345  
1    102  
2    101  
3    51  
Name: Dependents, dtype: int64
```

Approach 1 : In order to do Dependents column analysis we found "3+" value which not acceptable to do analysis so we replace with "3"

Approach 2 : Dependents column is categorical value so, we fill null value with mode

Self Employed columns

```
In [17]: df['Self_Employed'].value_counts()
```

```
Out[17]: No      500  
Yes      82  
Name: Self_Employed, dtype: int64
```

```
In [18]: df['Self_Employed'].dtype
```

```
Out[18]: dtype('O')
```

Approach : Self Employed column is categorical value so, we fill null value with mode

Loan Amount column

```
In [19]: df['LoanAmount'].value_counts()
```

```
Out[19]: 120.0    20  
110.0     17  
100.0     15  
160.0     12  
187.0     12  
..  
211.0      1  
250.0      1  
62.0       1  
85.0       1  
436.0      1  
Name: LoanAmount, Length: 203, dtype: int64
```

```
In [20]: df['LoanAmount'].dtype
```

```
Out[20]: dtype('float64')
```

Approach : Loan Amount column is continuous value so, we fill null value with mean

Loan Amount Term column

```
In [21]: df['Loan_Amount_Term'].value_counts()
```

```
Out[21]: 360.0    512
180.0     44
480.0     15
300.0     13
84.0      4
240.0     4
120.0     3
36.0      2
60.0      2
12.0      1
Name: Loan_Amount_Term, dtype: int64
```

```
In [22]: df['Loan_Amount_Term'].dtype
```

```
Out[22]: dtype('float64')
```

Approach : Loan Amount Term column is continuous value so, we fill null value with mean

Credit History column

```
In [23]: df['Credit_History'].value_counts()
```

```
Out[23]: 1.0    475
0.0     89
Name: Credit_History, dtype: int64
```

```
In [24]: df['Credit_History'].dtype
```

```
Out[24]: dtype('float64')
```

Approach : Credit History column is categorical value so, we fill null value with mode

Fill NaN

```
In [25]: df['Gender'] = df['Gender'].fillna(df['Gender'].mode()[0])
df['Married'] = df['Married'].fillna(df['Married'].mode()[0])
df['Dependents'] = df['Dependents'].fillna(df['Dependents'].mode()[0])
df['Self_Employed'] = df['Self_Employed'].fillna(df['Self_Employed'].mode()[0])
df['LoanAmount'] = df['LoanAmount'].fillna(df['LoanAmount'].mean())
df['Loan_Amount_Term'] = df['Loan_Amount_Term'].fillna(df['Loan_Amount_Term'].mode())
df['Credit_History'] = df['Credit_History'].fillna(df['Credit_History'].mode()[0])
```

```
In [26]: print('=====\\n')
print(df.isnull().sum())
print('\\n=====')
```

```
=====
```

```
Gender          0
Married         0
Dependents      0
Education        0
Self_Employed   0
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       0
Loan_Amount_Term 0
Credit_History    0
Property_Area     0
Loan_Status        0
dtype: int64
```

```
=====
```

There is no null value

Analysis of data respect to Loan Status

```
In [27]: df.head(2)
```

Out[27]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	Male	No	0	Graduate	No	5849	0.0
1	Male	Yes	1	Graduate	No	4583	1508.0

Gender column


```
In [28]: df['Gender'].value_counts()
```

Out[28]: Male 502
Female 112
Name: Gender, dtype: int64

```
In [29]: df['Loan_Status'].value_counts()
```


Out[29]: Y 422
N 192
Name: Loan_Status, dtype: int64

```
In [30]: sns.histplot(binwidth=0.5, x="Gender", hue="Gender", data=df, stat="count", multi  
plt.show()
```



Male is highest in number

```
In [31]: sns.histplot(binwidth=0.5, x="Loan_Status", hue="Loan_Status", data=df, stat="cou  
plt.show()
```




Repayment is high ↗↗

```
In [32]: gn = df.groupby('Gender')['Loan_Status'].value_counts()
gn
```

```
Out[32]: Gender  Loan_Status
Female    Y          75
           N          37
Male      Y         347
           N         155
Name: Loan_Status, dtype: int64
```

```
In [33]: gn.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Gender', c = 'purple', fontsize = 12)
plt.ylabel('Loan Status', c = 'purple', fontsize = 12 )
plt.title('Gender vs Loan Status', c = 'purple', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```




Above plot shows Repayment of Loan in male is higest 347 and in female is 75 which low as compare to male

Married column

```
In [34]: df['Married'].value_counts()
```

```
Out[34]: Yes     401
          No      213
Name: Married, dtype: int64
```

```
In [35]: sns.histplot(binwidth=0.5, x="Married", hue="Married", data=df, stat="count", multiple="stack")
```




Above plot shows no of Married person is highest

```
In [36]: m = df.groupby('Married')['Loan_Status'].value_counts()  
m
```

```
Out[36]: Married  Loan_Status  
No        Y        134  
          N         79  
Yes       Y        288  
          N        113  
Name: Loan_Status, dtype: int64
```

```
In [37]: m.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Married', c = 'k', fontsize = 12)
plt.ylabel('Loan Status', c = 'k', fontsize = 12 )
plt.title('Married vs Loan Status', c = 'k', fontsize = 12)
plt.xticks(c = 'g')
plt.yticks(c = 'g')
plt.show()
```



Above plot shows Repayment of Loan in Married person is higest 288 and In Not Married is 134 which low as compare to Married person

Dependents column


```
In [38]: df['Dependents'].value_counts()
```

```
Out[38]: 0    360
1    102
2    101
3     51
Name: Dependents, dtype: int64
```

```
In [39]: d = df.groupby('Dependents')['Loan_Status'].value_counts()  
d
```

```
Out[39]: Dependents  Loan_Status  
3           Y      33  
             N     18  
0           Y     247  
             N    113  
1           Y      66  
             N     36  
2           Y      76  
             N     25  
Name: Loan_Status, dtype: int64
```

```
In [40]: d.plot.bar(figsize = (10,5), rot = 360, color = ['g', 'r'])  
plt.xlabel('Dependents', c = 'r', fontsize = 12)  
plt.ylabel('Loan Status', c = 'r', fontsize = 12 )  
plt.title('Dependents vs Loan Status', c = 'r', fontsize = 12)  
plt.show()
```



Above plot shows Repayment of Loan in Dependents "0" is higest 247 and in Dependents "3" is 18 which low as compare to Dependents "0"

Education column


```
In [41]: df['Education'].value_counts()
```

```
Out[41]: Graduate      480  
Not Graduate   134  
Name: Education, dtype: int64
```

```
In [42]: g = df.groupby('Education')['Loan_Status'].value_counts()
g
```

```
Out[42]: Education    Loan_Status
Graduate      Y          340
              N          140
Not Graduate Y          82
              N          52
Name: Loan_Status, dtype: int64
```

```
In [43]: g.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Education', c = 'b', fontsize = 12)
plt.ylabel('Loan Status', c = 'b', fontsize = 12 )
plt.title('Education vs Loan Status', c = 'b', fontsize = 12)
plt.xticks(c = 'r')
plt.yticks(c = 'r')
plt.show()
```



Above plot shows Repayment of Loan in Graduate person is highest 340 and in Not Graduate person is 82 which low as compare to Graduate person

Self Employed column

```
In [44]: df['Self_Employed'].value_counts()
```

```
Out[44]: No      532
Yes     82
Name: Self_Employed, dtype: int64
```

```
In [45]: s = df.groupby('Self_Employed')['Loan_Status'].value_counts()  
s
```

```
Out[45]: Self_Employed  Loan_Status  
No           Y      366  
              N     166  
Yes          Y      56  
              N     26  
Name: Loan_Status, dtype: int64
```

```
In [46]: s.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])  
plt.xlabel('Self_Employed', c = 'k', fontsize = 12)  
plt.ylabel('Loan Status', c = 'k', fontsize = 12 )  
plt.title('Self_Employed vs Loan Status', c = 'k', fontsize = 12)  
plt.xticks(c = 'm')  
plt.yticks(c = 'm')  
plt.show()
```



Above plot shows Repayment of Loan in Self Employed person is 56 which low 🗔 😞

Applicant Income column


```
In [47]: df['ApplicantIncome'].value_counts()
```

```
Out[47]: 2500      9  
6000       6  
2600       6  
4583       6  
4166       5  
..  
5503       1  
3450       1  
2425       1  
2423       1  
4095       1  
Name: ApplicantIncome, Length: 505, dtype: int64
```

```
In [48]: a = df.groupby('ApplicantIncome')['Loan_Status'].value_counts()  
a
```

```
Out[48]: ApplicantIncome  Loan_Status  
150           N          1  
210           Y          1  
416           N          1  
645           Y          1  
674           Y          1  
..  
39147          Y          1  
39999          Y          1  
51763          Y          1  
63337          Y          1  
81000          N          1  
Name: Loan_Status, Length: 536, dtype: int64
```

```
In [49]: df.plot.scatter(x = 'Loan_Status', y = 'ApplicantIncome', c = 'r')  
plt.xlabel('Loan Status', c = 'g', fontsize = 12)  
plt.ylabel('Applicant Income', c = 'g', fontsize = 12 )  
plt.title('Applicant Income vs Loan Status', c = 'g', fontsize = 12)  
plt.show()
```



Above plot shows Applicant Income whoes earning in between (210 to 65000) higest ↗

Loan Amount column


```
In [50]: df['LoanAmount'].value_counts()
```

```
Out[50]: 146.412162    22
120.000000    20
110.000000    17
100.000000    15
160.000000    12
..
211.000000     1
250.000000     1
62.000000      1
85.000000      1
436.000000     1
Name: LoanAmount, Length: 204, dtype: int64
```

```
In [51]: am = df.groupby('LoanAmount')['Loan_Status'].value_counts()
am
```

```
Out[51]:   LoanAmount  Loan_Status
          9.0        N          1
         17.0       Y          1
         25.0       Y          2
         26.0       Y          1
         30.0       Y          2
..
        500.0       N          1
        570.0       N          1
        600.0       Y          2
        650.0       Y          1
        700.0       Y          1
Name: Loan_Status, Length: 282, dtype: int64
```

```
In [52]: df.plot.scatter(x = 'Loan_Status', y = 'LoanAmount', c = 'g')
plt.xlabel('Loan Status', c = 'k', fontsize = 12)
plt.ylabel('Loan Amount', c = 'k', fontsize = 12 )
plt.title('Loan Amount vs Loan Status', c = 'k', fontsize = 12)
plt.show()
```



Above plot shows Loan amount in between (17 to 700) higest

Property Area column


```
In [53]: df['Property_Area'].value_counts()
```

```
Out[53]: Semiurban    233
Urban        202
Rural        179
Name: Property_Area, dtype: int64
```

```
In [54]: pa = df.groupby('Property_Area')['Loan_Status'].value_counts()
pa
```

```
Out[54]: Property_Area  Loan_Status
Rural          Y          110
                  N          69
Semiurban      Y          179
                  N          54
Urban          Y          133
                  N          69
Name: Loan_Status, dtype: int64
```

```
In [55]: pa.plot.bar(figsize = (10,5), rot = 360, color = ['g','r'])
plt.xlabel('Property Area', c = 'purple', fontsize = 12)
plt.ylabel('Loan Status', c = 'purple', fontsize = 12 )
plt.title('Property Area vs Loan Status', c = 'purple', fontsize = 12)
plt.xticks(c = 'b')
plt.yticks(c = 'b')
plt.show()
```



Above plot shows Repayment of Loan in Semi urban is higest 179 and lowest is 54

Hypothesis : Applicant whoes gender is Male, Graduated, Married, Annual Income in between (210 to 65000) and lived in Semi Urban area higest chance of Loan Repayment

Encoding columns

```
In [56]: le = LabelEncoder()
```

```
In [57]: df['Gender'] = le.fit_transform(df['Gender'])
df['Married'] = le.fit_transform(df['Married'])
df['Education'] = le.fit_transform(df['Education'])
df['Self_Employed'] = le.fit_transform(df['Self_Employed'])
df['Property_Area'] = le.fit_transform(df['Property_Area'])
df['Loan_Status'] = le.fit_transform(df['Loan_Status'])
```

```
In [58]: df['Dependents'] = pd.to_numeric(df['Dependents'])
```

```
In [59]: print('=====\\n')
print(df.info())
print('=====')
```

```
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Gender            614 non-null    int32  
 1   Married           614 non-null    int32  
 2   Dependents        614 non-null    int64  
 3   Education         614 non-null    int32  
 4   Self_Employed     614 non-null    int32  
 5   ApplicantIncome   614 non-null    int64  
 6   CoapplicantIncome 614 non-null    float64 
 7   LoanAmount        614 non-null    float64 
 8   Loan_Amount_Term  614 non-null    float64 
 9   Credit_History    614 non-null    float64 
 10  Property_Area    614 non-null    int32  
 11  Loan_Status       614 non-null    int32  
dtypes: float64(4), int32(6), int64(2)
memory usage: 43.3 KB
None
=====
```

```
In [60]: df.head(2)
```

Out[60]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	1	0	0	0	0	5849	0.0
1	1	1	1	0	0	4583	1508.0


All columns are encoded

Data distribution and checking outliers and skewness

```
In [61]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in df:
    if plotnumber <=15:
        ax = plt.subplot(5,3, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 20)
        plt.ylabel('Loan Status', fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```



```
In [62]: df.skew()
```

```
Out[62]: Gender           -1.648795
Married          -0.644850
Dependents       1.015551
Education         1.367622
Self_Employed    2.159796
ApplicantIncome   6.539513
CoapplicantIncome 7.491531
LoanAmount        2.726601
Loan_Amount_Term -2.402112
Credit_History    -2.021971
Property_Area     -0.066196
Loan_Status        -0.809998
dtype: float64
```


Data has outliers and skewness

Corelation of Feature vs Label using Heat map

```
In [63]: print('-----')
print('Heat Map :-')
print('-----')
df_corr = df.corr().abs()

plt.figure(figsize = (11,8))
sns.heatmap(df_corr, vmin = -1, annot = True, square = True, center = 0, fmt = '.'
plt.tight_layout()
```

Heat Map :-



Loan Amount has highest corelation

Removing Outliers

```
In [64]: # we are removing the top 2% data from the Gender column
q = df['Gender'].quantile(0.99)
data_cleaned = df[df['Gender']<q]
# we are removing the top 2% data from the Dependents column
q = df['Dependents'].quantile(0.99)
data_cleaned = data_cleaned[data_cleaned['Dependents']<q]
# we are removing the top 2% data from the free Self_Employed column
q = df['Self_Employed'].quantile(0.99)
data_cleaned = data_cleaned[data_cleaned['Self_Employed']<q]
# we are removing the top 2% data from the ApplicantIncome column
q = df['ApplicantIncome'].quantile(0.99)
data_cleaned = data_cleaned[data_cleaned['ApplicantIncome']<q]
# we are removing the top 2% data from the CoapplicantIncome column
q = df['CoapplicantIncome'].quantile(0.99)
data_cleaned = data_cleaned[data_cleaned['CoapplicantIncome']<q]
# we are removing the top 2% data from the LoanAmount column
q = df['LoanAmount'].quantile(0.99)
data_cleaned = data_cleaned[data_cleaned['LoanAmount']<q]
# we are removing the top 2% data from the Loan_Amount_Term column
q = df['Loan_Amount_Term'].quantile(0.99)
data_cleaned = data_cleaned[data_cleaned['Loan_Amount_Term']<q]
# we are removing the top 2% data from the Credit_History column
q = df['Credit_History'].quantile(0.99)
data_cleaned = data_cleaned[data_cleaned['Credit_History']<q]
```

Checking Outliers and skewness removed or not


```
In [65]: # Let's see outliers are removed in columns or not.
```

```
print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in df:
    if plotnumber <=15:
        ax = plt.subplot(5,3, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 20)
        plt.ylabel('Loan Status', fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----
Distribution Plot :-
```



Outliers are removed

Splitting Dataset into features and label

```
In [66]: x = df.drop('Loan_Status', axis = 1)
y = df['Loan_Status']
print('Data has been splited')
```


Data has been splited

```
In [67]: # Let's see relation between features and label.
```

```
print('-----')
print('Scatter Plot :-')
print('-----')

plt.figure(figsize = (20,25), facecolor = 'yellow')
plotnumber = 1
for column in x:
    if plotnumber <=15:
        ax = plt.subplot(5,3, plotnumber)
        plt.scatter(x[column],y, c = 'g')
        plt.xlabel(column, fontsize = 18)
        plt.ylabel('Loan_Status', fontsize = 18)
    plotnumber += 1
plt.tight_layout()
```

```
-----
Scatter Plot :-
```



Checking for class imbalance

```
In [68]: df['Loan_Status'].value_counts()
```

```
Out[68]: 1    422  
0    192  
Name: Loan_Status, dtype: int64
```

Class are not balanced

Handling Class Imbalance

```
In [69]: sm = SMOTE()  
x_over, y_over = sm.fit_resample(x,y)
```

```
In [70]: print('-----')  
print('Class are balanced :-')  
print('-----')  
print(y_over.value_counts())  
print('-----')
```

```
-----  
Class are balanced :-  
-----  
0    422  
1    422  
Name: Loan_Status, dtype: int64  
-----
```

Data Scaling

```
In [71]: scaler = MinMaxScaler()
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
Out[71]: array([[1.        , 0.        , 0.        , ..., 0.74358974, 1.        ,
   1.        ],
 [1.        , 1.        , 0.33333333, ..., 0.74358974, 1.        ,
 0.        ],
 [1.        , 1.        , 0.        , ..., 0.74358974, 1.        ,
 1.        ],
 ...,
 [1.        , 1.        , 0.33333333, ..., 0.74358974, 1.        ,
 1.        ],
 [1.        , 1.        , 0.66666667, ..., 0.74358974, 1.        ,
 1.        ],
 [0.        , 0.        , 0.        , ..., 0.74358974, 0.        ,
 0.5       ]])
```

Data has been scaled

Split data into train and test. Model will be bulit on training data and tested on test data

```
In [72]: x_train, x_test, y_train, y_test = train_test_split(x_over, y_over, test_size = 6)
print('Data has been splited.')
```

Data has been splited.

Model Bulding

Decision Tree model instantiaing, training and evaluating

```
In [73]: bag_dt = BaggingClassifier(DecisionTreeClassifier(), n_estimators = 15, max_samples = 100, random_state= 3, oob_score = True)
```

```
In [74]: bag_dt.oob_score
```

```
Out[74]: True
```

```
In [75]: bag_dt.fit(x_train, y_train)
print('Bagging DT score ----->', bag_dt.score(x_test, y_test))
```

Bagging DT score -----> 0.8056872037914692

```
In [76]: y_pred = bag_dt.predict(x_test)
```

```
In [77]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

Classification Report:				
	precision	recall	f1-score	support
0	0.85	0.75	0.80	109
1	0.77	0.86	0.81	102
accuracy			0.81	211
macro avg	0.81	0.81	0.81	211
weighted avg	0.81	0.81	0.81	211

```
-----
```

Conclusion : Decision Tree model has 81% score

Cross Validation score to check if the model is overfitting


```
In [78]: cv = cross_val_score(bag_dt, x, y, cv = 5)
print('Cross Validation score of Decision Tree model --->', cv.mean())
```

Cross Validation score of Decision Tree model ---> 0.7720111955217913

Conclusion : Decision Tree model has 77% Cross Validation score

ROC, AUC Curve

```
In [79]: prob = bag_dt.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



XGBoost model instantiating, training and evaluating

```
In [80]: bag_xgb = BaggingClassifier(xgb.XGBClassifier(eval_metric = 'mlogloss'), n_estimators=100, random_state= 3, oob_score = True)
```

```
In [81]: bag_xgb.oob_score
```

```
Out[81]: True
```

```
In [82]: bag_xgb.fit(x_train, y_train)
print('Bagging XGBoost score ----->', bag_xgb.score(x_test, y_test))
```

Bagging XGBoost score -----> 0.8009478672985783

```
In [83]: y_pred = bag_xgb.predict(x_test)
```

```
In [84]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

	precision	recall	f1-score	support
0	0.85	0.75	0.80	109
1	0.76	0.85	0.81	102
accuracy			0.80	211
macro avg	0.80	0.80	0.80	211
weighted avg	0.81	0.80	0.80	211

Conclusion : XGBoost model has 80% score

Cross Validation score to check if the model is overfitting


```
In [85]: cv = cross_val_score(bag_xgb, x, y, cv = 5)
print('Cross Validation score of XGBoost model --->', cv.mean())
```

Cross Validation score of XGBoost model ---> 0.8013461282153805

Conclusion : XGBoost model has 80% Cross Validation score

ROC, AUC Curve

```
In [86]: prob = bag_xgb.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



Knn model instantiaing, training and evaluating

```
In [87]: bag_Knn = BaggingClassifier(KNeighborsClassifier(n_neighbors = 5), n_estimators =  
random_state= 3, oob_score = True)
```

```
In [88]: bag_Knn.oob_score
```

```
Out[88]: True
```

```
In [89]: bag_Knn.fit(x_train, y_train)  
print('Bagging KNN score ----->', bag_Knn.score(x_test, y_test))
```

```
Bagging KNN score -----> 0.5876777251184834
```

```
In [90]: y_pred = bag_dt.predict(x_test)
```

```
In [91]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

	precision	recall	f1-score	support
0	0.85	0.75	0.80	109
1	0.77	0.86	0.81	102
accuracy			0.81	211
macro avg	0.81	0.81	0.81	211
weighted avg	0.81	0.81	0.81	211

Conclusion : KNN model has 81% score

Cross Validation score to check if the model is overfitting


```
In [92]: cv = cross_val_score(bag_Knn, x, y, cv = 5)
print('Cross Validation score of Knn model --->', cv.mean())
```

Cross Validation score of Knn model ---> 0.6742636278821805

Conclusion : Knn model has 67% Cross Validation score

ROC, AUC Curve

```
In [93]: prob = bag_Knn.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



Random Forest model instantiaing, training and evaluating

```
In [94]: bag_Rn = BaggingClassifier(RandomForestClassifier(), n_estimators = 15, max_samples=200,  
random_state=3, oob_score = True)
```

```
In [95]: bag_Rn.oob_score
```

```
Out[95]: True
```

```
In [96]: bag_Rn.fit(x_train, y_train)
print('Bagging Random Forest score ----->', bag_Rn.score(x_test, y_test))
```

Bagging Random Forest score -----> 0.8009478672985783

```
In [97]: y_pred = bag_Rn.predict(x_test)
```

```
In [98]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

	precision	recall	f1-score	support
0	0.86	0.73	0.79	109
1	0.75	0.87	0.81	102
accuracy			0.80	211
macro avg	0.81	0.80	0.80	211
weighted avg	0.81	0.80	0.80	211

Conclusion : Random Forest model has 80% score

Cross Validation score to check if the model is overfitting


```
In [99]: cv = cross_val_score(bag_Rn, x, y, cv = 5)
print('Cross Validation score of Knn model --->', cv.mean())
```

Cross Validation score of Knn model ---> 0.8078501932560309

Conclusion : Random Forest model has 80% Cross Validation score

ROC, AUC Curve

```
In [100]: prob = bag_Rn.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob, figsize = (10,5))  
plt.show()
```



Logistic Regression model instantiating, training and evaluating

```
In [101]: bag_Lr = BaggingClassifier(LogisticRegression(), n_estimators = 15, max_samples =  
random_state= 3, oob_score = True)
```

```
In [102]: bag_Lr.oob_score
```

```
Out[102]: True
```

```
In [103]: bag_Lr.fit(x_train, y_train)  
print('Bagging Logistic Regression score ----->', bag_Lr.score(x_test, y_test))
```

```
Bagging Logistic Regression score -----> 0.7298578199052133
```

```
In [104]: y_pred = bag_Lr.predict(x_test)
```

```
In [105]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

```
-----
```

Classification Report:				
	precision	recall	f1-score	support
0	0.74	0.73	0.74	109
1	0.72	0.73	0.72	102
accuracy			0.73	211
macro avg	0.73	0.73	0.73	211
weighted avg	0.73	0.73	0.73	211

```
-----
```

Conclusion : Logistic Regression model has 73% score

Cross Validation score to check if the model is overfitting


```
In [113]: cv = cross_val_score(bag_Lr, x, y, cv = 5)
print('Cross Validation score of Logistic regression model --->', cv.mean())
```

Cross Validation score of Logistic regression model ---> 0.8062108489937359

Conclusion : Logistic Regression model has 80% Cross Validation score

ROC, AUC Curve

```
In [107]: prob = bag_lr.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred, prob, figsize = (10,5))  
plt.show()
```



Let's find ROC, AUC score

```
In [108]: # DecisionTreeClassifier  
roc_auc_score(y_test, bag_dt.predict(x_test))
```

Out[108]: 0.8075193380104335

```
In [109]: # XGBoostClassifier  
roc_auc_score(y_test, bag_xgb.predict(x_test))
```

Out[109]: 0.8026173772261198

```
In [110]: # KNeighborsClassifier  
roc_auc_score(y_test, bag_Knn.predict(x_test))
```


```
Out[110]: 0.5845475805000898
```

```
In [111]: # RandomForestClassifier  
roc_auc_score(y_test, bag_Rn.predict(x_test))
```

```
Out[111]: 0.8032469868681418
```

Let's check ROC, AUC Curve for the fitted model

```
In [112]: dis = plot_roc_curve(bag_dt, x_test, y_test)  
plot_roc_curve(bag_Rn, x_test, y_test, ax = dis.ax_) # ax_ = Axes with confusion  
plot_roc_curve(bag_Knn, x_test, y_test, ax = dis.ax_)  
plot_roc_curve(bag_xgb, x_test, y_test, ax = dis.ax_)  
plt.legend(prop = {'size':11}, loc = 'lower right')  
plt.show()
```



Looking CV score we found Random Forest has best model so we do Hyperparameter Tuning on it.

```
In [114]: param = {'n_estimators': [50,100,750], 'max_samples': [1.0], 'bootstrap': [True]}
```

```
In [115]: grid_search = GridSearchCV(estimator = bag_Rn, param_grid = param, cv = 5 , n_jobs
```

```
In [116]: grid_search.fit(x_train, y_train)
```

```
Out[116]: GridSearchCV(cv=5,
                       estimator=BaggingClassifier(base_estimator=RandomForestClassifier(),
                                                    max_samples=0.5, n_estimators=15,
                                                    oob_score=True, random_state=3),
                       n_jobs=-1,
                       param_grid={'bootstrap': [True], 'max_samples': [1.0],
                                   'n_estimators': [50, 100, 750]})
```

```
In [117]: best_parameters = grid_search.best_params_
print(best_parameters)

{'bootstrap': True, 'max_samples': 1.0, 'n_estimators': 50}
```

```
In [118]: hRn = BaggingClassifier(base_estimator=RandomForestClassifier(), max_samples = 1.0)
hRn.fit(x_train, y_train)
hRn.score(x_test, y_test)
```

```
Out[118]: 0.8056872037914692
```

```
In [119]: y_pred = hRn.predict(x_test)
```

```
In [120]: print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----\n')
```

	precision	recall	f1-score	support
0	0.85	0.75	0.80	109
1	0.77	0.86	0.81	102
accuracy			0.81	211
macro avg	0.81	0.81	0.81	211
weighted avg	0.81	0.81	0.81	211

Saving The Model

```
In [129]: # saving the model to the Local file system
filename = 'Loan Application Status Project.pickle'
pickle.dump(hRn, open(filename, 'wb'))
```

Predict Loan Repayment

```
In [130]: model = pickle.load(open('Loan Application Status Project.pickle', 'rb'))
```

```
In [131]: a = model.predict(x)  
a
```

```
In [132]: pred = pd.DataFrame(a)
pred
```

Out[132]:

	0
0	1
1	0
2	1
3	1
4	1
...	...
609	1
610	1
611	1
612	1
613	0

614 rows × 1 columns

```
In [ ]:
```