

Problem Statement:

Data Set Information: This data set includes descriptions of hypothetical samples corresponding to 23 species of gilled mushrooms in the Agaricus and Lepiota Family (pp. 500-525). Each species is identified as definitely edible, definitely poisonous, or of unknown edibility and not recommended. This latter class was combined with the poisonous one. The Guide clearly states that there is no simple rule for determining the edibility of a mushroom; no rule like ``leaflets three, let it be'' for Poisonous Oak and Ivy.

Importing required libraries

```
In [32]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import pickle
import seaborn as sns
from category_encoders import BinaryEncoder
from scipy.stats import zscore
import statsmodels.api as sm
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import power_transform, StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

Display Maximum columns and rows

```
In [2]: pd.set_option("display.max_columns", None)
pd.set_option("display.max_rows", None)
```

Reading Data

```
In [3]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\Data Analysis With Python\ML Files\mushrooms.csv")
df.head()
```

Out[3]:

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	stalk-shape	stalk-root	su
0	p	x	s	n	t	p	f	c	n	k	e	e	o
1	e	x	s	y	t	a	f	c	b	k	e	c	o
2	e	b	s	w	t	l	f	c	b	n	e	c	o
3	p	x	y	w	t	p	f	c	n	n	e	e	o
4	e	x	s	g	f	n	f	w	b	k	t	e	o

```
In [4]: # Here we use shape command to know total no of rows and columns present in our dataset
print('Rows and Columns in Dataset : ', df.shape)
```

Rows and Columns in Dataset : (8124, 23)

In [5]: # Here we use info command to know all details about dataset i.e, size, type etc.

```
print('-----')
print('\nInformations of dataset :-\n')
print(df.info())
print('\n-----')
```

Informations of dataset :-

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8124 entries, 0 to 8123
Data columns (total 23 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   class            8124 non-null    object  
 1   cap-shape        8124 non-null    object  
 2   cap-surface      8124 non-null    object  
 3   cap-color        8124 non-null    object  
 4   bruises          8124 non-null    object  
 5   odor             8124 non-null    object  
 6   gill-attachment  8124 non-null    object  
 7   gill-spacing     8124 non-null    object  
 8   gill-size        8124 non-null    object  
 9   gill-color       8124 non-null    object  
 10  stalk-shape     8124 non-null    object  
 11  stalk-root       8124 non-null    object  
 12  stalk-surface-above-ring 8124 non-null    object  
 13  stalk-surface-below-ring 8124 non-null    object  
 14  stalk-color-above-ring 8124 non-null    object  
 15  stalk-color-below-ring 8124 non-null    object  
 16  veil-type        8124 non-null    object  
 17  veil-color       8124 non-null    object  
 18  ring-number      8124 non-null    object  
 19  ring-type        8124 non-null    object  
 20  spore-print-color 8124 non-null    object  
 21  population        8124 non-null    object  
 22  habitat           8124 non-null    object  
dtypes: object(23)
memory usage: 1.4+ MB
None
```

Our dataset set has all object value we need to apply encoding technique.

```
In [6]: # Here we use isna() command to identify of nan in our dataset.
```

```
print('-----')
print('\nNaN in dataset :-\n')
print(df.isna().sum())
print('\n-----')
```

NaN in dataset :-

```
class                  0
cap-shape              0
cap-surface             0
cap-color               0
bruises                0
odor                   0
gill-attachment         0
gill-spacing            0
gill-size               0
gill-color               0
stalk-shape             0
stalk-root               0
stalk-surface-above-ring 0
stalk-surface-below-ring 0
stalk-color-above-ring      0
stalk-color-below-ring      0
veil-type                0
veil-color               0
ring-number              0
ring-type                 0
spore-print-color         0
population                0
habitat                  0
dtype: int64
```

There is no null values in dataset.

Checking class imbalance

```
In [7]: df['class'].value_counts()
```

```
Out[7]: e    4208
p    3916
Name: class, dtype: int64
```

Dataset is balance.

Label Encoder

```
In [8]: lab_enc = LabelEncoder()
y = lab_enc.fit_transform(df['class'])
y
```

```
Out[8]: array([1, 0, 0, ..., 0, 1, 0])
```

Binary Encoder

```
In [9]: bi_enc = BinaryEncoder()
x = bi_enc.fit_transform(df)
x.head()
```

```
Out[9]:
```

	class_0	class_1	cap-shape_0	cap-shape_1	cap-shape_2	cap-shape_3	cap-surface_0	cap-surface_1	cap-surface_2	cap-color_
0	0	1	0	0	0	1	0	0	1	
1	1	0	0	0	0	1	0	0	0	1
2	1	0	0	0	1	0	0	0	0	1
3	0	1	0	0	0	1	0	1	0	
4	1	0	0	0	0	1	0	0	1	

Above we encode our categorical data into binary.

```
In [10]: # Here we use shape command to know total no of rows and columns present in our dataset
print('Rows and Columns in Dataset : ', x.shape)
```

```
Rows and Columns in Dataset : (8124, 78)
```

Drop repeated class

```
In [11]: x = x.drop(columns = ['class_0', 'class_1'], axis = 1)  
x.head()
```

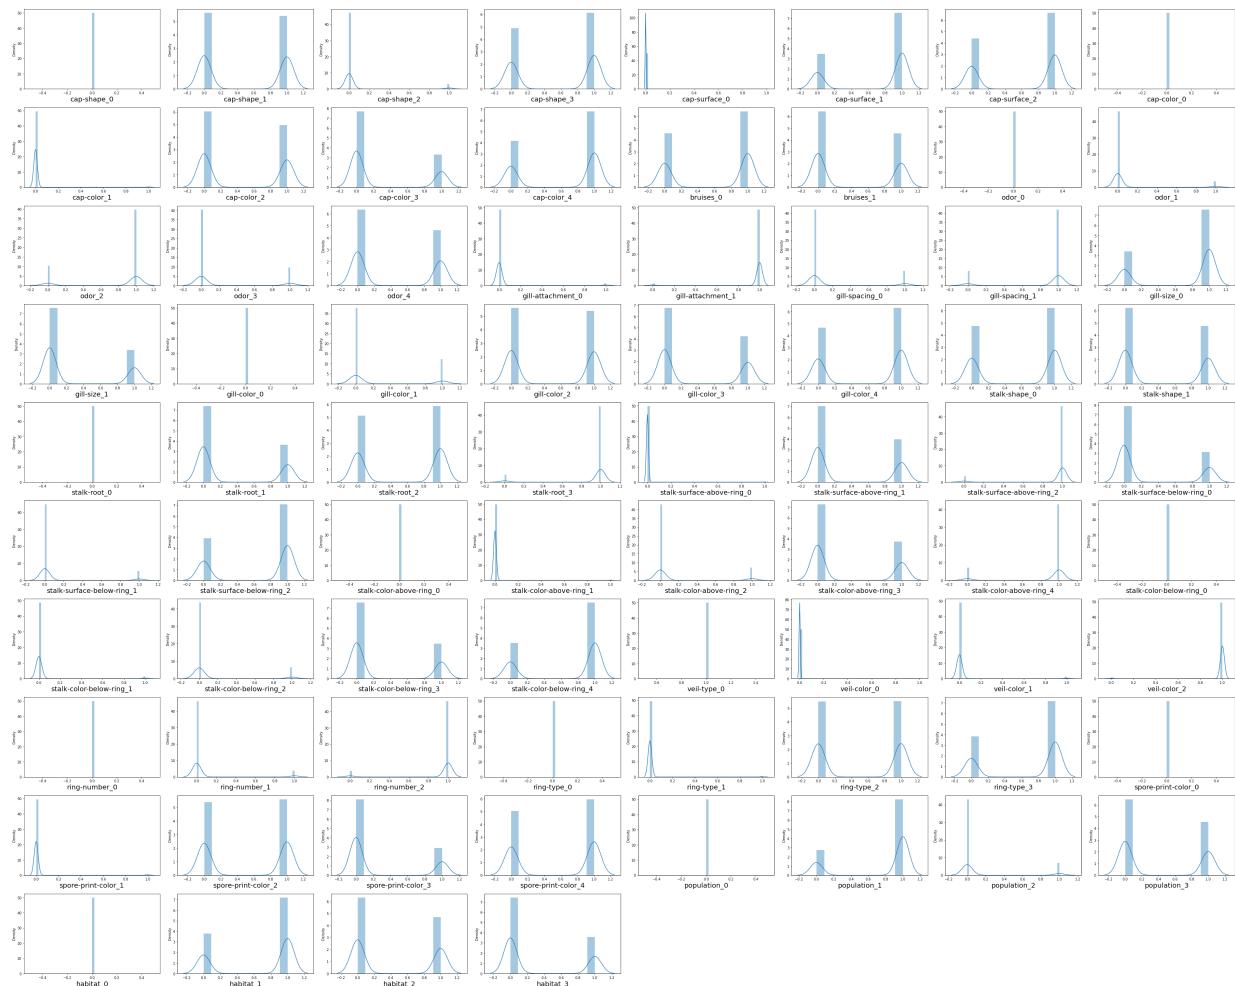
Out[11]:

	cap-shape_0	cap-shape_1	cap-shape_2	cap-shape_3	cap-surface_0	cap-surface_1	cap-surface_2	cap-color_0	cap-color_1	cap-color_2
0	0	0	0	1	0	0	1	0	0	0
1	0	0	0	1	0	0	1	0	0	0
2	0	0	1	0	0	0	1	0	0	0
3	0	0	0	1	0	1	0	0	0	0
4	0	0	0	1	0	0	1	0	0	0

Checking for outliers

```
In [12]: # Checking outlier
plt.figure(figsize = (50,40))
plotnumber = 1
```

```
for column in x:
    if plotnumber <=80:
        ax = plt.subplot(10,8, plotnumber)
        sns.distplot(x[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```



There is no outliers.

Split data into train and test. Model will be bulit on training data and tested on test data.

```
In [13]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25, random_state = 42)
print('Data has been splitted.')
```

Data has been splitted.

Model Building.

Logistic Regression model instantiating, training and evaluating

```
In [14]: Lr = LogisticRegression()
Lr.fit(x_train, y_train)
y_pred = Lr.predict(x_test)
```

```
In [15]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

```
-----  
Confusion Matrix :  
[[1056   2]  
 [   3 970]]  
-----
```

```
Classification Report:  
      precision    recall  f1-score   support  
  
          0       1.00     1.00     1.00     1058  
          1       1.00     1.00     1.00      973  
  
    accuracy                           1.00     2031  
   macro avg       1.00     1.00     1.00     2031  
weighted avg       1.00     1.00     1.00     2031  
-----
```

Conclusion : Logistic Regression model has 100% score .

Cross Validation score to check if the model is overfitting

```
In [16]: cv = cross_val_score(Lr, x, y, cv = 5)
print('Cross Validation score of Logistic Regression model --->', cv.mean())
```

```
Cross Validation score of Logistic Regression model ---> 0.877243956043956
```

Conclusion : Logistic Regression model has 87% Cross Validation score .

Decision Tree model instantiaing, training and evaluating

```
In [17]: DT = DecisionTreeClassifier()
DT.fit(x_train, y_train)
y_pred = DT.predict(x_test)
```

```
In [18]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

```
-----  
Confusion Matrix :  
[[1058  0]  
 [ 0 973]]  
-----
```

```
Classification Report:  
      precision    recall   f1-score   support  
  
       0          1.00     1.00     1.00      1058  
       1          1.00     1.00     1.00      973  
  
accuracy                           1.00      2031  
macro avg       1.00     1.00     1.00      2031  
weighted avg     1.00     1.00     1.00      2031  
-----
```

Conclusion : Decision Tree model has 100% score .

Cross Validation score to check if the model is overfitting

```
In [19]: cv = cross_val_score(DT, x, y, cv = 5)
print('Cross Validation score of Decision Tree model --->', cv.mean())
```

```
Cross Validation score of Decision Tree model ---> 0.9280909435392195
```

Conclusion : Decision Tree model has 93% Cross Validation score

Knn model instantiaing, training and evaluating

```
In [20]: Knn = KNeighborsClassifier()
Knn.fit(x_train, y_train)
y_pred = Knn.predict(x_test)
```

```
In [21]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

Confusion Matrix :

```
[[1058    0]
 [  0  973]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1058
1	1.00	1.00	1.00	973
accuracy			1.00	2031
macro avg	1.00	1.00	1.00	2031
weighted avg	1.00	1.00	1.00	2031

Conclusion : Knn model has 100% score .

Cross Validation score to check if the model is overfitting

```
In [22]: cv = cross_val_score(Knn, x, y, cv = 5)
print('Cross Validation score of Knn model --->', cv.mean())
```

Cross Validation score of Knn model ---> 0.9308065176203109

Conclusion : Knn model has 93% Cross Validation score .

Random Forest model instantiaing, training and evaluating

```
In [23]: Rn = RandomForestClassifier()
Rn.fit(x_train, y_train)
y_pred = Rn.predict(x_test)
```

```
In [24]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

Confusion Matrix :

```
[[1058    0]
 [    0  973]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1058
1	1.00	1.00	1.00	973
accuracy			1.00	2031
macro avg	1.00	1.00	1.00	2031
weighted avg	1.00	1.00	1.00	2031

Conclusion : Random Forest model has 100% score .

Cross Validation score to check if the model is overfitting

```
In [25]: cv = cross_val_score(Rn, x, y, cv = 5)
print('Cross Validation score of Knn model --->', cv.mean())
```

```
Cross Validation score of Knn model ---> 0.9209553618794999
```

Conclusion : Random Forest model has 90% Cross Validation score

SVM model instantiating, training and evaluating

```
In [26]: svc = SVC()
svc.fit(x_train, y_train)
y_pred = svc.predict(x_test)
```

```
In [27]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('Classification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

Confusion Matrix :

```
[[1058    0]
 [  0  973]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1058
1	1.00	1.00	1.00	973
accuracy			1.00	2031
macro avg	1.00	1.00	1.00	2031
weighted avg	1.00	1.00	1.00	2031

Conclusion : SVM model has 100% score .

Cross Validation score to check if the model is overfitting

```
In [28]: cv = cross_val_score(svc, x, y, cv = 5)
print('Cross Validation score of Knn model --->', cv.mean())
```

```
Cross Validation score of Knn model ---> 0.8893054945054946
```

Conclusion : SVM model has 88% Cross Validation score .

Looking CV score we found Decision Tree has best model so we do Hyperparameter Tuning on it.

```
In [29]: # we are tuning hyperparameter, we are passing the different values for both parameters
grid_param = {'criterion': ['gini', 'entropy'], 'max_depth': range(2, 20, 3), 'min_samples_leaf': range(1, 50, 2)}
```

```
In [34]: grid_search = GridSearchCV(estimator = DT, param_grid = grid_param, cv = 5 , n_jobs=-1)
```

```
In [35]: grid_search.fit(x_train, y_train)
```

```
Out[35]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(), n_jobs=-1,
param_grid={'criterion': ['gini', 'entropy'],
'max_depth': range(2, 20, 3),
'min_samples_leaf': range(1, 50, 2),
'min_samples_split': range(2, 50, 2)})
```

```
In [36]: best_parameters = grid_search.best_params_
print(best_parameters)
```

```
{'criterion': 'gini', 'max_depth': 8, 'min_samples_leaf': 1, 'min_samples_split': 2}
```

```
In [38]: hdt = DecisionTreeClassifier(criterion = 'gini', max_depth = 8 , min_samples_leaf = 1)
hdt.fit(x_train, y_train)
hdt.score(x_test, y_test)
```

```
Out[38]: 1.0
```

```
In [39]: y_pred = hdt.predict(x_test)
```

```
In [40]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

Confusion Matrix :

```
[[1058    0]
 [    0  973]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1058
1	1.00	1.00	1.00	973
accuracy			1.00	2031
macro avg	1.00	1.00	1.00	2031
weighted avg	1.00	1.00	1.00	2031

After Hyperparameter Tuning model accuracy score increase to 100%

Saving The Model

```
In [42]: # saving the model to the Local file system
filename = 'hdt_model.pickle'
pickle.dump(hdt, open(filename, 'wb'))
```

Final Conclusion : Decision Tree is our best model.

```
In [ ]:
```