

## Problem Statement:

The Titanic Problem is based on the sinking of the ‘Unsinkable’ ship Titanic in early 1912. It gives you information about multiple people like their ages, sexes, sibling counts, embarkment points, and whether or not they survived the disaster. Based on these features, you have to predict if an arbitrary passenger on Titanic would survive the sinking or not.

## Importing required libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pickle
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from scipy.stats import zscore
import scikitplot as skplt
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

## Reading data

```
In [2]: df = pd.read_csv(r"C:\Users\Kushal Arya\Desktop\Data Analysis With Python\ML Files\titanic.csv")
df.head()
```

Out[2]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	Na
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833	C8
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	Na
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C12
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	Na

## Check no of row and column

```
In [3]: print('No of Rows and Columns ----->', df.shape )
```

No of Rows and Columns -----> (891, 12)

## Checking for Null values

```
In [4]: print('-----\n')
print(df.isnull().sum())
print('\n-----')
```

```
-----  
PassengerId      0  
Survived         0  
Pclass           0  
Name             0  
Sex              0  
Age            177  
SibSp           0  
Parch           0  
Ticket          0  
Fare            0  
Cabin          687  
Embarked        2  
dtype: int64  
-----
```

## Information about dataset

```
In [5]: print('-----\n')
print(df.info())
print('\n-----')
```

```
-----  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 891 entries, 0 to 890  
Data columns (total 12 columns):  
 #   Column      Non-Null Count  Dtype     
---  --    
 0   PassengerId  891 non-null    int64    
 1   Survived     891 non-null    int64    
 2   Pclass       891 non-null    int64    
 3   Name         891 non-null    object    
 4   Sex          891 non-null    object    
 5   Age          714 non-null    float64   
 6   SibSp        891 non-null    int64    
 7   Parch        891 non-null    int64    
 8   Ticket       891 non-null    object    
 9   Fare         891 non-null    float64   
 10  Cabin        204 non-null    object    
 11  Embarked     889 non-null    object    
dtypes: float64(2), int64(5), object(5)  
memory usage: 83.7+ KB  
None  
-----
```

There is Sex, Name, Ticket, Cabin and Embarked feature which is object type we have convert into int.

## Dropping Unwanted Column

```
In [6]: col = ['PassengerId', 'Name', 'Ticket', 'Cabin']
```

```
In [7]: df = df.drop(columns = col, axis = 1)
df.head()
```

Out[7]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	71.2833	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S

We drop Cabin because it has more than 50% data is missing or other words more than 50% NaN present.

## Handling missing data

```
In [8]: df['Age'] = df['Age'].fillna(df['Age'].median())
df['Embarked'] = df['Embarked'].fillna(df['Embarked'].mode()[0])
```

```
In [9]: # check nan remove or not
df.isna().sum()
```

Out[9]:

Survived	0
Pclass	0
Sex	0
Age	0
SibSp	0
Parch	0
Fare	0
Embarked	0
dtype: int64	

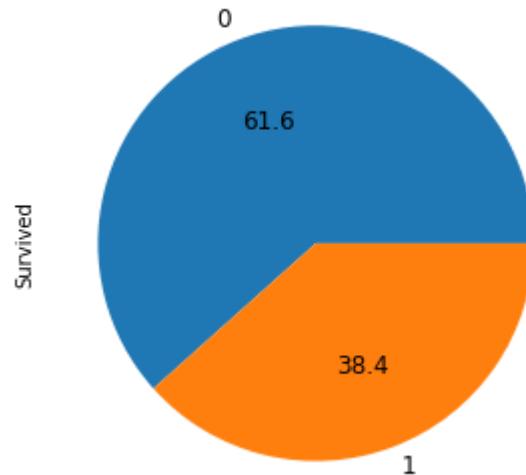
Nan removed from our dataset.

## Relation feature and feature

```
In [10]: s = df['Survived'].value_counts()  
s
```

```
Out[10]: 0    549  
1    342  
Name: Survived, dtype: int64
```

```
In [11]: s.plot.pie( fontsize = 12, autopct = '%.1f', figsize = (10,5))  
plt.show()
```



**Observation : 61.6% people died and 38.4% people survived.**

```
In [12]: df['Pclass'].value_counts()
```

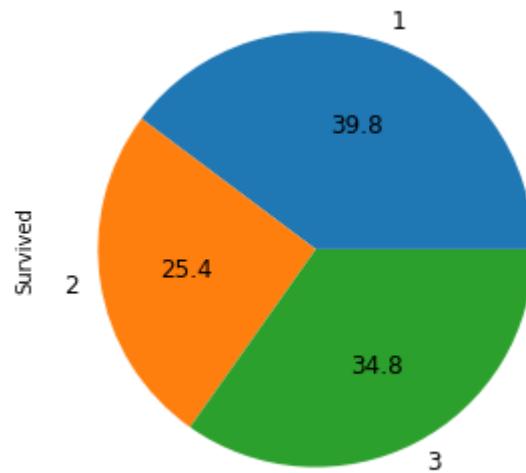
```
Out[12]: 3    491  
1    216  
2    184  
Name: Pclass, dtype: int64
```

```
In [13]: print('-----')  
print(df.groupby('Pclass')['Survived'].sum())  
print('-----')
```

```
Pclass  
1    136  
2     87  
3    119  
Name: Survived, dtype: int64
```

```
In [14]: p = df.groupby('Pclass')['Survived'].sum()
```

```
In [15]: p.plot.pie( fontsize = 12, autopct = '%.1f',figsize = (10,5))  
plt.show()
```



**Observation : 1st class passanger Survived more compare to other class and 2nd class passanger Survived least.**

```
In [16]: df['Fare'].value_counts()
```

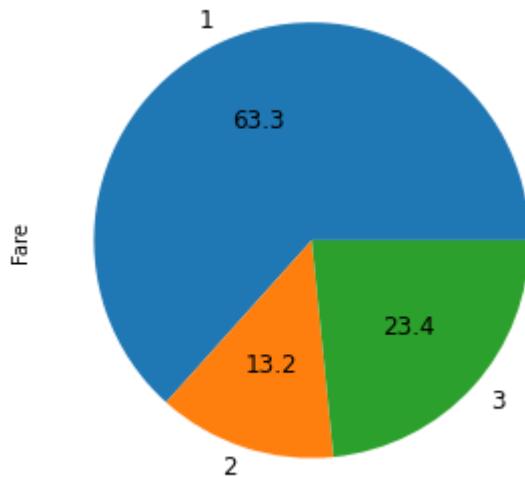
```
Out[16]: 8.0500    43  
13.0000    42  
7.8958    38  
7.7500    34  
26.0000    31  
..  
50.4958    1  
13.8583    1  
8.4583    1  
7.7250    1  
7.5208    1  
Name: Fare, Length: 248, dtype: int64
```

```
In [17]: print('-----')  
print(df.groupby('Pclass')['Fare'].sum())  
print('-----')
```

```
-----  
Pclass  
1    18177.4125  
2    3801.8417  
3    6714.6951  
Name: Fare, dtype: float64  
-----
```

```
In [18]: f = df.groupby('Pclass')['Fare'].sum()
```

```
In [19]: f.plot.pie(fontsize = 12, autopct = '%.1f', figsize = (10,5))  
plt.show()
```



**Observation : 1st class passanger spent more on tickets.**

```
In [20]: df['Age'].value_counts()
```

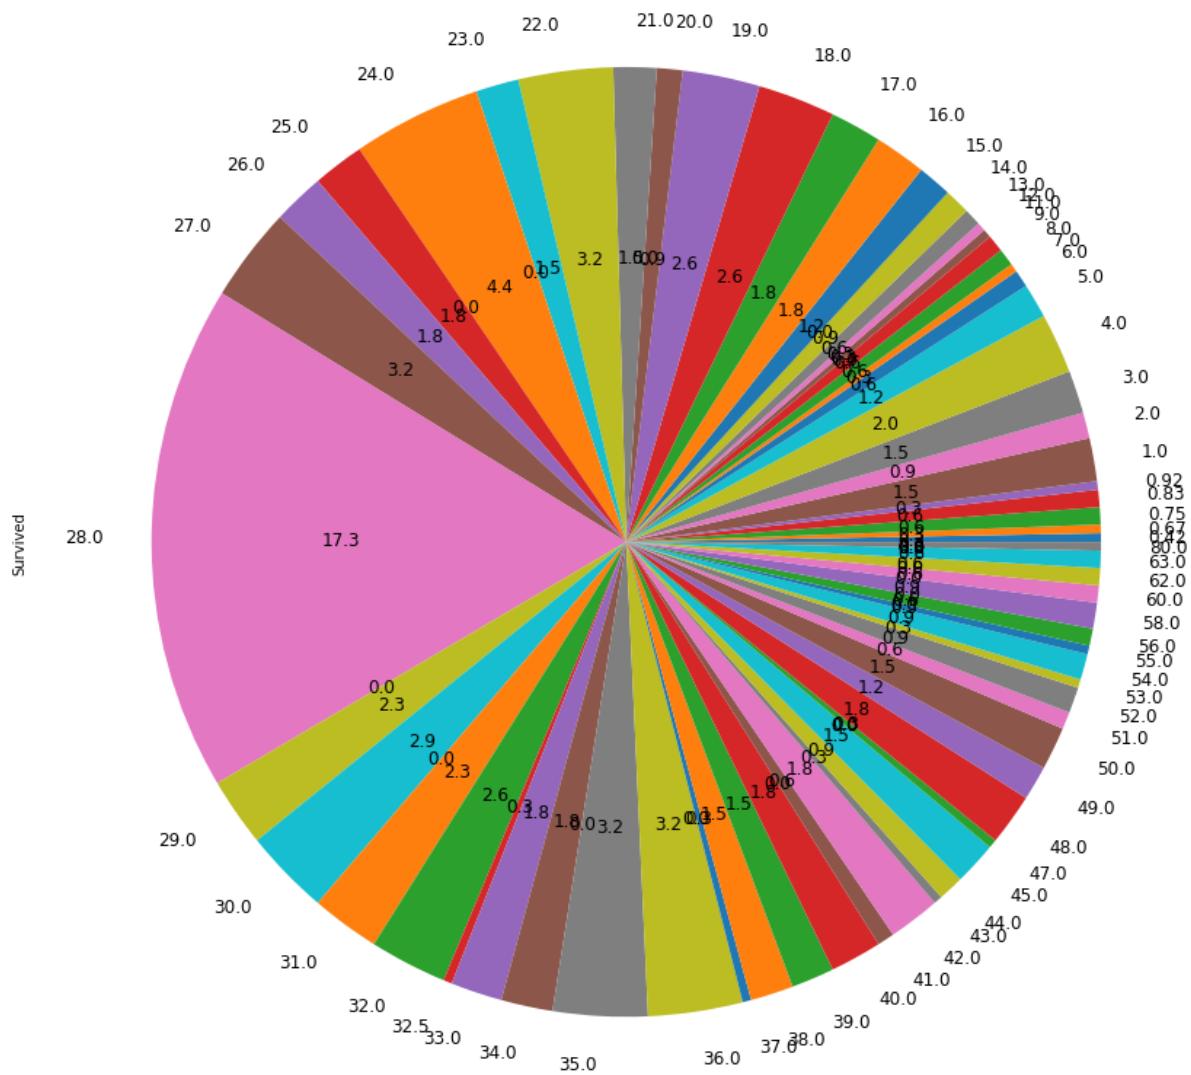
```
Out[20]: 28.00    202  
24.00     30  
22.00     27  
18.00     26  
19.00     25  
...  
55.50      1  
74.00      1  
0.92      1  
70.50      1  
12.00      1  
Name: Age, Length: 88, dtype: int64
```

```
In [21]: print('-----')
print(df.groupby('Age')['Survived'].sum())
print('-----')
```

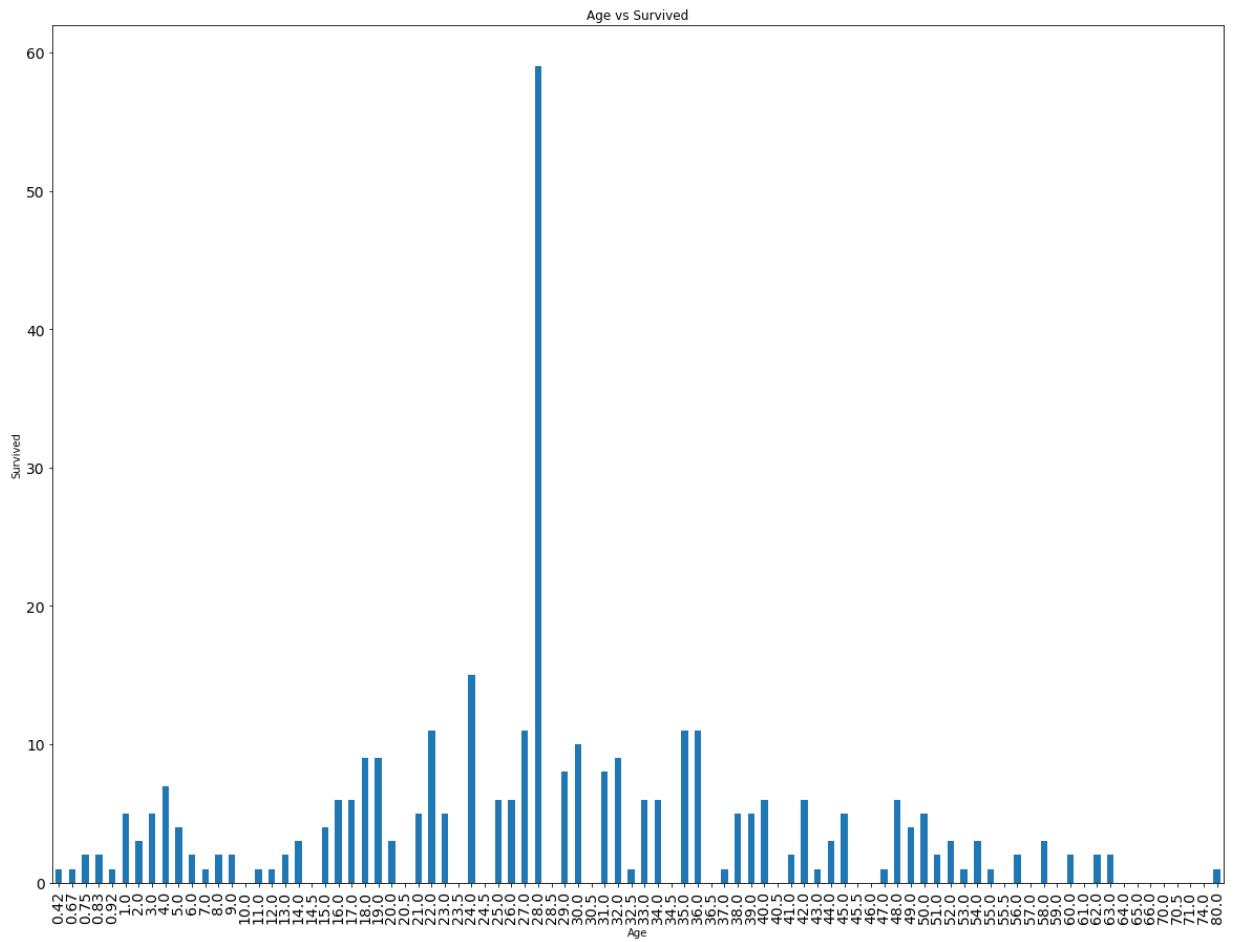
```
-----
Age
0.42      1
0.67      1
0.75      2
0.83      2
0.92      1
..
70.00     0
70.50     0
71.00     0
74.00     0
80.00     1
Name: Survived, Length: 88, dtype: int64
-----
```

```
In [22]: a = df.groupby('Age')['Survived'].sum()
```

```
In [23]: a.plot.pie( fontsize = 12, autopct = '%.1f',figsize = (20,15))  
plt.show()
```



```
In [24]: a.plot.bar(figsize = (20,15), fontsize = 14)
plt.xlabel('Age')
plt.ylabel('Survived')
plt.title('Age vs Survived')
plt.show()
```



**Observation :**Above both plot shows age group of 24 and 36 is more Survived.

## Checking for Imbalance Label

```
In [25]: print('-----')
print('No of Survived present in Titanic :')
print('-----')
print(df['Survived'].value_counts())
print('-----')
```

```
-----
No of Survived present in Titanic :
-----
0    549
1    342
Name: Survived, dtype: int64
-----
```

**Observation : Class is balanced.**

## Label Encoder

```
In [26]: lab_enc = LabelEncoder()
df['Sex'] = lab_enc.fit_transform(df['Sex'])
```

```
In [27]: lab_enc = LabelEncoder()
df['Embarked'] = lab_enc.fit_transform(df['Embarked'])
```

```
In [28]: df.head()
```

Out[28]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	1	22.0	1	0	7.2500	2
1	1	1	0	38.0	1	0	71.2833	0
2	1	3	0	26.0	0	0	7.9250	2
3	1	1	0	35.0	1	0	53.1000	2
4	0	3	1	35.0	0	0	8.0500	2

```
In [29]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 8 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   Survived   891 non-null   int64  
 1   Pclass     891 non-null   int64  
 2   Sex        891 non-null   int32  
 3   Age        891 non-null   float64 
 4   SibSp      891 non-null   int64  
 5   Parch      891 non-null   int64  
 6   Fare        891 non-null   float64 
 7   Embarked   891 non-null   int32  
dtypes: float64(2), int32(2), int64(4)
memory usage: 48.9 KB
```

Converting categorical into numerical value is done.

## Statistic of Dataset

```
In [71]: # We use describe command to extracte statistical infomation about dataset.
df.describe()
```

Out[71]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Emba
<b>count</b>	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000	891.00
<b>mean</b>	0.383838	2.308642	0.647587	29.361582	0.523008	0.381594	32.204208	1.53
<b>std</b>	0.486592	0.836071	0.477990	13.019697	1.102743	0.806057	49.693429	0.79
<b>min</b>	0.000000	1.000000	0.000000	0.420000	0.000000	0.000000	0.000000	0.00
<b>25%</b>	0.000000	2.000000	0.000000	22.000000	0.000000	0.000000	7.910400	1.00
<b>50%</b>	0.000000	3.000000	1.000000	28.000000	0.000000	0.000000	14.454200	2.00
<b>75%</b>	1.000000	3.000000	1.000000	35.000000	1.000000	0.000000	31.000000	2.00
<b>max</b>	1.000000	3.000000	1.000000	80.000000	8.000000	6.000000	512.329200	2.00

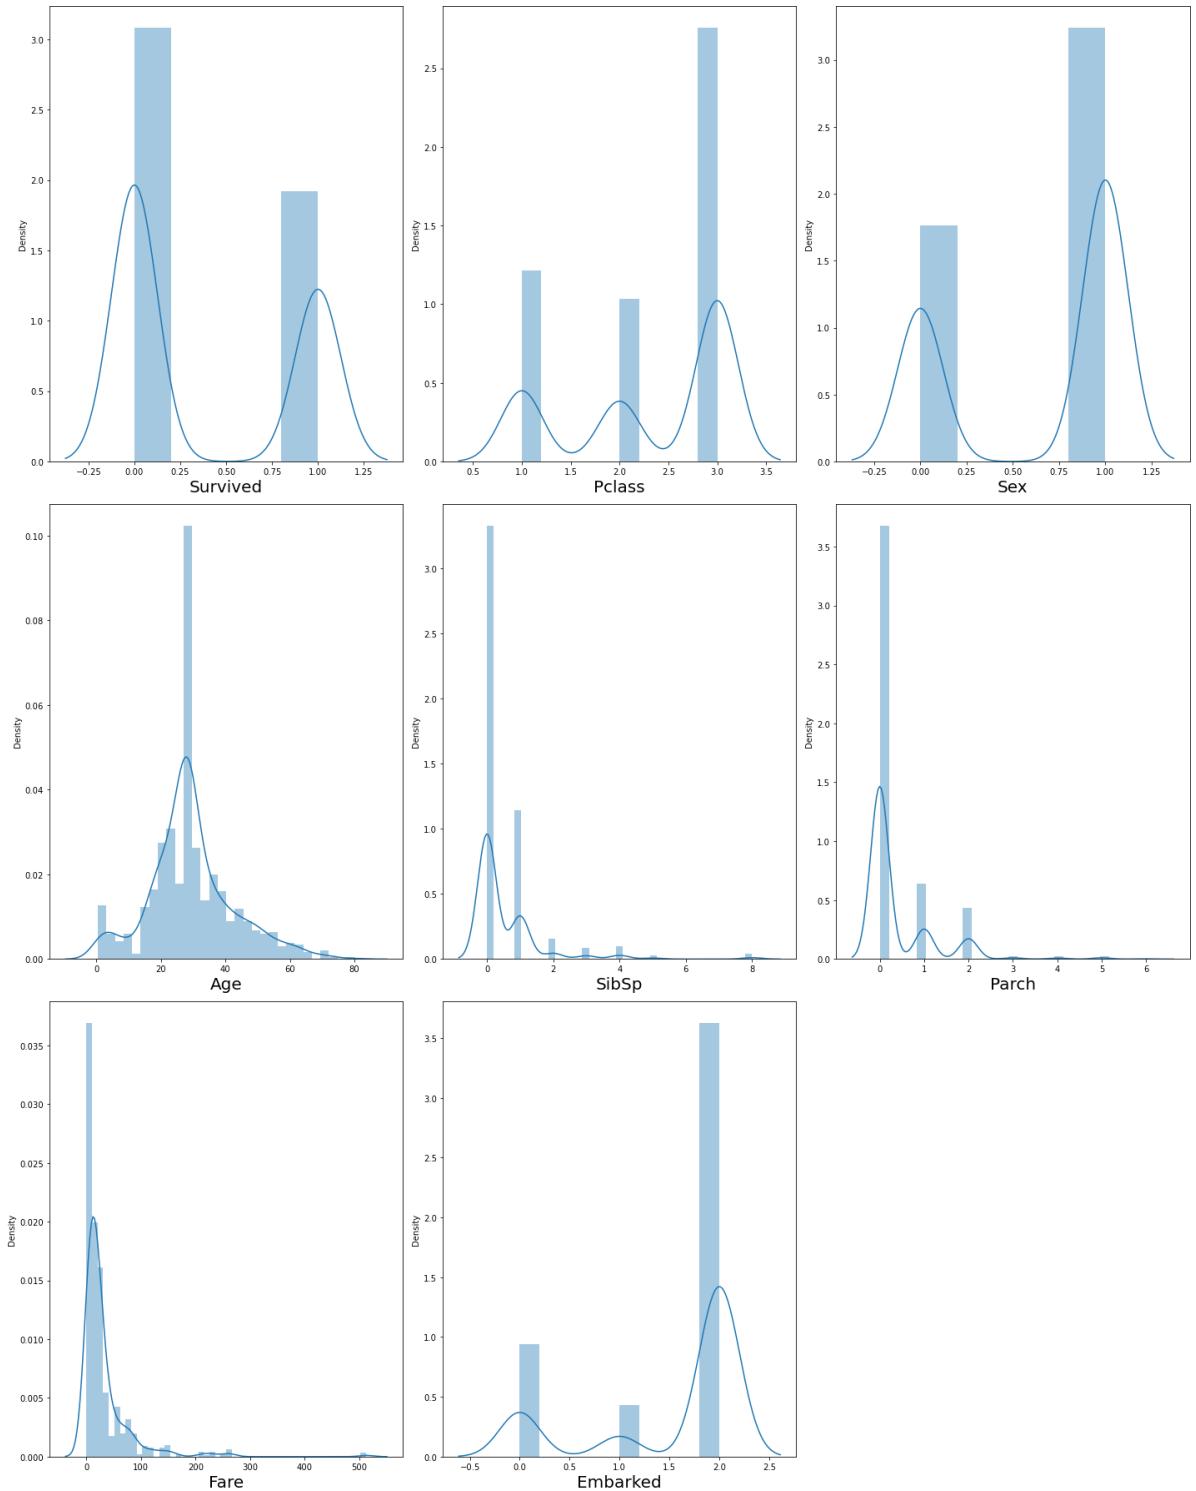
## Plot how data is distributed

```
In [30]: print('-----')
print('Distribution Plot :- ')
print('-----')

plt.figure(figsize = (20,25))
plotnumber = 1

for column in df:
    if plotnumber <=9:
        ax = plt.subplot(3,3, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column, fontsize = 20)
    plotnumber +=1
plt.tight_layout()
```

```
-----
Distribution Plot :-
```

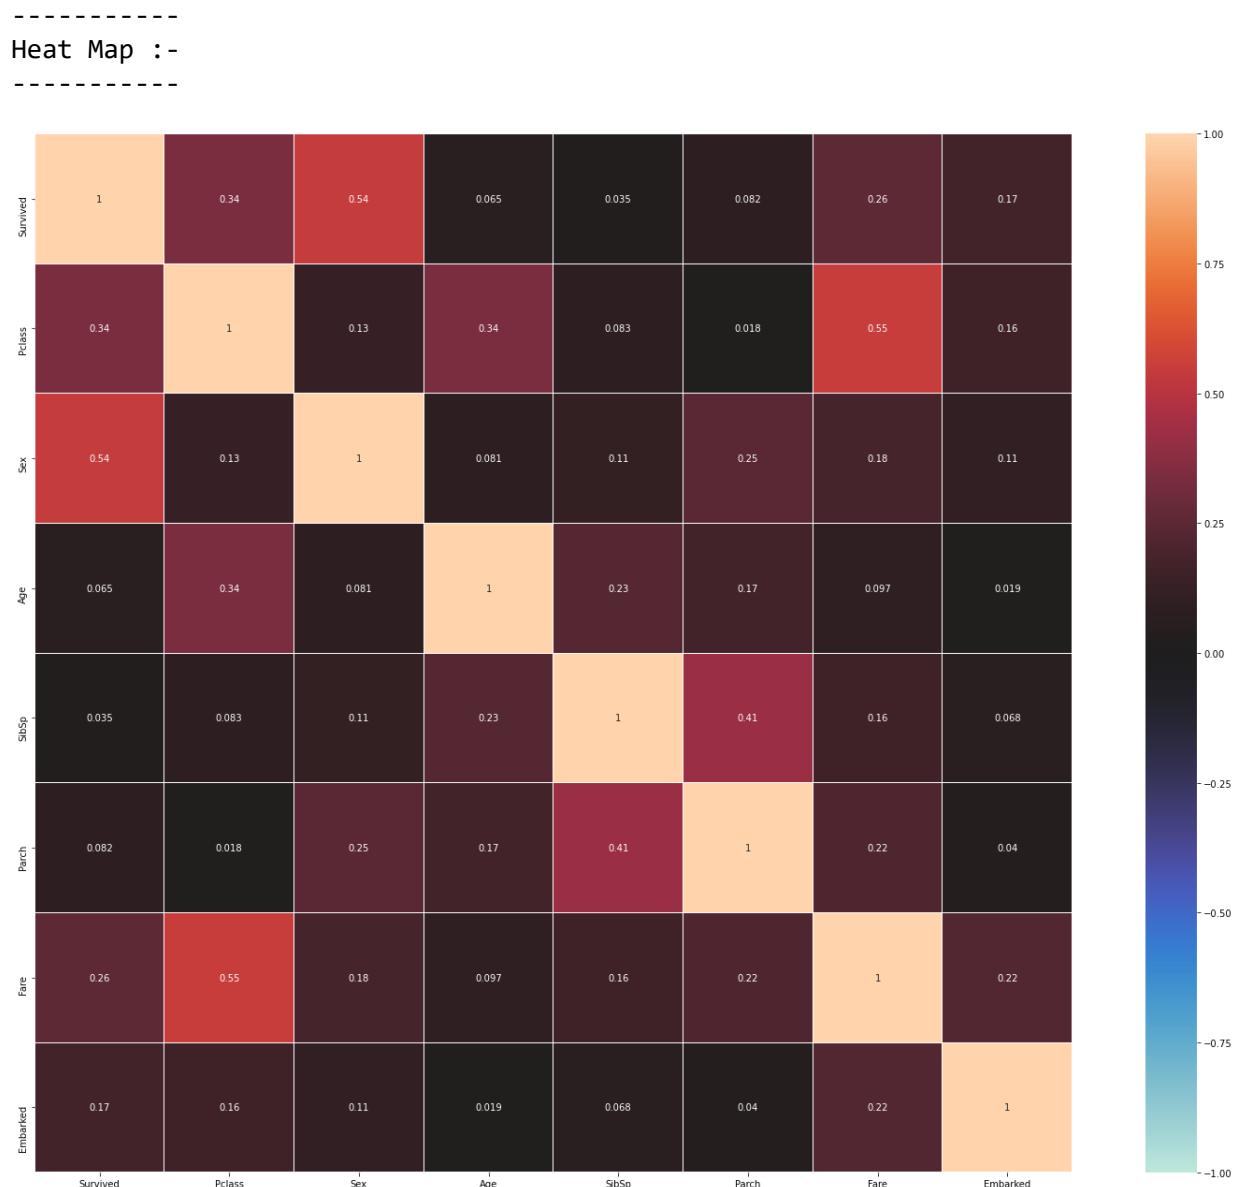


**Observation :**Above both plot shows their is no outliers.

**Corelation of Feature vs Label using Heat map**

```
In [31]: print('-----')
print('Heat Map :-')
print('-----')
df_corr = df.corr().abs()

plt.figure(figsize = (22,16))
sns.heatmap(df_corr, vmin = -1, annot = True, square = True, center = 0, fmt = '.2f')
plt.tight_layout()
```



**Observation : Pclass and Fare show maximum relation.**

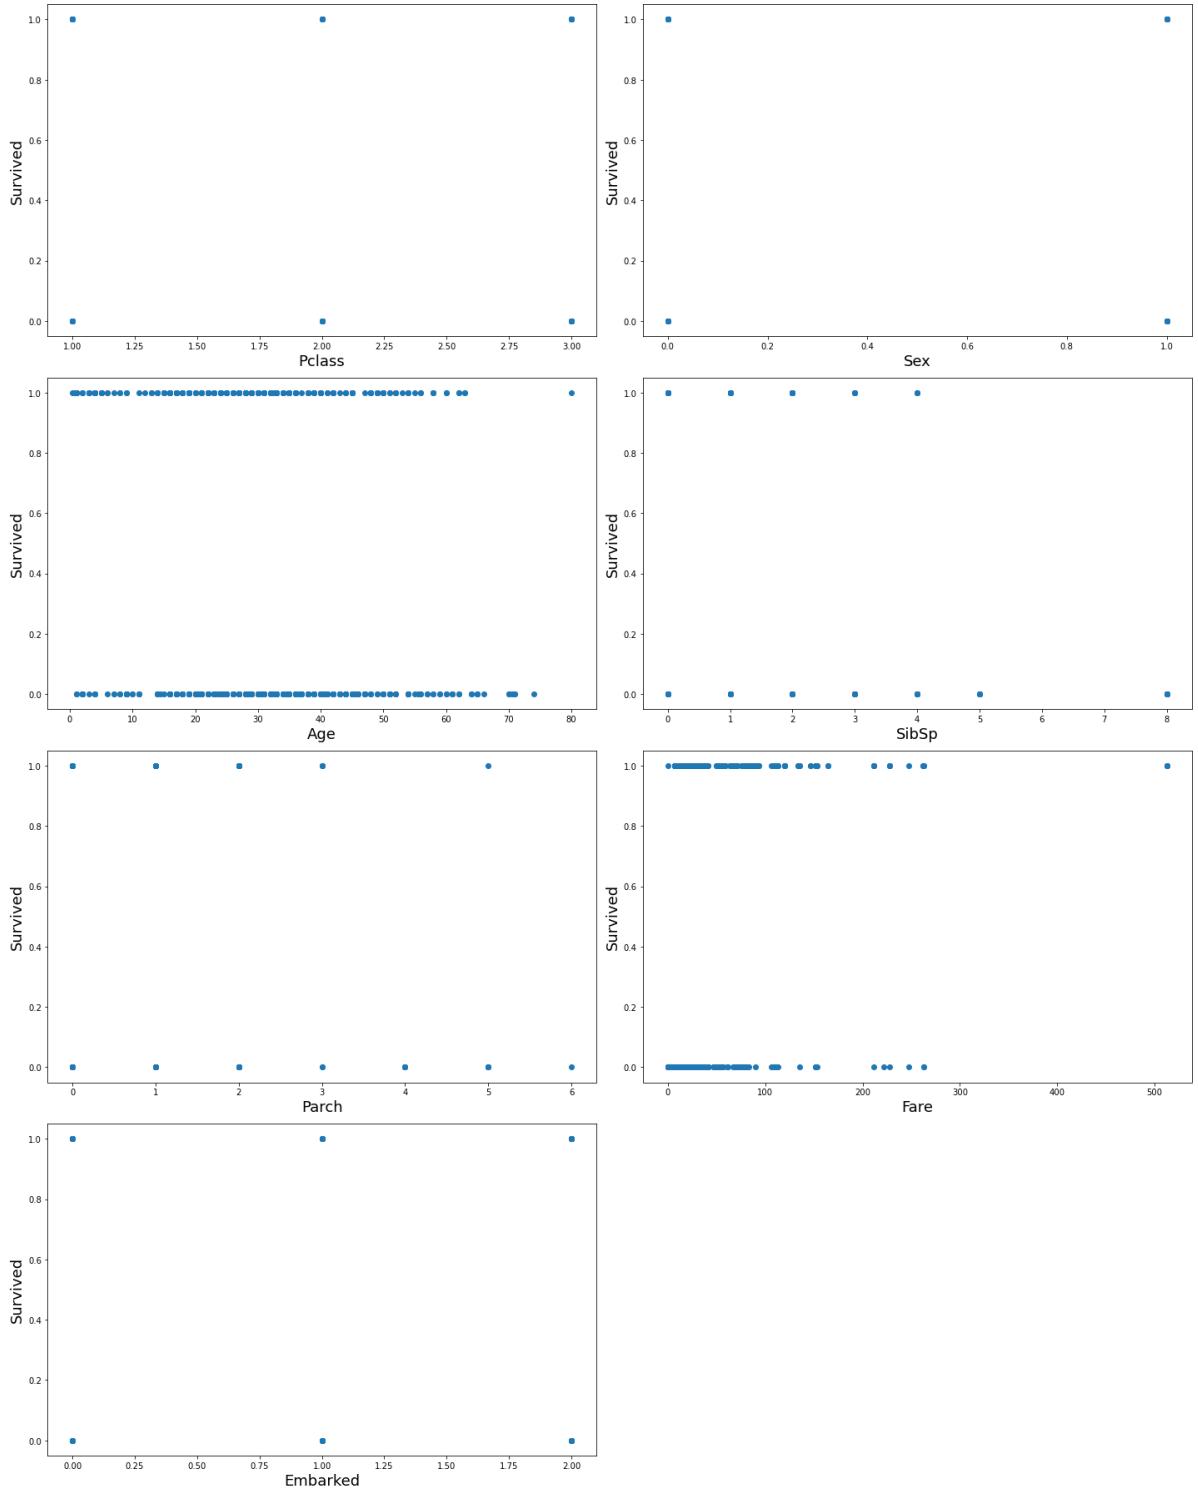
## Splitting Dataset into features and labels

```
In [32]: x = df.drop('Survived', axis = 1)
y = df.Survived
print('Data has been splited')
```

Data has been splited

```
In [33]: # Let's see relation between features and labels.  
print('-----')  
print('Distribution Plot :-')  
print('-----')  
  
plt.figure(figsize = (20,25), facecolor = 'white')  
plotnumber = 1  
for column in x:  
    if plotnumber <=8:  
        ax = plt.subplot(4,2, plotnumber)  
        plt.scatter(x[column],y)  
        plt.xlabel(column, fontsize = 18)  
        plt.ylabel('Survived', fontsize = 18)  
    plotnumber += 1  
plt.tight_layout()
```

-----  
Distribution Plot :-  
-----



**Features are related to class**

**Checking skewnessm**

```
In [34]: x.skew()
```

```
Out[34]: Pclass      -0.630548
Sex         -0.618921
Age         0.510245
SibSp       3.695352
Parch       2.749117
Fare        4.787317
Embarked    -1.264823
dtype: float64
```

Some feature has skewed

## Data Scaling

```
In [35]: scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
Out[35]: array([[ 0.82737724,  0.73769513, -0.56573646, ... , -0.47367361,
   -0.50244517,  0.58595414],
 [-1.56610693, -1.35557354,  0.66386103, ... , -0.47367361,
   0.78684529, -1.9423032 ],
 [ 0.82737724, -1.35557354, -0.25833709, ... , -0.47367361,
   -0.48885426,  0.58595414],
 ... ,
 [ 0.82737724, -1.35557354, -0.1046374 , ... ,  2.00893337,
   -0.17626324,  0.58595414],
 [-1.56610693,  0.73769513, -0.25833709, ... , -0.47367361,
   -0.04438104, -1.9423032 ],
 [ 0.82737724,  0.73769513,  0.20276197, ... , -0.47367361,
   -0.49237783, -0.67817453]])
```

Split data into train and test. Model will be bulit on training data and tested on test data

```
In [36]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25, random_state=42)
print('Data has been splited.')
```

Data has been splited.

## Model Building

Logistic Regression model instantiaing, training and evaluating

```
In [37]: Lr = LogisticRegression()
Lr.fit(x_train, y_train)
y_pred = Lr.predict(x_test)
```

```
In [38]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

-----

Confusion Matrix :

```
[[117  28]
 [ 24  54]]
```

-----

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.81	0.82	145
1	0.66	0.69	0.68	78
accuracy			0.77	223
macro avg	0.74	0.75	0.75	223
weighted avg	0.77	0.77	0.77	223

-----

**Conclusion : Logistic Regression model has 77% score**

**Cross Validation score to check if the model is overfitting**

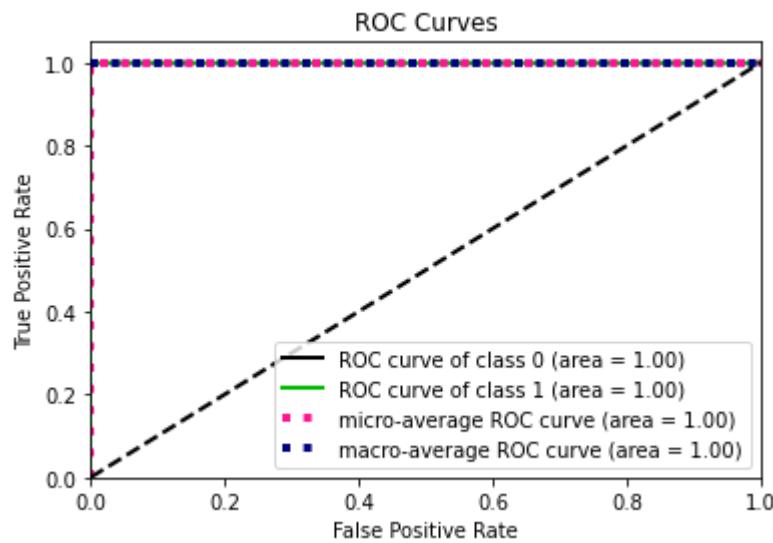
```
In [39]: cv = cross_val_score(Lr, x, y, cv = 5)
print('Cross Validation score of Logistic Regression model --->', cv.mean())
```

```
Cross Validation score of Logistic Regression model ---> 0.7890025735986442
```

**Conclusion : Logistic Regression model has 78% Cross Validation score**

**ROC, AUC Curve**

```
In [40]: prob = Lr.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob)  
plt.show()
```



## Decision Tree model instantiaing, training and evaluating

```
In [41]: DT = DecisionTreeClassifier()  
DT.fit(x_train, y_train)  
y_pred = DT.predict(x_test)
```

```
In [42]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

```
-----  
Confusion Matrix :  
[[106  39]  
 [ 23  55]]  
-----
```

```
Classification Report:  
      precision    recall  f1-score   support  
  
       0          0.82     0.73      0.77      145  
       1          0.59     0.71      0.64       78  
  
accuracy                           0.72      223  
macro avg       0.70     0.72      0.71      223  
weighted avg     0.74     0.72      0.73      223  
-----
```

**Conclusion : Decision Tree model has 72% score .**

**Cross Validation score to check if the model is overfitting**

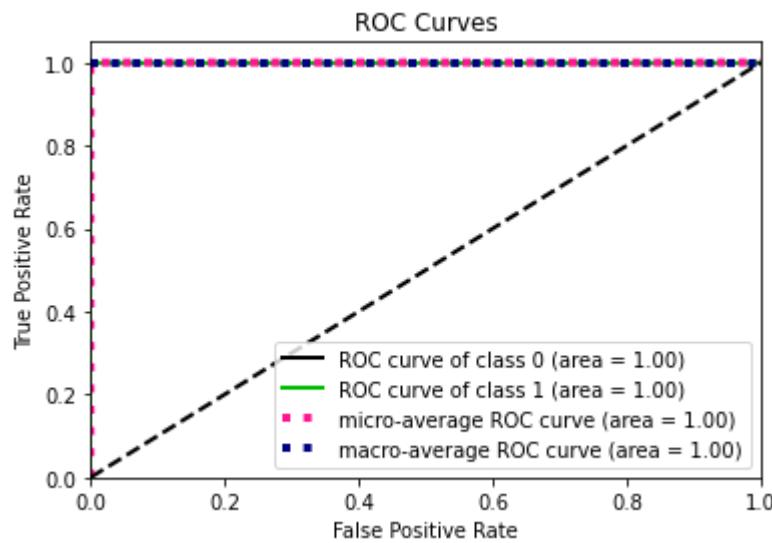
```
In [43]: cv = cross_val_score(DT, x, y, cv = 5)
print('Cross Validation score of Decision Tree model --->', cv.mean())
```

```
Cross Validation score of Decision Tree model ---> 0.7822923859142553
```

**Conclusion : Decision Tree model has 77% Cross Validation score**

**ROC, AUC Curve**

```
In [44]: prob = DT.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob)  
plt.show()
```



## Knn model instantiaing, training and evaluating

```
In [45]: Knn = KNeighborsClassifier()  
Knn.fit(x_train, y_train)  
y_pred = Knn.predict(x_test)
```

```
In [46]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

```
-----  
Confusion Matrix :  
[[110  35]  
 [ 36  42]]  
-----
```

```
Classification Report:  
      precision    recall  f1-score   support  
  
       0          0.75     0.76     0.76      145  
       1          0.55     0.54     0.54       78  
  
accuracy                           0.68      223  
macro avg       0.65     0.65     0.65      223  
weighted avg     0.68     0.68     0.68      223  
-----
```

**Conclusion : Knn model has 68% score**

**Cross Validation score to check if the model is overfitting**

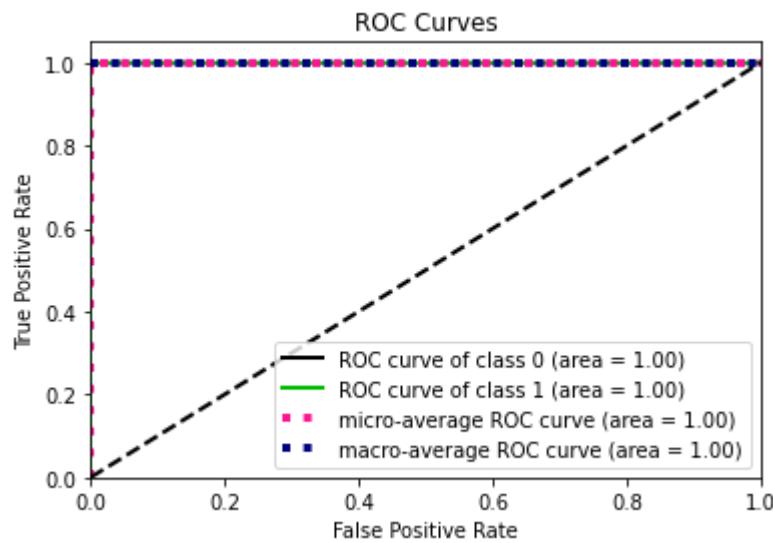
```
In [47]: cv = cross_val_score(Knn, x, y, cv = 5)
print('Cross Validation score of Knn model --->', cv.mean())
```

```
Cross Validation score of Knn model ---> 0.6947649237336011
```

**Conclusion : Knn model has 69% Cross Validation score**

**ROC, AUC Curve**

```
In [48]: prob = Knn.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred, prob)  
plt.show()
```



## Random Forest model instantiating, training and evaluating

```
In [49]: Rn = RandomForestClassifier()  
Rn.fit(x_train, y_train)  
y_pred = Rn.predict(x_test)
```

```
In [50]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

-----

Confusion Matrix :

```
[[111  34]
 [ 24  54]]
```

-----

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.77	0.79	145
1	0.61	0.69	0.65	78
accuracy			0.74	223
macro avg	0.72	0.73	0.72	223
weighted avg	0.75	0.74	0.74	223

-----

**Conclusion : Random Forest model has 75% score**

**Cross Validation score to check if the model is overfitting**

```
In [51]: cv = cross_val_score(Rn, x, y, cv = 5)
print('Cross Validation score of Random Forest model --->', cv.mean())
```

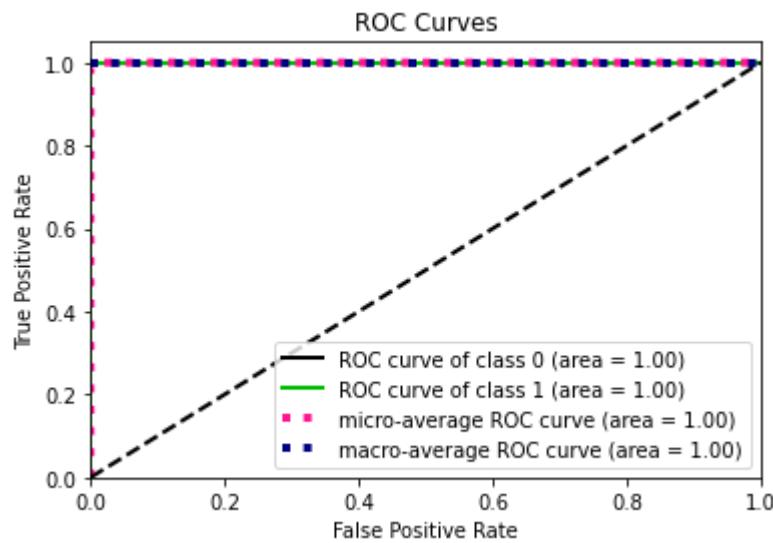
Cross Validation score of Random Forest model ---> 0.807005209967987

**Conclusion : Random Forest model has 80% Cross Validation score**

.

**ROC, AUC Curve**

```
In [52]: prob = Rn.predict_proba(x_test) # calculating probability  
skplt.metrics.plot_roc(y_pred,prob)  
plt.show()
```



## SVM model instantiating, training and evaluating

```
In [53]: svc = SVC(probability=True)  
svc.fit(x_train, y_train)  
y_pred = svc.predict(x_test)
```

```
In [54]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

-----

Confusion Matrix :

```
[[126  19]
 [ 57  21]]
```

-----

Classification Report:

	precision	recall	f1-score	support
0	0.69	0.87	0.77	145
1	0.53	0.27	0.36	78
accuracy			0.66	223
macro avg	0.61	0.57	0.56	223
weighted avg	0.63	0.66	0.62	223

-----

**Conclusion : SVM model has 66% score**

**Cross Validation score to check if the model is overfitting**

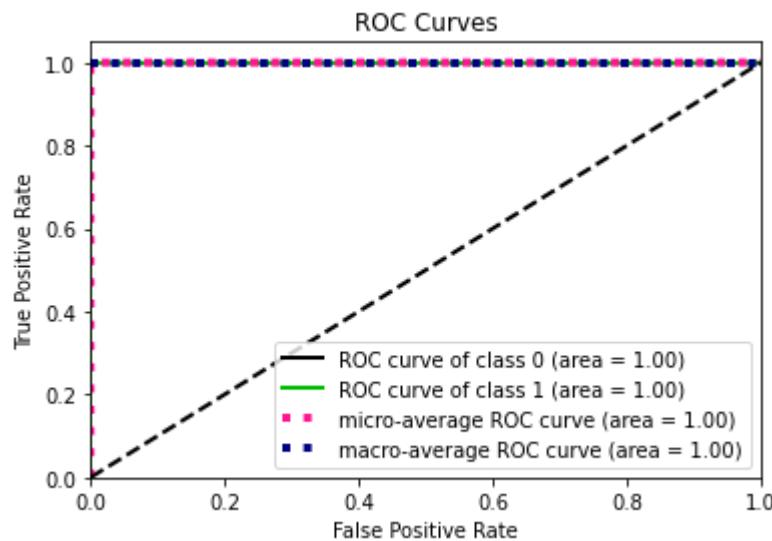
```
In [55]: cv = cross_val_score(svc, x, y, cv = 5)
print('Cross Validation score of svc model --->', cv.mean())
```

```
Cross Validation score of svc model ---> 0.674615529470843
```

**Conclusion : SVM model has 67% Cross Validation score**

**ROC, AUC Curve**

```
In [56]: prob = svc.predict_proba(x_test) # calculating probability
skplt.metrics.plot_roc(y_pred, prob)
plt.show()
```



## Let's find ROC, AUC score

```
In [57]: # LogisticRegression
roc_auc_score(y_test, Lr.predict(x_test))
```

Out[57]: 0.7496021220159151

```
In [58]: # DecisionTreeClassifier
roc_auc_score(y_test, DT.predict(x_test))
```

Out[58]: 0.718081343943413

```
In [59]: # RandomForestClassifier
roc_auc_score(y_test, Rn.predict(x_test))
```

Out[59]: 0.7289124668435013

```
In [60]: # KNeighborsClassifier
roc_auc_score(y_test, Knn.predict(x_test))
```

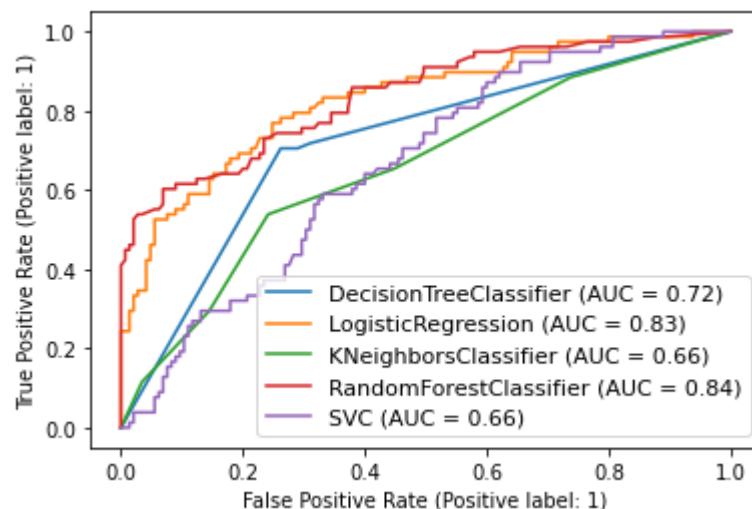
Out[60]: 0.6485411140583554

```
In [61]: # SVMClassifier  
roc_auc_score(y_test, svc.predict(x_test))
```

```
Out[61]: 0.5690981432360742
```

## Let's check ROC, AUC Curve for the fitted model

```
In [62]: disp = plot_roc_curve(DT, x_test, y_test)  
plot_roc_curve(Lr, x_test, y_test, ax = disp.ax_) # ax_ = Axes with confusion mat  
plot_roc_curve(Knn, x_test, y_test, ax = disp.ax_)  
plot_roc_curve(Rn, x_test, y_test, ax = disp.ax_)  
plot_roc_curve(svc, x_test, y_test, ax = disp.ax_)  
plt.legend(prop = {'size':11}, loc = 'lower right')  
plt.show()
```



Looking CV score we found Logistic Regression has best model so we do Hyperparameter Tuning on it.

```
In [63]: param = { 'max_iter':[100], 'penalty': ['l1','l2'], 'fit_intercept':[True,False]}
```

```
In [64]: grid_search = GridSearchCV(estimator = Lr, param_grid = param, cv = 5,n_jobs = -1)
```

```
In [65]: grid_search.fit(x_train, y_train)
```

```
Out[65]: GridSearchCV(cv=5, estimator=LogisticRegression(), n_jobs=-1,  
param_grid={'fit_intercept': [True, False], 'max_iter': [100],  
'penalty': ['l1', 'l2']})
```

```
In [66]: best_parameters = grid_search.best_params_  
print(best_parameters)
```

```
{'fit_intercept': False, 'max_iter': 100, 'penalty': 'l2'}
```

```
In [67]: hlr = LogisticRegression(max_iter = 100, penalty = 'l2')
hlr.fit(x_train, y_train)
hlr.score(x_test, y_test)
```

```
Out[67]: 0.7668161434977578
```

```
In [68]: y_pred = hlr.predict(x_test)
```

```
In [69]: print('-----\n')
print('Confusion Matrix :')
cfm = confusion_matrix(y_test, y_pred)
print(cfm)
print('\n-----')
print('\nClassification Report:')
print(classification_report(y_test, y_pred, digits = 2))
print('-----')
```

```
Confusion Matrix :
[[117  28]
 [ 24  54]]
```

```
Classification Report:
      precision    recall  f1-score   support

          0       0.83     0.81     0.82      145
          1       0.66     0.69     0.68       78

   accuracy                           0.77      223
    macro avg       0.74     0.75     0.75      223
weighted avg       0.77     0.77     0.77      223
```

After Hyperparameter Tuning model accuracy score 77%

## Saving The Model

```
In [70]: # saving the model to the Local file system
filename = 'Titanic project.pickle'
pickle.dump(hlr, open(filename, 'wb'))
```

Final Conclusion : Logistic Regression is our best model.

In [ ]: